

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>

<https://github.com/tensorflow/tfjs-models>

README.md

## Pre-trained TensorFlow.js models

This repository hosts a set of pre-trained models that have been ported to TensorFlow.js.

The models are hosted on NPM and unpkg so they can be used in any project out of the box. They can be used directly or used in a transfer learning setting with TensorFlow.js.

To find out about APIs for models, look at the README in each of the respective directories. In general, we try to hide tensors so the API can be used by non-machine learning experts.

For those intested in contributing a model, please file a [GitHub issue on tfjs](#) to gauge interest. We are trying to add models that complement the existing set of models and can be used as building blocks in other apps.

## Models

### Image

- [MobileNet](#) - Classify images with labels from the [ImageNet](#) database.
  - `npm i @tensorflow-models/mobilenet`
- [PoseNet](#) - Realtime pose detection. Blog post [here](#).
  - `npm i @tensorflow-models/posenet`
- [Coco SSD](#) - Object detection based on the [TensorFlow object detection API](#).
  - `npm i @tensorflow-models/coco-ssd`

### Audio

- [Speech commands](#) - Classify 1 second audio snippets from the [speech commands dataset](#).
  - `npm i @tensorflow-models/speech-commands`

### General utilities

- [KNN Classifier](#) - Create a custom k-nearest neighbors classifier. Can be used for transfer learning.
  - `npm i @tensorflow-models/knn-classifier`

## Toxicity classifier

The toxicity model detects whether text contains toxic content such as threatening language, insults, obscenities, identity-based hate, or sexually explicit language. The model was trained on the civil comments dataset:

[https://figshare.com/articles/data\\_json/7376747](https://figshare.com/articles/data_json/7376747) which contains ~2 million comments labeled for toxicity. The model is built on top of the Universal Sentence Encoder ([Cer et al., 2018](#)).

More information about how the toxicity labels were calibrated can be found [here](#).

text	identity attack	insult	obscene	severe toxicity	sexual explicit	threat	toxicity
We're dudes on computers, moron. You are quite astonishingly stupid.	false	true	false	false	false	false	true
Please stop. If you continue to vandalize Wikipedia, as you did to Kmart, you will be blocked from editing.	false	false	false	false	false	false	false
I respect your point of view, and when this discussion originated on 8th April I would have tended to agree with you.	false	false	false	false	false	false	false

Enter text below and click 'Classify' to add it to the table.

i.e. 'you suck'









CLASSIFY

Check out our [demo](#), which uses the toxicity model to predict the toxicity of several sentences taken from this [Kaggle dataset](#). Users can also input their own text for classification.









```
<script src="https://cdn.jsdelivr.net/npm/@tensorflow/tfjs@latest"></script>  
<script src="https://cdn.jsdelivr.net/npm/@tensorflow-models/toxicity"></script>
```

```
<script src="https://cdn.jsdelivr.net/npm/@tensorflow/tfjs@latest"></script>  
<script src="https://cdn.jsdelivr.net/npm/@tensorflow-models/toxicity"></script>
```

<https://github.com/tensorflow/tfjs-models>

 <a href="#">body-pix</a>	Fix typo in source code (#183)	5 days ago
 <a href="#">coco-ssd</a>	Update bodypix, cocossd, knn-classifier, posenet to depend on tfjs 1.0 (	a month ago
 <a href="#">knn-classifier</a>	Update bodypix, cocossd, knn-classifier, posenet to depend on tfjs 1.0 (	a month ago
 <a href="#">mobilenet</a>	Update versions of tfjs and mobilenet in the example code (#174)	18 days ago
 <a href="#">posenet</a>	Update bodypix, cocossd, knn-classifier, posenet to depend on tfjs 1.0 (	a month ago
 <a href="#">speech-commands</a>	[speech-commands] Fix incorrect metadata field for word labels; v0.3.4 (	an hour ago
 <a href="#">toxicity</a>	Update toxicity demo per reviewer feedback. (#172)	22 days ago
 <a href="#">universal-sentence-encoder</a>	Depend on tfjs 1.0 in USE. (#164)	a month ago

<https://github.com/tensorflow/tfjs-models>

 <a href="#">body-pix</a>	Fix typo in source code (#183)	5 days ago
 <a href="#">coco-ssd</a>	Update bodypix, cocossd, knn-classifier, posenet to depend on tfjs 1.0 (	a month ago
 <a href="#">knn-classifier</a>	Update bodypix, cocossd, knn-classifier, posenet to depend on tfjs 1.0 (	a month ago
 <a href="#">mobilenet</a>	Update versions of tfjs and mobilenet in the example code (#174)	18 days ago
 <a href="#">posenet</a>	Update bodypix, cocossd, knn-classifier, posenet to depend on tfjs 1.0 (	a month ago
 <a href="#">speech-commands</a>	[speech-commands] Fix incorrect metadata field for word labels; v0.3.4 (	an hour ago
 <a href="#">toxicity</a>	Update toxicity demo per reviewer feedback. (#172)	22 days ago
 <a href="#">universal-sentence-encoder</a>	Depend on tfjs 1.0 in USE. (#164)	a month ago

```
<html>
<head>
<script src="https://cdn.jsdelivr.net/npm/@tensorflow/tfjs@latest"></script>
<script src="https://cdn.jsdelivr.net/npm/@tensorflow-models/toxicity"></script>
</head>
<body></body>
</html>
```



```
<html>
<head>
<script src="https://cdn.jsdelivr.net/npm/@tensorflow/tfjs@latest"></script>
<script src="https://cdn.jsdelivr.net/npm/@tensorflow-models/toxicity"></script>
<script>
// Your code here
</script>
</head>
<body></body>
</html>
```

```
const threshold = 0.9;
```

Value for not-insult

```
const threshold = 0.9;
```

```
"label": "insult",
```

```
"results": [{
```

```
  "probabilities": [0.9187529683113098, 0.08124706149101257],
```

```
  "match": false
```

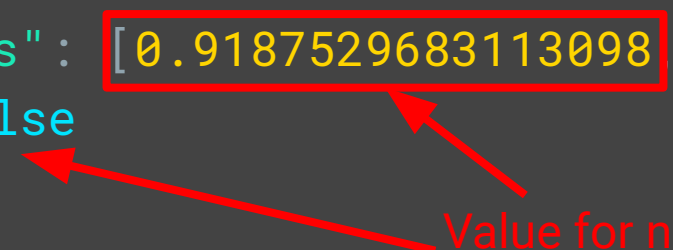
```
}]
```

Value for insult

```
const threshold = 0.9;
```

```
"label": "insult",  
  "results": [{  
    "probabilities": [0.9187529683113098 0.08124706149101257],  
    "match": false  
  }]
```

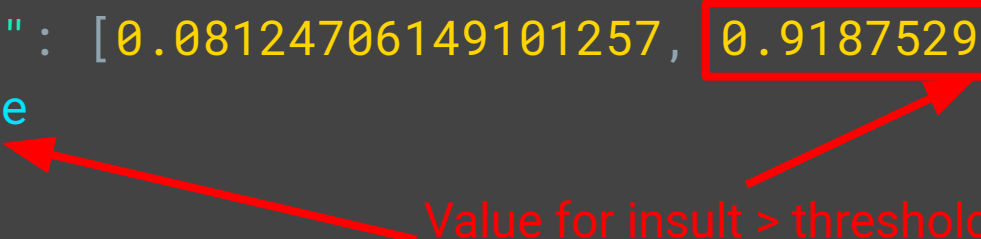
Value for not-insult > threshold  
So match = false



```
const threshold = 0.9;
```

```
"label": "insult",  
  "results": [{  
    "probabilities": [0.08124706149101257, 0.9187529683113098],  
    "match": true  
  }]
```

Value for insult > threshold  
So match = true



```
const threshold = 0.9;
```

```
"label": "insult",  
  "results": [{  
    "probabilities": [0.5, 0.5],  
    "match": null  
  }]
```

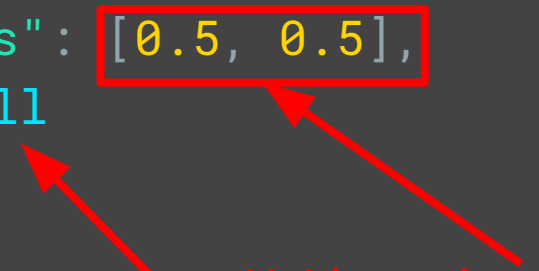


Diagram illustrating the logic flow: Two red arrows originate from the array `[0.5, 0.5]` in the JSON object. One arrow points to the `"match": null` line, and the other points to the explanatory text below.

Neither value > threshold  
So match = null

```
toxicity.load(threshold).then(model => {  
  const sentences = ['you suck!'];  
  model.classify(sentences).then(predictions => {  
    // Handle Results  
  });  
});
```

```
toxicity.load(threshold).then(model => {  
  const sentences = ['you suck!'];  
  model.classify(sentences).then(predictions => {  
    // Handle Results  
  });  
});
```



```
toxicity.load(threshold).then(model => {  
  const sentences = ['you suck!'];  
  model.classify(sentences).then(predictions => {  
    // Handle Results  
  });  
});
```

```
toxicity.load(threshold).then(model => {  
  const sentences = ['you suck!'];  
  model.classify(sentences).then(predictions => {  
    // Handle Results  
  });  
});
```

```
toxicity.load(threshold).then(model => {  
  const sentences = ['you suck!'];  
  model.classify(sentences).then(predictions => {  
    // Handle Results  
  });  
});
```

```
toxicity.load(threshold).then(model => {  
  const sentences = ['you suck!'];  
  model.classify(sentences).then(predictions => {  
    // Handle Results  
  });  
});
```

```
▼ (7) [{...}, {...}, {...}, {...}, {...}, {...}, {...}] ⓘ  
  ► 0: {label: "identity_attack", results: Array(1)}  
  ► 1: {label: "insult", results: Array(1)}  
  ► 2: {label: "obscene", results: Array(1)}  
  ► 3: {label: "severe_toxicity", results: Array(1)}  
  ► 4: {label: "sexual_explicit", results: Array(1)}  
  ► 5: {label: "threat", results: Array(1)}  
  ► 6: {label: "toxicity", results: Array(1)}  
    length: 7  
  ► __proto__: Array(0)
```

```
▼ 1:
  label: "insult"
  ▼ results: Array(1)
    ▼ 0:
      match: true
      ► probabilities: Float32Array(2) [0.05890671908855438, 0.94109326601028...]
      ► __proto__: Object
      length: 1
      ► __proto__: Array(0)
      ► proto : Object
```

```
predictions[1].label;
predictions[1].results[0].probabilities[0];
predictions[1].results[0].probabilities[1];
predictions[1].results[0].match;
```

The diagram consists of four red arrows pointing from code attributes to JSON object keys:

- From `predictions[1].label;` to `"label": "insult",`
- From `predictions[1].results[0].probabilities[0];` to `"results": [{`
- From `predictions[1].results[0].probabilities[1];` to `"probabilities": [0.5, 0.5],`
- From `predictions[1].results[0].match;` to `"match": null`


```
    "label": "insult",
    "results": [{
      "probabilities": [0.5, 0.5],
      "match": null
    }]
```

```
for(i=0; i<7; i++){  
    if(predictions[i].results[0].match){  
        console.log(predictions[i].label + " was found with a probability of " +  
                    predictions[i].results[0].probabilities[1]);  
    }  
}
```



insult was found with a probability of 0.9410932660102844  
toxicity was found with a probability of 0.9766321778297424

```
predictions[1].label;  
predictions[1].results[n].probabilities[0];  
predictions[1].results[n].probabilities[1];  
predictions[1].results[n].match;
```



<http://bit.ly/mobilenet-labels>

- 00: background
- 01: tench
- 02: goldfish
- 03: great white shark
- 04: tiger shark
- 05: hammerhead
- 06: electric ray
- 07: stingray
- 08: cock
- 09: hen
- 10: ostrich
- 11: brambling
- 12: goldfinch
- 13: house finch
- ...

```
<script src="https://cdn.jsdelivr.net/npm/@tensorflow/tfjs@1.0.1"> </script>  
<script src="https://cdn.jsdelivr.net/npm/@tensorflow-models/mobilenet@1.0.0"> </script>
```

```
<body>  
  </img>  
  <div id="output" style="font-family:courier;font-size:24px;height=300px"></div>  
</body>
```

```
<body>  
  </img>  
  <div id="output" style="font-family:courier;font-size:24px;height=300px"></div>  
</body>
```

```
const img = document.getElementById('img');  
mobilenet.load().then(model => {  
  model.classify(img).then(predictions => {  
    console.log(predictions);  
  });  
});
```



```
const img = document.getElementById('img');  
mobilenet.load().then(model => {  
  model.classify(img).then(predictions => {  
    console.log(predictions);  
  });  
});
```

```
const img = document.getElementById('img');  
mobilenet.load().then(model => {  
  model.classify(img).then(predictions => {  
    console.log(predictions);  
  });  
});
```

```
const img = document.getElementById('img');  
mobilenet.load().then(model => {  
  model.classify(img).then(predictions => {  
    console.log(predictions);  
  });  
});
```

```
const img = document.getElementById('img');  
mobilenet.load().then(model => {  
  model.classify(img).then(predictions => {  
    console.log(predictions);  
  });  
});
```

```
const img = document.getElementById('img');  
mobilenet.load().then(model => {  
  model.classify(img).then(predictions => {  
    console.log(predictions);  
  });  
});
```

```
const img = document.getElementById('img');  
mobilenet.load().then(model => {  
  model.classify(img).then(predictions => {  
    console.log(predictions);  
  });  
});
```



Elements Console Sources Network Performance Memory

Page Filesystem >> tfjs@1.0.1 x

top  
127.0.0.1:8887  
mobilenet.html  
coffee.jpg  
cdn.jsdelivr.net

Pretty-print this minified file?  
Source Map detected.

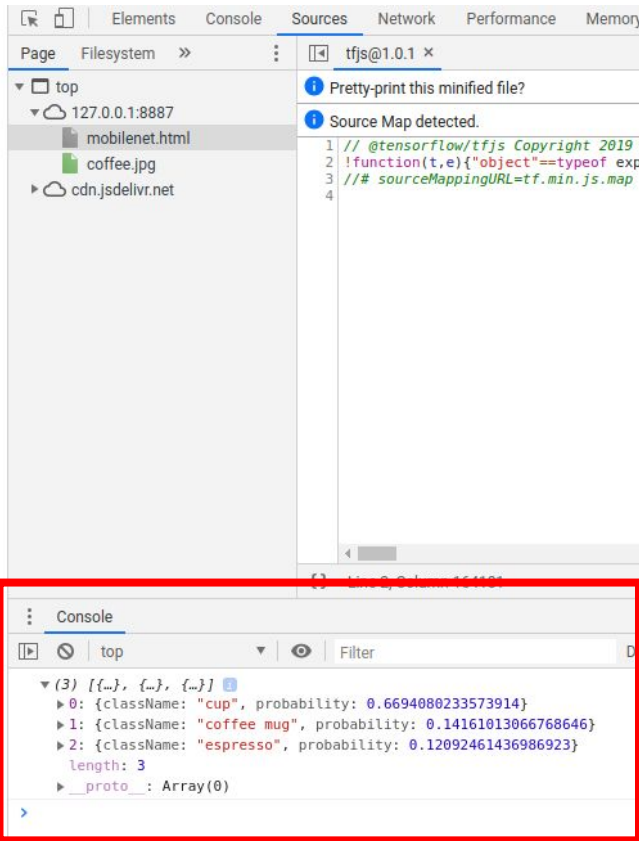
```
1 // @tensorflow/tfjs Copyright 2019  
2 !function(t,e){"object"===typeof exp  
3   //# sourceMappingURL=tf.min.js.map  
4
```

{ } Line 2, Column 164181

Console

top

(3) [{...}, {...}, {...}]  
 0: {className: "cup", probability: 0.6694080233573914}  
 1: {className: "coffee mug", probability: 0.14161013066768646}  
 2: {className: "espresso", probability: 0.12092461436986923}  
 length: 3  
 \_\_proto\_\_: Array(0)





▼ (3) [{...}, {...}, {...}] ⓘ

- ▶ 0: {className: "cup", probability: 0.6694080233573914}
  - ▶ 1: {className: "coffee mug", probability: 0.14161013066768646}
  - ▶ 2: {className: "espresso", probability: 0.12092461436986923}
- length: 3
- ▶ \_\_proto\_\_: Array(0)
-

```
const img = document.getElementById('img');
const outp = document.getElementById('output');
mobilenet.load().then(model => {
  model.classify(img).then(predictions => {
    console.log(predictions);
    for(var i = 0; i<predictions.length; i++){
      outp.innerHTML += "<br/>" + predictions[i].className
                      + " : " + predictions[i].probability;
    }
  });
});
```

```
const img = document.getElementById('img');  
const outp = document.getElementById('output');  
mobilenet.load().then(model => {  
  model.classify(img).then(predictions => {  
    console.log(predictions);  
    for(var i = 0; i<predictions.length; i++){  
      outp.innerHTML += "<br/>" + predictions[i].className  
        + " : " + predictions[i].probability;  
    }  
  });  
});
```

```
const img = document.getElementById('img');
const outp = document.getElementById('output');
mobilenet.load().then(model => {
  model.classify(img).then(predictions => {
    console.log(predictions);
    for(var i = 0; i<predictions.length; i++){
      outp.innerHTML += "<br/>" + predictions[i].className
        + " : " + predictions[i].probability;
    }
  });
});
```

```
const img = document.getElementById('img');
const outp = document.getElementById('output');
mobilenet.load().then(model => {
  model.classify(img).then(predictions => {
    console.log(predictions);
    for(var i = 0; i<predictions.length; i++){
      outp.innerHTML += "<br/>" + predictions[i].className
        + " : " + predictions[i].probability;
    }
  });
});
```



tiger cat : 0.44101303815841675

leopard, *Panthera pardus* : 0.18026068806648254

Madagascar cat, ring-tailed lemur, *Lemur catta* : 0.14320671558380127

```
!pip install tensorflowjs
```

```
import numpy as np
import tensorflow as tf
from tensorflow import keras
print(tf.__version__)
model = tf.keras.Sequential([keras.layers.Dense(units=1, input_shape=[1])])
model.compile(optimizer='sgd', loss='mean_squared_error')
xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)
model.fit(xs, ys, epochs=500)
```



```
print(model.predict([10.0]))
```

```
import time
saved_model_path = "/tmp/saved_models/{}".format(int(time.time()))

# For TensorFlow 2.0 use this:
# tf.keras.experimental.export_saved_model(model, saved_model_path)

# For TensorFlow 1.x use this:
tf.contrib.saved_model.save_keras_model(model, saved_model_path)
```

```
import time
saved_model_path = "/tmp/saved_models/{}".format(int(time.time()))
```

```
# For TensorFlow 2.0 use this:
```

```
# tf.keras.experimental.export_saved_model(model, saved_model_path)
```

```
# For TensorFlow 1.x use this:
```

```
tf.contrib.saved_model.save_keras_model(model, saved_model_path)
```

```
import time
saved_model_path = "/tmp/saved_models/{}".format(int(time.time()))

# For TensorFlow 2.0 use this:
# tf.keras.experimental.export_saved_model(model, saved_model_path)

# For TensorFlow 1.x use this:
tf.contrib.saved_model.save_keras_model(model, saved_model_path)
```

```
import time
saved_model_path = "/tmp/saved_models/{}".format(int(time.time()))

# For TensorFlow 2.0 use this:
# tf.keras.experimental.export_saved_model(model, saved_model_path)

# For TensorFlow 1.x use this:
tf.contrib.saved_model.save_keras_model(model, saved_model_path)
```

```
INFO:tensorflow:SavedModel written to /tmp/saved_models/1554528640/1554528642/saved_model.pb  
b'/tmp/saved_models/1554528640/1554528642'
```

```
!tensorflowjs_converter \  
  --input_format=keras_saved_model \  
  /tmp/saved_models/1554528640/1554528642 \  
  /tmp/linear
```

```
!tensorflowjs_converter \  
  --input_format=keras_saved_model \  
  /tmp/saved_models/1554528640/1554528642 \  
  /tmp/linear
```





```
!tensorflowjs_converter \  
  --input_format=keras_saved_model \  
  /tmp/saved_models/1554528640/1554528642 \  
  /tmp/linear
```

```
!tensorflowjs_converter \  
  --input_format=keras_saved_model \  
  /tmp/saved_models/1554528640/1554528642 \  
  /tmp/linear
```

▼  tmp

▼  linear

 group1-shard1of1.bin

 model.json

```
<html>
<head>
<script src="https://cdn.jsdelivr.net/npm/@tensorflow/tfjs@latest"></script>
<script>
  async function run(){
    const MODEL_URL = 'http://127.0.0.1:8887/model.json';
    const model = await tf.loadLayersModel(MODEL_URL);
    console.log(model.summary());
    const input = tf.tensor2d([10.0], [1, 1]);
    const result = model.predict(input);
    alert(result);
  }
  run();
</script>
<body>
</body>
</html>
```

```
<html>
<head>
<script src="https://cdn.jsdelivr.net/npm/@tensorflow/tfjs@latest"></script>
<script>
  async function run(){
    const MODEL_URL = 'http://127.0.0.1:8887/model.json';
    const model = await tf.loadLayersModel(MODEL_URL);
    console.log(model.summary());
    const input = tf.tensor2d([10.0], [1, 1]);
    const result = model.predict(input);
    alert(result);
  }
  run();
</script>
<body>
</body>
</html>
```

```
<html>
<head>
<script src="https://cdn.jsdelivr.net/npm/@tensorflow/tfjs@latest"></script>
<script>
  async function run(){
    const MODEL_URL = 'http://127.0.0.1:8887/model.json';
    const model = await tf.loadLayersModel(MODEL_URL);
    console.log(model.summary());
    const input = tf.tensor2d([10.0], [1, 1]);
    const result = model.predict(input);
    alert(result);
  }
  run();
</script>
<body>
</body>
</html>
```

```
<html>
<head>
<script src="https://cdn.jsdelivr.net/npm/@tensorflow/tfjs@latest"></script>
<script>
  async function run(){
    const MODEL_URL = 'http://127.0.0.1:8887/model.json';
    const model = await tf.loadLayersModel(MODEL_URL);
    console.log(model.summary());
    const input = tf.tensor2d([10.0], [1, 1]);
    const result = model.predict(input);
    alert(result);
  }
  run();
</script>
<body>
</body>
</html>
```

```
<html>
<head>
<script src="https://cdn.jsdelivr.net/npm/@tensorflow/tfjs@latest"></script>
<script>
  async function run(){
    const MODEL_URL = 'http://127.0.0.1:8887/model.json';
    const model = await tf.loadLayersModel(MODEL_URL);
    console.log(model.summary());
    const input = tf.tensor2d([10.0], [1, 1]);
    const result = model.predict(input);
    alert(result);
  }
  run();
</script>
<body>
</body>
</html>
```



```
<html>
<head>
<script src="https://cdn.jsdelivr.net/npm/@tensorflow/tfjs@latest"></script>
<script>
  async function run(){
    const MODEL_URL = 'http://127.0.0.1:8887/model.json';
    const model = await tf.loadLayersModel(MODEL_URL);
    console.log(model.summary());
    const input = tf.tensor2d([10.0], [1, 1]);
    const result = model.predict(input);
    alert(result);
  }
  run();
</script>
<body>
</body>
</html>
```

```
<html>
<head>
<script src="https://cdn.jsdelivr.net/npm/@tensorflow/tfjs@latest"></script>
<script>
  async function run(){
    const MODEL_URL = 'http://127.0.0.1:8887/model.json';
    const model = await tf.loadLayersModel(MODEL_URL);
    console.log(model.summary());
    const input = tf.tensor2d([10.0], [1, 1]);
    const result = model.predict(input);
    alert(result);
  }
  run();
</script>
<body>
</body>
</html>
```