

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>

Features

Lightweight

Low-latency

Privacy

Improved power
consumption

Efficient model
format

Pre-trained models

Features

Lightweight

Low-latency

Privacy

Improved power
consumption

Efficient model
format

Pre-trained models

Features

Lightweight

Low-latency

Privacy

Improved power
consumption

Efficient model
format

Pre-trained models

Features

Lightweight

Low-latency

Privacy

Improved power
consumption

Efficient model
format

Pre-trained models

Features

Lightweight

Low-latency

Privacy

Improved power
consumption

Efficient model
format

Pre-trained models

Features

Lightweight

Low-latency

Privacy

Improved power
consumption

Efficient model
format

Pre-trained models

Features

Lightweight

Low-latency

Privacy

Improved power
consumption

Efficient model
format

Pre-trained models

Components in TensorFlow Lite

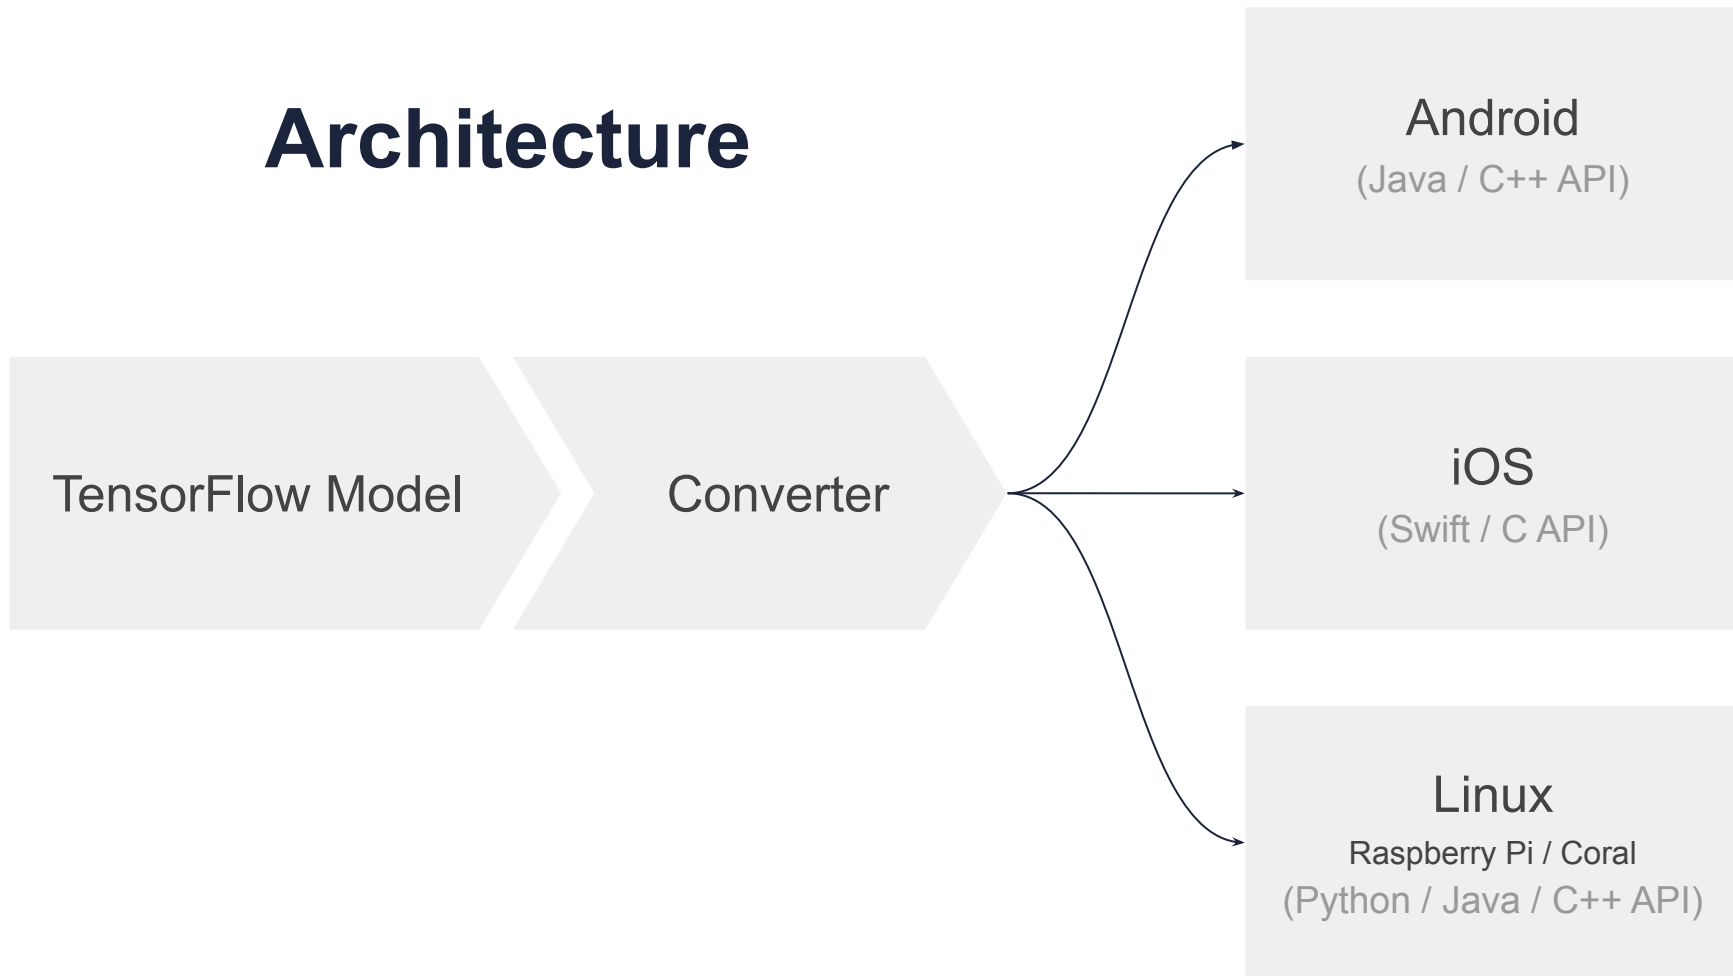
Converter
(to TensorFlow Lite
format)

- Transforms TensorFlow models into a form efficient for reading by the interpreter
- Introduces optimizations to improve binary size model performance and/or reduce model size.

Interpreter
Core

- Diverse platform support, (Android, iOS, embedded Linux and microcontrollers)
- Platform APIs for accelerated inference

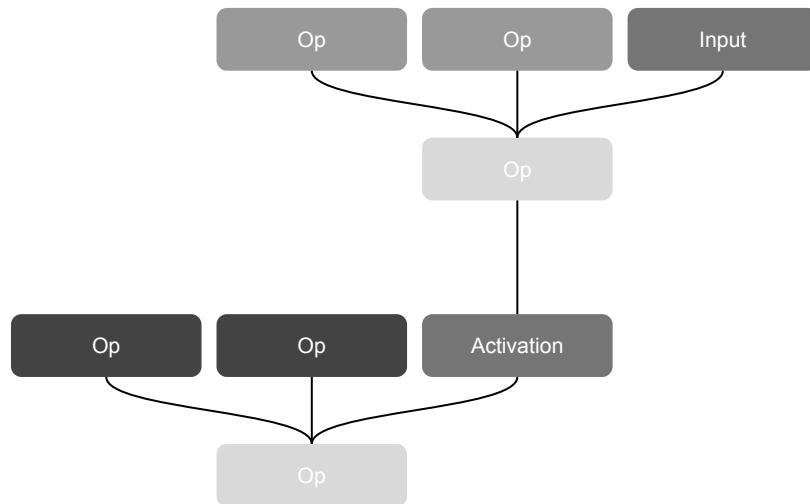

Architecture







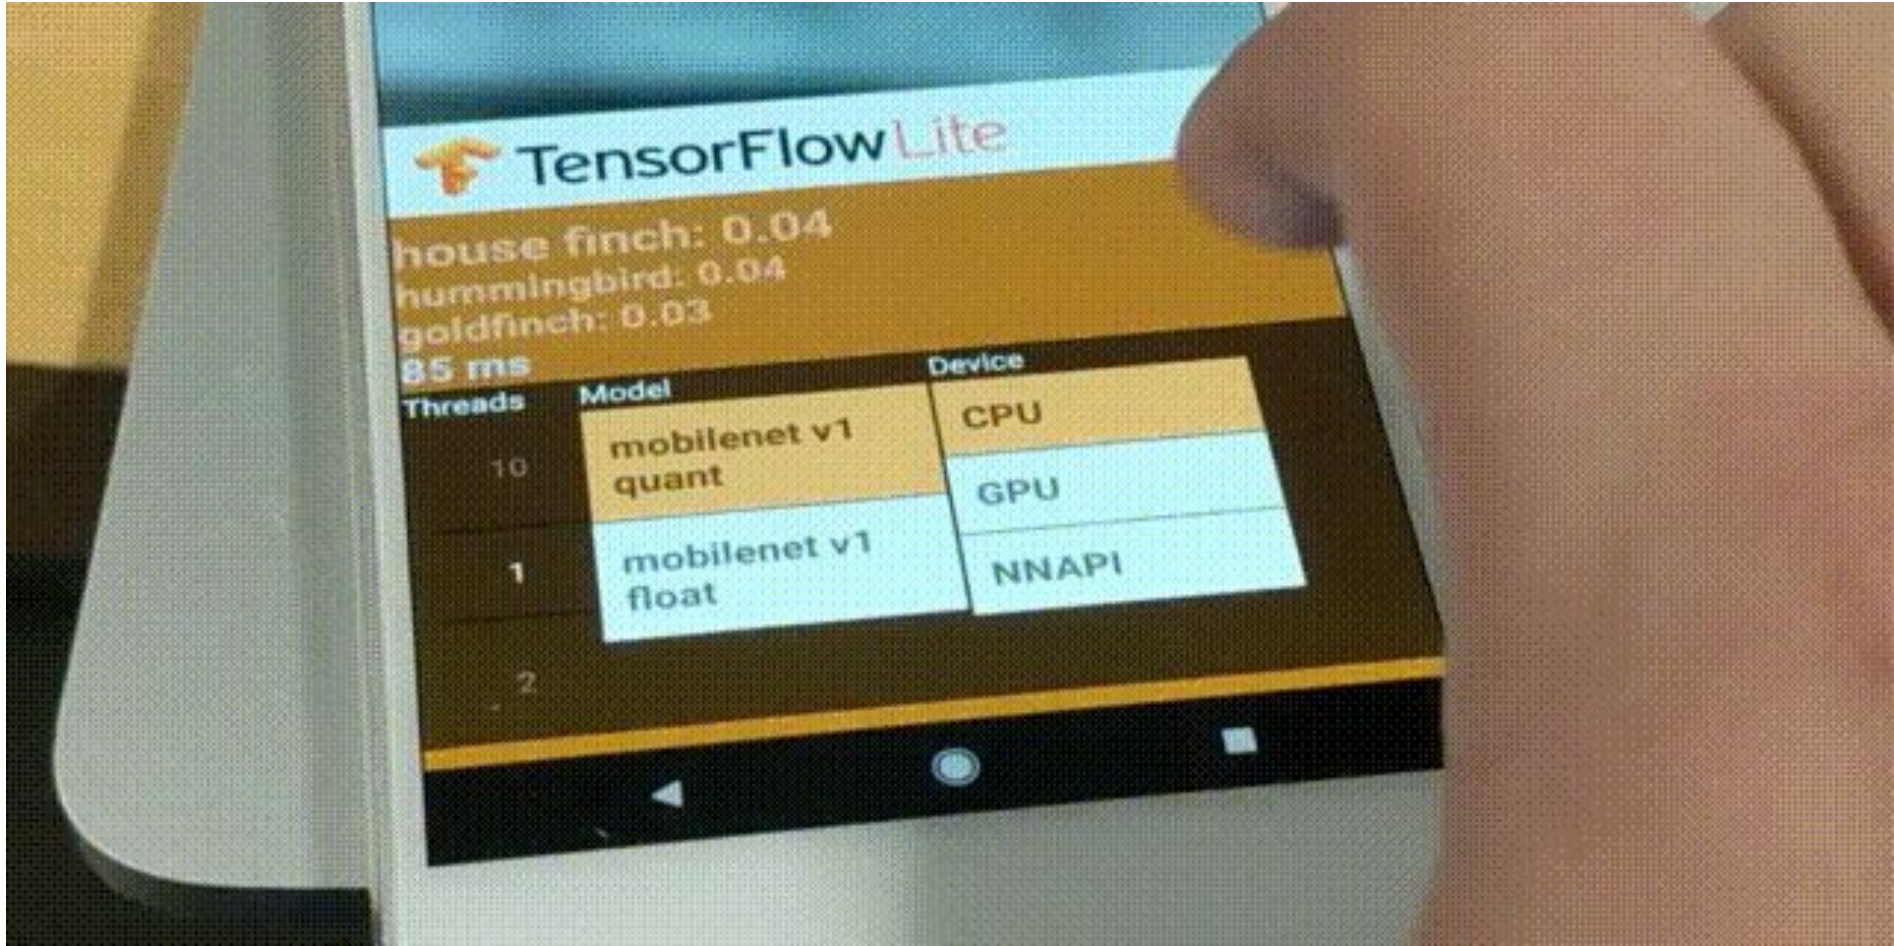
Performance

Acceleration	Available
Software	NN API (also a delegate)
	Edge TPU
	GPU
	CPU Optimizations (ARM and x86)
Hardware	

Delegates







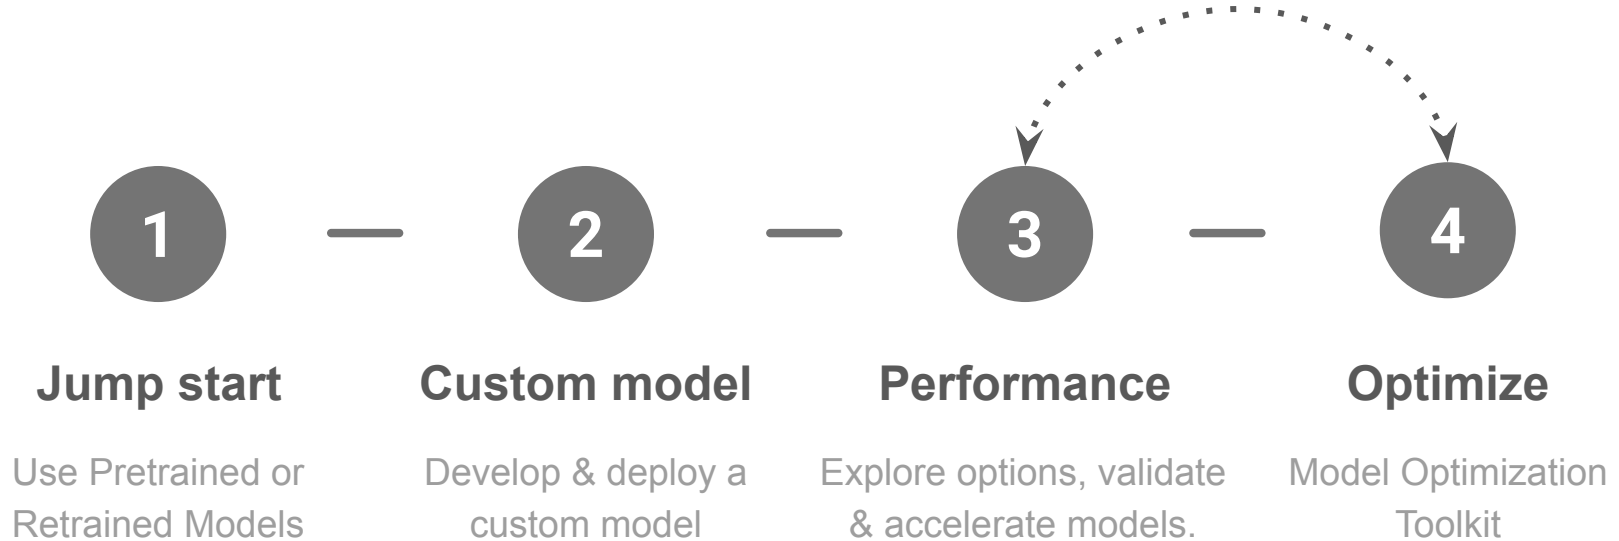
Techniques

- Quantization
- Weight pruning
- Model topology transforms
 - Tensor Decomposition
 - Distillation

Why Quantize?

- All available CPU platforms are supported
- Reducing latency and inference cost
- Low memory footprint
- Allow execution on hardware restricted-to or optimized-for fixed-point operations
- Optimized models for special purpose HW accelerators (TPUs)

Putting it all together





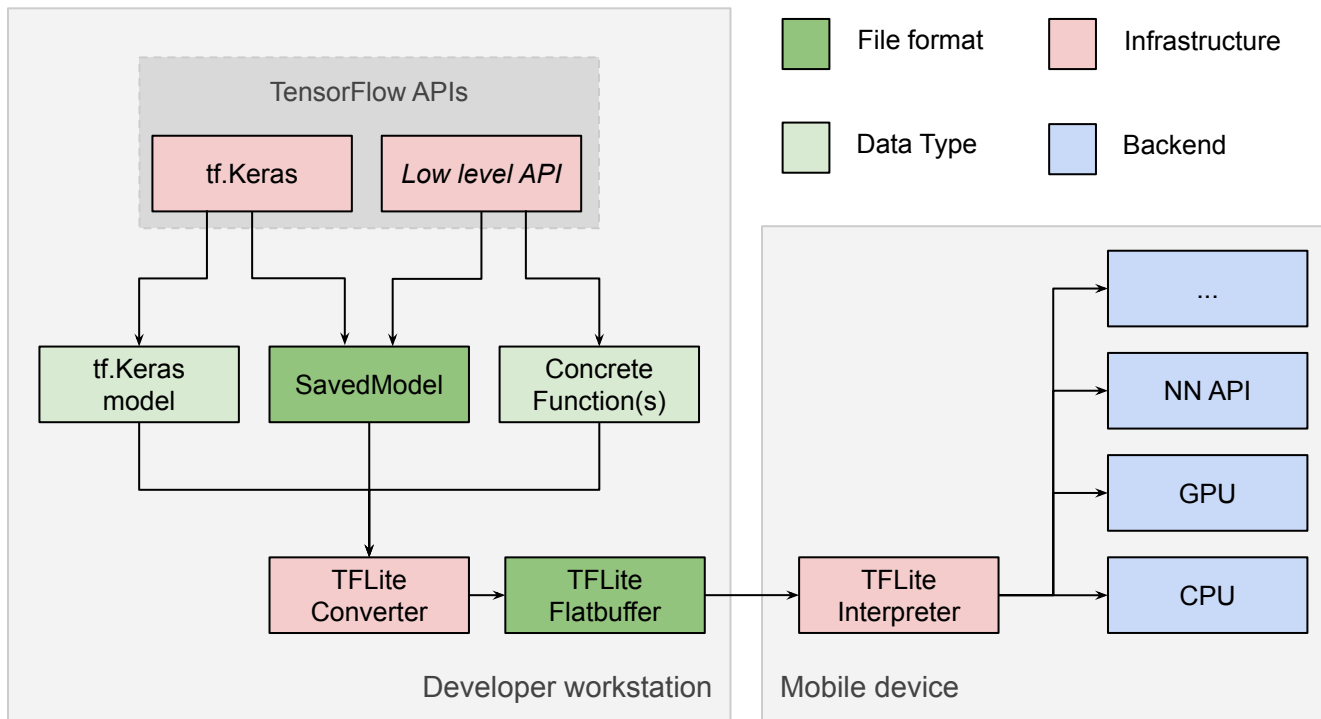
TensorFlow

Saved
Model

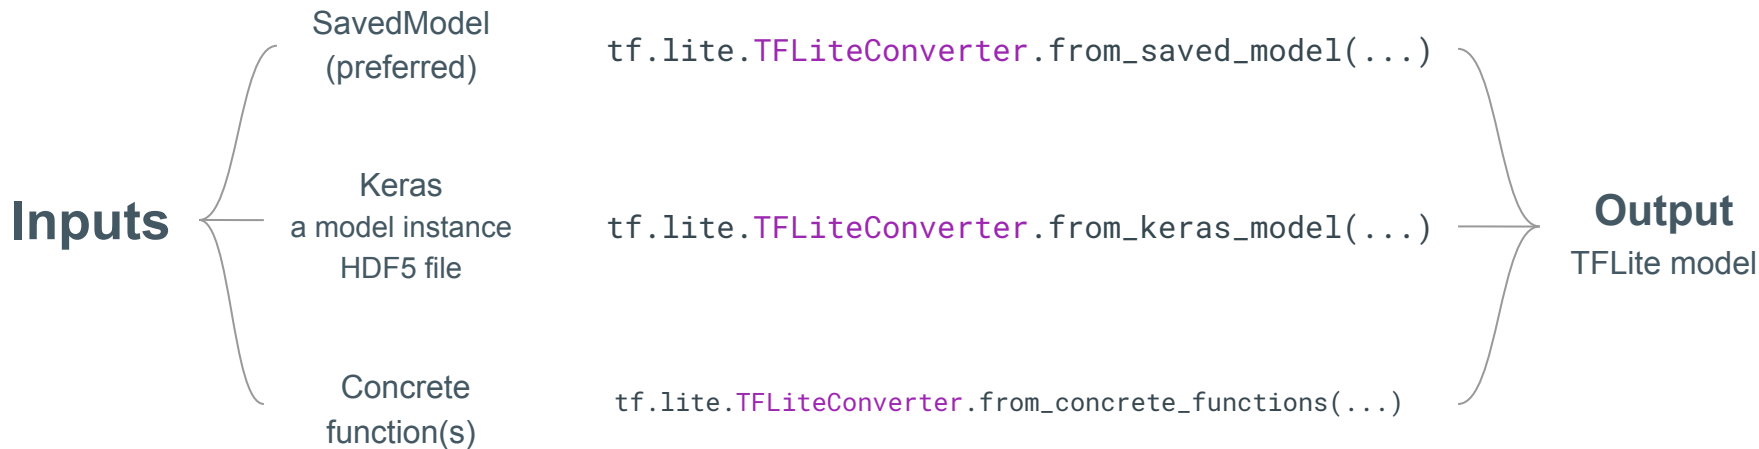
TF Lite
Converter

TF Lite
Model

TensorFlow Lite Converter



Parameters for conversion



Python API (preferred)

SavedModel

- The standard for serializing a TensorFlow model
- A MetaGraph to hold metadata
- Holds snapshot of the trained model (with model weights and computation)
- No model building code required
- Supports model versioning

Inspecting with SavedModel's CLI

```
!saved_model_cli show --dir /tmp/mobilenet/1 \  
    --tag_set serve \  
    --signature_def serving_default
```

Inspecting with SavedModel's CLI

The given `SavedModel SignatureDef` contains the following input(s):

```
inputs['input_1'] tensor_info:  
  dtype: DT_FLOAT  
  shape: (-1, 224, 224, 3)  
  name: serving_default_input_1:0
```

The given `SavedModel SignatureDef` contains the following output(s):

```
outputs['act_softmax'] tensor_info:  
  dtype: DT_FLOAT  
  shape: (-1, 1000)  
  name: StatefulPartitionedCall:0
```

Method name `is`: tensorflow/serving/predict

Exporting a SavedModel from Keras

```
pretrained_model = tf.keras.applications.MobileNet()  
tf.saved_model.save(pretrained_model, '/tmp/saved_model/1/')
```



```
import tensorflow as tf
```

```
# Store data for x and y
```

```
x = [-1, 0, 1, 2, 3, 4]
```

```
y = [-3, -1, 1, 3, 5, 7]
```

```
# Create a simple Keras model.
```


```
model = tf.keras.models.Sequential(
```

```
    [tf.keras.layers.Dense(units=1, input_shape=[1])])
```

```
model.compile(optimizer='sgd', loss='mean_squared_error')
```

```
model.fit(x, y, epochs=500)
```

$$y = 2x - 1$$



```
import pathlib

# Export the SavedModel
export_dir = '/tmp/saved_model'
tf.saved_model.save(model, export_dir)

# Convert the model
converter = tf.lite.TFLiteConverter.from_saved_model(export_dir)
tflite_model = converter.convert()

# Save the model
tflite_model_file = pathlib.Path('/tmp/foo.tflite')
tflite_model_file.write_bytes(tflite_model)
```

```
import pathlib
```

```
# Export the SavedModel
```

```
export_dir = '/tmp/saved_model'
```

```
tf.saved_model.save(model, export_dir)
```

```
# Convert the model
```

```
converter = tf.lite.TFLiteConverter.from_saved_model(export_dir)
```

```
tflite_model = converter.convert()
```

```
# Save the model
```

```
tflite_model_file = pathlib.Path('/tmp/foo.tflite')
```

```
tflite_model_file.write_bytes(tflite_model)
```

```
import pathlib

# Export the SavedModel
export_dir = '/tmp/saved_model'
tf.saved_model.save(model, export_dir)

# Convert the model
converter = tf.lite.TFLiteConverter.from_saved_model(export_dir)
tflite_model = converter.convert()

# Save the model
tflite_model_file = pathlib.Path('/tmp/foo.tflite')
tflite_model_file.write_bytes(tflite_model)
```

```
import pathlib

# Export the SavedModel
export_dir = '/tmp/saved_model'
tf.saved_model.save(model, export_dir)

# Convert the model
converter = tf.lite.TFLiteConverter.from_saved_model(export_dir)
tflite_model = converter.convert()

# Save the model
tflite_model_file = pathlib.Path('/tmp/foo.tflite')
tflite_model_file.write_bytes(tflite_model)
```

```
import tensorflow as tf
import pathlib

# Load the MobileNet tf.keras model.
model = tf.keras.applications.MobileNetV2(weights="imagenet", input_shape=(224, 224, 3))
# Saving the model for later use by tflite_convert
model.save('model.h5')

# Convert the model.
converter = tf.lite.TFLiteConverter.from_keras_model(model)
tflite_model = converter.convert()

# Save the model
tflite_model_file = pathlib.Path('/tmp/foo.tflite')
tflite_model_file.write_bytes(tflite_model)
```

```
import tensorflow as tf
import pathlib
```

```
# Load the MobileNet tf.keras model.
model = tf.keras.applications.MobileNetV2(weights="imagenet", input_shape=(224, 224, 3))
# Saving the model for later use by tflite_convert
model.save('model.h5')
```

```
# Convert the model.
converter = tf.lite.TFLiteConverter.from_keras_model(model)
tflite_model = converter.convert()
```

```
# Save the model
tflite_model_file = pathlib.Path('/tmp/foo.tflite')
tflite_model_file.write_bytes(tflite_model)
```

```
import tensorflow as tf
import pathlib

# Load the MobileNet tf.keras model.
model = tf.keras.applications.MobileNetV2(weights="imagenet", input_shape=(224, 224, 3))
# Saving the model for later use by tflite_convert
model.save('model.h5')

# Convert the model.
converter = tf.lite.TFLiteConverter.from_keras_model(model)
tflite_model = converter.convert()

# Save the model
tflite_model_file = pathlib.Path('/tmp/foo.tflite')
tflite_model_file.write_bytes(tflite_model)
```



```
import tensorflow as tf
import pathlib

# Load the MobileNet tf.keras model.
model = tf.keras.applications.MobileNetV2(weights="imagenet", input_shape=(224, 224, 3))
# Saving the model for later use by tflite_convert
model.save('model.h5')

# Convert the model.
converter = tf.lite.TFLiteConverter.from_keras_model(model)
tflite_model = converter.convert()

# Save the model
tflite_model_file = pathlib.Path('/tmp/foo.tflite')
tflite_model_file.write_bytes(tflite_model)
```

Command-line usage

```
#!/usr/bin/env bash
```

```
# Saving with the command-line from a SavedModel
```

```
tflite_convert --output_file=model.tflite --saved_model_dir=/tmp/saved_model
```

```
# Saving with the command-line from a Keras model
```

```
tflite_convert --output_file=model.tflite --keras_model_file=model.h5
```

Command-line usage

```
#!/usr/bin/env bash
```

```
# Saving with the command-line from a SavedModel
```

```
tflite_convert --output_file=model.tflite --saved_model_dir=/tmp/saved_model
```

```
# Saving with the command-line from a Keras model
```

```
tflite_convert --output_file=model.tflite --keras_model_file=model.h5
```

Command-line usage

```
#!/usr/bin/env bash
```

```
# Saving with the command-line from a SavedModel
```


```
tflite_convert --output_file=model.tflite --saved_model_dir=/tmp/saved_model
```

```
# Saving with the command-line from a Keras model
```

```
tflite_convert --output_file=model.tflite --keras_model_file=model.h5
```

Post-training quantization

- Reduced precision representation with
3x lower latency
- Little degradation in model accuracy
- Optimization modes
 - ↕ Default (both size and latency)
 - ↓ Size
 - ↓ Latency
- Efficiently represents an arbitrary
magnitude of ranges
- Quantization target specification
(FP32/INT8)



```
import tensorflow as tf

converter = tf.lite.TFLiteConverter.from_saved_model(saved_model_dir)

converter.optimizations = [tf.lite.Optimize.OPTIMIZE_FOR_SIZE]

tflite_quant_model = converter.convert()
```

Post-training integer quantization



```
# Define the generator
```

```
def generator():
```

```
    data = tfds.load(...)
```

```
    for _ in range(num_calibration_steps):
```

```
        image, = data.take(1)
```

```
        yield [image]
```

```
converter = tf.lite.TFLiteConverter.from_saved_model(saved_model_dir)
```

```
# Set the optimization mode
```

```
converter.optimizations = [tf.lite.Optimize.DEFAULT]
```

```
# Pass the representative dataset to the converter
```

```
converter.representative_dataset = tf.lite.RepresentativeDataset(generator)
```



```
# Define the generator
```

```
def generator():
```

```
    data = tfds.load(...)
```

```
    for _ in range(num_calibration_steps):
```

```
        image, = data.take(1)
```

```
        yield [image]
```

```
converter = tf.lite.TFLiteConverter.from_saved_model(saved_model_dir)
```

```
# Set the optimization mode
```

```
converter.optimizations = [tf.lite.Optimize.DEFAULT]
```

```
# Pass the representative dataset to the converter
```

```
converter.representative_dataset = tf.lite.RepresentativeDataset(generator)
```

```
# Define the generator
```

```
def generator():
```

```
    data = tfds.load(...)
```

```
    for _ in range(num_calibration_steps):
```

```
        image, = data.take(1)
```

```
        yield [image]
```

```
converter = tf.lite.TFLiteConverter.from_saved_model(saved_model_dir)
```

```
# Set the optimization mode
```

```
converter.optimizations = [tf.lite.Optimize.DEFAULT]
```

```
# Pass the representative dataset to the converter
```

```
converter.representative_dataset = tf.lite.RepresentativeDataset(generator)
```

```
# Define the generator
```

```
def generator():
```

```
    data = tfds.load(...)
```

```
    for _ in range(num_calibration_steps):
```

```
        image, = data.take(1)
```

```
        yield [image]
```

```
converter = tf.lite.TFLiteConverter.from_saved_model(saved_model_dir)
```

```
# Set the optimization mode
```

```
converter.optimizations = [tf.lite.Optimize.DEFAULT]
```

```
# Pass the representative dataset to the converter
```

```
converter.representative_dataset = tf.lite.RepresentativeDataset(generator)
```

```
...  
# Set the optimization mode  
converter.optimizations = [tf.lite.Optimize.OPTIMIZE_FOR_LATENCY]  
  
# Pass the representative dataset to the converter  
converter.representative_dataset = tf.lite.RepresentativeDataset(generator)  
  
# Restricting supported target op specification to INT8  
converter.target_spec.supported_ops = [tf.lite.OpsSet.TFLITE_BUILTINS_INT8]
```

Learn more about supported ops:

https://www.tensorflow.org/lite/guide/ops_compatibility

TF-Select to overcome unsupported ops

TF-Select

https://www.tensorflow.org/lite/guide/ops_select

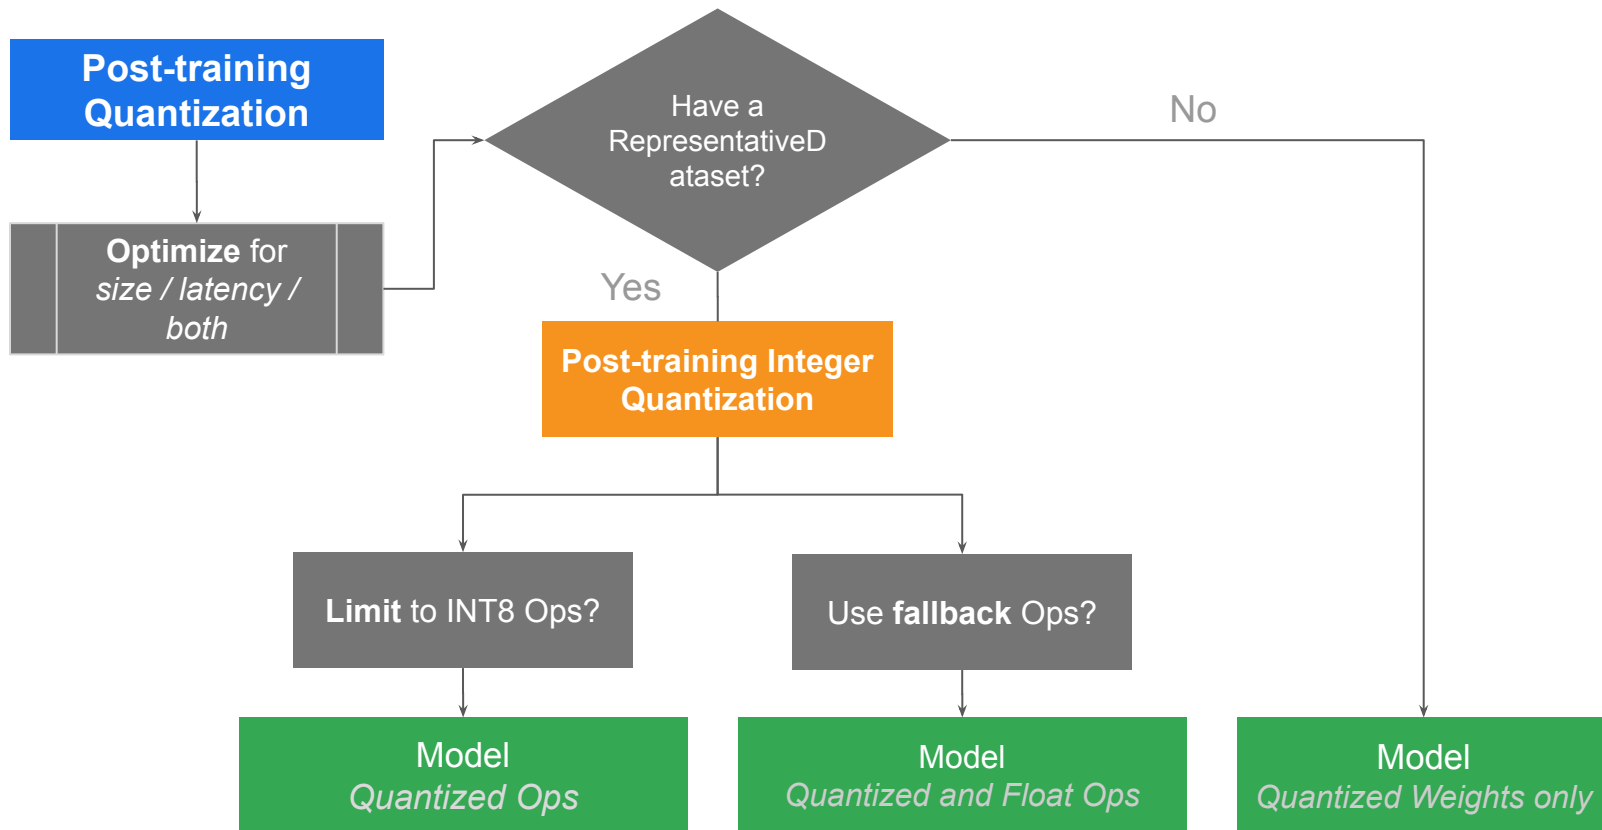
```
import tensorflow as tf

converter = tf.lite.TFLiteConverter.from_saved_model(saved_model_dir)

converter.target_ops = [tf.lite.OpsSet.TFLITE_BUILTINS,
                        tf.lite.OpsSet.SELECT_TF_OPS]

tflite_model = converter.convert()
```

Optimizing your models in a nutshell



TensorFlow Lite Interpreter in Python

```
# Load TFLite model and allocate tensors
```

```
interpreter = tf.lite.Interpreter(model_content=tflite_model)
```

```
interpreter.allocate_tensors()
```

```
# Get input and output tensors.
```

```
input_details = interpreter.get_input_details()
```

```
output_details = interpreter.get_output_details()
```

```
# Point the data to be used for testing and run the interpreter
```

```
interpreter.set_tensor(input_details[0]['index'], input_data)
```

```
interpreter.invoke()
```

```
tflite_results = interpreter.get_tensor(output_details[0]['index'])
```


TensorFlow Lite Interpreter in Python

```
# Load TFLite model and allocate tensors
```

```
interpreter = tf.lite.Interpreter(model_content=tflite_model)
```

```
interpreter.allocate_tensors()
```

```
# Get input and output tensors.
```

```
input_details = interpreter.get_input_details()
```

```
output_details = interpreter.get_output_details()
```

```
# Point the data to be used for testing and run the interpreter
```

```
interpreter.set_tensor(input_details[0]['index'], input_data)
```

```
interpreter.invoke()
```

```
tflite_results = interpreter.get_tensor(output_details[0]['index'])
```

TensorFlow Lite Interpreter in Python

```
# Load TFLite model and allocate tensors
```

```
interpreter = tf.lite.Interpreter(model_content=tflite_model)
```

```
interpreter.allocate_tensors()
```

```
# Get input and output tensors.
```

```
input_details = interpreter.get_input_details()
```

```
output_details = interpreter.get_output_details()
```

```
# Point the data to be used for testing and run the interpreter
```

```
interpreter.set_tensor(input_details[0]['index'], input_data)
```

```
interpreter.invoke()
```

```
tflite_results = interpreter.get_tensor(output_details[0]['index'])
```

TensorFlow Lite Interpreter in Python

```
# Load TFLite model and allocate tensors
```

```
interpreter = tf.lite.Interpreter(model_content=tflite_model)  
interpreter.allocate_tensors()
```

```
# Get input and output tensors.
```

```
input_details = interpreter.get_input_details()  
output_details = interpreter.get_output_details()
```

```
# Point the data to be used for testing and run the interpreter
```

```
interpreter.set_tensor(input_details[0]['index'], input_data)  
interpreter.invoke()
```

```
tflite_results = interpreter.get_tensor(output_details[0]['index'])
```

TensorFlow Lite Interpreter in Python

```
# Load TFLite model and allocate tensors
```

```
interpreter = tf.lite.Interpreter(model_content=tflite_model)  
interpreter.allocate_tensors()
```

```
# Get input and output tensors.
```

```
input_details = interpreter.get_input_details()  
output_details = interpreter.get_output_details()
```

```
# Point the data to be used for testing and run the interpreter
```

```
interpreter.set_tensor(input_details[0]['index'], input_data)  
interpreter.invoke()
```

```
tflite_results = interpreter.get_tensor(output_details[0]['index'])
```

Running models

Pretrained models

Image classification
Object detection
Smart reply
Pose estimation
Segmentation

TensorFlow Hub

Classification modules
Feature vector modules
Embedding modules

Not what you're
looking for?

Build a custom
model!

Getting a basic model running

Get started

1

Build a model

Create a simple model
for $(y = 2x - 1)$
from simulated data
and train it

2

Export & Convert

Generate the
SavedModel and
convert it to TFLite

3

Verify

Perform evaluation on
random data to verify the
results

4

Deploy

Deploy the converted
model on a mobile
device (Android/iOS)

Getting a basic model running

Get started

1

Build a model

Create a simple model
for $(y = 2x - 1)$
from simulated data
and train it

2

Export & Convert

Generate the
SavedModel and
convert it to TFLite

3

Verify

Perform evaluation on
random data to verify the
results

4

Deploy

Deploy the converted
model on a mobile
device (Android/iOS)

Getting a basic model running

Get started

1

Build a model

Create a simple model
for $(y = 2x - 1)$
from simulated data
and train it

2

Export & Convert

Generate the
SavedModel and
convert it to TFLite

3

Verify

Perform evaluation on
random data to verify the
results

4

Deploy

Deploy the converted
model on a mobile
device (Android/iOS)

Getting a basic model running

Get started

1

Build a model

Create a simple model
for $(y = 2x - 1)$
from simulated data
and train it

2

Export & Convert

Generate the
SavedModel and
convert it to TFLite

3

Verify

Perform evaluation on
random data to verify the
results

4

Deploy

Deploy the converted
model on a mobile
device (Android/iOS)

Getting a basic model running

Get started

1

Build a model

Create a simple model
for $(y = 2x - 1)$
from simulated data
and train it

2

Export & Convert

Generate the
SavedModel and
convert it to TFLite

3

Verify


Perform evaluation on
random data to verify the
results

4

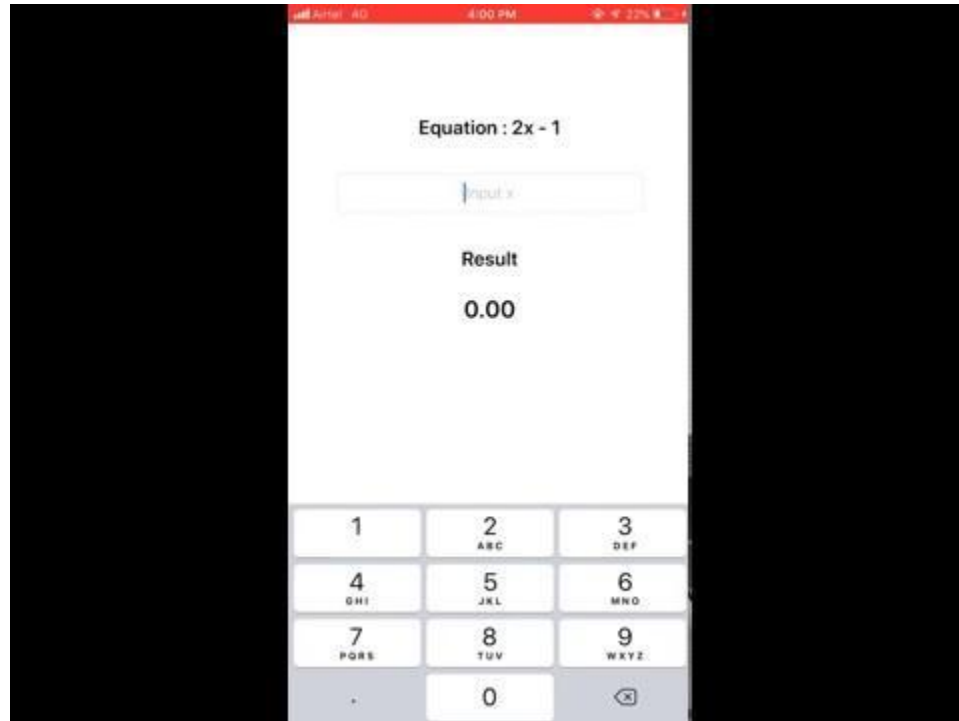
Deploy

Deploy the converted
model on a mobile
device (Android/iOS)


Android

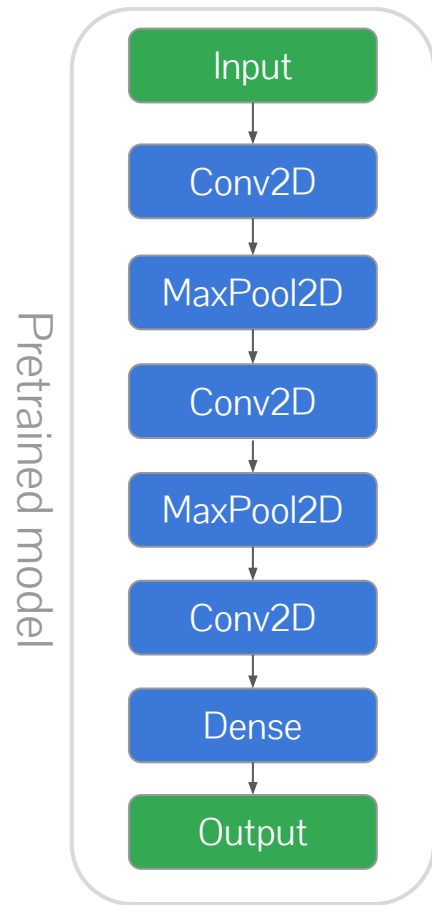


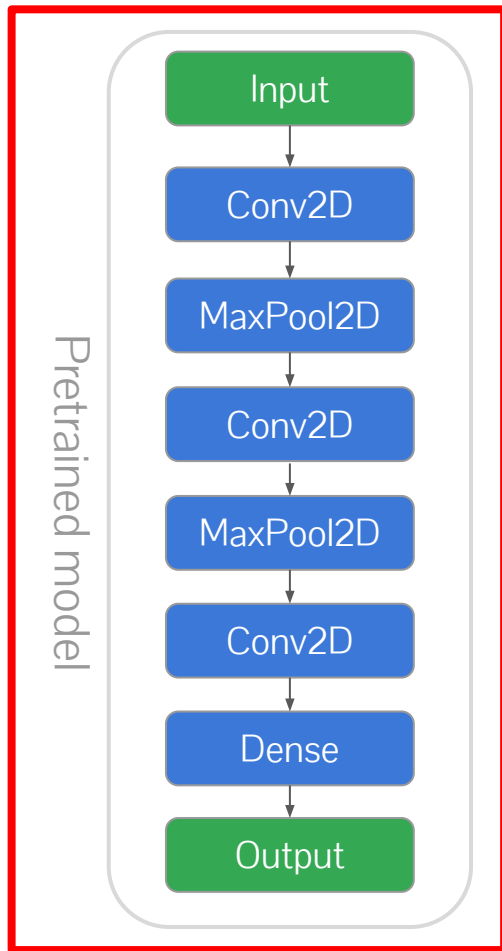
iOS

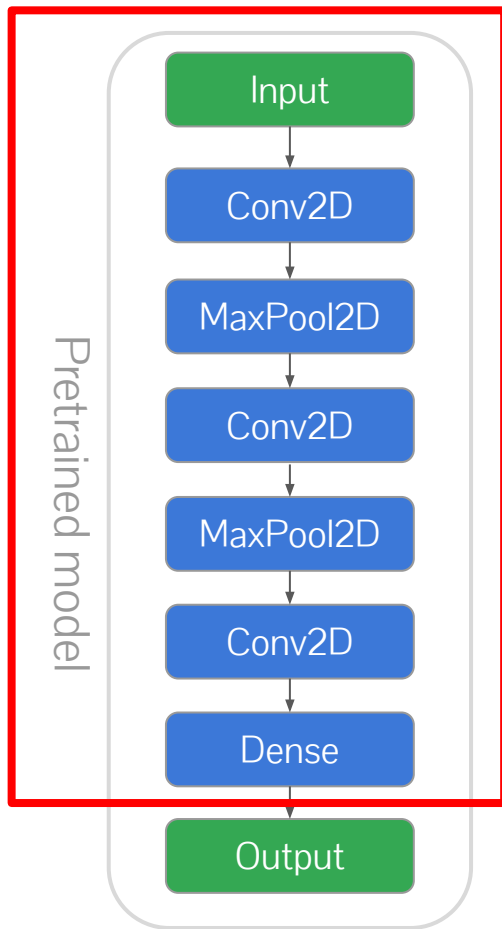


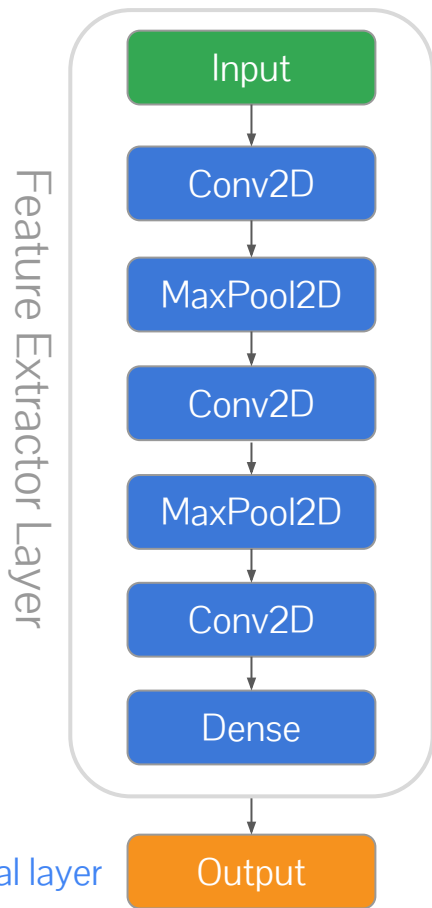
Transfer Learning











Train only the final layer

Transfer Learning on Cats vs Dogs with TensorFlow Hub

Get started

1

Prepare the dataset

Download the Cats vs. Dogs tf.Dataset, split into sets (train, val, test), and preprocess

2

Transfer Learning

Choose a feature vector module (MobileNet V2) from TFHub and perform transfer learning

3

Export & Convert

Export the trained model to SavedModel and convert it to TFLite

4

Deploy

Deploy the converted model on a mobile device (Android/iOS)

Transfer Learning on Cats vs Dogs with TensorFlow Hub

Get started

1

Prepare the dataset

Download the Cats vs. Dogs tf.Dataset, split into sets (train, val, test), and preprocess

2

Transfer Learning

Choose a feature vector module (MobileNet V2) from TFHub and perform transfer learning

3

Export & Convert

Export the trained model to SavedModel and convert it to TFLite

4

Deploy

Deploy the converted model on a mobile device (Android/iOS)

Transfer Learning on Cats vs Dogs with TensorFlow Hub

Get started

1

Prepare the dataset

Download the Cats vs. Dogs tf.Dataset, split into sets (train, val, test), and preprocess

2

Transfer Learning

Choose a feature vector module (MobileNet V2) from TFHub and perform transfer learning

3

Export & Convert

Export the trained model to SavedModel and convert it to TFLite

4

Deploy

Deploy the converted model on a mobile device (Android/iOS)

Transfer Learning on Cats vs Dogs with TensorFlow Hub

Get started

1

Prepare the dataset

Download the Cats vs. Dogs tf.Dataset, split into sets (train, val, test), and preprocess

2

Transfer Learning

Choose a feature vector module (MobileNet V2) from TFHub and perform transfer learning

3

Export & Convert

Export the trained model to SavedModel and convert it to TFLite

4

Deploy

Deploy the converted model on a mobile device (Android/iOS)

Transfer Learning on Cats vs Dogs with TensorFlow Hub

Get started

1

Prepare the dataset

Download the Cats vs. Dogs tf.Dataset, split into sets (train, val, test), and preprocess

2

Transfer Learning

Choose a feature vector module (MobileNet V2) from TFHub and perform transfer learning

3

Export & Convert

Export the trained model to SavedModel and convert it to TFLite

4

Deploy

Deploy the converted model on a mobile device (Android/iOS)