

# Copyright Notice

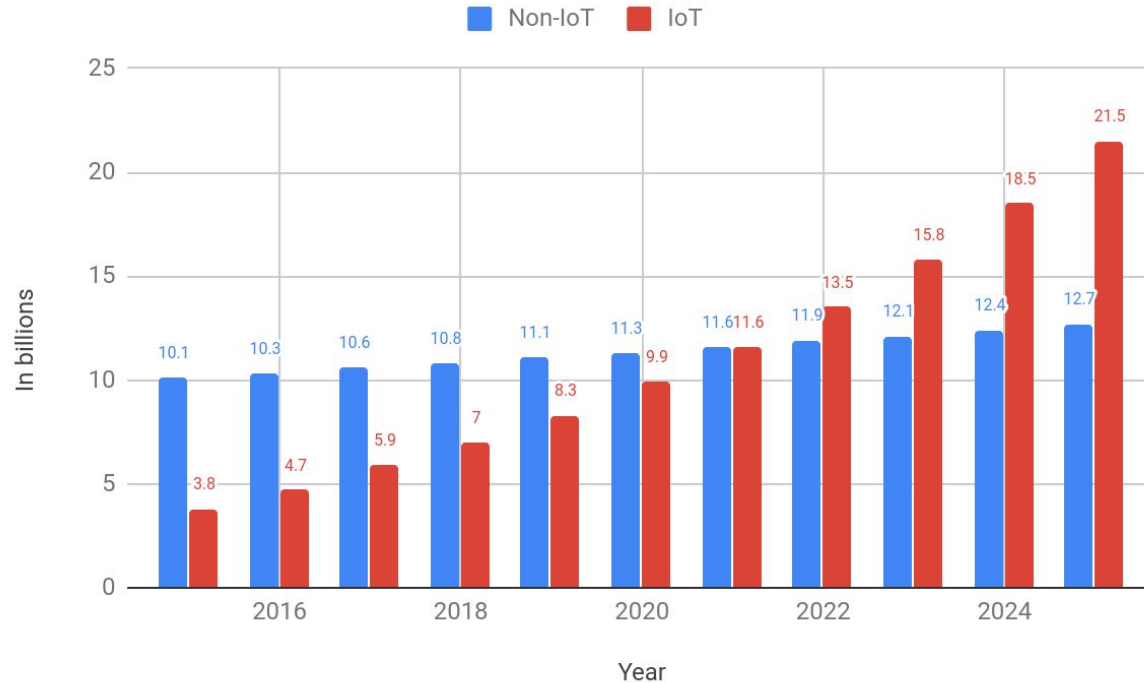
These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>

# Trends in adoption of Smart devices



IoT Analytics - State of the IoT 2018

# Factors driving this trend

- Smaller, faster models
- On device accelerators
- Demands move ML capability from cloud to on-device
- Smart device growth demands bring ML to the edge


# Benefits of on-device ML



High-performance



Better privacy




Local data accessibility



Works offline

# Coral



Coral Dev Board



Coral USB Accelerator

# Software for Coral

- Mendel OS
- Edge TPU Compiler
- Mendel Development Tool (mdt)
- Edge TPU models : <https://coral.withgoogle.com/models/>

# Raspberry Pi


- Small sized
- Low cost
- Just like a computer
- Accessibility
- Raspbian (OS)



[Image credits](#)

# Micro-controllers

- Low-powered
- Small form-factor
- Some specially designed for ML tasks
- No reliance on network connectivity



**SparkFun Edge**

[Image credits](#)



# Options

- Compile TensorFlow from Source
- Install TensorFlow with pip
- Use TensorFlow Lite Interpreter directly

# Build From Source

[https://www.tensorflow.org/install/source\\_rpi](https://www.tensorflow.org/install/source_rpi)

```
curl -sSL https://get.docker.com | sh
```

```
sudo docker run hello-world
```

# Build From Source

[https://www.tensorflow.org/install/source\\_rpi](https://www.tensorflow.org/install/source_rpi)

```
git clone https://github.com/tensorflow/tensorflow.git
```

```
cd tensorflow
```

# Build From Source

[https://www.tensorflow.org/install/source\\_rpi](https://www.tensorflow.org/install/source_rpi)

```
sudo CI_DOCKER_EXTRA_PARAMS= \  
    "-e CI_BUILD_PYTHON=python3 \  
    -e CROSSTOOL_PYTHON_INCLUDE_PATH=/usr/include/python3.4" \  
tensorflow/tools/ci_build/ci_build.sh PI-PYTHON3 \  
tensorflow/tools/ci_build/pi/build_raspberry_pi.sh
```

# Build From Source

[https://www.tensorflow.org/install/source\\_rpi](https://www.tensorflow.org/install/source_rpi)

```
pip install <your wheel name>
```

# Use Pre-built Packages

<https://www.tensorflow.org/install/pip>

```
sudo apt update  
sudo apt install python3-dev python3-pip  
sudo apt install libatlas-base-dev  
pip install --upgrade tensorflow
```

# Use Interpreter Only

<https://www.tensorflow.org/lite/guide/python>

```
pip3 install tf-lite-runtime-1.14.0-cp37-cp37m-linux_armv7l.whl
```



dog	0.91
rabbit	0.07
hamster	0.02

Identify classes of different objects in the image



# Model

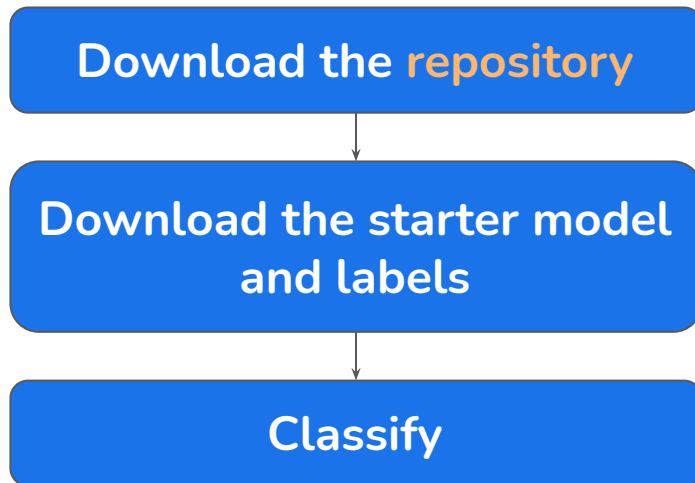
Pre-quantized MobileNet trained on ImageNet

1000 different classes of objects

More details on the model can be found at

**[https://www.tensorflow.org/lite/models/image\\_classification/overview](https://www.tensorflow.org/lite/models/image_classification/overview)**

# Quickstart



1

## **Initialize the Interpreter**

Load the interpreter with the model and make it ready for inference

2

## **Preprocess input**

Preprocess by resizing and normalizing the image data

3

## **Perform Inference**

Pass input to the Interpreter and invoke it

4

## **Obtain results and map**

Extract the resulting scores for each class and map them

# Initializing the Interpreter

```
# Load the model and allocate tensors  
interpreter = tf.lite.Interpreter(model_path='mobilenet_v2_1.0_224.tflite')  
interpreter.allocate_tensors()
```

# Get the model's tensors

```
# Load the model and allocate tensors
```

```
interpreter = tf.lite.Interpreter(model_content=tflite_model)
```

```
interpreter.allocate_tensors()
```

```
# Get input and output tensors.
```

```
input_details = interpreter.get_input_details()
```

```
output_details = interpreter.get_output_details()
```

```
...
```

1



2



3



4

## **Initialize the Interpreter**

Load the interpreter with the model and make it ready for inference

## **Preprocess input**

Preprocess by resizing and normalizing the image data

## **Perform Inference**

Pass input to the Interpreter and invoke it

## **Obtain results and map**

Extract the resulting scores for each class and map them

# Preprocess the image

```
# Read image and decode
```

```
img = tf.io.read_file(filename)
```

```
img_tensor = tf.image.decode_image(img)
```

```
# Preprocess image
```

```
img_tensor = tf.image.resize(img_tensor, size)
```

```
img_tensor = tf.cast(img_tensor, tf.uint8)
```

```
# Add a batch dimension
```

```
input_data = tf.expand_dims(img_tensor, axis=0)
```



# Preprocess the image

```
# Read image and decode
```

```
img = tf.io.read_file(filename)
```

```
img_tensor = tf.image.decode_image(img)
```

```
# Preprocess image
```

```
img_tensor = tf.image.resize(img_tensor, size)
```

```
img_tensor = tf.cast(img_tensor, tf.uint8)
```

```
# Add a batch dimension
```

```
input_data = tf.expand_dims(img_tensor, axis=0)
```





# Preprocess the image

```
# Read image and decode
```

```
img = tf.io.read_file(filename)
```

```
img_tensor = tf.image.decode_image(img)
```

```
# Preprocess image
```

```
img_tensor = tf.image.resize(img_tensor, size)
```

```
img_tensor = tf.cast(img_tensor, tf.uint8)
```

```
# Add a batch dimension
```

```
input_data = tf.expand_dims(img_tensor, axis=0)
```



# Preprocess the image

```
# Read image and decode
```

```
img = tf.io.read_file(filename)
```

```
img_tensor = tf.image.decode_image(img)
```

```
# Preprocess image
```

```
img_tensor = tf.image.resize(img_tensor, size)
```

```
img_tensor = tf.cast(img_tensor, tf.uint8)
```

```
# Add a batch dimension
```

```
input_data = tf.expand_dims(img_tensor, axis=0)
```



1

—

2

—

3

—

4

## **Initialize the Interpreter**

Load the interpreter with the model and make it ready for inference

## **Preprocess input**

Preprocess by resizing and normalizing the image data

## **Perform Inference**

Pass input to the Interpreter and invoke it

## **Obtain results and map**

Extract the resulting scores for each class and map them

# Perform inference

```
# Point data to be used for testing the interpreter and run it
interpreter.set_tensor(input_details[0]['index'], input_data)
interpreter.invoke()
```

1

—

2

—

3

—

4

## **Initialize the Interpreter**

Load the interpreter with the model and make it ready for inference

## **Preprocess input**

Preprocess by resizing and normalizing the image data

## **Perform Inference**

Pass input to the Interpreter and invoke it

## **Obtain results and map**

Extract the resulting scores for each class and map them

# Report only the top results

```
# Obtain results
predictions = interpreter.get_tensor(output_details[0]['index'])

# Get indices of the top k results
_, top_k_indices = tf.math.top_k(predictions, k=top_k_results)
top_k_indices = np.array(top_k_indices)[0]

for i in range(top_k_results):
    print(labels[top_k_indices[i]],
          predictions[top_k_indices[i]] / 255.0)
```

Label	Probability
dog	0.91
rabbit	0.07
hamster	0.02

# Report only the top results

```
# Obtain results
predictions = interpreter.get_tensor(output_details[0]['index'])

# Get indices of the top k results
_, top_k_indices = tf.math.top_k(predictions, k=top_k_results)
top_k_indices = np.array(top_k_indices)[0]

for i in range(top_k_results):
    print(labels[top_k_indices[i]],
          predictions[top_k_indices[i]] / 255.0)
```

Label	Probability
dog	0.91
rabbit	0.07
hamster	0.02

# Report only the top results

```
# Obtain results
```

```
predictions = interpreter.get_tensor(output_details[0]['index'])
```

```
# Get indices of the top k results
```

```
_, top_k_indices = tf.math.top_k(predictions, k=top_k_results)
```

```
top_k_indices = np.array(top_k_indices)[0]
```

```
for i in range(top_k_results):
```

```
    print(labels[top_k_indices[i]],
```

```
          predictions[top_k_indices[i]] / 255.0)
```

Label	Probability
dog	0.91
rabbit	0.07
hamster	0.02



# Report only the top results

```
# Obtain results
```

```
predictions = interpreter.get_tensor(output_details[0]['index'])
```

```
# Get indices of the top k results
```

```
_, top_k_indices = tf.math.top_k(predictions, k=top_k_results)
```

```
top_k_indices = np.array(top_k_indices)[0]
```

```
for i in range(top_k_results):  
    print(labels[top_k_indices[i]],  
          predictions[top_k_indices[i]] / 255.0)
```

Label	Probability
dog	0.91
rabbit	0.07
hamster	0.02

**Detect multiple objects  
within an image**

**Recognize different  
classes of objects**

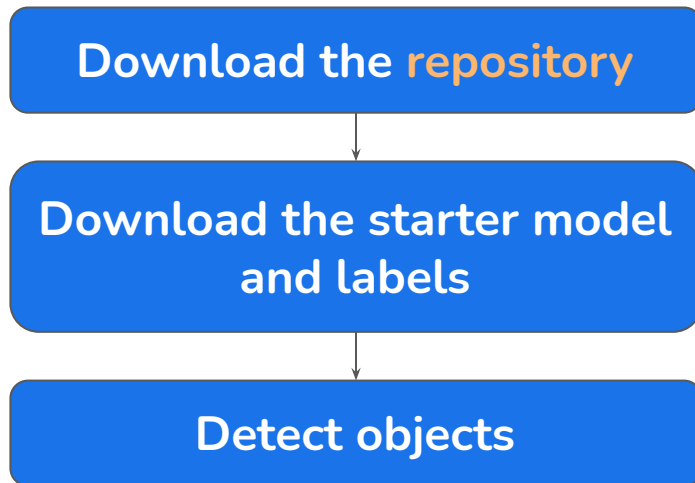


# Model

- Pre-optimized MobileNet SSD trained on COCO dataset
- COCO dataset has 80 common object categories
- A labels file to map the model's outputs
- You can download the model (.tflite) and labels (.txt)

[http://storage.googleapis.com/download.tensorflow.org/models/tflite/coco\\_ssd\\_mobilenet\\_v1\\_1.0\\_quant\\_2018\\_06\\_29.zip](http://storage.googleapis.com/download.tensorflow.org/models/tflite/coco_ssd_mobilenet_v1_1.0_quant_2018_06_29.zip)

# Quickstart



1

## **Initialize the Interpreter**

Load the interpreter with the model and make it ready for inference by getting the input and output tensors

2

## **Preprocess input**

Center crop the input image so that the model can accept the input

3

## **Perform Inference**

Pass the preprocessed input to the Interpreter and invoke it

4

## **Fetch the outputs**

Extract the outputs as locations, classes, scores and number of detections

# Initializing the Interpreter

```
# Load the model and allocate tensors
```

```
interpreter = tf.lite.Interpreter(model_path='detect.tflite')
```

```
interpreter.allocate_tensors()
```

```
# Get input and output tensors.
```

```
input_details = interpreter.get_input_details()
```

```
output_details = interpreter.get_output_details()
```

1

—

2

—

3

—

4

## **Initialize the Interpreter**

Load the interpreter with the model and make it ready for inference

## **Get cam feed and preprocess**

Gather images from the camera feed and center crop the input image so that the model can accept the input

## **Perform Inference**

Pass the preprocessed input to the Interpreter and invoke it

## **Fetch the outputs**

Extract the outputs as locations, classes, scores and number of detections

# Raspberry Pi Camera

```
with picamera.PiCamera() as camera:  
    camera.resolution = (640, 480)  
    while True:  
        # Input image  
        image = np.empty((480, 640, 3), dtype=np.uint8)  
        camera.capture(image, 'rgb')  
        # Use the frame captured from the stream
```



# Preprocessing




# Center and Crop

```
# Get the dimensions
height, width, _ = frame.shape # Image shape
new_width, new_height = (300, 300) # Target shape

# Calculate offsets between heights and widths
offset_height = (height - new_height) // 2
offset_width = (width - new_width) // 2

# Crop to the biggest square in the center
image = tf.image.crop_to_bounding_box(frame,
                                     offset_height,
                                     offset_width,
                                     new_height,
                                     new_width)
```



1



2



3



4

## **Initialize the Interpreter**

Load the interpreter with the model and make it ready for inference

## **Get cam feed and preprocess**

Gather images from the camera feed and center crop the input image so that the model can accept the input

## **Perform Inference**

Pass the preprocessed input to the Interpreter and invoke it

## **Fetch the outputs**

Extract the outputs as locations, classes, scores and number of detections

# Perform inference

```
def detect(self, image, threshold=0.1):  
    # Add a batch dimension  
    frame = np.expand_dims(image, axis=0)  
  
    # run model  
    self.interpreter.set_tensor(self.input_details[0]['index'], frame)  
    self.interpreter.invoke()
```

1



2



3



4

## **Initialize the Interpreter**

Load the interpreter with the model and make it ready for inference

## **Get cam feed and preprocess**

Gather images from the camera feed and center crop the input image so that the model can accept the input

## **Perform Inference**


Pass the preprocessed input to the Interpreter and invoke it

## **Fetch the outputs**

Extract the outputs as locations, classes, scores and number of detections

# How a detected object is represented

(top, left)




Class	Score	Location
<b>Dog</b>	0.92	top, left, bottom, right

(bottom, right)

# How the COCO model sees outputs

index	name	
0	locations	A list of floats in [0, 1] representing <i>normalized bounding boxes</i> [top left bottom right]
1	classes	A list of integers (output as float) each indicating the <i>index of a class label</i> from the labels file
2	scores	Array of floats in [0, 1] representing <i>probability</i> that a class was detected
3	number of detections	Floating point value expressing <i>total number of results</i>

# Interpreting the number of results



Class	Score	Location (top, left, bottom, right)
Apple	0.96	275, 257, 407, 379
Glasses	0.89	4, 257, 224, 356
Computer Keyboard	0.77	0, 2, 292, 80
Banana	0.67	345, 0, 417, 284



# Fetching the outputs

```
# Normalized coordinates of the detected objects
```

```
boxes = interpreter.get_tensor(output_details[0]['index'])[0]
```

```
# Recognized classes of objects
```

```
classes = interpreter.get_tensor(output_details[1]['index'])[0]
```

```
# Probabilities of the detected classes
```

```
scores = interpreter.get_tensor(output_details[2]['index'])[0]
```

```
# Maximum number of results
```

```
num_detections = interpreter.get_tensor(output_details[3]['index'])[0]
```

# Discarding less relevant results

```
min_score_thresh = 0.6
number_boxes = boxes.shape[0]
detected_boxes = []
probabilities = []
categories = []

for i in range(number_boxes):
    if scores is None or scores[i] > min_score_thresh:
        box = tuple(boxes[i].tolist()) # [top, left, bottom, right]
        detected_boxes.append(box)
        probabilities.append(scores[i])
        categories.append(self.category_index[classes[i]])
```

# Reporting results

```
# Convert normalized boxes to regular bounding boxes
```

```
(top, bottom, left, right) = (top * im_height, bottom * im_height,  
                               left * im_width, right * im_width)
```

```
# Draw lines for the detected boxes
```

```
draw.line([(left, top), (left, bottom), (right, bottom), (right, top),  
          (left, top)], width=thickness, fill=color)
```

```
# Draw the display string (predicted class)
```

```
draw.text(
```


```
    # Calculate position of text to be placed at the top-left corner
```

```
    (left + margin, text_bottom - text_height - margin),
```

```
    display_str,
```

```
    fill='black',
```

```
    font=font)
```





<https://www.youtube.com/watch?v=e-Gu6Sp4OUQ>

Detecting




<https://www.sparkfun.com/products/15170>







<https://www.sparkfun.com/products/15096>











Red LED  
("no")

Blue  
LED

Yellow LED  
("yes")

1 Introduction

2 Set up your hardware

3 Set up your software

4 Build and prepare the binary

5 Get ready to flash the binary

6 Flash the binary

7 Verify flashing was successful

8 Extend the code

9 Read the debug output

10 Next steps

## 1. Introduction

Machine learning helps developers build software that can understand our world. We can use it to create intelligent tools that make users' lives easier, like the [Google Assistant](#), and fun experiences that let users express their creativity, like [Google Pixel's portrait mode](#).

But often, these experiences require a lot of computation. Machine learning often needs to run on a powerful cloud server, or at least a powerful mobile phone.

With [TensorFlow Lite](#), it's possible to run machine learning inference on tiny, low-powered hardware, like microcontrollers. This means you can build amazing experiences that add intelligence to the smallest devices, bringing machine learning closer to the world around us.

In this codelab, we'll learn to deploy a machine learning model to the [SparkFun Edge](#), a microcontroller designed by Google and SparkFun to help developers experiment with ML on tiny devices.

### What is the SparkFun Edge?

The SparkFun Edge is a microcontroller-based platform: a tiny computer on a single circuit board. It has a processor, memory, and I/O hardware that allows it to send and receive digital signals to other devices. It also has four software-controllable LEDs, in your favorite Google colors.



[https://www.tensorflow.org/lite/microcontrollers/get\\_started](https://www.tensorflow.org/lite/microcontrollers/get_started)

