

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.


For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>


Boosting tf.data pipeline

- Rise of accelerators and parallelism
- High-performant tf.data input pipelines
- Adopt pipelines to different scenarios
- Learn better ways of using tf.data operations

ETL Revisited



Extract



Transform



Load

Local (HDD/SSD)

Remote (GCS/HDFS)

Shuffling & Batching

Decompression

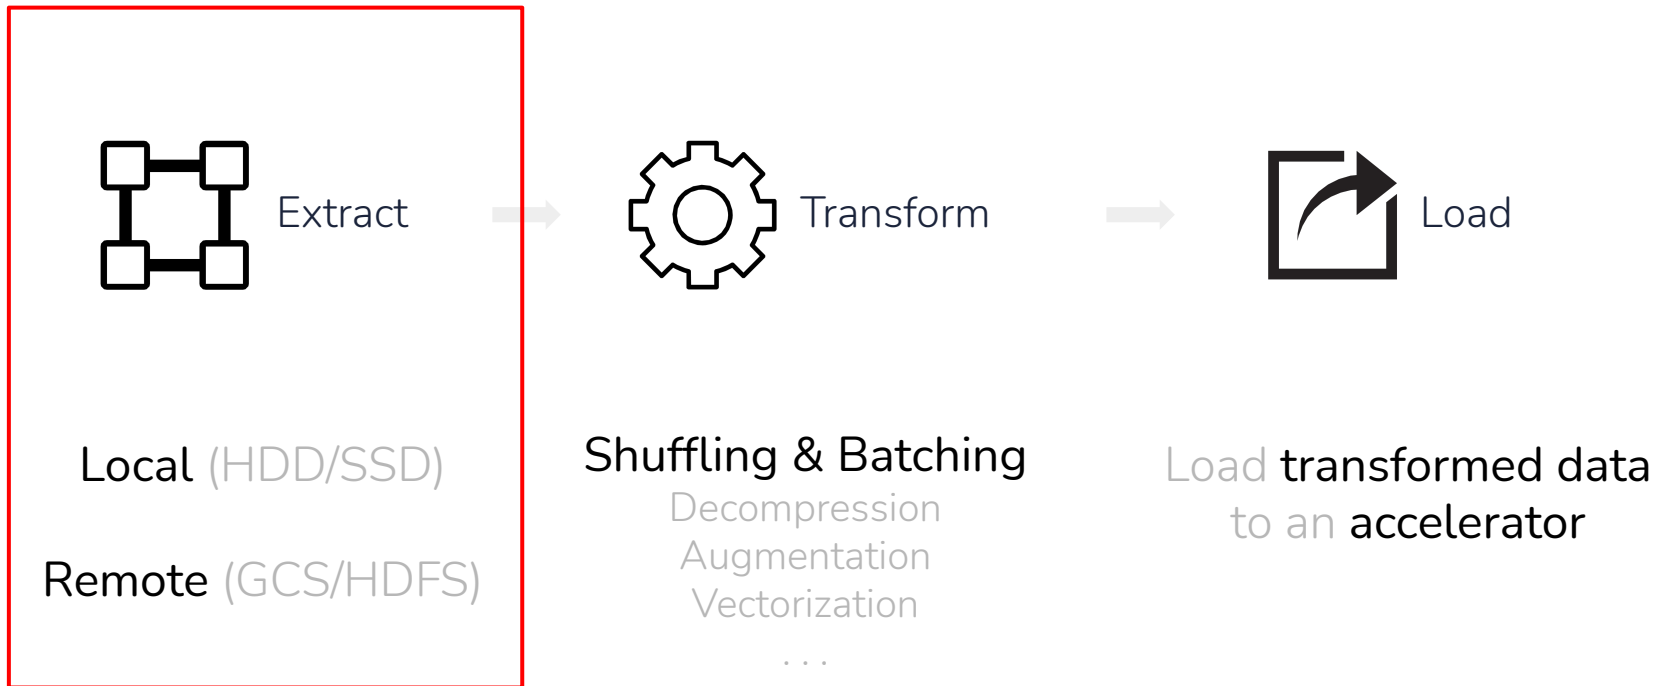
Augmentation

Vectorization

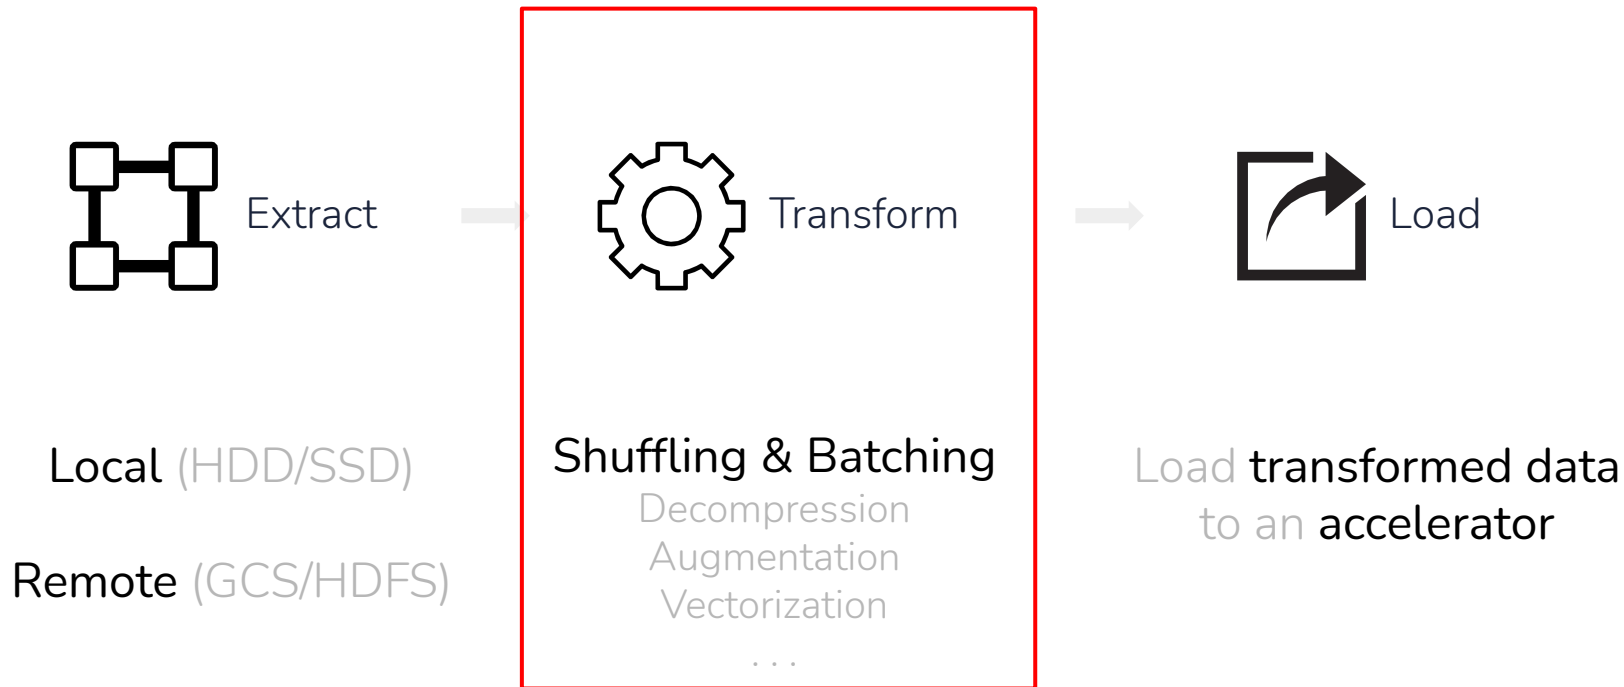
...

Load transformed data
to an accelerator


ETL Revisited




ETL Revisited



ETL Revisited



Extract



Transform



Load

Local (HDD/SSD)

Remote (GCS/HDFS)

Shuffling & Batching

Decompression

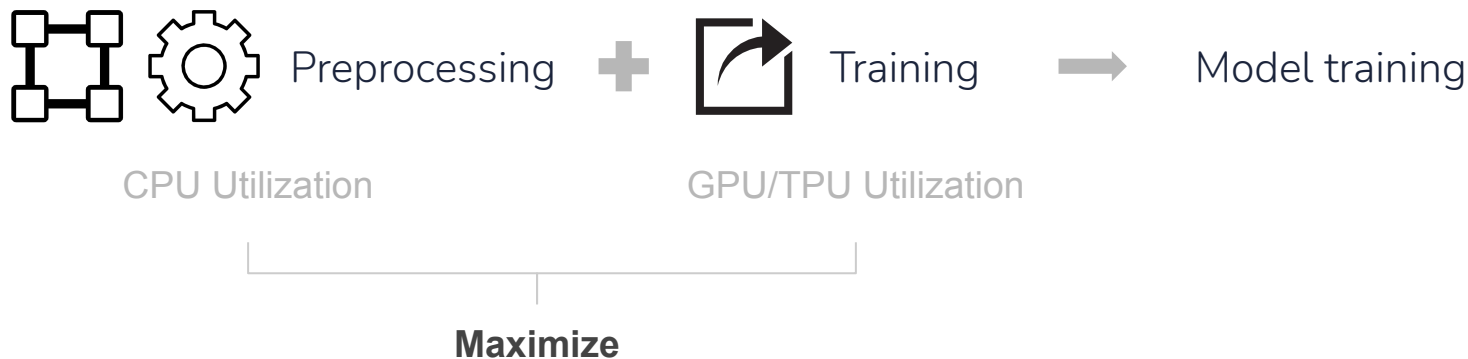
Augmentation

Vectorization

...

Load transformed data
to an accelerator

What happens when you train a model?



Data and its problems

- Bound to come across fitting input data locally
- When **distributed training** expects **distributed data**
- Avoid having the same data on every machine

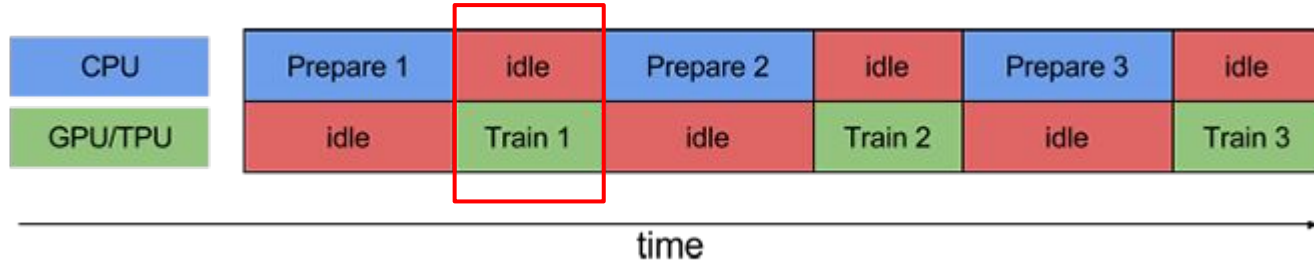
Data and models



Data and models



Data and models



Data and models



Data and models



Pipelining

Without
pipelining



With
pipelining



Improve training time with caching

- In-memory

```
tf.data.Dataset.cache()
```

- Disk

```
tf.data.Dataset.cache(filename=...)
```

Caching with tf.data

```
dataset = tfds.load('cats_vs_dogs', split=tfds.Split.TRAIN)
```

```
# In-memory caching
```

```
train_dataset = dataset.cache()
```

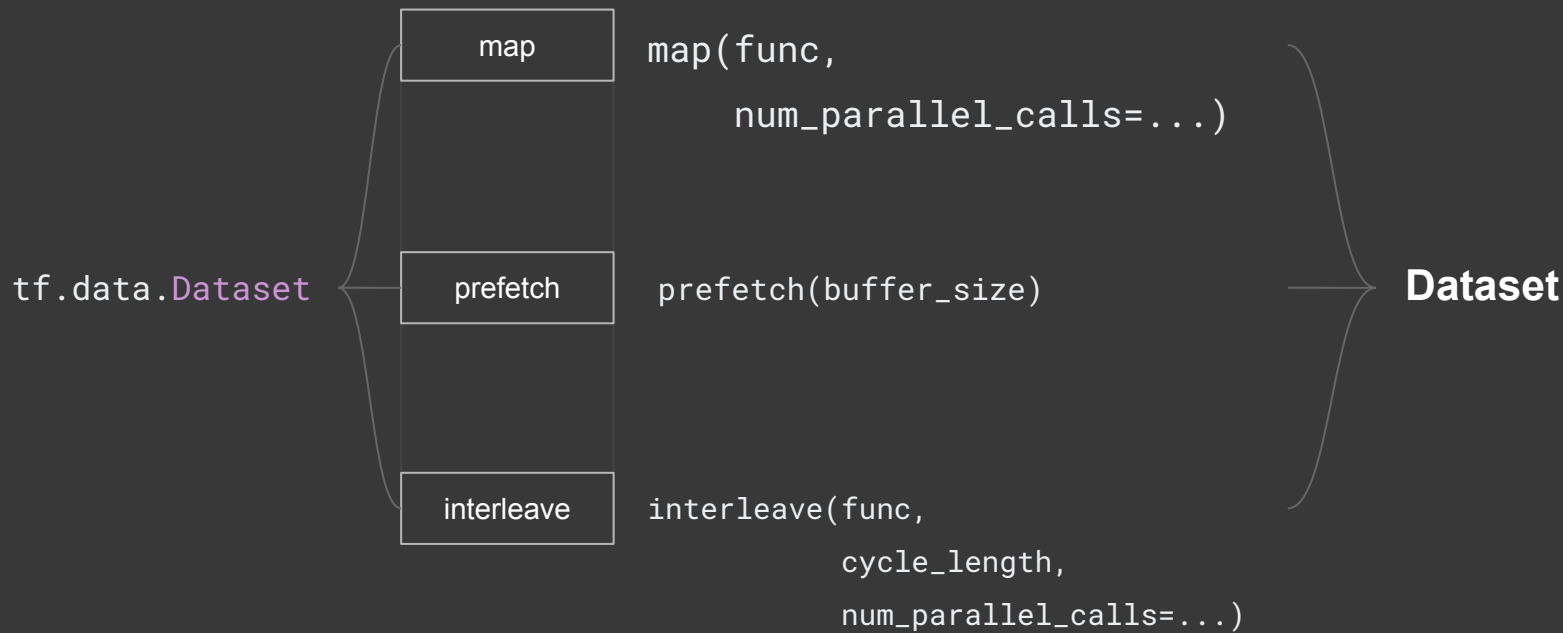
```
model.fit(train_dataset, epochs=...)
```

```
# Disk caching
```

```
train_dataset = dataset.cache(filename='cache')
```

```
model.fit(train_dataset, epochs=...)
```


Parallelism with tf.data



Data transformations

- Transformations can be expensive
- Time-consuming as CPU is not fully utilized

e.g., Resizing, preprocessing, augmentation in images

Consider the following transformation

```
def augment(features):  
    X = tf.image.random_flip_left_right(features['image'])  
    X = tf.image.random_flip_up_down(X)  
    X = tf.image.random_brightness(X, max_delta=0.1)  
    X = tf.image.random_saturation(X, lower=0.75, upper=1.5)  
    X = tf.image.random_hue(X, max_delta=0.15)  
    X = tf.image.random_contrast(X, lower=0.75, upper=1.5)  
    X = tf.image.resize(X, (224, 224))  
    image = X / 255.0  
    return image, features['label']
```

Whats happens when you map that transformation?

```
dataset = tfds.load('cats_vs_dogs',  
                    split=tfds.Split.TRAIN)
```

```
augmented_dataset = dataset.map(augment)
```

Parallelizing data transformation

```
map(func, num_parallel_calls=...)
```

```
augmented_dataset = dataset.map(augment, num_parallel_calls=1)
```

Parallelizing data transformation

```
map(func, num_parallel_calls=...)
```

```
augmented_dataset = dataset.map(augment, num_parallel_calls=1)
```

Maximizing the utilization of CPU cores

```
# Get the number of available cpu cores
num_cores = multiprocessing.cpu_count()

# Set num_parallel_calls with 'num_cores'
augmented_dataset = dataset.map(augment, num_parallel_calls=num_cores)
```

Autotuning

- **`tf.data.experimental.AUTOTUNE`**
- Tunes the value dynamically at runtime
- Decides on the level of parallelism
- Tweaks values of parameters in transformations (`tf.data`)
 - Buffer size (`map`, `prefetch`, `shuffle`, ...)
 - CPU budget (`num_parallel_calls`)
 - I/O (`num_parallel_reads`)

Autotune in practice

```
from tensorflow.data.experimental import AUTOTUNE
```

```
augmented_dataset = dataset.map(  
    augment,  
    num_parallel_calls=AUTOTUNE)
```

Maximizing utilization

With
prefetch



Parallelizing data loading

`prefetch(buffer_size)`

```
dataset = tfds.load('cats_vs_dogs', split=tfds.Split.TRAIN)
```

```
# With prefetch
```

```
train_dataset = dataset.map(format_image).prefetch(tf.data.experimental.AUTOTUNE)
```

Parallelizing data loading


`prefetch(buffer_size)`

```
dataset = tfds.load('cats_vs_dogs', split=tfds.Split.TRAIN)
```

```
# With prefetch
```

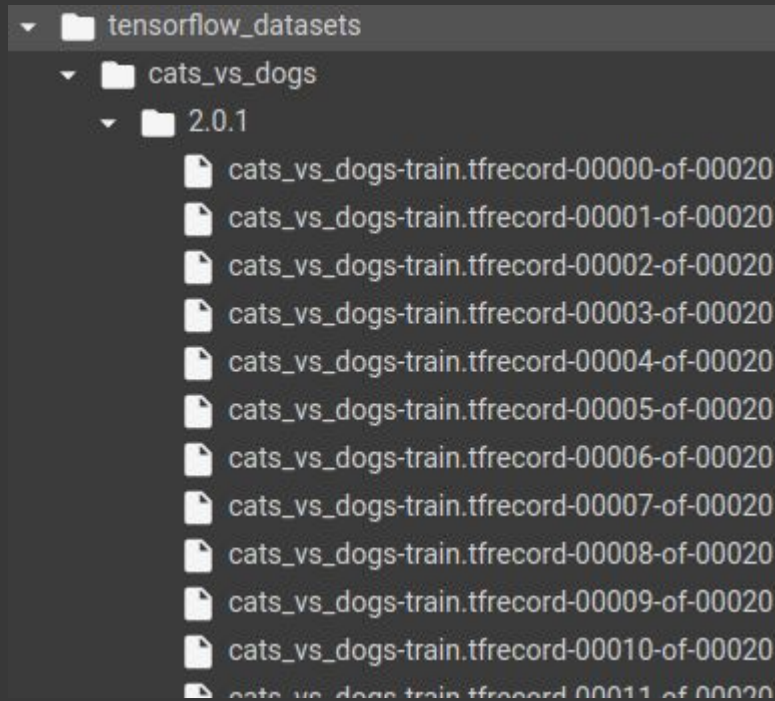
```
train_dataset = dataset.map(format_image).prefetch(tf.data.experimental.AUTOTUNE)
```

Maximizing I/O utilization



Let's inspect TFRecords of a TFDS

```
dataset = tfds.load(name='cats_vs_dogs',  
                    split=tfds.Split.TRAIN)
```



Parallelizing data extraction

```
TFRECORDS_DIR = '/root/tensorflow_datasets/cats_vs_dogs/<dataset-version>/'
files = tf.data.Dataset.list_files(TFRECORDS_DIR +
                                   "cats_vs_dogs-train.tfrecord-*")

num_parallel_reads = 4

dataset = files.interleave(
    tf.data.TFRecordDataset, # map function
    cycle_length=num_parallel_reads, # ...
    num_parallel_calls=tf.data.experimental.AUTOTUNE) # ...
```

Parallelizing data extraction

```
TFRECORDS_DIR = '/root/tensorflow_datasets/cats_vs_dogs/<dataset-version>/'
files = tf.data.Dataset.list_files(TFRECORDS_DIR +
                                   "cats_vs_dogs-train.tfrecord-*")

num_parallel_reads = 4

dataset = files.interleave(
    tf.data.TFRecordDataset, # map function
    cycle_length=num_parallel_reads, # ...
    num_parallel_calls=tf.data.experimental.AUTOTUNE) # ...
```


Performance considerations

- The Dataset APIs are designed to be flexible
- Most operations are commutative
- Order transformations accordingly

e.g., map, batch, shuffle, repeat, interleave, prefetch, etc.,

Map and Batch

The map transformation has overhead in terms of

- Scheduling
- Executing the user-defined function

Map and Batch

Solution: Vectorize the user-defined function

```
dataset = dataset.batch(BATCH_SIZE).map(func)
```

or

```
options = tf.data.Options()  
options.experimental_optimization.map_vectorization.enabled = True  
dataset= dataset.with_options(options)
```

Map and Batch

Solution: Vectorize the user-defined function

```
dataset = dataset.batch(BATCH_SIZE).map(func)
```

or

```
options = tf.data.Options()  
options.experimental_optimization.map_vectorization.enabled = True  
dataset= dataset.with_options(options)
```

Map and Batch

Solution: Vectorize the user-defined function

```
dataset = dataset.batch(BATCH_SIZE).map(func)
```

or

```
options = tf.data.Options()  
options.experimental_optimization.map_vectorization.enabled = True  
dataset= dataset.with_options(options)
```

Map and Cache

```
# Use map before cache when the transformation is expensive  
transformed_dataset = dataset.map(transforms).cache()
```

Shuffle and Repeat

- Shuffling the dataset before applying repeat can cause slow downs
- `shuffle.repeat` for ordering guarantees
- `repeat.shuffle` for better performance

Map and (Interleave / Prefetch / Shuffle)

- All transformations maintain an internal buffer
- Memory footprint is affected if map affects the size of elements
- Generally, have order that affects the memory usage the least