

# Copyright Notice

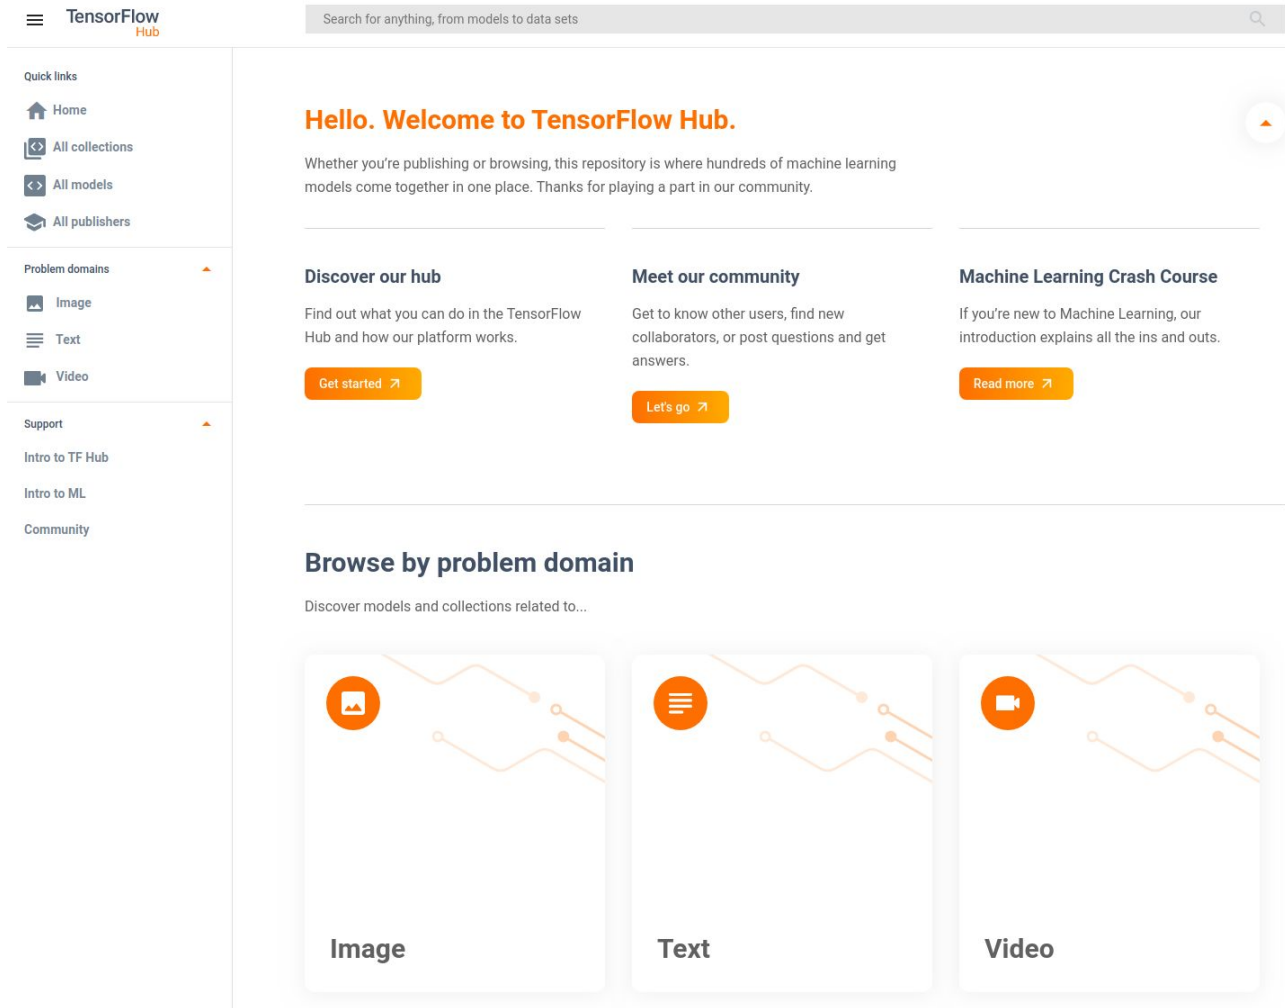
These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>

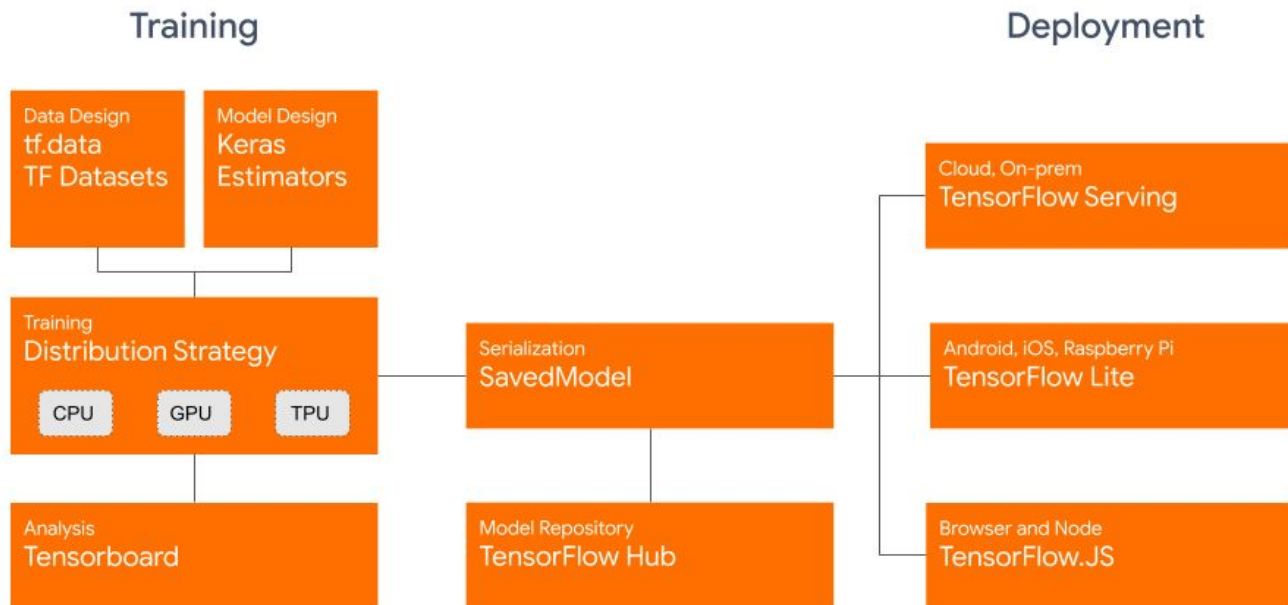
# tfhub.dev



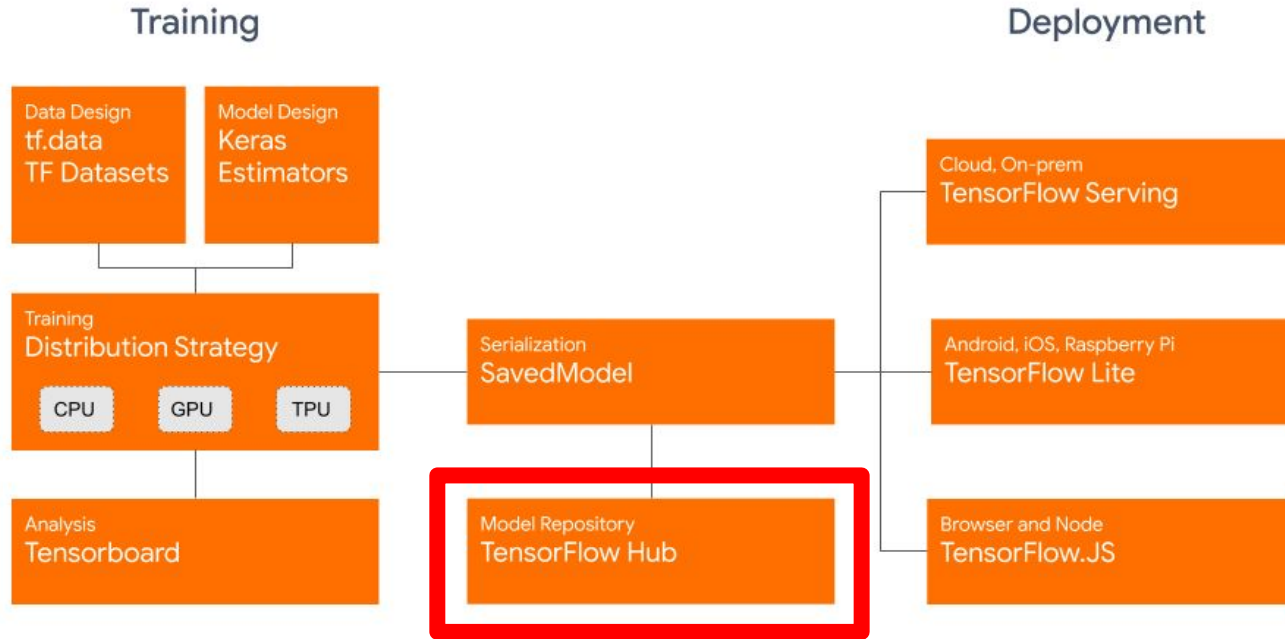
# TensorFlow Hub

- An extensive library of existing modules
- Well-tested modules designed for better accuracy
- Easy to use and integrate with your model APIs

# Where does Hub fit in?



# Where does Hub fit in?



# Who publishes the modules?



BigGAN

Inflated 3D

Wide ResNets



MobileNet

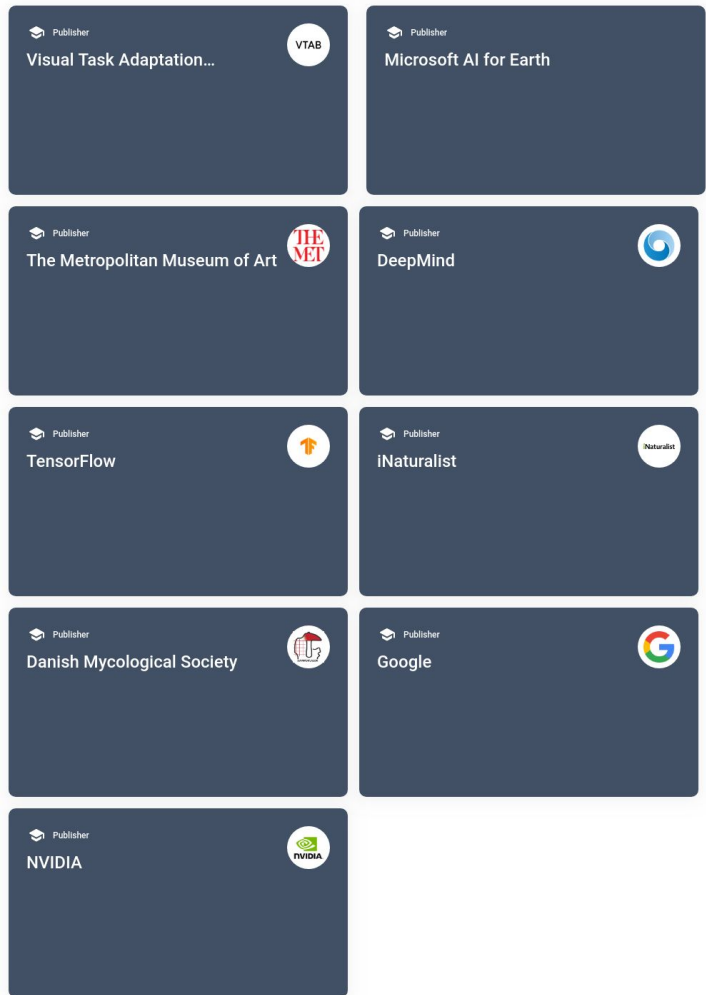
Inception

NASNet



TF-GAN

<https://tfhub.dev/s?subtype=publisher>



# Transfer Learning with TF Hub

- Train models with less data
- Improve generalization
- Speed up training



# Problem domains

## Text

Embedding

## Image

Classification

Feature vector

Augmentation

Object Detection

Generator

Style Transfer

## Video

Classification

# Installation

```
pip install tensorflow_hub
```

```
import tensorflow_hub as hub
```

# Loading a module

```
MODULE_HANDLE = 'https://tfhub.dev/google/tf2-preview/mobilenet_v2/classification/4'  
module = hub.load(MODULE_HANDLE) # Load a the MobileNet image classification module
```

Problem domain

Classification

Publisher

Google

Architecture

V2

Data set

ilsvrc 2012 cls

Fine tunable: No

Last updated: 10/17/2019

Imagenet (ILSVRC-2012-CLS) classification with MobileNet V2 (depth multiplier 1.00).

V4

## TF2 SavedModel

This is a [SavedModel in TensorFlow 2 format](#). Using it requires TensorFlow 2 (or 1.15) and TensorFlow Hub 0.5.0 or newer.

## Overview

MobileNet V2 is a family of neural network architectures for efficient on-device image classification and related tasks, originally published by

- Mark Sandler, Andrew Howard, Menglong Zhu, Andrey Zhmoginov, Liang-Chieh Chen: ["Inverted Residuals and Linear Bottlenecks: Mobile Networks for Classification, Detection and Segmentation"](#), 2018.

Mobilenets come in various sizes controlled by a multiplier for the depth (number of features) in the convolutional layers. They can also be trained for various sizes of input images to control inference speed.

This TF Hub model uses the TF-Slim implementation of `mobilenet_v2` with a depth multiplier of 1.0 and an input size of 160x160 pixels.

The model contains a trained instance of the network, packaged to do the [image classification](#) that the network was trained on. If you merely want to transform images into feature vectors, use `google/imagenet/mobilenet_v2_100_160/feature_vector/4` instead, and save the space occupied by the classification layer.

## Training

The checkpoint exported into this model was `mobilenet_v2_1.0_160/mobilenet_v2_1.0_160.cpkt` downloaded from [MobileNet V2 pre-trained models](#). Its weights were originally obtained by training on the ILSVRC-2012-CLS dataset for image classification ("Imagenet").

## Usage

This model can be used with the `hub.KerasLayer` as follows. It cannot be used with the `hub.Module` API for TensorFlow 1.

```
m = tf.keras.Sequential([
    hub.KerasLayer("https://tfhub.dev/google/imagenet/mobilenet_v2_100_160/classification/4")
])
m.build([None, 160, 160, 3]) # Batch input shape.
```

# Running inference on the module

```
MODULE_HANDLE = 'https://tfhub.dev/google/tf2-preview/mobilenet_v2/classification/4'
```

```
module = hub.load(MODULE_HANDLE)
```

```
images = ... # Batches of images
```

```
predictions = tf.nn.softmax(module(images))
```



Label	Probability
Labrador retriever	0.58
kuvasz	0.08
Great Pyrenees	0.07

# Using a module with Keras

```
MODULE_HANDLE = 'https://tfhub.dev/google/tf2-preview/mobilenet_v2/classification/4'  
OUTPUT_SIZE = 1001  
IMAGE_SIZE = (224, 224)  
  
hub.KerasLayer(MODULE_HANDLE,  
                output_shape=[OUTPUT_SIZE],  
                input_shape=IMAGE_SIZE + (3,))
```

# Using a module with Keras

```
MODULE_HANDLE = 'https://tfhub.dev/google/tf2-preview/mobilenet_v2/classification/4'
OUTPUT_SIZE = 1001
IMAGE_SIZE = (224, 224)

model = tf.keras.Sequential([
    hub.KerasLayer(MODULE_HANDLE,
                    output_shape=[OUTPUT_SIZE], input_shape=IMAGE_SIZE + (3,)),
    tf.keras.layers.Activation('softmax')
])

images = ... # Batches of images
predictions = model.predict(images)
```

# Using a Feature Vector

```
MODULE_HANDLE = "https://tfhub.dev/google/tf2-preview/mobilenet_v2/feature_vector/4"
FV_SIZE = 1280
IMAGE_SIZE = (224, 224)


model = tf.keras.Sequential([
    hub.KerasLayer(MODULE_HANDLE,
                   output_shape=[FV_SIZE],
                   input_shape=IMAGE_SIZE + (3,)),
    tf.keras.layers.Dense(NUM_CLASSES, activation='softmax')
])

images = ... # Batches of images
predictions = model.predict(images)
```



# Where do the modules get stored?

```
module = hub.load(...)
```



# Saving a module for local use

```
MODULE_HANDLE = 'https://tfhub.dev/...?tf-hub-format=compressed'  
!wget $MODULE_HANDLE  
  
# Untar the tarball and load it with hub  
hub_module = hub.load('/path/to/saved_model')
```

# Inspecting the Hub module

```
SAVED_MODEL_DIR = ...
```

```
model = tf.saved_model.load(SAVED_MODEL_DIR, tags='serve')
```

```
images = ... # Batch of images
```

```
predictions = model(images)
```

```
>>> print(predictions)
```

```
Tensor("StatefulPartitionedCall_2:0", shape=(1, 1001), dtype=float32)
```


# Relocating TF Hub modules

```
import os
os.environ['TFHUB_CACHE_DIR'] = '/home/hub_cache_dir'

export TFHUB_CACHE_DIR='/home/hub_cache_dir'
```

# Relocating TF Hub modules

```
module = hub.load(...)
```



# IMDB Reviews

```
train_validation_split = tfds.Split.TRAIN.subsplit([6, 4])

(train_data, validation_data), test_data = tfds.load(
    name="imdb_reviews",
    split=(train_validation_split, tfds.Split.TEST),
    as_supervised=True)
```

# Explore the dataset

```
train_examples_batch, train_labels_batch = next(iter(train_data.batch(10)))
```

```
>>> train_examples_batch.numpy()
```

```
array([b"As a lifelong fan of Dickens, I have invariably been disappointed by ...",  
      b"I absolutely LOVED this movie when I was a kid. I cried every time I ...",  
      ...  
      dtype=object])
```

# Building a model for text classification

## Architectural decisions

- How to **represent** the text?
- How many **layers** to use in the model?
- How many **hidden units** to use for each **layer**?



# Embeddings to choose from

- [google/tf2-preview/gnews-swivel-20dim\(-with-oov\)/1](#)
- [google/tf2-preview/nnlm-en-dim50/1](#)
- [google/tf2-preview/nnlm-en-dim128/a1](#)

# Experimenting with embeddings

```
embedding = "https://tfhub.dev/google/tf2-preview/gnews-swivel-20dim/1"

hub_layer = hub.KerasLayer(embedding, input_shape=[], dtype=tf.string, trainable=True)

>>> hub_layer(train_examples_batch[:3])

<tf.Tensor: id=910, shape=(3, 20), dtype=float32, numpy=
array([[ 3.9819887 , -4.4838037 ,  5.177359  , -2.3643482 , -3.2938678 ,
        -3.5364532 , -2.4786978 ,  2.5525482 ,  6.688532  , -2.3076782 ,
        -1.9807833 ,  1.1315885 , -3.0339816 , -0.7604128 , -5.743445  ,
         3.4242578 ,  4.790099  , -4.03061  , -5.992149  , -1.7297493 ],
       ...
       dtype=float32)>
```

# Experimenting with embeddings

```
embedding = "https://tfhub.dev/google/tf2-preview/gnews-swivel-20dim/1"
```

```
hub_layer = hub.KerasLayer(embedding, input_shape=[], dtype=tf.string, trainable=True)
```

```
>>> hub_layer(train_examples_batch[:3])
```

```
<tf.Tensor: id=910, shape=(3, 20), dtype=float32, numpy=
```

```
array([[ 3.9819887, -4.4838037,  5.177359, -2.3643482, -3.2938678,
        -3.5364532, -2.4786978,  2.5525482,  6.688532, -2.3076782,
        -1.9807833,  1.1315885, -3.0339816, -0.7604128, -5.743445,
         3.4242578,  4.790099, -4.03061, -5.992149, -1.7297493 ],
```

```
...
```

```
dtype=float32)>
```

# Experimenting with embeddings

```
embedding = "https://tfhub.dev/google/tf2-preview/gnews-swivel-20dim/1"
```

```
hub_layer = hub.KerasLayer(embedding, input_shape=[], dtype=tf.string, trainable=True)
```

```
>>> hub_layer(train_examples_batch[:3])
```

```
<tf.Tensor: id=910, shape=(3, 20), dtype=float32, numpy=
```

```
array([[ 3.9819887, -4.4838037,  5.177359, -2.3643482, -3.2938678,
        -3.5364532, -2.4786978,  2.5525482,  6.688532, -2.3076782,
        -1.9807833,  1.1315885, -3.0339816, -0.7604128, -5.743445,
         3.4242578,  4.790099, -4.03061, -5.992149, -1.7297493 ],
```

```
...
```

```
dtype=float32)>
```

# Experimenting with embeddings

```
embedding = "https://tfhub.dev/google/tf2-preview/gnews-swivel-20dim/1"
```

```
hub_layer = hub.KerasLayer(embedding, input_shape=[], dtype=tf.string, trainable=True)
```

```
>>> hub_layer(train_examples_batch[:3])  
<tf.Tensor: id=910, shape=(3, 20), dtype=float32, numpy=  
array([[ 3.9819887, -4.4838037,  5.177359, -2.3643482, -3.2938678,  
        -3.5364532, -2.4786978,  2.5525482,  6.688532, -2.3076782,  
        -1.9807833,  1.1315885, -3.0339816, -0.7604128, -5.743445,  
         3.4242578,  4.790099, -4.03061, -5.992149, -1.7297493],  
        ...  
        dtype=float32)>
```

# Create the model

```
model = tf.keras.Sequential([  
    hub_layer,  
    tf.keras.layers.Dense(16, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')])
```

# Train the model

```
# Compile the model
```

```
model.compile(optimizer='adam',  
              loss='binary_crossentropy',  
              metrics=['accuracy'])
```

```
# Start training
```

```
model.fit(train_data.shuffle(10000).batch(512),  
          epochs=20,  
          validation_data=validation_data.batch(512),  
          verbose=1)
```

# Time to evaluate our model

```
results = model.evaluate(test_data.batch(512), verbose=2)
```

```
>>> for name, value in zip(model.metrics_names, results):  
    print("%s: %.3f" % (name, value))
```

```
loss: 0.318
```

```
accuracy: 0.865
```



# Classify Cats and Dogs

```
splits = tfds.Split.ALL.subsplit(weighted=(80, 10, 10))
```

```
splits, info = tfds.load('cats_vs_dogs', with_info=True, as_supervised=True, split=splits)
```

```
(train_examples, validation_examples, test_examples) = splits
```

```
num_examples = info.splits['train'].num_examples
```

```
num_classes = info.features['label'].num_classes
```

# Prepare datasets

```
def format_image(image, label):  
    # Resizes and normalizes the image  
    image = tf.image.resize(image, IMAGE_SIZE) / 255.0  
    return image, label  
  
# Crate train and validation datasets  
train_batches = train_examples.shuffle(SHUFFLE_SIZE)  
    .map(format_image)  
    .batch(BATCH_SIZE)  
    .prefetch(tf.data.experimental.AUTOTUNE)  
validation_batches = validation_examples.map(format_image)  
    .batch(BATCH_SIZE)  
    .prefetch(tf.data.experimental.AUTOTUNE)
```

# Select feature vector module

```
# Pick a feature vector module. e.g., InceptionV3, MobileNetV2, ...
handle_base = "mobilenet_v2"
MODULE_HANDLE = "https://tfhub.dev/google/tf2-preview/{}/feature_vector/4"
                .format(handle_base)

# This is size of the feature vector
FV_SIZE = 1280 # For MobileNetV2

IMAGE_SIZE = (224, 224)
```

[https://tfhub.dev/google/imagenet/mobilenet\\_v2\\_050\\_96/feature\\_vector/4](https://tfhub.dev/google/imagenet/mobilenet_v2_050_96/feature_vector/4)

feature\_vector

Publisher

Google

Architecture

V2

Data set

ILSVRC 2012 cls

Fine tunable: No

Last updated: 10/17/2019

V4

## TF2 SavedModel

This is a [SavedModel in TensorFlow 2 format](#). Using it requires TensorFlow 2 (or 1.15) and TensorFlow Hub 0.5.0 or newer.

## Overview

MobileNet V2 is a family of neural network architectures for efficient on-device image classification and related tasks, originally published by

- Mark Sandler, Andrew Howard, Menglong Zhu, Andrey Zhmoginov, Liang-Chieh Chen: "[Inverted Residuals and Linear Bottlenecks: Mobile Networks for Classification, Detection and Segmentation](#)", 2018.

Mobilenets come in various sizes controlled by a multiplier for the depth (number of features) in the convolutional layers. They can also be trained for various sizes of input images to control inference speed.

This TF Hub model uses the TF-Slim implementation of `mobilenet_v2` with a depth multiplier of 0.5 and an input size of 96x96 pixels. This implementation of Mobilenet V2 rounds feature depths to multiples of 8 (an optimization *not* described in the paper). Depth multipliers less than 1.0 are not applied to the last convolutional layer (from which the module takes the image feature vector).

The model contains a trained instance of the network, packaged to get [feature vectors from images](#). If you want the full model including the classification it was originally trained for, use `google/imagenet/mobilenet_v2_050_96/classification/4` instead.

## Training

The checkpoint exported into this model was `mobilenet_v2_0.5_96/mobilenet_v2_0.5_96.ckpt` downloaded from [MobileNet V2 pre-trained models](#). Its weights were originally obtained by training on the ILSVRC-2012-CLS dataset for image classification ("Imagenet").

## Usage

This model can be used with the `hub.KerasLayer` as follows. It *cannot* be used with the `hub.Module` API for TensorFlow 1.

```
m = tf.keras.Sequential([
    hub.KerasLayer("https://tfhub.dev/google/imagenet/mobilenet_v2_050_96/feature_vector/4",
                  trainable=False), # Can be True, see below.
    tf.keras.layers.Dense(num_classes, activation='softmax')
])
m.build([None, 96, 96, 3]) # Batch input shape.
```



# Transfer Learning with TF Hub

```
MODULE_HANDLE = 'https://tfhub.dev/google/tf2-preview/mobilenet_v2/feature_vector/4'
```

```
feature_extractor = hub.KerasLayer(MODULE_HANDLE,  
                                   input_shape=IMAGE_SIZE + (3,),  
                                   output_shape=[FV_SIZE],  
                                   trainable=True)
```

```
>>> len(feature_extractor.variables)
```

```
260
```

```
>>> len(feature_extractor.trainable_variables)
```

```
156
```

```
>>> len(feature_extractor.non_trainable_variables)
```

```
104
```

# Build model with TF Hub module

```
# Fine-tuning switch
```

```
do_fine_tuning = True
```

```
feature_extractor = hub.KerasLayer(MODULE_HANDLE,  
                                   input_shape=IMAGE_SIZE + (3,),  
                                   output_shape=[FV_SIZE],  
                                   trainable=do_fine_tuning)
```

```
# Build model with the chosen module handle
```

```
model = tf.keras.Sequential([  
    feature_extractor,  
    tf.keras.layers.Dense(2, activation='softmax') # (cat, dog)  
])
```

# Model summary

```
>>> model.summary()
```

```
Model: "sequential"
```

```
-----  
Layer (type)                 Output Shape                 Param #  
=====
```

```
keras_layer (KerasLayer)     (None, 1280)                 2257984  
-----
```

```
dense (Dense)                 (None, 2)                   2562  
=====
```

```
Total params: 2,260,546
```

```
Trainable params: 2,226,434
```

```
Non-trainable params: 34,112  
-----
```

# Choose the right training configuration

```
# Define loss and metrics
```

```
loss = 'sparse_categorical_crossentropy'
```

```
metrics = ['accuracy']
```

```
# Choose an appropriate optimizer when fine-tuning
```

```
if do_fine_tuning:
```

```
    optimizer = tf.keras.optimizers.SGD(lr=0.002, momentum=0.9)
```

```
else:
```

```
    optimizer = 'adam'
```



# Train the model

```
# Compile the model and fit
```

```
model.compile(optimizer=optimizer, loss=loss, metrics=metrics)
```

```
model.fit(train_batches, epochs=5,  
          validation_data=validation_batches)
```