

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>

# Introduction to Callbacks

# Callbacks

- Provides some functionality at various stages of training
- Subclasses `tf.keras.callbacks.Callback`
- Useful in understanding a model's state during training
  - internal states
  - statistics e.g., losses and metrics

# Training specific methods

```
class Callback(object):  
    def __init__(self):  
        self.validation_data = None  
        self.model = None  
  
    def on_epoch_begin(self, epoch, logs=None):  
        """Called at the beginning of an epoch during training."""  
  
    def on_epoch_end(self, epoch, logs=None):  
        """Called at the end of an epoch during training."""
```

# Training specific methods

```
class Callback(object):  
    def __init__(self):  
        self.validation_data = None  
        self.model = None
```

```
    def on_epoch_begin(self, epoch, logs=None):  
        """Called at the beginning of an epoch during training."""
```

```
    def on_epoch_end(self, epoch, logs=None):  
        """Called at the end of an epoch during training."""
```

# Training specific methods

```
class Callback(object):  
    def __init__(self):  
        self.validation_data = None  
        self.model = None  
  
    def on_epoch_begin(self, epoch, logs=None):  
        """Called at the beginning of an epoch during training."""  
  
    def on_epoch_end(self, epoch, logs=None):  
        """Called at the end of an epoch during training."""
```

# Common methods for training/testing/predicting

```
class Callback(object):  
    ...  
    def on_(train|test|predict)_begin(self, logs=None):  
        """Called at the begin of fit/evaluate/predict."""  
  
    def on_(train|test|predict)_end(self, logs=None):  
        """Called at the end of fit/evaluate/predict."""  
  
    def on_(train|test|predict)_batch_begin(self, batch, logs=None):  
        """Called right before processing a batch during training/testing/predicting."""  
  
    def on_(train|test|predict)_batch_end(self, batch, logs=None):  
        """Called at the end of training/testing/predicting a batch."""
```

# Common methods for training/testing/predicting

```
class Callback(object):
```

```
...
```

```
def on_(train|test|predict)_begin(self, logs=None):
```

```
    """Called at the begin of fit/evaluate/predict."""
```

```
def on_(train|test|predict)_end(self, logs=None):
```

```
    """Called at the end of fit/evaluate/predict."""
```

```
def on_(train|test|predict)_batch_begin(self, batch, logs=None):
```

```
    """Called right before processing a batch during training/testing/predicting."""
```

```
def on_(train|test|predict)_batch_end(self, batch, logs=None):
```

```
    """Called at the end of training/testing/predicting a batch."""
```



# Common methods for training/testing/predicting

```
class Callback(object):  
    ...  
    def on_(train|test|predict)_begin(self, logs=None):  
        """Called at the begin of fit/evaluate/predict."""  
  
    def on_(train|test|predict)_end(self, logs=None):  
        """Called at the end of fit/evaluate/predict."""  
  
    def on_(train|test|predict)_batch_begin(self, batch, logs=None):  
        """Called right before processing a batch during training/testing/predicting."""  
  
    def on_(train|test|predict)_batch_end(self, batch, logs=None):  
        """Called at the end of training/testing/predicting a batch."""
```

# Where can you use them?

Model methods that take callbacks

- `fit(..., callbacks=[...])`
- `fit_generator(..., callbacks=[...])`
- `evaluate(..., callbacks=[...])`
- `evaluate_generator(..., callbacks=[...])`
- `predict(..., callbacks=[...])`
- `predict_generator(..., callbacks=[...])`

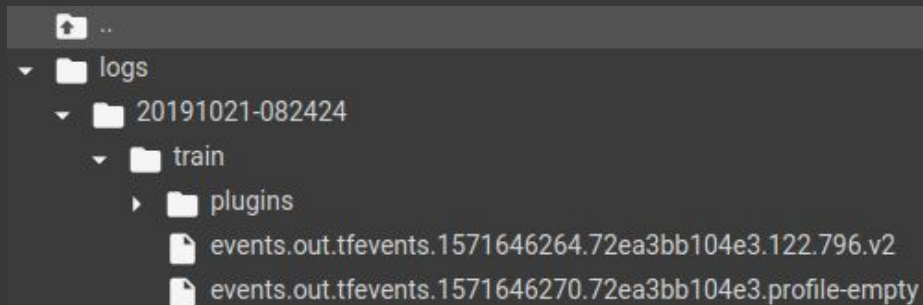
# TensorBoard Callback

- Visualize machine learning experiments
- Track metrics (e.g., loss, accuracy)
- View the model graph

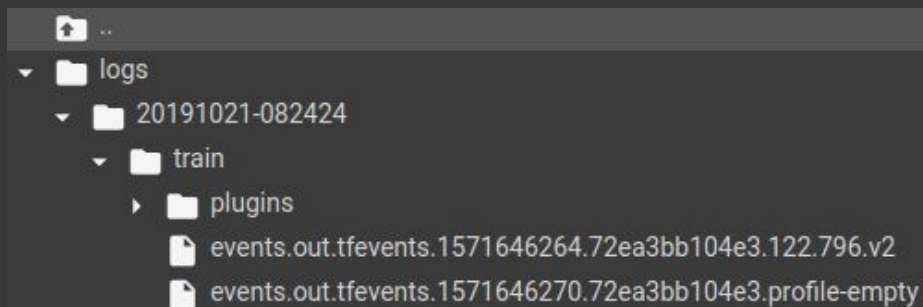
```
TensorBoard(log_dir='./logs', update_freq='epoch', **kwargs)
```

<https://www.tensorflow.org/tensorboard>

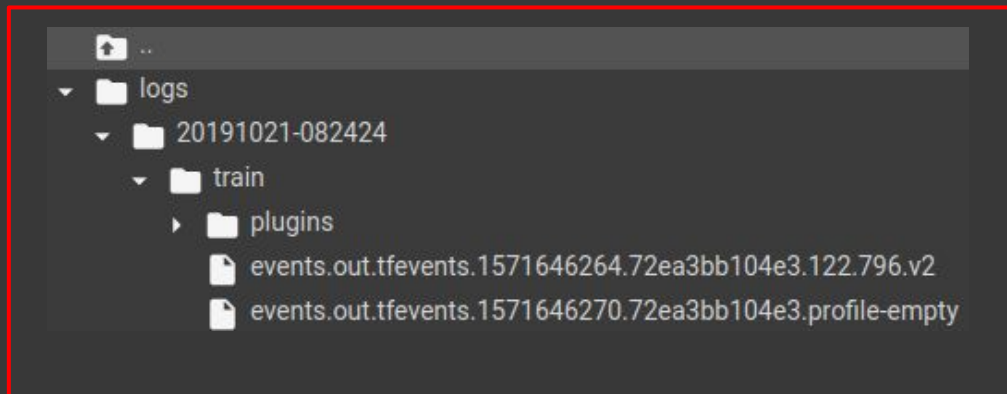
```
log_dir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))  
tensorboard = tf.keras.callbacks.TensorBoard(log_dir=log_dir)  
model.fit(train_batches, epochs=10, callbacks=[tensorboard])
```



```
log_dir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))  
tensorboard = tf.keras.callbacks.TensorBoard(log_dir=log_dir)  
model.fit(train_batches, epochs=10, callbacks=[tensorboard])
```



```
log_dir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))  
tensorboard = tf.keras.callbacks.TensorBoard(log_dir=log_dir)  
model.fit(train_batches, epochs=10, callbacks=[tensorboard])
```

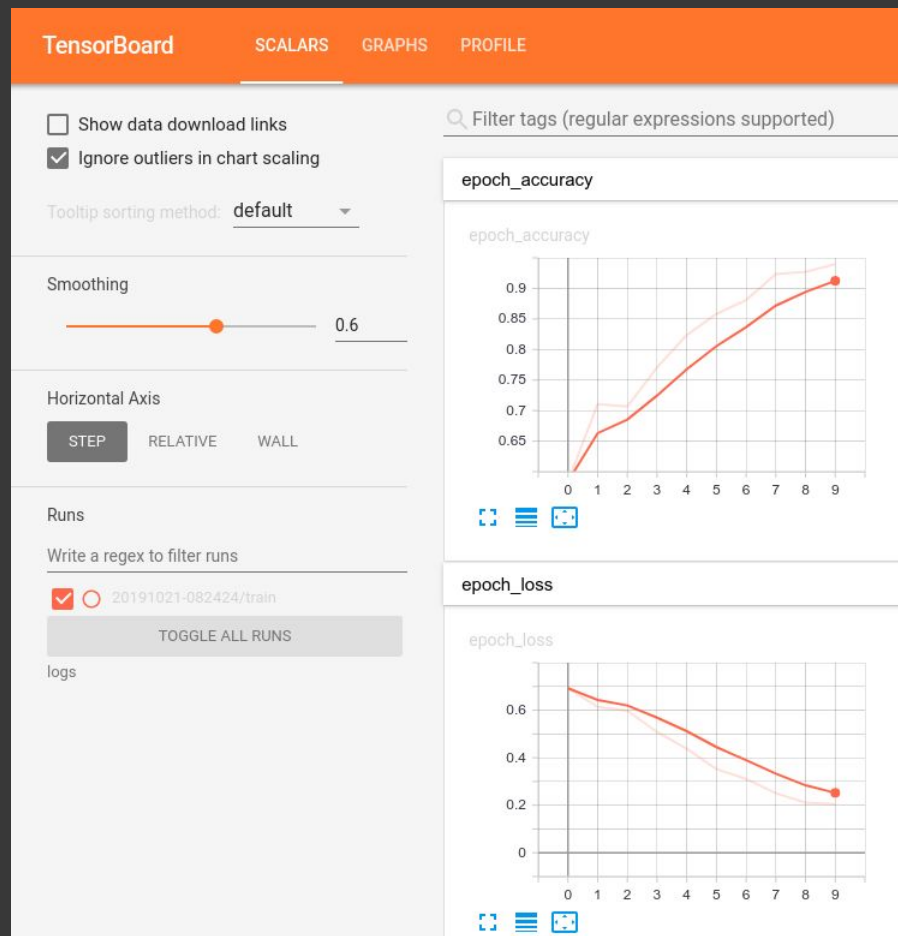


```
# Load the extension
```

```
%load_ext tensorboard
```

```
# Run TensorBoard
```

```
%tensorboard --logdir logs
```



# Model Checkpoints



# ModelCheckpoint

- Saves the model every so often
- Choose to save only the best checkpoints / weights

```
ModelCheckpoint(filepath, monitor='val_loss', mode='auto',  
                save_best_only=False, save_weights_only=False,  
                verbose=0, save_freq=1, **kwargs)
```

```
model.fit(train_batches, epochs=5, validation_data=validation_batches,  
          callbacks=[ModelCheckpoint('model.h5', verbose=1)])
```

Epoch 1/5

Epoch 00001: saving model to model.h5

33/33 - 7s - loss: 0.6879 - accuracy: 0.6702 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/5

Epoch 00002: saving model to model.h5

33/33 - 6s - loss: 0.6721 - accuracy: 0.8447 - val\_loss: 0.6608 - val\_accuracy: 0.8667

Epoch 3/5

Epoch 00003: saving model to model.h5

33/33 - 6s - loss: 0.6435 - accuracy: 0.8840 - val\_loss: 0.6217 - val\_accuracy: 0.9417

Epoch 4/5

Epoch 00004: saving model to model.h5

33/33 - 6s - loss: 0.5920 - accuracy: 0.8849 - val\_loss: 0.5591 - val\_accuracy: 0.8667

Epoch 5/5

Epoch 00005: saving model to model.h5

33/33 - 6s - loss: 0.5047 - accuracy: 0.9089 - val\_loss: 0.4485 - val\_accuracy: 0.8583

<tensorflow.python.keras.callbacks.History at 0x7f09ccef97f0>

```
model.fit(train_batches, epochs=5, validation_data=validation_batches,  
         callbacks=[ModelCheckpoint('model.h5', verbose=1)])
```

Epoch 1/5

Epoch 00001: saving model to model.h5

33/33 - 7s - loss: 0.6879 - accuracy: 0.6702 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/5

Epoch 00002: saving model to model.h5

33/33 - 6s - loss: 0.6721 - accuracy: 0.8447 - val\_loss: 0.6608 - val\_accuracy: 0.8667

Epoch 3/5

Epoch 00003: saving model to model.h5

33/33 - 6s - loss: 0.6435 - accuracy: 0.8840 - val\_loss: 0.6217 - val\_accuracy: 0.9417

Epoch 4/5

Epoch 00004: saving model to model.h5

33/33 - 6s - loss: 0.5920 - accuracy: 0.8849 - val\_loss: 0.5591 - val\_accuracy: 0.8667

Epoch 5/5

Epoch 00005: saving model to model.h5

33/33 - 6s - loss: 0.5047 - accuracy: 0.9089 - val\_loss: 0.4485 - val\_accuracy: 0.8583

<tensorflow.python.keras.callbacks.History at 0x7f09ccef97f0>

```
model.fit(train_batches, epochs=5, validation_data=validation_batches,  
         callbacks=[ModelCheckpoint('model.h5', verbose=1)])
```

Epoch 1/5

```
Epoch 00001: saving model to model.h5  
33/33 - 7s - loss: 0.6879 - accuracy: 0.6702 - val_loss: 0.0000e+00 - val_accuracy: 0.0000e+00  
Epoch 2/5
```

```
Epoch 00002: saving model to model.h5  
33/33 - 6s - loss: 0.6721 - accuracy: 0.8447 - val_loss: 0.6608 - val_accuracy: 0.8667  
Epoch 3/5
```

```
Epoch 00003: saving model to model.h5  
33/33 - 6s - loss: 0.6435 - accuracy: 0.8840 - val_loss: 0.6217 - val_accuracy: 0.9417  
Epoch 4/5
```

```
Epoch 00004: saving model to model.h5  
33/33 - 6s - loss: 0.5920 - accuracy: 0.8849 - val_loss: 0.5591 - val_accuracy: 0.8667  
Epoch 5/5
```

```
Epoch 00005: saving model to model.h5  
33/33 - 6s - loss: 0.5047 - accuracy: 0.9089 - val_loss: 0.4485 - val_accuracy: 0.8583  
<tensorflow.python.keras.callbacks.History at 0x7f09ccef97f0>
```

```
model.fit(train_batches, epochs=5, validation_data=validation_batches, verbose=2,  
          callbacks=[ModelCheckpoint('model.h5', save_weights_only=True, verbose=1)])
```

Epoch 1/2

Epoch 00001: saving model to model.h5

33/33 - 7s - loss: 0.6493 - accuracy: 0.6184 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/2

Epoch 00002: saving model to model.h5

33/33 - 6s - loss: 0.5684 - accuracy: 0.7507 - val\_loss: 0.5183 - val\_accuracy: 0.7083

<tensorflow.python.keras.callbacks.History at 0x7f09cb5547f0>

```
model.fit(train_batches, epochs=5, validation_data=validation_batches, verbose=2,  
         callbacks=[ModelCheckpoint('model.h5', monitor='val_loss',  
                                   save_best_only=True, verbose=1)])
```

Epoch 1/5

Epoch 00001: val\_loss improved from inf to 0.65278, saving model to model.h5  
33/33 - 7s - loss: 0.6753 - accuracy: 0.5772 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00  
Epoch 2/5

Epoch 00002: val\_loss improved from 0.65278 to 0.62279, saving model to model.h5  
33/33 - 6s - loss: 0.6219 - accuracy: 0.7584 - val\_loss: 0.6228 - val\_accuracy: 0.5417  
Epoch 3/5

Epoch 00003: val\_loss improved from 0.62279 to 0.47633, saving model to model.h5  
33/33 - 6s - loss: 0.5448 - accuracy: 0.7977 - val\_loss: 0.4763 - val\_accuracy: 0.8750  
Epoch 4/5

Epoch 00004: val\_loss improved from 0.47633 to 0.44497, saving model to model.h5  
33/33 - 6s - loss: 0.4673 - accuracy: 0.8054 - val\_loss: 0.4450 - val\_accuracy: 0.8000  
Epoch 5/5

Epoch 00005: val\_loss improved from 0.44497 to 0.30997, saving model to model.h5  
33/33 - 6s - loss: 0.4030 - accuracy: 0.8677 - val\_loss: 0.3100 - val\_accuracy: 0.9000  
<tensorflow.python.keras.callbacks.History at 0x7f09cc9b7128>

```
model.fit(train_batches, epochs=5, validation_data=validation_batches, verbose=2,  
          callbacks=[ModelCheckpoint('model.h5', monitor='val_loss',  
                                     save_best_only=True, verbose=1)])
```

Epoch 1/5

Epoch 00001: val\_loss improved from inf to 0.65278, saving model to model.h5  
33/33 - 7s - loss: 0.6753 - accuracy: 0.5772 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00  
Epoch 2/5

Epoch 00002: val\_loss improved from 0.65278 to 0.62279, saving model to model.h5  
33/33 - 6s - loss: 0.6219 - accuracy: 0.7584 - val\_loss: 0.6228 - val\_accuracy: 0.5417  
Epoch 3/5

Epoch 00003: val\_loss improved from 0.62279 to 0.47633, saving model to model.h5  
33/33 - 6s - loss: 0.5448 - accuracy: 0.7977 - val\_loss: 0.4763 - val\_accuracy: 0.8750  
Epoch 4/5

Epoch 00004: val\_loss improved from 0.47633 to 0.44497, saving model to model.h5  
33/33 - 6s - loss: 0.4673 - accuracy: 0.8054 - val\_loss: 0.4450 - val\_accuracy: 0.8000  
Epoch 5/5

Epoch 00005: val\_loss improved from 0.44497 to 0.30997, saving model to model.h5  
33/33 - 6s - loss: 0.4030 - accuracy: 0.8677 - val\_loss: 0.3100 - val\_accuracy: 0.9000  
<tensorflow.python.keras.callbacks.History at 0x7f09cc9b7128>

```
model.fit(..., callbacks=[ModelCheckpoint('saved_model', ...)])
```

Epoch 1/2

Epoch 00001: saving model to model.h5

33/33 - 7s - loss: 0.6714 - accuracy: 0.5695 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/2

Epoch 00002: saving model to model.h5

33/33 - 6s - loss: 0.6238 - accuracy: 0.6366 - val\_loss: 0.6459 - val\_accuracy: 0.5417

```
model.fit(..., callbacks=[ModelCheckpoint('model.h5', ...)])
```

Epoch 1/2

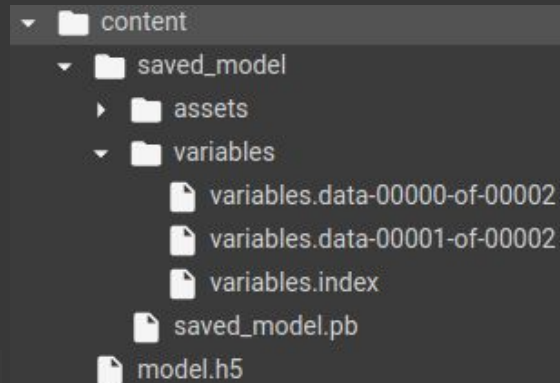
Epoch 00001: saving model to model.h5

33/33 - 7s - loss: 0.6714 - accuracy: 0.5695 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/2

Epoch 00002: saving model to model.h5

33/33 - 6s - loss: 0.6238 - accuracy: 0.6366 - val\_loss: 0.6459 - val\_accuracy: 0.5417





```
model.fit(..., callbacks=[ModelCheckpoint('weights.{epoch:02d}-{val_loss:.2f}.h5', verbose=1)])
```

Epoch 1/5

Epoch 00001: saving model to weights.01-0.63.h5

33/33 - 6s - loss: 0.6709 - accuracy: 0.6098 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/5

Epoch 00002: saving model to weights.02-0.60.h5

33/33 - 6s - loss: 0.6088 - accuracy: 0.7124 - val\_loss: 0.6046 - val\_accuracy: 0.5917

Epoch 3/5

Epoch 00003: saving model to weights.03-0.46.h5

33/33 - 6s - loss: 0.5354 - accuracy: 0.7613 - val\_loss: 0.4602 - val\_accuracy: 0.8500

Epoch 4/5

Epoch 00004: saving model to weights.04-0.38.h5

33/33 - 6s - loss: 0.4769 - accuracy: 0.7891 - val\_loss: 0.3848 - val\_accuracy: 0.9250

Epoch 5/5

Epoch 00005: saving model to weights.05-0.33.h5

33/33 - 6s - loss: 0.3961 - accuracy: 0.8600 - val\_loss: 0.3263 - val\_accuracy: 0.8667

content

weights.01-0.63.h5

weights.02-0.60.h5

weights.03-0.46.h5

weights.04-0.38.h5

weights.05-0.33.h5

```
model.fit(..., callbacks=[ModelCheckpoint('weights.{epoch:02d}-{val_loss:.2f}.h5', verbose=1)])
```

Epoch 1/5

Epoch 00001: saving model to weights.01-0.63.h5

33/33 - 6s - loss: 0.6709 - accuracy: 0.6098 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/5

Epoch 00002: saving model to weights.02-0.60.h5

33/33 - 6s - loss: 0.6088 - accuracy: 0.7124 - val\_loss: 0.6046 - val\_accuracy: 0.5917

Epoch 3/5

Epoch 00003: saving model to weights.03-0.46.h5

33/33 - 6s - loss: 0.5354 - accuracy: 0.7613 - val\_loss: 0.4602 - val\_accuracy: 0.8500

Epoch 4/5

Epoch 00004: saving model to weights.04-0.38.h5

33/33 - 6s - loss: 0.4769 - accuracy: 0.7891 - val\_loss: 0.3848 - val\_accuracy: 0.9250

Epoch 5/5

Epoch 00005: saving model to weights.05-0.33.h5

33/33 - 6s - loss: 0.3961 - accuracy: 0.8600 - val\_loss: 0.3263 - val\_accuracy: 0.8667

content

weights.01-0.63.h5

weights.02-0.60.h5

weights.03-0.46.h5

weights.04-0.38.h5

weights.05-0.33.h5

# EarlyStopping

- Helps you keep track of a certain metric/loss and change training behavior accordingly
- Stops training when there's no improvement observed

```
EarlyStopping(monitor='val_loss', verbose=0,  
              min_delta=0, patience=0, mode='auto',  
              baseline=None, restore_best_weights=False,  
              **kwargs)
```

```
model.fit(train_batches, epochs=50, validation_data=validation_batches,  
          callbacks=[EarlyStopping(patience=3, monitor='val_loss')])
```

```
Epoch 11/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2423 - accuracy: 0.9319 - val_loss: 0.1474 - val_accuracy: 0.9500  
Epoch 12/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1521 - accuracy: 0.9607 - val_loss: 0.1990 - val_accuracy: 0.9000  
Epoch 13/50  
33/33 [=====] - 6s 182ms/step - loss: 0.1571 - accuracy: 0.9511 - val_loss: 0.1176 - val_accuracy: 0.9500  
Epoch 14/50  
33/33 [=====] - 6s 186ms/step - loss: 0.1409 - accuracy: 0.9569 - val_loss: 0.1071 - val_accuracy: 0.9583  
Epoch 15/50  
33/33 [=====] - 6s 185ms/step - loss: 0.1450 - accuracy: 0.9626 - val_loss: 0.0953 - val_accuracy: 0.9583  
Epoch 16/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2023 - accuracy: 0.9396 - val_loss: 0.1413 - val_accuracy: 0.9583  
Epoch 17/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1443 - accuracy: 0.9655 - val_loss: 0.1771 - val_accuracy: 0.9167  
Epoch 18/50  
33/33 [=====] - 6s 183ms/step - loss: 0.1442 - accuracy: 0.9521 - val_loss: 0.1201 - val_accuracy: 0.9333  
Epoch 00018: early stopping
```

```
model.fit(train_batches, epochs=50, validation_data=validation_batches,  
          callbacks=[EarlyStopping(patience=3, monitor='val_loss')])
```

```
Epoch 11/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2423 - accuracy: 0.9319 - val_loss: 0.1474 - val_accuracy: 0.9500  
Epoch 12/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1521 - accuracy: 0.9607 - val_loss: 0.1990 - val_accuracy: 0.9000  
Epoch 13/50  
33/33 [=====] - 6s 182ms/step - loss: 0.1571 - accuracy: 0.9511 - val_loss: 0.1176 - val_accuracy: 0.9500  
Epoch 14/50  
33/33 [=====] - 6s 186ms/step - loss: 0.1409 - accuracy: 0.9569 - val_loss: 0.1071 - val_accuracy: 0.9583  
Epoch 15/50  
33/33 [=====] - 6s 185ms/step - loss: 0.1450 - accuracy: 0.9626 - val_loss: 0.0953 - val_accuracy: 0.9583  
Epoch 16/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2023 - accuracy: 0.9396 - val_loss: 0.1413 - val_accuracy: 0.9583  
Epoch 17/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1443 - accuracy: 0.9655 - val_loss: 0.1771 - val_accuracy: 0.9167  
Epoch 18/50  
33/33 [=====] - 6s 183ms/step - loss: 0.1442 - accuracy: 0.9521 - val_loss: 0.1201 - val_accuracy: 0.9333  
Epoch 00018: early stopping
```

```
model.fit(train_batches, epochs=50, validation_data=validation_batches,  
          callbacks=[EarlyStopping(patience=3, monitor='val_loss')])
```

```
Epoch 11/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2423 - accuracy: 0.9319 - val_loss: 0.1474 - val_accuracy: 0.9500  
Epoch 12/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1521 - accuracy: 0.9607 - val_loss: 0.1990 - val_accuracy: 0.9000  
Epoch 13/50  
33/33 [=====] - 6s 182ms/step - loss: 0.1571 - accuracy: 0.9511 - val_loss: 0.1176 - val_accuracy: 0.9500  
Epoch 14/50  
33/33 [=====] - 6s 186ms/step - loss: 0.1409 - accuracy: 0.9569 - val_loss: 0.1071 - val_accuracy: 0.9583  
Epoch 15/50  
33/33 [=====] - 6s 185ms/step - loss: 0.1450 - accuracy: 0.9626 - val_loss: 0.0953 - val_accuracy: 0.9583  
Epoch 16/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2023 - accuracy: 0.9396 - val_loss: 0.1413 - val_accuracy: 0.9583  
Epoch 17/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1443 - accuracy: 0.9655 - val_loss: 0.1771 - val_accuracy: 0.9167  
Epoch 18/50  
33/33 [=====] - 6s 183ms/step - loss: 0.1442 - accuracy: 0.9521 - val_loss: 0.1201 - val_accuracy: 0.9333  
Epoch 00018: early stopping
```

```
model.fit(train_batches, epochs=50, validation_data=validation_batches,  
          callbacks=[EarlyStopping(patience=3, monitor='val_loss')])
```

```
Epoch 11/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2423 - accuracy: 0.9319 - val_loss: 0.1474 - val_accuracy: 0.9500  
Epoch 12/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1521 - accuracy: 0.9607 - val_loss: 0.1990 - val_accuracy: 0.9000  
Epoch 13/50  
33/33 [=====] - 6s 182ms/step - loss: 0.1571 - accuracy: 0.9511 - val_loss: 0.1176 - val_accuracy: 0.9500  
Epoch 14/50  
33/33 [=====] - 6s 186ms/step - loss: 0.1409 - accuracy: 0.9569 - val_loss: 0.1071 - val_accuracy: 0.9583  
Epoch 15/50  
33/33 [=====] - 6s 185ms/step - loss: 0.1450 - accuracy: 0.9626 - val_loss: 0.0953 - val_accuracy: 0.9583  
Epoch 16/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2023 - accuracy: 0.9396 - val_loss: 0.1413 - val_accuracy: 0.9583  
Epoch 17/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1443 - accuracy: 0.9655 - val_loss: 0.1771 - val_accuracy: 0.9167  
Epoch 18/50  
33/33 [=====] - 6s 183ms/step - loss: 0.1442 - accuracy: 0.9521 - val_loss: 0.1201 - val_accuracy: 0.9333  
Epoch 00018: early stopping
```

```
model.fit(...,  
          callbacks=[EarlyStopping(patience=3, restore_best_weights=True,  
                                   monitor='val_loss', verbose=1)])
```

```
Epoch 11/50  
33/33 - 6s - loss: 0.1380 - accuracy: 0.9616 - val_loss: 0.0968 - val_accuracy: 0.9750  
Epoch 12/50  
33/33 - 6s - loss: 0.1202 - accuracy: 0.9655 - val_loss: 0.0741 - val_accuracy: 0.9917  
Epoch 13/50  
33/33 - 6s - loss: 0.1716 - accuracy: 0.9434 - val_loss: 0.1083 - val_accuracy: 0.9750  
Epoch 14/50  
33/33 - 6s - loss: 0.1331 - accuracy: 0.9626 - val_loss: 0.0861 - val_accuracy: 0.9667  
Epoch 15/50  
Restoring model weights from the end of the best epoch.  
33/33 - 6s - loss: 0.1393 - accuracy: 0.9578 - val_loss: 0.0771 - val_accuracy: 0.9750  
Epoch 00015: early stopping
```



```
model.fit(...,  
           callbacks=[EarlyStopping(  
                                   patience=3,  
                                   min_delta=0.05,  
                                   baseline=0.8,  
                                   mode='min',  
                                   monitor='val_loss',  
                                   verbose=1  
                               )])
```

```
model.fit(..., callbacks=[CSVLogger('training.csv')])
```

epoch	accuracy	loss	val_accuracy	val_loss
0	0.574305	0.682536	0.775000	0.655427
1	0.760307	0.633610	0.675000	0.595201
2	0.758389	0.573186	0.850000	0.503174
3	0.835091	0.472031	0.808333	0.416691
4	0.854267	0.419491	0.916667	0.309128

## Multiple callbacks

```
model.fit(..., callbacks=[EarlyStopping(...),
model.evaluate(...      ModelCheckpoint(...),
model.predict(...      TensorBoard(...),
                        ...
                        ])
```

# Build a simple model

```
model = tf.keras.Sequential()
```

```
model.add(tf.keras.layers.Dense(units=1,  
                                  activation='linear',  
                                  input_dim=(784,)))
```

```
model.compile(optimizer=tf.keras.optimizers.RMSprop(lr=0.1),  
              loss='mean_squared_error', metrics=['mae'])
```

# How a custom callback looks

```
import datetime

class MyCustomCallback(tf.keras.callbacks.Callback):
    def on_train_batch_begin(self, batch, logs=None):
        print('Training: batch {} begins at {}'.format(batch, datetime.datetime.now().time()))

    def on_train_batch_end(self, batch, logs=None):
        print('Training: batch {} ends at {}'.format(batch, datetime.datetime.now().time()))
```

# How a custom callback looks

```
import datetime
```

```
class MyCustomCallback(tf.keras.callbacks.Callback):
```

```
    def on_train_batch_begin(self, batch, logs=None):
```

```
        print('Training: batch {} begins at {}'.format(batch, datetime.datetime.now().time()))
```

```
    def on_train_batch_end(self, batch, logs=None):
```

```
        print('Training: batch {} ends at {}'.format(batch, datetime.datetime.now().time()))
```

```
my_custom_callback = MyCustomCallback()
```

```
model.fit(x_train, y_train, batch_size=64, epochs=1, verbose=0,  
          callbacks=[my_custom_callback])
```

```
my_custom_callback = MyCustomCallback()
```

```
model.fit(x_train, y_train, batch_size=64, epochs=1, verbose=0,  
          callbacks=[my_custom_callback])
```



```
class DetectOverfittingCallback(tf.keras.callbacks.Callback):  
    def __init__(self, threshold):  
        super(DetectOverfittingCallback, self).__init__()  
        self.threshold = threshold  
  
    def on_epoch_end(self, epoch, logs=None):  
        ratio = logs["val_loss"] / logs["loss"]  
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))  
  
        if ratio>threshold:  
            print("Stopping training...")  
            self.model.stop_training = True  
  
model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])
```

```
class DetectOverfittingCallback(tf.keras.callbacks.Callback):  
    def __init__(self, threshold):  
        super(DetectOverfittingCallback, self).__init__()  
        self.threshold = threshold  
  
    def on_epoch_end(self, epoch, logs=None):  
        ratio = logs["val_loss"] / logs["loss"]  
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))  
  
        if ratio > threshold:  
            print("Stopping training...")  
            self.model.stop_training = True  
  
model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])
```

```
class DetectOverfittingCallback(tf.keras.callbacks.Callback):
```

```
    def __init__(self, threshold):
```

```
        super(DetectOverfittingCallback, self).__init__()
```

```
        self.threshold = threshold
```

```
    def on_epoch_end(self, epoch, logs=None):
```

```
        ratio = logs["val_loss"] / logs["loss"]
```

```
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))
```

```
        if ratio>threshold:
```

```
            print("Stopping training...")
```

```
            self.model.stop_training = True
```

```
model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])
```

```
class DetectOverfittingCallback(tf.keras.callbacks.Callback):  
    def __init__(self, threshold):  
        super(DetectOverfittingCallback, self).__init__()  
        self.threshold = threshold
```

```
    def on_epoch_end(self, epoch, logs=None):  
        ratio = logs["val_loss"] / logs["loss"]  
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))
```

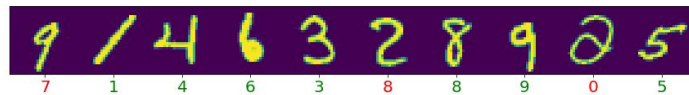
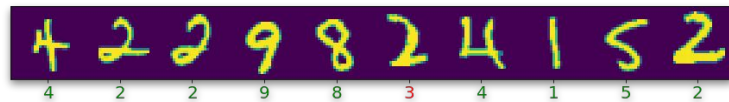
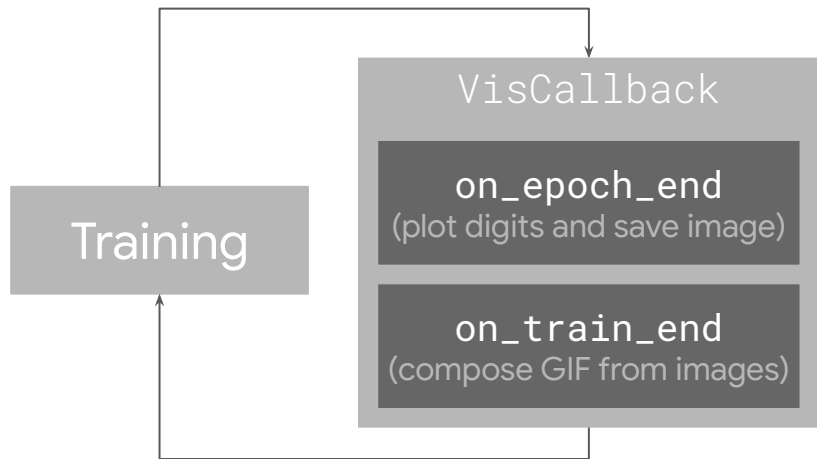
```
    if ratio>threshold:  
        print("Stopping training...")  
        self.model.stop_training = True
```

```
model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])
```

```
class DetectOverfittingCallback(tf.keras.callbacks.Callback):  
    def __init__(self, threshold):  
        super(DetectOverfittingCallback, self).__init__()  
        self.threshold = threshold  
  
    def on_epoch_end(self, epoch, logs=None):  
        ratio = logs["val_loss"] / logs["loss"]  
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))  
  
        if ratio > threshold:  
            print("Stopping training...")  
            self.model.stop_training = True
```

```
model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])
```

```
class DetectOverfittingCallback(tf.keras.callbacks.Callback):  
    def __init__(self, threshold):  
        super(DetectOverfittingCallback, self).__init__()  
        self.threshold = threshold  
  
    def on_epoch_end(self, epoch, logs=None):  
        ratio = logs["val_loss"] / logs["loss"]  
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))  
  
        if ratio>threshold:  
            print("Stopping training...")  
            self.model.stop_training = True  
  
model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])
```



```
class VisCallback(tf.keras.callbacks.Callback):  
    def __init__(self, inputs, ground_truth, display_freq=10,  
                 n_samples=10):  
        self.inputs = inputs  
        self.ground_truth = ground_truth  
        self.images = []  
        self.display_freq = display_freq  
        self.n_samples = n_samples
```



```
class VisCallback(tf.keras.callbacks.Callback):  
    def __init__(self, inputs, ground_truth, display_freq=10,  
                 n_samples=10):  
        self.inputs = inputs  
        self.ground_truth = ground_truth  
        self.images = []  
        self.display_freq = display_freq  
        self.n_samples = n_samples
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    def __init__(self, inputs, ground_truth, display_freq=10,  
                 n_samples=10):  
        self.inputs = inputs  
        self.ground_truth = ground_truth  
        self.images = []  
        self.display_freq = display_freq  
        self.n_samples = n_samples
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    def __init__(self, inputs, ground_truth, display_freq=10,  
                 n_samples=10):  
  
        self.inputs = inputs  
        self.ground_truth = ground_truth  
        self.images = []  
        self.display_freq = display_freq  
        self.n_samples = n_samples
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    def __init__(self, inputs, ground_truth, display_freq=10,  
                 n_samples=10):  
        self.inputs = inputs  
        self.ground_truth = ground_truth  
        self.images = []  
        self.display_freq = display_freq  
        self.n_samples = n_samples
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    def __init__(self, inputs, ground_truth, display_freq=10,  
                 n_samples=10):  
  
        self.inputs = inputs  
        self.ground_truth = ground_truth  
        self.images = []  
        self.display_freq = display_freq  
        self.n_samples = n_samples
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    ...  
    def on_epoch_end(self, epoch, logs=None):  
        # Randomly sample data  
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)  
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]  
        predictions = np.argmax(self.model.predict(X_test), axis=1)
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    ...  
    def on_epoch_end(self, epoch, logs=None):  
        # Randomly sample data  
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)  
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]  
        predictions = np.argmax(self.model.predict(X_test), axis=1)
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    ...  
    def on_epoch_end(self, epoch, logs=None):  
        # Randomly sample data  
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)  
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]  
        predictions = np.argmax(self.model.predict(X_test), axis=1)
```



```
class VisCallback(tf.keras.callbacks.Callback):  
    ...  
    def on_epoch_end(self, epoch, logs=None):  
        # Randomly sample data  
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)  
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]  
        predictions = np.argmax(self.model.predict(X_test), axis=1)
```

```
class VisCallback(tf.keras.callbacks.Callback):
    ...
    def on_epoch_end(self, epoch, logs=None):
        # Randomly sample data
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]
        predictions = np.argmax(self.model.predict(X_test), axis=1)

        # Plot the digits
        display_digits(X_test, predictions, y_test, epoch, n=self.display_freq)
        ...
```

```
class VisCallback(tf.keras.callbacks.Callback):
```

```
...
```

```
def on_epoch_end(self, epoch, logs=None):
```

```
...
```

```
# Save the figure
```

```
buf = io.BytesIO()
```

```
plt.savefig(buf, format='png')
```

```
buf.seek(0)
```

```
image = Image.open(buf)
```

```
self.images.append(np.array(image))
```

```
# Display the digits every now and then
```

```
if epoch % self.display_freq == 0:
```

```
    plt.show()
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    ...  
    def on_epoch_end(self, epoch, logs=None):  
        ...  
        # Save the figure  
        buf = io.BytesIO()  
        plt.savefig(buf, format='png')  
        buf.seek(0)  
        image = Image.open(buf)  
        self.images.append(np.array(image))  
  
        # Display the digits every now and then  
        if epoch % self.display_freq == 0:  
            plt.show()
```

```
import imageio
```

```
class VisCallback(tf.keras.callbacks.Callback):
```

```
...
```

```
def on_train_end(self, logs=None):
```

```
    imageio.mimsave('animation.gif', self.images, fps=1)
```

```
# Train the model
```

```
model.fit(..., callbacks=[VisCallback(x_test, y_test)])
```

```
import imageio
```

```
class VisCallback(tf.keras.callbacks.Callback):
```

```
    ...
```

```
    def on_train_end(self, logs=None):
```

```
        imageio.mimsave('animation.gif', self.images, fps=1)
```

```
# Train the model
```

```
model.fit(..., callbacks=[VisCallback(x_test, y_test)])
```

9 1 4 6 3 2 8 9 0 5