

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>

Data Parallelism

Different Data Slices





Data 1

Data 2


Data 3

Data 4

Same Model
Architecture



Aggregate and Update model
variables



New
Master
Model

Data Parallelism

Different Data Slices

Data 1

Data 2






Data 3

Data 4

Same Model Architecture

Aggregate and Update model variables

New Master Model



Data Parallelism

Different Data Slices





Data 1

Data 2


Data 3

Data 4

Same Model Architecture



Aggregate and Update model variables



New Master Model

`tf.distribute.Strategy`

- High-level APIs
- Custom training loops
- TensorFlow 2: eager mode & graph mode
- Supported on multiple configurations.
- Convenient to use with little to no code changes

Commonly used terms

- Device
- Replica
- Worker
- Mirrored variable

Commonly used terms

- Device



CPU

Accelerator: GPU, TPU

- Replica
- Worker
- Mirrored variable

Commonly used terms


- Device



CPU

Accelerator: GPU, TPU

- Replica



- Worker

- Mirrored variable

Commonly used terms


- Device



CPU

Accelerator: GPU, TPU

- Replica



- Worker



- Mirrored variable

Commonly used terms



- Device
- Replica
- Worker

- Mirrored variable



CPU

Accelerator: GPU, TPU



Classifying strategies

Hardware platforms

- Single-machine multi-device
- Multi-machine (with 0 or more accelerators)

Training

- Synchronous (All-reduce)
- Asynchronous (Parameter Server)

Classifying strategies

Hardware platforms

- Single-machine multi-device
- Multi-machine (with 0 or more accelerators)



Training

- Synchronous (All-reduce)
- Asynchronous (Parameter Server)

Classifying strategies

Hardware platforms

- Single-machine multi-device
- Multi-machine (with 0 or more accelerators)



Training

- Synchronous (All-reduce)
- Asynchronous (Parameter Server)

Classifying strategies

Hardware platforms

- Single-machine multi-device
- Multi-machine (with 0 or more accelerators)



Training

- Synchronous (All-reduce)
- Asynchronous (Parameter Server)

Classifying strategies

Hardware platforms

- Single-machine multi-device
- Multi-machine (with 0 or more accelerators)



Training

- Synchronous (All-reduce)
- Asynchronous (Parameter Server)

MirroredStrategy

MultiWorkerMirroredStrategy

ParameterServerStrategy

DefaultStrategy

TPUStrategy

CentralStorageStrategy

OneDeviceStrategy

MirroredStrategy

- Single-machine multi-GPU
- Creates a replica per **GPU**
- Each variable is ***mirrored***
- All-reduce ***across devices***

MultiWorkerMirroredStrategy

ParameterServerStrategy

DefaultStrategy

TPUStrategy

CentralStorageStrategy

OneDeviceStrategy

MirroredStrategy

- Single-machine multi-GPU
- Creates a replica per **GPU**
- Each variable is ***mirrored***
- All-reduce ***across devices***

MultiWorkerMirroredStrategy

ParameterServerStrategy

DefaultStrategy

TPUStrategy

- Same as **MirroredStrategy**
- All-reduce across ***TPU cores***

CentralStorageStrategy

OneDeviceStrategy

MirroredStrategy

- Single-machine multi-GPU
- Creates a replica per **GPU**
- Each variable is **mirrored**
- All-reduce **across devices**

MultiWorkerMirroredStrategy

- Multi-machine multi-GPU
- Replicates variables per device **across workers**
- All-reduce based on
 - hardware
 - network topology
 - tensor sizes

ParameterServerStrategy

DefaultStrategy

TPUStrategy

- Same as **MirroredStrategy**
- All-reduce across **TPU cores**

CentralStorageStrategy

OneDeviceStrategy

MirroredStrategy

- Single-machine multi-GPU
- Creates a replica per **GPU**
- Each variable is **mirrored**
- All-reduce **across devices**

MultiWorkerMirroredStrategy

- Multi-machine multi-GPU
- Replicates variables per device **across workers**
- All-reduce based on
 - hardware
 - network topology
 - tensor sizes

ParameterServerStrategy

DefaultStrategy

TPUStrategy

- Same as **MirroredStrategy**
- All-reduce across **TPU cores**

CentralStorageStrategy

- Variables are **not mirrored**
(instead placed on the CPU)
- Done in-memory on a device

OneDeviceStrategy

MirroredStrategy

- Single-machine multi-GPU
- Creates a replica per **GPU**
- Each variable is **mirrored**
- All-reduce **across devices**

MultiWorkerMirroredStrategy

- Multi-machine multi-GPU
- Replicates variables per device **across workers**
- All-reduce based on
 - hardware
 - network topology
 - tensor sizes

ParameterServerStrategy

- Some machines designated as **workers**
- Some others as **parameter servers**

DefaultStrategy

TPUStrategy

- Same as **MirroredStrategy**
- All-reduce across **TPU cores**

CentralStorageStrategy

- Variables are **not mirrored**
(instead placed on the CPU)
- Done in-memory on a device

OneDeviceStrategy

MirroredStrategy

- Single-machine multi-GPU
- Creates a replica per **GPU**
- Each variable is **mirrored**
- All-reduce **across devices**

MultiWorkerMirroredStrategy

- Multi-machine multi-GPU
- Replicates variables per device **across workers**
- All-reduce based on
 - hardware
 - network topology
 - tensor sizes

ParameterServerStrategy

- Some machines designated as **workers**
- Some others as **parameter servers**

TPUStrategy

- Same as **MirroredStrategy**
- All-reduce across **TPU cores**

CentralStorageStrategy

- Variables are **not mirrored**
(instead placed on the CPU)
- Done in-memory on a device


DefaultStrategy

- Simple Passthrough

OneDeviceStrategy

- Single device

Mirrored Strategy



Mirrored Strategy

- Model declaration
- Data preprocessing


```
model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(10)
])
```

```
model.compile(
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    optimizer=tf.keras.optimizers.Adam(),
    metrics=['accuracy'])
```

```
model = tf.keras.Sequential([  
    tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),  
    tf.keras.layers.MaxPooling2D(),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(64, activation='relu'),  
    tf.keras.layers.Dense(10)  
])
```

```
model.compile(  
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
    optimizer=tf.keras.optimizers.Adam(),  
    metrics=['accuracy'])
```

```
model = tf.keras.Sequential([  
    tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),  
    tf.keras.layers.MaxPooling2D(),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(64, activation='relu'),  
    tf.keras.layers.Dense(10)  
])
```

```
model.compile(  
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
    optimizer=tf.keras.optimizers.Adam(),  
    metrics=['accuracy'])
```

```
def scale(image, label):  
    image = tf.cast(image, tf.float32)  
    image /= 255  
    return image, label
```

```
num_train_examples = info.splits['train'].num_examples  
num_test_examples = info.splits['test'].num_examples
```

```
BUFFER_SIZE = 10000
```

```
BATCH_SIZE = 64
```

```
train_dataset = mnist_train.map(scale).cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)  
eval_dataset = mnist_test.map(scale).batch(BATCH_SIZE)
```

```
def scale(image, label):  
    image = tf.cast(image, tf.float32)  
    image /= 255  
    return image, label
```

```
num_train_examples = info.splits['train'].num_examples  
num_test_examples = info.splits['test'].num_examples
```

```
BUFFER_SIZE = 10000
```

```
BATCH_SIZE = 64
```

```
train_dataset = mnist_train.map(scale).cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)  
eval_dataset = mnist_test.map(scale).batch(BATCH_SIZE)
```

```
def scale(image, label):  
    image = tf.cast(image, tf.float32)  
    image /= 255  
    return image, label
```

```
num_train_examples = info.splits['train'].num_examples  
num_test_examples = info.splits['test'].num_examples
```

```
BUFFER_SIZE = 10000
```

```
BATCH_SIZE = 64
```

```
train_dataset = mnist_train.map(scale).cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)  
eval_dataset = mnist_test.map(scale).batch(BATCH_SIZE)
```

```
strategy = tf.distribute.MirroredStrategy()
```

```
print('Number of devices: {}'.format(strategy.num_replicas_in_sync))
```

```
strategy = tf.distribute.MirroredStrategy()
```

```
print('Number of devices: {}'.format(strategy.num_replicas_in_sync))
```



```
def scale(image, label):  
    image = tf.cast(image, tf.float32)  
    image /= 255  
  
    return image, label  
  
num_train_examples = info.splits['train'].num_examples  
num_test_examples = info.splits['test'].num_examples  
  
BUFFER_SIZE = 10000  
  
BATCH_SIZE_PER_REPLICA = 64  
BATCH_SIZE = BATCH_SIZE_PER_REPLICA * strategy.num_replicas_in_sync  
  
train_dataset =  
mnist_train.map(scale).cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)  
eval_dataset = mnist_test.map(scale).batch(BATCH_SIZE)
```

```
def scale(image, label):  
    image = tf.cast(image, tf.float32)  
    image /= 255  
  
    return image, label  
  
num_train_examples = info.splits['train'].num_examples  
num_test_examples = info.splits['test'].num_examples  
  
BUFFER_SIZE = 10000  
  
BATCH_SIZE_PER_REPLICA = 64  
BATCH_SIZE = BATCH_SIZE_PER_REPLICA * strategy.num_replicas_in_sync  
  
train_dataset =  
mnist_train.map(scale).cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)  
eval_dataset = mnist_test.map(scale).batch(BATCH_SIZE)
```

```
def scale(image, label):  
    image = tf.cast(image, tf.float32)  
    image /= 255  
  
    return image, label  
  
num_train_examples = info.splits['train'].num_examples  
num_test_examples = info.splits['test'].num_examples  
  
BUFFER_SIZE = 10000  
  
BATCH_SIZE_PER_REPLICA = 64  
BATCH_SIZE = BATCH_SIZE_PER_REPLICA * strategy.num_replicas_in_sync  
  
train_dataset =  
mnist_train.map(scale).cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)  
eval_dataset = mnist_test.map(scale).batch(BATCH_SIZE)
```

```
with strategy.scope():  
    model = tf.keras.Sequential([  
        tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),  
        tf.keras.layers.MaxPooling2D(),  
        tf.keras.layers.Flatten(),  
        tf.keras.layers.Dense(64, activation='relu'),  
        tf.keras.layers.Dense(10)  
    ])  
  
model.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
              optimizer=tf.keras.optimizers.Adam(),  
              metrics=['accuracy'])
```

```
with strategy.scope():
```

```
    model = tf.keras.Sequential([  
        tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),  
        tf.keras.layers.MaxPooling2D(),  
        tf.keras.layers.Flatten(),  
        tf.keras.layers.Dense(64, activation='relu'),  
        tf.keras.layers.Dense(10)  
    ])
```

```
model.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
              optimizer=tf.keras.optimizers.Adam(),  
              metrics=['accuracy'])
```

```
with strategy.scope():  
    model = tf.keras.Sequential([  
        tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),  
        tf.keras.layers.MaxPooling2D(),  
        tf.keras.layers.Flatten(),  
        tf.keras.layers.Dense(64, activation='relu'),  
        tf.keras.layers.Dense(10)  
    ])
```

```
model.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
              optimizer=tf.keras.optimizers.Adam(),  
              metrics=['accuracy'])
```

Epoch 1/12

INFO:tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadcast to
INFO:tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadcast to
INFO:tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadcast to
INFO:tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadcast to

Training across local GPUs

`tf.distribute.MirroredStrategy`

- Each variable in the model is mirrored across all replicas
- Variables are treated as `MirroredVariable`
- **Synchronization** done with NVIDIA NCCL


```
# Create Datasets from the batches
```

```
train_dataset = tf.data.Dataset.from_tensor_slices((train_images, train_labels))  
                .shuffle(BUFFER_SIZE).batch(GLOBAL_BATCH_SIZE)
```

```
test_dataset = tf.data.Dataset.from_tensor_slices((test_images, test_labels))  
                .batch(GLOBAL_BATCH_SIZE)
```

```
# Create Distributed Datasets from the datasets
```

```
train_dist_dataset = strategy.experimental_distribute_dataset(train_dataset)  
test_dist_dataset = strategy.experimental_distribute_dataset(test_dataset)
```

```
# Create Datasets from the batches
```

```
train_dataset = tf.data.Dataset.from_tensor_slices((train_images, train_labels))  
                .shuffle(BUFFER_SIZE).batch(GLOBAL_BATCH_SIZE)
```

```
test_dataset = tf.data.Dataset.from_tensor_slices((test_images, test_labels))  
                .batch(GLOBAL_BATCH_SIZE)
```

```
# Create Distributed Datasets from the datasets
```

```
train_dist_dataset = strategy.experimental_distribute_dataset(train_dataset)
```

```
test_dist_dataset = strategy.experimental_distribute_dataset(test_dataset)
```

```
# Create Datasets from the batches
```

```
train_dataset = tf.data.Dataset.from_tensor_slices((train_images, train_labels))  
                .shuffle(BUFFER_SIZE).batch(GLOBAL_BATCH_SIZE)
```

```
test_dataset = tf.data.Dataset.from_tensor_slices((test_images, test_labels))  
                .batch(GLOBAL_BATCH_SIZE)
```

```
# Create Distributed Datasets from the datasets
```

```
train_dist_dataset = strategy.experimental_distribute_dataset(train_dataset)  
test_dist_dataset = strategy.experimental_distribute_dataset(test_dataset)
```

```
# Create Datasets from the batches
```

```
train_dataset = tf.data.Dataset.from_tensor_slices((train_images, train_labels))  
                .shuffle(BUFFER_SIZE).batch(GLOBAL_BATCH_SIZE)
```

```
test_dataset = tf.data.Dataset.from_tensor_slices((test_images, test_labels))  
                .batch(GLOBAL_BATCH_SIZE)
```

```
# Create Distributed Datasets from the datasets
```

```
train_dist_dataset = strategy.experimental_distribute_dataset(train_dataset)  
test_dist_dataset = strategy.experimental_distribute_dataset(test_dataset)
```

```
EPOCHS = 10
for epoch in range(EPOCHS):
    # Do Training
    total_loss = 0.0
    num_batches = 0
    for batch in train_dist_dataset:
        total_loss += distributed_train_step(batch)
        num_batches += 1
    train_loss = total_loss / num_batches
```

```
EPOCHS = 10
for epoch in range(EPOCHS):
    # Do Training
    total_loss = 0.0
    num_batches = 0
    for batch in train_dist_dataset:
        total_loss += distributed_train_step(batch)
        num_batches += 1
    train_loss = total_loss / num_batches
```

```
EPOCHS = 10
for epoch in range(EPOCHS):
    # Do Training
    total_loss = 0.0
    num_batches = 0
    for batch in train_dist_dataset:
        total_loss += distributed_train_step(batch)
        num_batches += 1
    train_loss = total_loss / num_batches
```

```
EPOCHS = 10
for epoch in range(EPOCHS):
    # Do Training
    total_loss = 0.0
    num_batches = 0
    for batch in train_dist_dataset:
        total_loss += distributed_train_step(batch)
        num_batches += 1
    train_loss = total_loss / num_batches
```



```
EPOCHS = 10
for epoch in range(EPOCHS):
    # Do Training
    total_loss = 0.0
    num_batches = 0
    for batch in train_dist_dataset:
        total_loss += distributed_train_step(batch)
        num_batches += 1
    train_loss = total_loss / num_batches
```



```
def train_step(inputs):  
    images, labels = inputs  
    with tf.GradientTape() as tape:  
        predictions = model(images, training=True)  
        loss = compute_loss(labels, predictions)  
  
    gradients = tape.gradient(loss, model.trainable_variables)  
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))  
  
    train_accuracy.update_state(labels, predictions)  
    return loss
```

```
def train_step(inputs):  
    images, labels = inputs  
    with tf.GradientTape() as tape:  
        predictions = model(images, training=True)  
        loss = compute_loss(labels, predictions)  
  
    gradients = tape.gradient(loss, model.trainable_variables)  
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))  
  
    train_accuracy.update_state(labels, predictions)  
    return loss
```



```
def train_step(inputs):  
    images, labels = inputs  
    with tf.GradientTape() as tape:  
        predictions = model(images, training=True)  
        loss = compute_loss(labels, predictions)  
  
    gradients = tape.gradient(loss, model.trainable_variables)  
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))  
  
    train_accuracy.update_state(labels, predictions)  
    return loss
```

```
def train_step(inputs):  
    images, labels = inputs  
    with tf.GradientTape() as tape:  
        predictions = model(images, training=True)  
        loss = compute_loss(labels, predictions)  
  
        gradients = tape.gradient(loss, model.trainable_variables)  
        optimizer.apply_gradients(zip(gradients, model.trainable_variables))  
  
    train_accuracy.update_state(labels, predictions)  
    return loss
```

```
def train_step(inputs):  
    images, labels = inputs  
    with tf.GradientTape() as tape:  
        predictions = model(images, training=True)  
        loss = compute_loss(labels, predictions)  
  
    gradients = tape.gradient(loss, model.trainable_variables)  
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))  
  
    train_accuracy.update_state(labels, predictions)  
    return loss
```

```
def train_step(inputs):  
    images, labels = inputs  
    with tf.GradientTape() as tape:  
        predictions = model(images, training=True)  
        loss = compute_loss(labels, predictions)  
  
    gradients = tape.gradient(loss, model.trainable_variables)  
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))  
  
    train_accuracy.update_state(labels, predictions)  
    return loss
```


Notebook settings

Hardware accelerator

TPU



To get the most out of Colab Pro, avoid using a TPU unless you need one. [Learn more](#)

Runtime shape

Standard



☐ Omit code cell output when saving this notebook

CANCEL

SAVE

```
# Detect hardware
```

```
try:
```

```
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
```

```
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address)
```

```
    tf.config.experimental_connect_to_cluster(tpu)
```

```
    tf.tpu.experimental.initialize_tpu_system(tpu)
```

```
    strategy = tf.distribute.experimental.TPUStrategy(tpu)
```

```
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
```

```
    print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
except ValueError:
```

```
    print('TPU failed to initialize.')
```

```
# Detect hardware
```

```
try:
```

```
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']  
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address)  
    tf.config.experimental_connect_to_cluster(tpu)  
    tf.tpu.experimental.initialize_tpu_system(tpu)  
    strategy = tf.distribute.experimental.TPUStrategy(tpu)  
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])  
    print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
except ValueError:
```

```
    print('TPU failed to initialize.')
```



```
# Detect hardware
```

```
try:
```

```
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
```

```
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address)
```

```
    tf.config.experimental_connect_to_cluster(tpu)
```

```
    tf.tpu.experimental.initialize_tpu_system(tpu)
```

```
    strategy = tf.distribute.experimental.TPUStrategy(tpu)
```

```
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
```

```
    print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
except ValueError:
```

```
    print('TPU failed to initialize.')
```

```
# Detect hardware
```

```
try:
```

```
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
```

```
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address)
```

```
    tf.config.experimental_connect_to_cluster(tpu)
```

```
    tf.tpu.experimental.initialize_tpu_system(tpu)
```

```
    strategy = tf.distribute.experimental.TPUStrategy(tpu)
```

```
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
```

```
    print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
except ValueError:
```

```
    print('TPU failed to initialize.')
```

```
# Detect hardware
```

```
try:
```

```
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
```

```
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address)
```

```
    tf.config.experimental_connect_to_cluster(tpu)
```

```
    tf.tpu.experimental.initialize_tpu_system(tpu)
```

```
    strategy = tf.distribute.experimental.TPUStrategy(tpu)
```

```
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
```

```
    print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
except ValueError:
```

```
    print('TPU failed to initialize.')
```

```
# Detect hardware
```

```
try:
```

```
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
```

```
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address)
```

```
    tf.config.experimental_connect_to_cluster(tpu)
```

```
    tf.tpu.experimental.initialize_tpu_system(tpu)
```

```
    strategy = tf.distribute.experimental.TPUStrategy(tpu)
```

```
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
```

```
    print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
except ValueError:
```

```
    print('TPU failed to initialize.')
```

```
# Detect hardware
```

```
try:
```

```
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
```

```
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address)
```

```
    tf.config.experimental_connect_to_cluster(tpu)
```

```
    tf.tpu.experimental.initialize_tpu_system(tpu)
```

```
    strategy = tf.distribute.experimental.TPUStrategy(tpu)
```

```
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
```

```
    print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
except ValueError:
```

```
    print('TPU failed to initialize.')
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:TPU_SYSTEM:0,  
TPU_SYSTEM, 0, 0)
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:TPU_SYSTEM:0,  
TPU_SYSTEM, 0, 0)
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:XLA_CPU:0,  
XLA_CPU, 0, 0)
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:XLA_CPU:0,  
XLA_CPU, 0, 0)
```

```
Running on TPU ['10.109.132.10:8470']  
Number of accelerators: 8
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:TPU_SYSTEM:0,  
TPU_SYSTEM, 0, 0)
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:TPU_SYSTEM:0,  
TPU_SYSTEM, 0, 0)
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:XLA_CPU:0,  
XLA_CPU, 0, 0)
```

```
INFO:tensorflow:*** Available Device:  
_DeviceAttributes(/job:worker/replica:0/task:0/device:XLA_CPU:0,  
XLA_CPU, 0, 0)
```

```
Running on TPU ['10.109.132.10:8470']  
Number of accelerators: 8
```

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in `per-replica-losses` structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in `per-replica-losses` structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in `per-replica-losses` structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in per-replica-losses structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in `per-replica-losses` structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in `per-replica-losses` structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in `per-replica-losses` structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

Training with TPU strategy

- Use a custom training loop
- Call the distributed training function within the loop
 - Use `strategy.run` to call your usual training function across all replicas
 - Results will be in `per-replica-losses` structure
 - Use `strategy.reduce` to reduce losses
- Call the distributed testing function within the loop
 - Use `strategy.run` to call your usual testing function across all replicas

```
@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses, axis=None)

def train_step(inputs):
    images, labels = inputs
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    train_accuracy.update_state(labels, predictions)
    return loss
```



```
@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses, axis=None)
```

```
def train_step(inputs):
    images, labels = inputs
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    train_accuracy.update_state(labels, predictions)
    return loss
```

```
@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses, axis=None)

def train_step(inputs):
    images, labels = inputs
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    train_accuracy.update_state(labels, predictions)
    return loss
```

```
@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses, axis=None)

def train_step(inputs):
    images, labels = inputs
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    train_accuracy.update_state(labels, predictions)
    return loss
```

```
@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses, axis=None)

def train_step(inputs):
    images, labels = inputs
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    train_accuracy.update_state(labels, predictions)
    return loss
```

```
@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses, axis=None)

def train_step(inputs):
    images, labels = inputs
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)
        gradients = tape.gradient(loss, model.trainable_variables)
        optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    train_accuracy.update_state(labels, predictions)
    return loss
```

```
@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses, axis=None)

def train_step(inputs):
    images, labels = inputs
    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    train_accuracy.update_state(labels, predictions)
    return loss
```

Training on a single device

```
tf.distribute.OneDeviceStrategy
```



Input data is distributed

```
strategy = tf.distribute.OneDeviceStrategy(device="/gpu:0")
```


Training across many machines

`tf.distribute.experimental.MultiWorkerMirroredStrategy`

- Done across multiple workers, each with multiple GPUs
- Variables are replicated on each device across workers
- Fault tolerance with `tf.keras.callbacks.ModelCheckpoint`
- **Synchronization** done with `CollectiveOps`



```
multiworker_strategy = tf.distribute.experimental.MultiWorkerMirroredStrategy()
```

Multi-worker training

- Run **workers** in a cluster
- **Tasks** (training/input pipelines)
- **Roles** (chief, worker, ps, evaluator)
- Configuring the cluster (next..)

Cluster specification

```
os.environ["TF_CONFIG"] = json.dumps({  
    "cluster": {  
        "worker": ["host1:port", "host2:port", ...],  
    },  
    "task": {"type": "worker", "index": 0}  
})
```

https://www.tensorflow.org/tutorials/distribute/multi_worker_with_keras

Other strategies

CentralStorageStrategy

Variables - not mirrored, but placed on CPU

Computation - replicated across local GPUs

ParameterServerStrategy

Some machines are designated as workers

... some as parameter servers

Variables - placed on one parameter server (ps)

Computation - replicated across GPUs of all the workers

Other strategies

CentralStorageStrategy

Variables - not mirrored, but placed on CPU

Computation - replicated across local GPUs

ParameterServerStrategy

Some machines are designated as workers

... some as parameter servers

Variables - placed on one parameter server (ps)

Computation - replicated across GPUs of all the workers