

Copyright Notice


These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.


For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>

Image Segmentation




[source](#)



Zombie


Door




Zombie

Door





Semantic Segmentation



Instance Segmentation



Classes indices = [0 = Background 1 = People]

Image Segmentation Basic Architecture

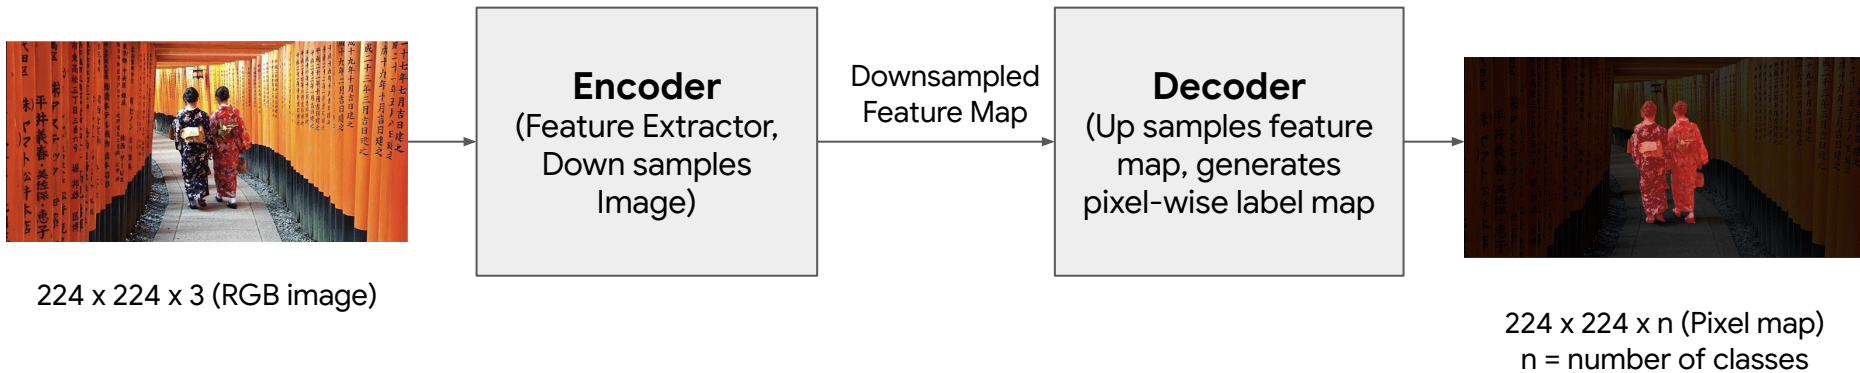




Image Segmentation Basic Architecture



224 x 224 x 3 (RGB image)



Downsampled
Feature Map



224 x 224 x n (Pixel map)
n = number of classes

Image Segmentation Basic Architecture




Image Segmentation Basic Architecture

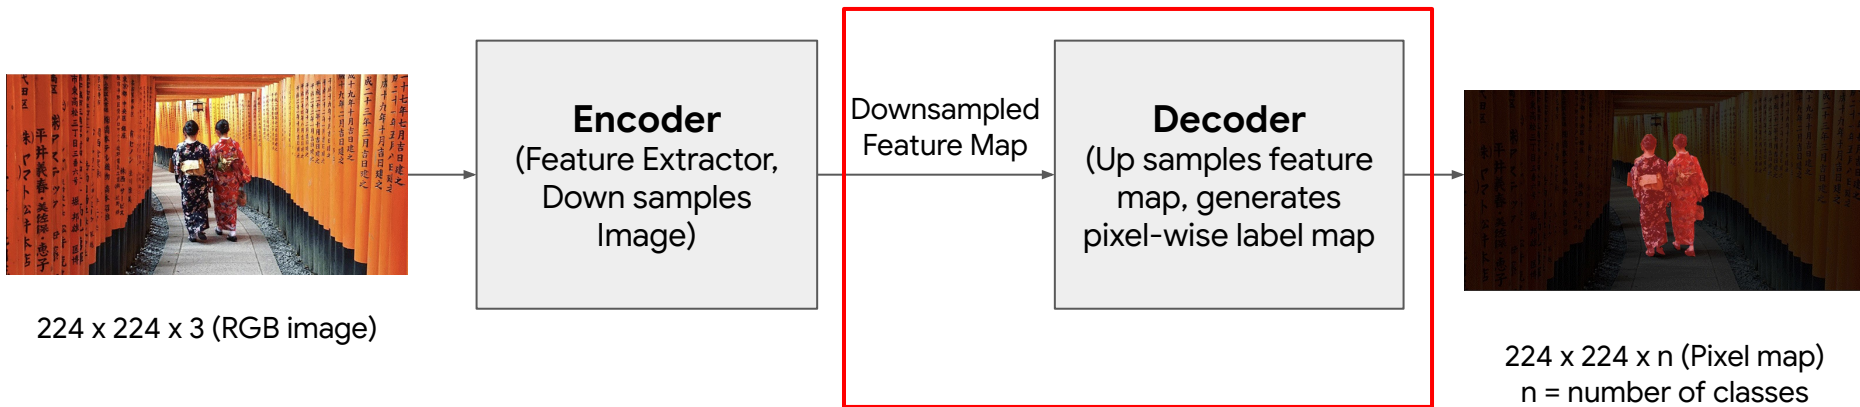


Image Segmentation Basic Architecture

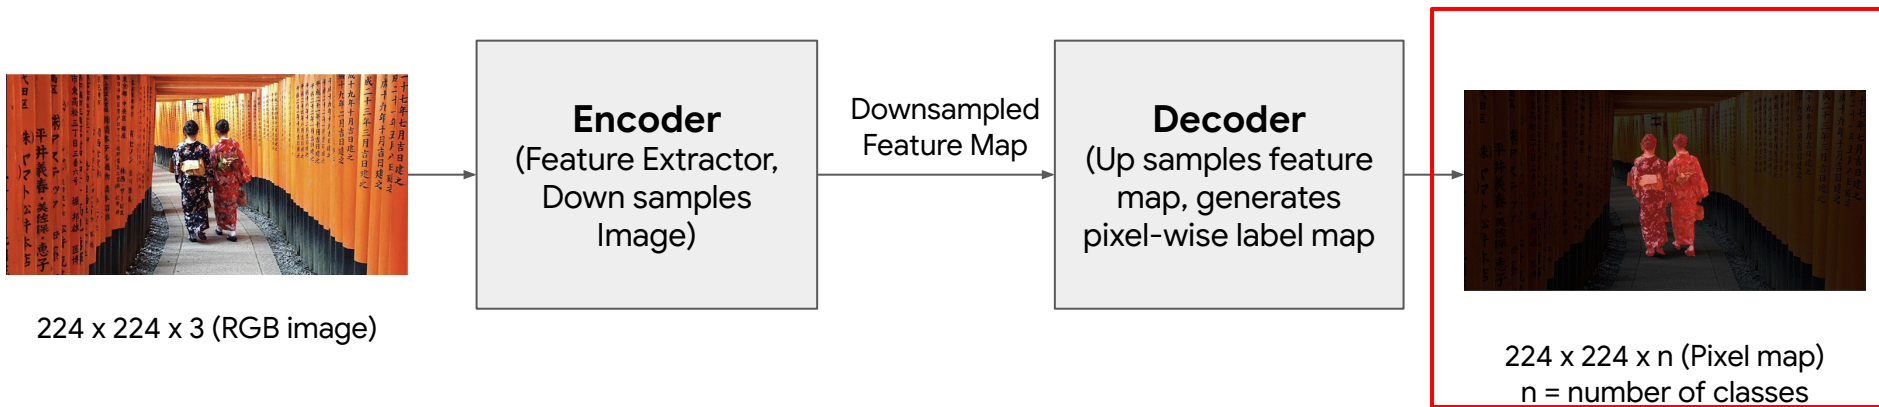


Image Segmentation

- Encoder
 - CNN without fully connected layers
 - Aggregates low level features to high level features
- Decoder
 - Replaces fully connected layers in a CNN
 - Up samples image to original size to generate a pixel mask

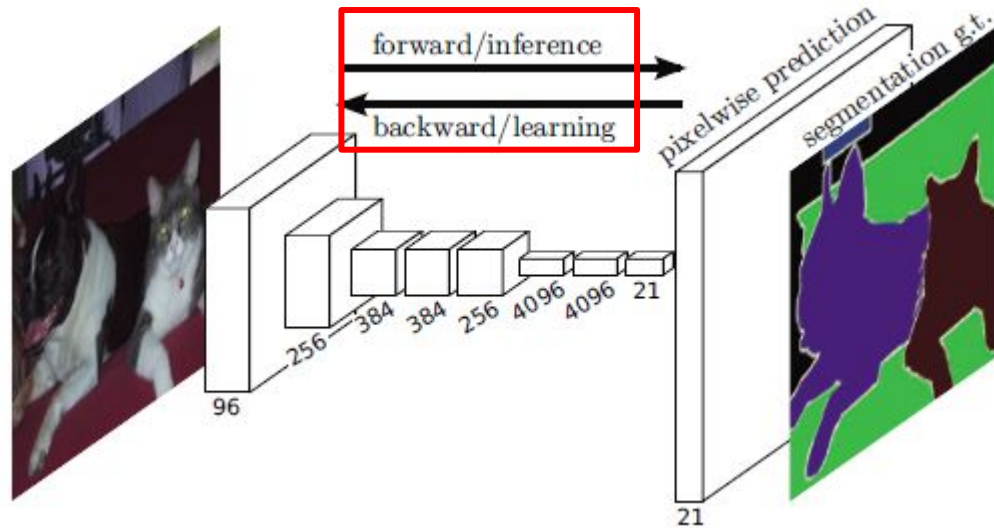
Popular Architectures

- Fully Convolutional Neural Networks.
 - SegNet
 - UNet
 - PSPNet
 - Mask-RCNN

Fully Convolutional Neural Networks

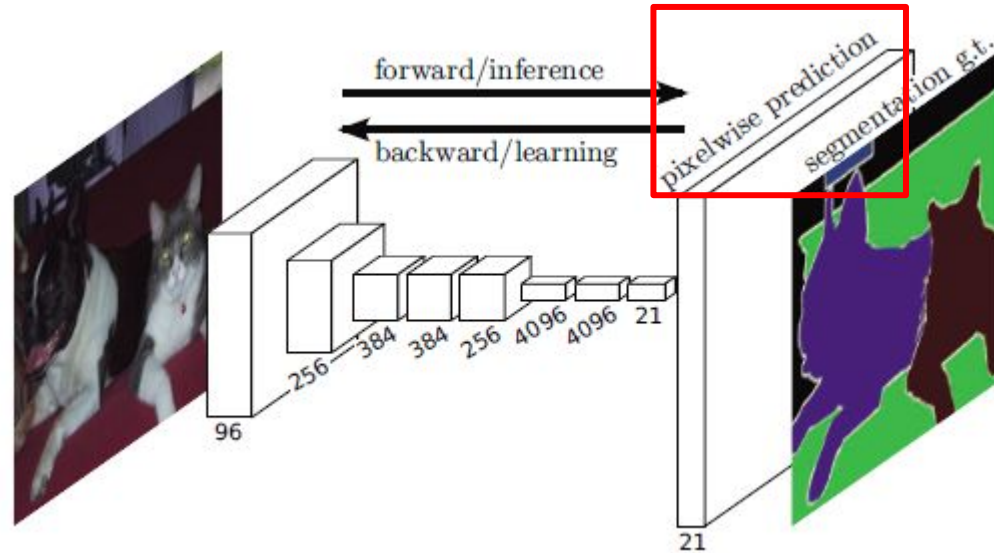
- “Fully Convolutional Networks for Semantic Segmentation”
<https://arxiv.org/abs/1411.4038>
- Replace the fully connected layers with convolutional layers
- Earlier conv layers: Feature extraction and down sampling
- Later conv layers: up sample and pixel-wise labelmap.

Fully Convolutional Neural Networks



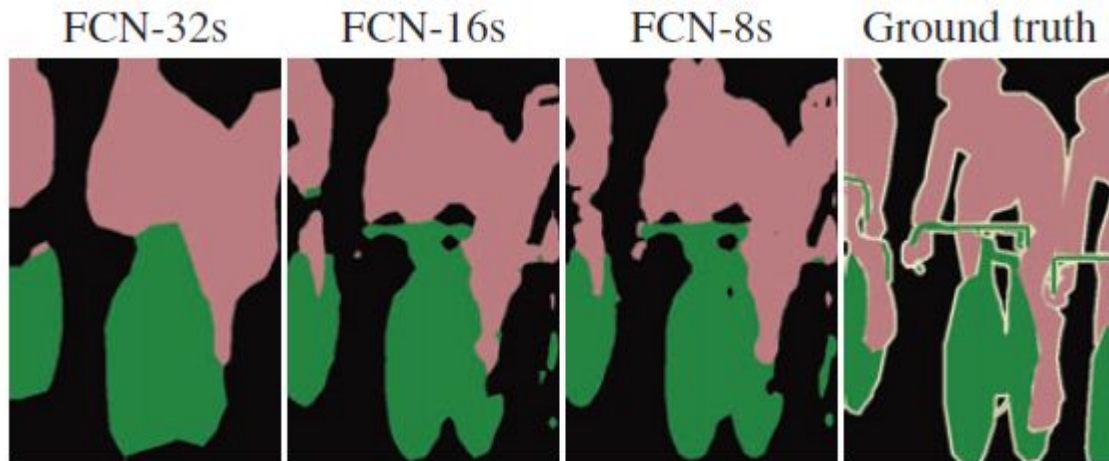
https://people.eecs.berkeley.edu/~jonlong/long_shelhamer_fcn.pdf

Fully Convolutional Neural Networks



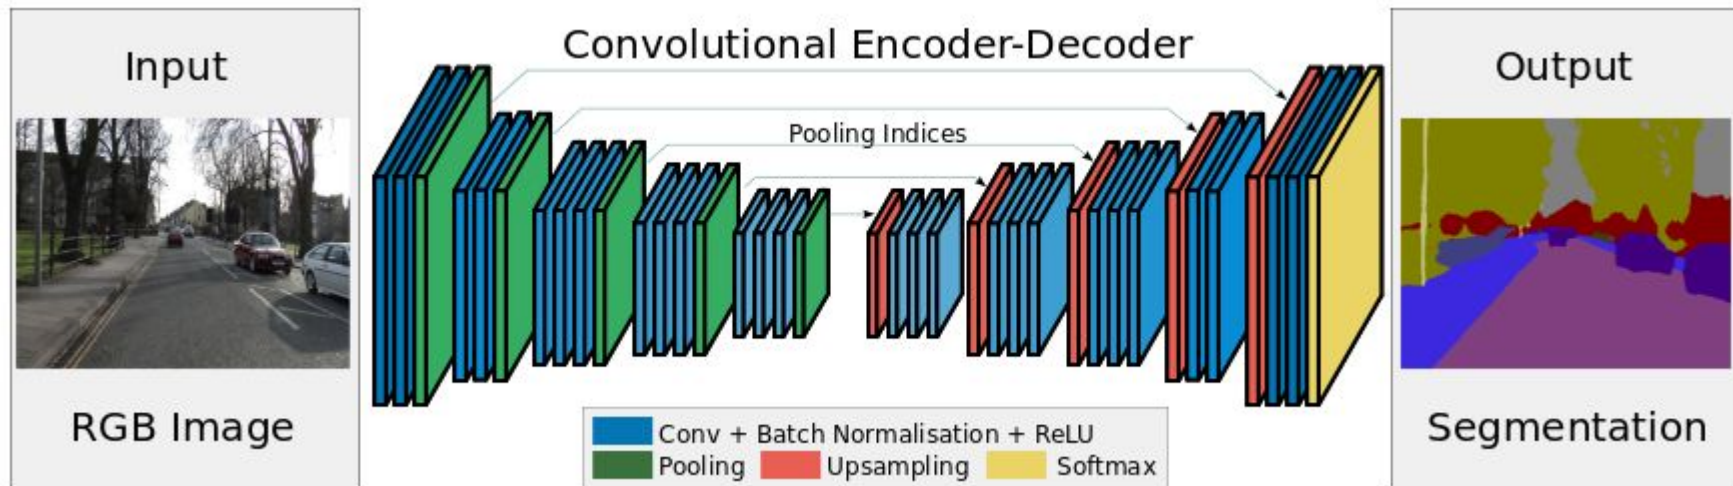
https://people.eecs.berkeley.edu/~jonlong/long_shelhamer_fcn.pdf

Comparison of Different FCNs




<https://arxiv.org/pdf/1411.4038.pdf>


SegNet




SegNet




SegNet




SegNet




SegNet



UNet




Mask R-CNN



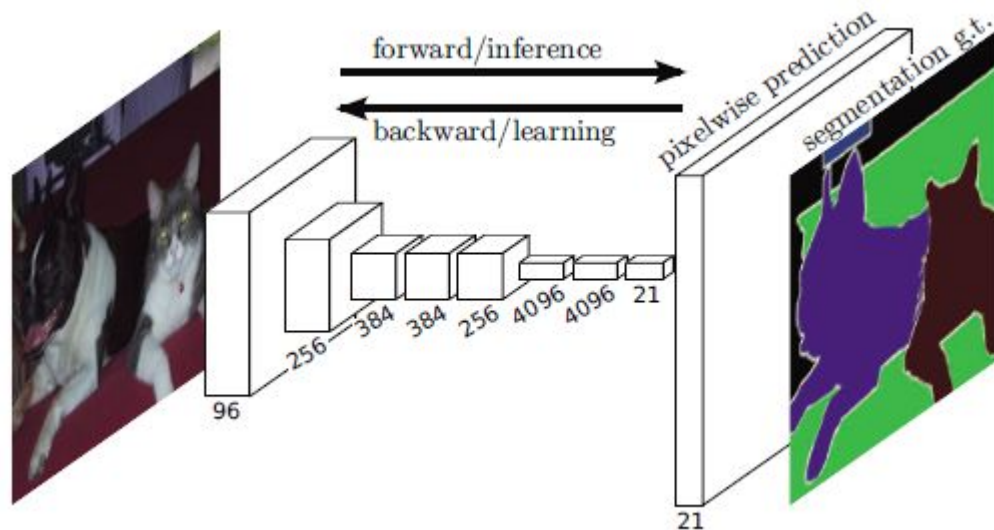
<https://arxiv.org/abs/1703.06870>

Mask R-CNN



<https://arxiv.org/abs/1703.06870>

Fully Convolutional Neural Networks



Fully Convolutional Networks for Semantic Segmentation

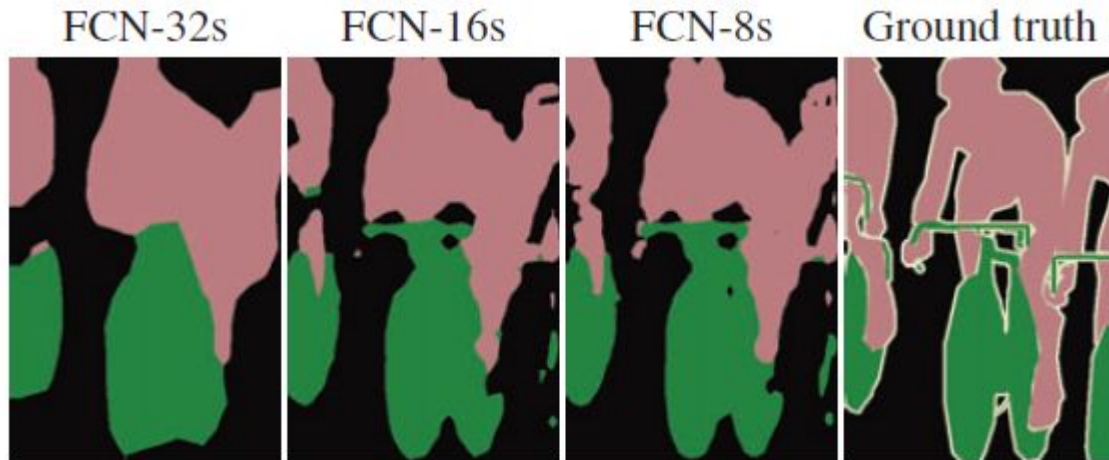
By Jonathan Long, Evan Shelhamer, Trevor Darrell

<https://arxiv.org/pdf/1411.4038.pdf>

Encoder

- Popular encoder architectures:
 - VGG-16
 - ResNet-50
 - MobileNet
- Reuse convolutional layers for feature extraction.
 - Do not reuse fully connected layers


Decoder



Image

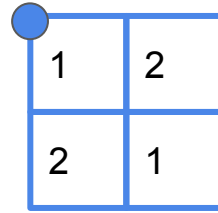
| | |
|---|---|
| 1 | 2 |
| 2 | 1 |
| 0 | 1 |
| 0 | 1 |

Image




| | |
|---|---|
| 1 | 2 |
| 2 | 1 |
| 0 | 1 |
| 0 | 1 |

Pooling window
2 x 2



| | |
|---|---|
| 1 | 2 |
| 2 | 1 |

Pooled result



| |
|-----|
| 1.5 |
|-----|

Image

| | |
|---|---|
| 1 | 2 |
| 2 | 1 |
| 0 | 1 |
| 0 | 1 |

Pooling window
 2×2

Stride
 2×2



| | |
|---|---|
| 1 | 2 |
| 2 | 1 |



Pooled result

| |
|-----|
| 1.5 |
|-----|

Image

| | |
|---|---|
| 1 | 2 |
| 2 | 1 |
| 0 | 1 |
| 0 | 1 |

Pooling window
 2×2

Stride
 2×2



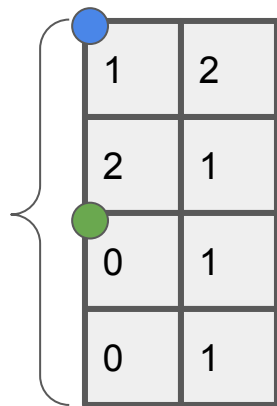
| | |
|---|---|
| 1 | 2 |
| 2 | 1 |
| 0 | 1 |
| 0 | 1 |



Pooled result

| |
|-----|
| 1.5 |
| 0.5 |

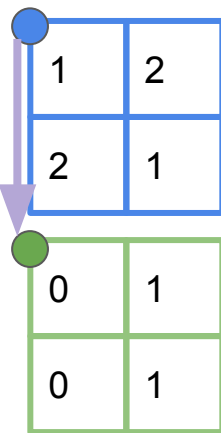
Image



| | |
|---|---|
| 1 | 2 |
| 2 | 1 |
| 0 | 1 |
| 0 | 1 |

Pooling window
 2×2

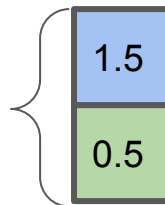
Stride
 2×2



| | |
|---|---|
| 1 | 2 |
| 2 | 1 |
| 0 | 1 |
| 0 | 1 |

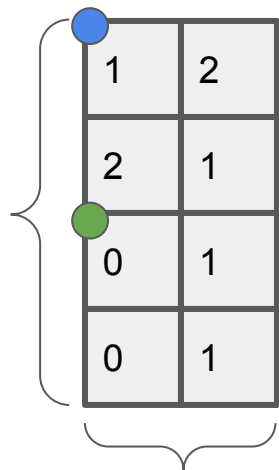


Pooled result



| |
|-----|
| 1.5 |
| 0.5 |

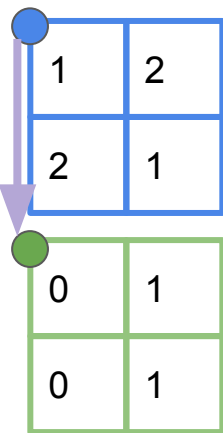
Image



| | |
|---|---|
| 1 | 2 |
| 2 | 1 |
| 0 | 1 |
| 0 | 1 |

Pooling window
 2×2

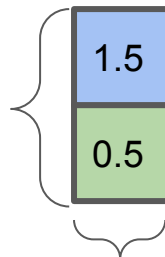
Stride
 2×2



| | |
|---|---|
| 1 | 2 |
| 2 | 1 |
| 0 | 1 |
| 0 | 1 |

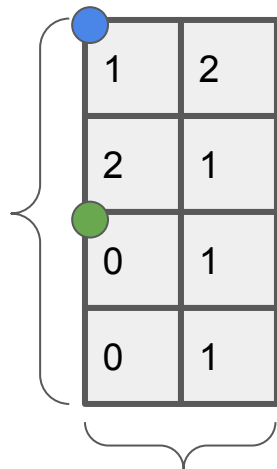


Pooled result



| |
|-----|
| 1.5 |
| 0.5 |

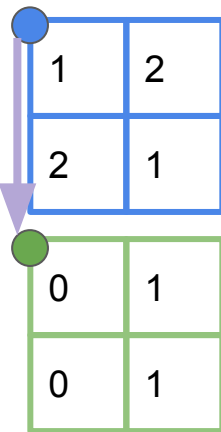
Image




| | |
|---|---|
| 1 | 2 |
| 2 | 1 |
| 0 | 1 |
| 0 | 1 |

Pooling window
 2×2

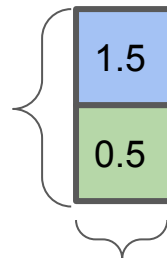
Stride
 2×2



| | |
|---|---|
| 1 | 2 |
| 2 | 1 |
| 0 | 1 |
| 0 | 1 |



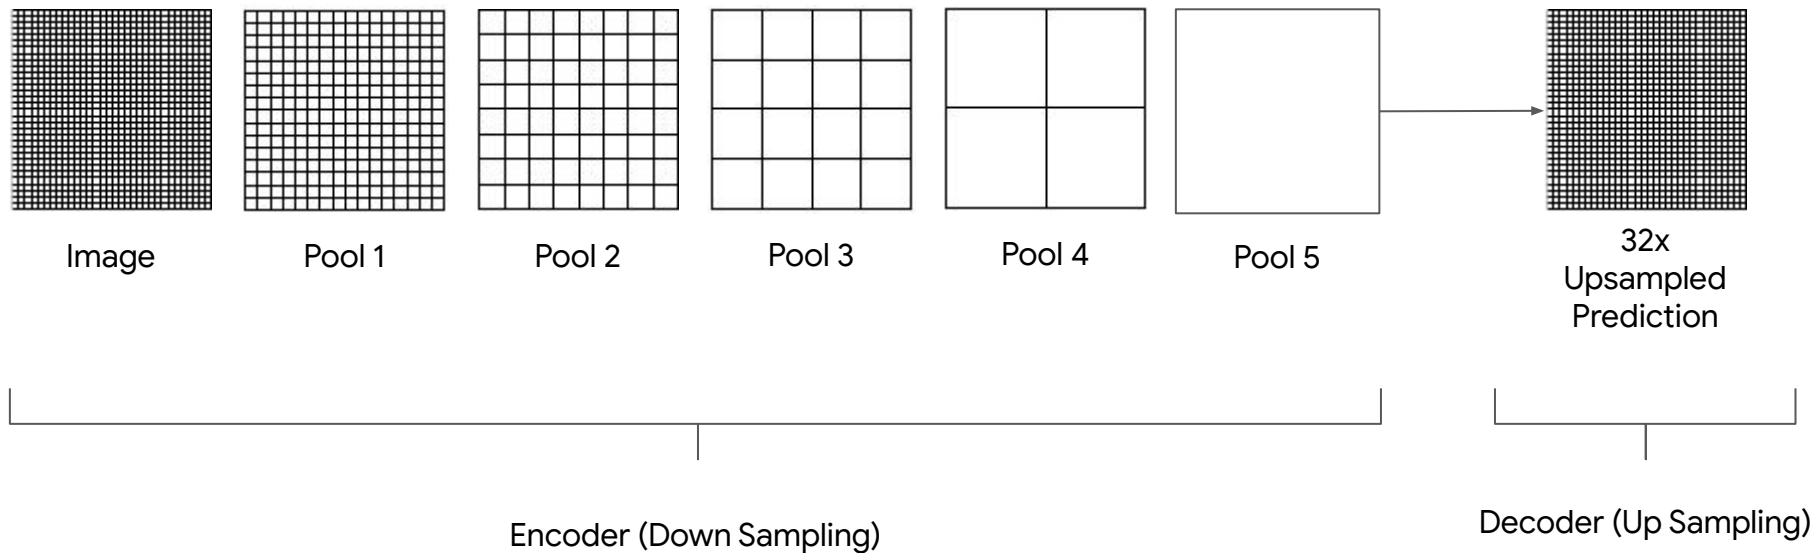
Pooled result



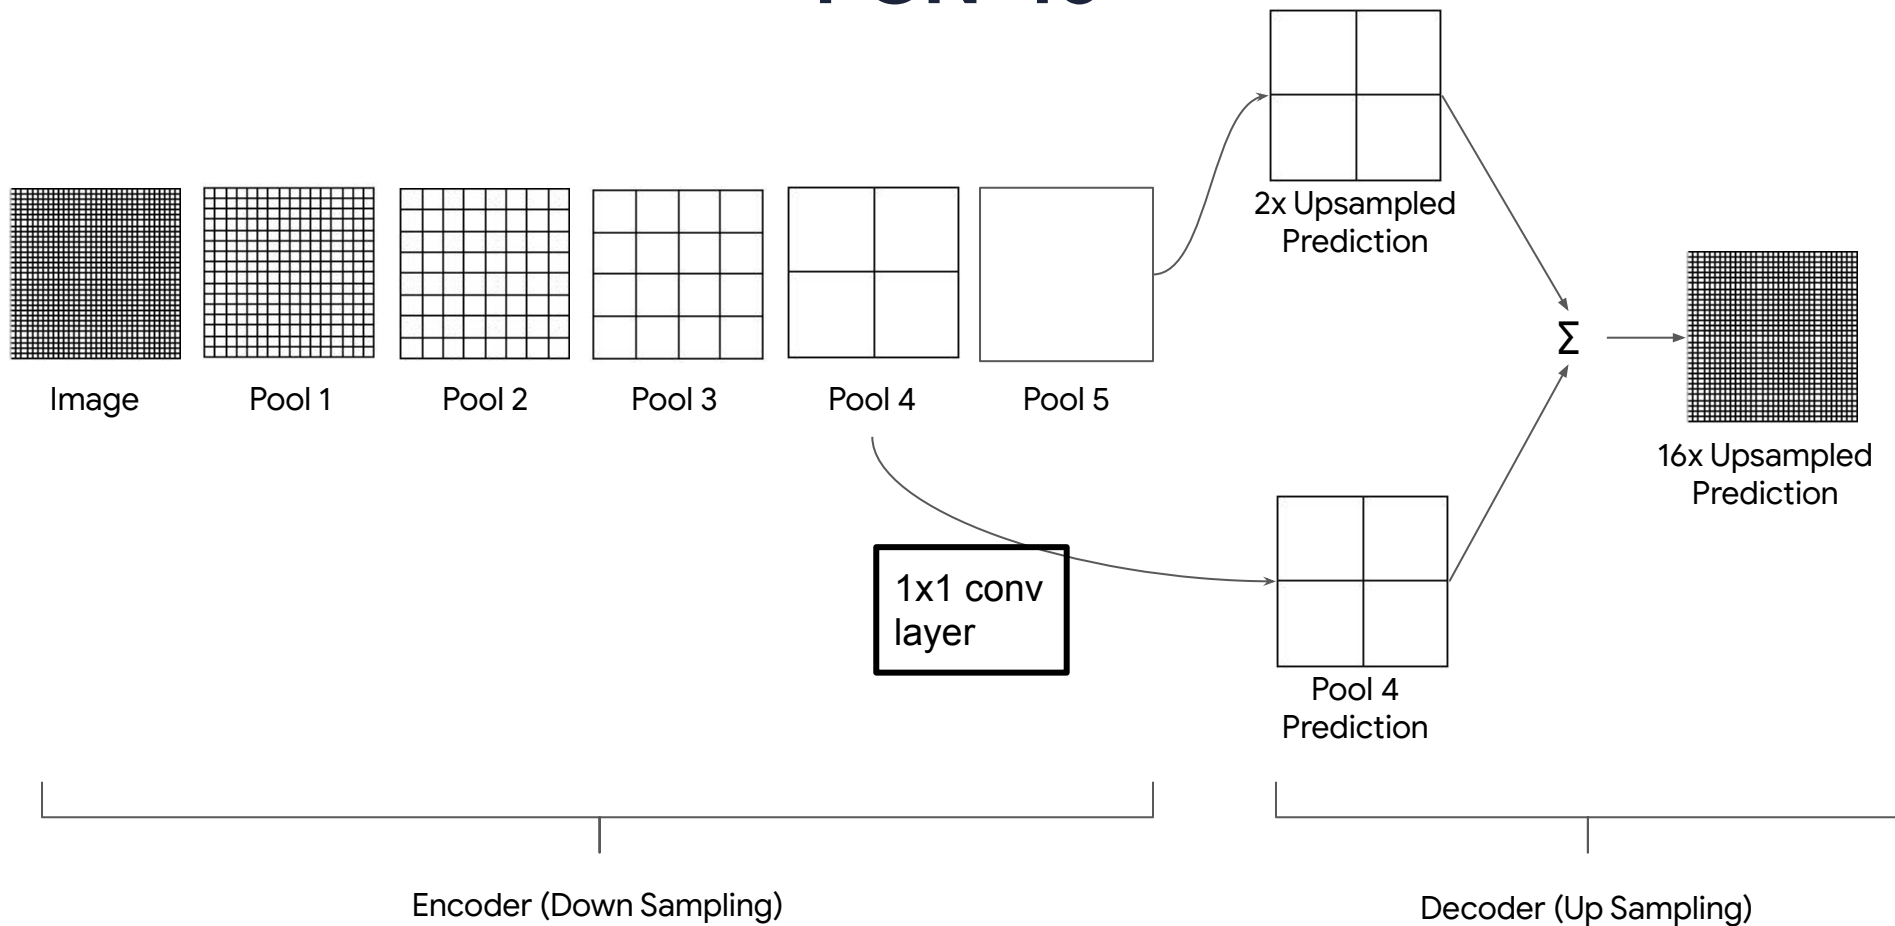
| |
|-----|
| 1.5 |
| 0.5 |

Height and width halved


FCN-32




FCN-16



FCN-8



Comparison of Different FCNs



Upsampling

- Upsampling is increasing height and width of the feature map.
- Two types of layers used in TensorFlow:
 - Simple Scaling - UpSampling2D
 - Transposed Convolution(Deconvolution) - Conv2DTranspose

Simple Scaling - UpSampling2D

- Upsampling2D scales up the image
- Two Types of scaling:
 - Nearest
 - Copies value from nearest pixel.
 - Bilinear
 - linear interpolation from nearby pixels.

UpSampling2D - Usage


```
x = UpSampling2D(  
    size=(2, 2),  
    data_format=None,  
    interpolation='nearest')(x)
```

size: `int` or `tuple` of two `ints`


data_format: `'channels_first'`, `'channels_last'` or `None`

interpolation: `'nearest'` or `'bilinear'`

Transposed Convolution



Convolution->




<- Transposed Convolution

Transposed Convolution

| | | |
|-----|-----|-----|
| 0 | 64 | 128 |
| 48 | 192 | 144 |
| 142 | 226 | 168 |




| | | |
|-----|-----|------|
| -1 | 0 | -2 |
| .5 | 4.5 | -1.5 |
| 1.5 | 2 | -3 |



Convolution->

| | | |
|-----|-----|-----|
| 32 | 64 | 64 |
| 48 | 192 | 128 |
| 128 | 192 | 168 |




| | | |
|-----|-----|------|
| -1 | 0 | -2 |
| .5 | 4.5 | -1.5 |
| 1.5 | 2 | -3 |




<- Transposed Convolution

Transposed Convolution



Convolution->




Conv2DTranspose

- Reverse of Convolution.
- Applied to output of a convolution operation.
- Uses a kernel of a specified size and stride in order to recreate the original input before the convolution operation.

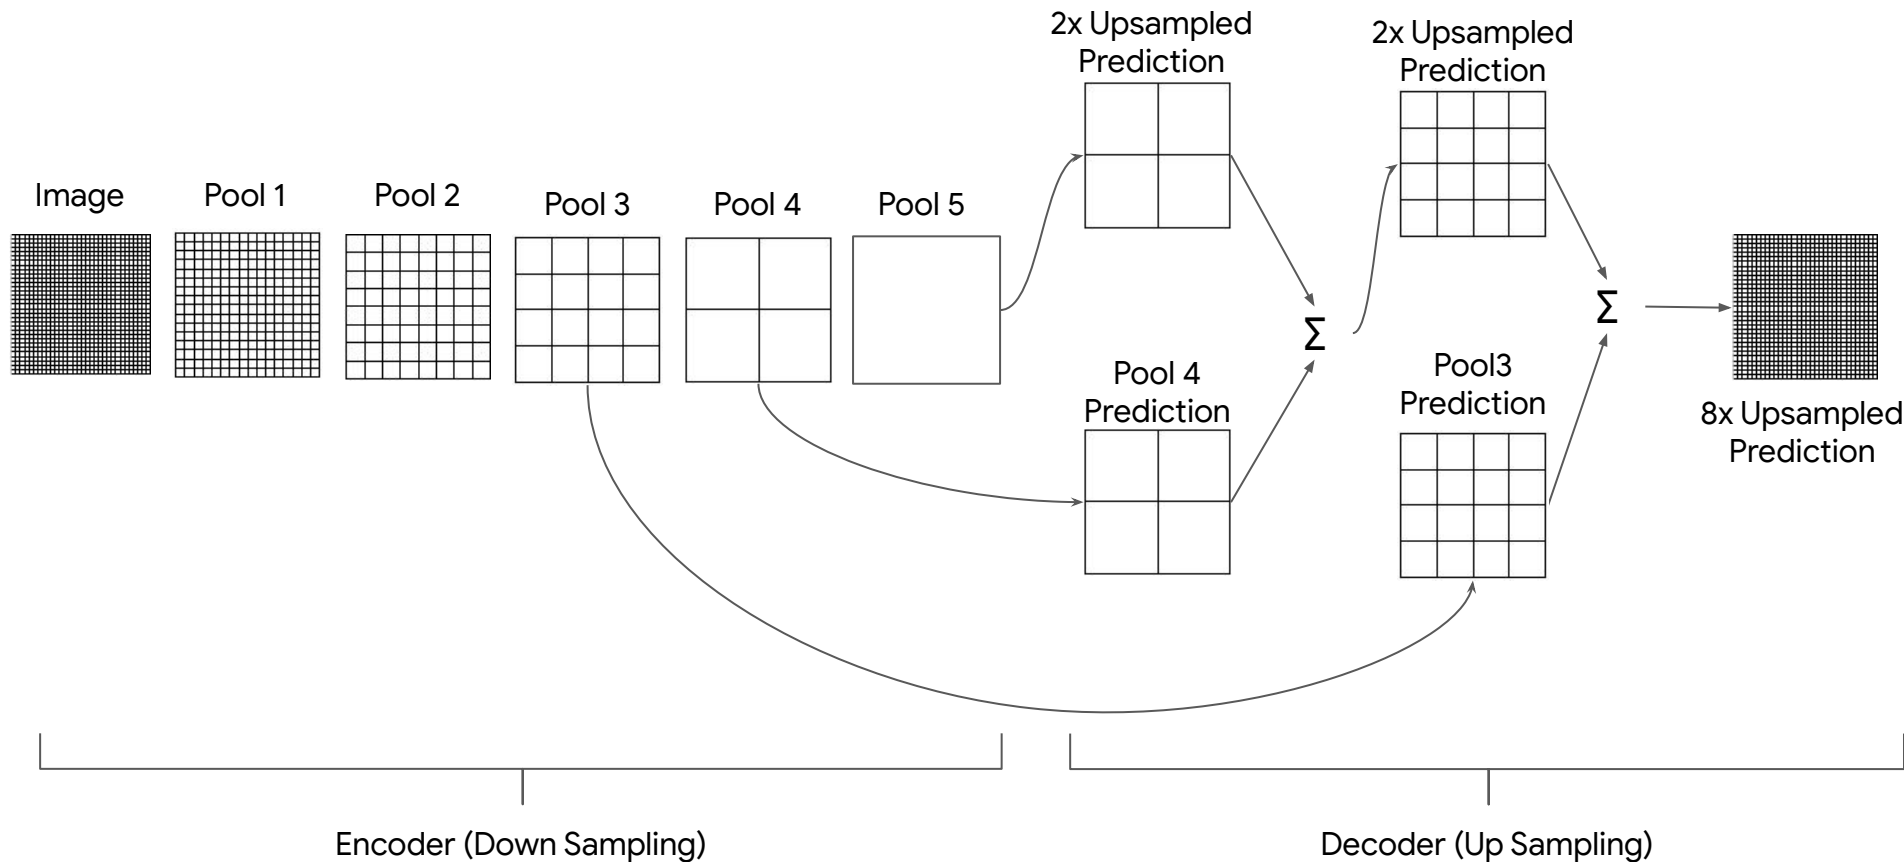
Conv2DTranspose - Usage

```
Conv2DTranspose(  
    filters=32,  
    kernel_size=(3, 3)  
)
```


- The Cambridge Driving, labelled video database (aka CamVid) contains 10 minutes of 30fps video, segmented and labelled with 32 classes
- GitHub account [divamgupta](#) has taken a subsample of the CamVid dataset to create a smaller dataset.



FCN-8



| |
|----------|
| conv1_1 |
| conv1_2 |
| pooling2 |

| |
|----------|
| conv2_1 |
| conv2_2 |
| pooling2 |

| |
|----------|
| conv3_1 |
| conv3_2 |
| conv3_3 |
| pooling3 |

| |
|----------|
| conv4_1 |
| conv4_2 |
| conv4_3 |
| pooling4 |

| |
|----------|
| conv5_1 |
| conv5_2 |
| conv5_3 |
| pooling5 |

```
def block(x, n_convs, filters, kernel_size, activation, pool_size, pool_stride, block_name):  
    for i in range(n_convs):  
        x = tf.keras.layers.Conv2D(filters=filters,  
                                     kernel_size=kernel_size, activation=activation,  
                                     padding='same',  
                                     name="{}_conv{}".format(block_name, i + 1))(x)  
  
    x = tf.keras.layers.MaxPooling2D(pool_size=pool_size, strides=pool_stride,  
                                     name="{}_pool{}".format(block_name, i+1 ))(x)  
  
    return x
```

| |
|----------|
| conv4_1 |
| conv4_2 |
| conv4_3 |
| pooling4 |

```
def block(x, n_convs, filters, kernel_size, activation, pool_size, pool_stride, block_name):  
    for i in range(n_convs):  
        x = tf.keras.layers.Conv2D(filters=filters,  
                                     kernel_size=kernel_size, activation=activation,  
                                     padding='same',  
                                     name="{}_conv{}".format(block_name, i + 1))(x)  
  
    x = tf.keras.layers.MaxPooling2D(pool_size=pool_size, strides=pool_stride,  
                                     name="{}_pool{}".format(block_name, i+1 ))(x)  
  
    return x
```

conv4_1

conv4_2

conv4_3

pooling4

```
def block(x, n_convs, filters, kernel_size, activation, pool_size, pool_stride, block_name):  
    for i in range(n_convs):  
        x = tf.keras.layers.Conv2D(filters=filters,  
                                     kernel_size=kernel_size, activation=activation,  
                                     padding='same',  
                                     name="{}_conv{}".format(block_name, i + 1))(x)  
  
x = tf.keras.layers.MaxPooling2D(pool_size=pool_size, strides=pool_stride,  
                                  name="{}_pool{}".format(block_name, i+1 ))(x)  
  
return x
```

| |
|----------|
| conv4_1 |
| conv4_2 |
| conv4_3 |
| pooling4 |

```
def VGG_16(image_input):  
    x = block(image_input, n_convs=2, filters=64, kernel_size=(3,3),  
              activation='relu', pool_size=(2,2), pool_stride=(2,2),  
              block_name='block1')  
  
    p1= x  
  
    x = block(x, n_convs=2, filters=128, kernel_size=(3,3),  
             activation='relu', pool_size=(2,2), pool_stride=(2,2),  
             block_name='block2')  
  
    p2 = x  
  
    ...
```

```
def VGG_16(image_input):  
    x = block(image_input, n_convs=2, filters=64, kernel_size=(3,3),  
              activation='relu', pool_size=(2,2), pool_stride=(2,2),  
              block_name='block1')  
    p1= x
```

```
    x = block(x, n_convs=2, filters=128, kernel_size=(3,3),  
              activation='relu', pool_size=(2,2), pool_stride=(2,2),  
              block_name='block2')  
    p2 = x
```

...

| |
|----------|
| conv1_1 |
| conv1_2 |
| pooling2 |


```
def VGG_16(image_input):  
    x = block(image_input, n_convs=2, filters=64, kernel_size=(3,3),  
              activation='relu', pool_size=(2,2), pool_stride=(2,2),  
              block_name='block1')  
  
    p1= x
```

```
    x = block(x, n_convs=2, filters=128, kernel_size=(3,3),  
              activation='relu', pool_size=(2,2), pool_stride=(2,2),  
              block_name='block2')  
  
    p2 = x
```

...

| |
|----------|
| conv1_1 |
| conv1_2 |
| pooling2 |

| |
|----------|
| conv2_1 |
| conv2_2 |
| pooling2 |

```
x = block(x, n_convs=3, filters=256, kernel_size=(3,3), activation='relu', pool_size=(2,2),  
          pool_stride=(2,2), block_name='block3')
```

```
p3 = x
```

| |
|----------|
| conv3_1 |
| conv3_2 |
| conv3_3 |
| pooling3 |

```
x = block(x, n_convs=3, filters=512, kernel_size=(3,3), activation='relu', pool_size=(2,2),  
          pool_stride=(2,2), block_name='block4')
```

```
p4 = x
```

| |
|----------|
| conv4_1 |
| conv4_2 |
| conv4_3 |
| pooling4 |


```
x = block(x, n_convs=3, filters=512, kernel_size=(3,3), activation='relu', pool_size=(2,2),  
          pool_stride=(2,2), block_name='block5')
```

```
p5 = x
```

| |
|----------|
| conv5_1 |
| conv5_2 |
| conv5_3 |
| pooling5 |

```
...
```

2x Upsampling



Define Decoder - 2x UpSampling

```
def fcn8_decoder(convs, n_classes):  
    f1, f2, f3, f4, f5 = convs  
  
    o = tf.keras.layers.Conv2DTranspose(n_classes, kernel_size=(4,4),  
                                         strides=(2,2), use_bias=False )(f5)  
  
    o = tf.keras.layers.Cropping2D(cropping=(1,1))(o)  
  
    o2 = f4  
    o2 = ( tf.keras.layers.Conv2D(n_classes, (1,1),  
                                   activation='relu', padding='same'))(o2)  
  
    o = tf.keras.layers.Add()(o, o2)  
    ...
```

Define Decoder - 2x UpSampling

```
def fcn8_decoder(convs, n_classes):  
    f1, f2, f3, f4, f5 = convs  
  
    o = tf.keras.layers.Conv2DTranspose(n_classes, kernel_size=(4,4),  
                                         strides=(2,2), use_bias=False )(f5)  
  
    o = tf.keras.layers.Cropping2D(cropping=(1,1))(o)  
  
    o2 = f4  
    o2 = ( tf.keras.layers.Conv2D(n_classes, (1,1),  
                                   activation='relu', padding='same'))(o2)  
  
    o = tf.keras.layers.Add()(o, o2)  
    ...
```

Define Decoder - 2x UpSampling

```
def fcn8_decoder(convs, n_classes):  
    f1, f2, f3, f4, f5 = convs  
  
    o = tf.keras.layers.Conv2DTranspose(n_classes, kernel_size=(4,4),  
                                         strides=(2,2), use_bias=False )(f5)  
  
    o = tf.keras.layers.Cropping2D(cropping=(1,1))(o)  
  
    o2 = f4  
    o2 = ( tf.keras.layers.Conv2D(n_classes, (1,1),  
                                   activation='relu', padding='same'))(o2)  
  
    o = tf.keras.layers.Add()( [o, o2])  
    ...
```

Define Decoder - 2x UpSampling

```
def fcn8_decoder(convs, n_classes):  
    f1, f2, f3, f4, f5 = convs  
  
    o = tf.keras.layers.Conv2DTranspose(n_classes, kernel_size=(4,4),  
                                         strides=(2,2), use_bias=False )(f5)  
  
    o = tf.keras.layers.Cropping2D(cropping=(1,1))(o)  
  
    o2 = f4  
    o2 = ( tf.keras.layers.Conv2D(n_classes, (1,1),  
                                   activation='relu', padding='same'))(o2)  
  
    o = tf.keras.layers.Add()(o, o2)  
    ...
```

Define Decoder - 2x UpSampling

```
def fcn8_decoder(convs, n_classes):  
    f1, f2, f3, f4, f5 = convs  
  
    o = tf.keras.layers.Conv2DTranspose(n_classes, kernel_size=(4,4),  
                                         strides=(2,2), use_bias=False )(f5)  
  
    o = tf.keras.layers.Cropping2D(cropping=(1,1))(o)  
  
    o2 = f4  
    o2 = ( tf.keras.layers.Conv2D(n_classes, (1,1),  
                                   activation='relu', padding='same'))(o2)  
  
    o = tf.keras.layers.Add()(o, o2)  
    ...
```


1 x 1 Convolutions

(**B, F, H, W**) - *B = # batches; F= # filters, H, W = Height/Width*

1 x 1 Convolutions

(**B**, **F**, **H**, **W**) - *B* = # *batches*; *F* = # *filters*, *H*, *W* = *Height/Width*

Apply a layer with N 1x1 Convolutions with stride of 1:

1 x 1 Convolutions

(**B, F, H, W**) - B = # batches; F = # filters, H, W = Height/Width

Apply a layer with N 1x1 Convolutions with stride of 1:

(**B, N, H, W**) - B = # batches; N = # filters, H, W = Height/Width

1 x 1 Convolutions


(**B**, **F**, H, W) - B = # batches; F = # filters, H , W = Height/Width

(**B**, **N**, H, W) - B = # batches; N = # filters, H , W = Height/Width

1 x 1 Convolutions

(**B**, **F**, H, W) - B = # batches; F = # filters, H , W = Height/Width

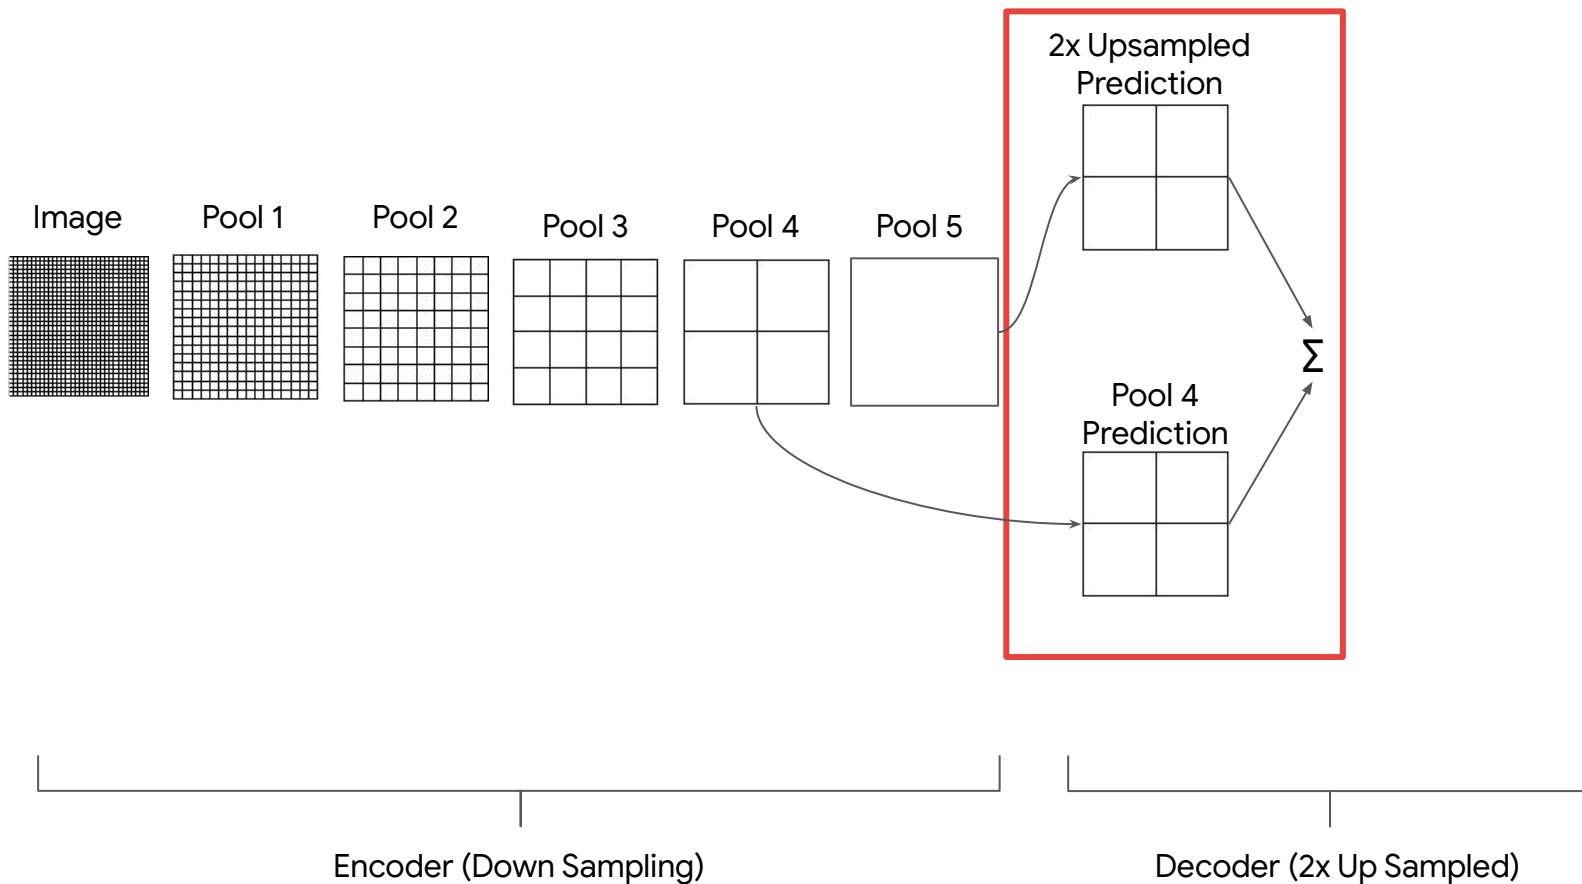
(**B**, **N**, H, W) - B = # batches; N = # filters, H , W = Height/Width



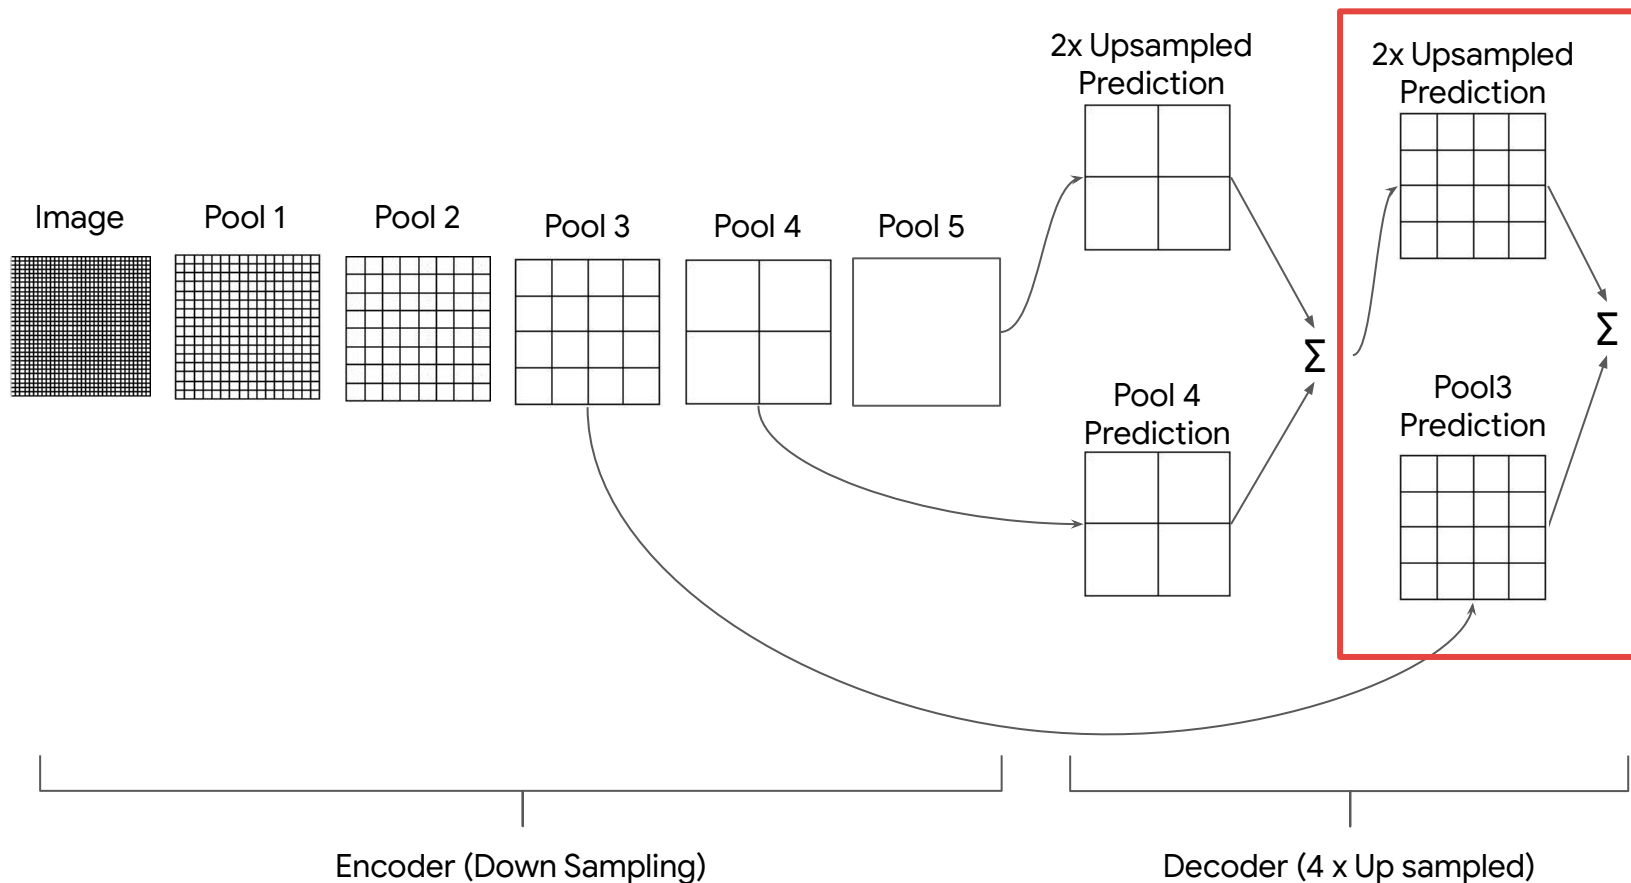
Define Decoder - 2x UpSampling

```
def fcn8_decoder(convs, n_classes):  
    f1, f2, f3, f4, f5 = convs  
  
    o = tf.keras.layers.Conv2DTranspose(n_classes, kernel_size=(4,4),  
                                         strides=(2,2), use_bias=False )(f5)  
  
    o = tf.keras.layers.Cropping2D(cropping=(1,1))(o)  
  
    o2 = f4  
    o2 = ( tf.keras.layers.Conv2D(n_classes, (1,1),  
                                   activation='relu', padding='same'))(o2)  
  
    o = tf.keras.layers.Add()( [o, o2])  
    ...
```

2x Upsampling



2 x 2 Upsampled




```
def fcn8_decoder(convs, n_classes):  
    ...  
    o = (tf.keras.layers.Conv2DTranspose( n_classes, kernel_size=(4,4),  
                                           strides=(2,2)))(o)  
  
    o = tf.keras.layers.Cropping2D(cropping=(1, 1))(o)  
  
    o2 = ( tf.keras.layers.Conv2D(n_classes,(1,1), activation='relu',  
                                   padding='same'))(f3)  
  
    o = tf.keras.layers.Add()([o, o2])  
  
    ...
```

```
def fcn8_decoder(convs, n_classes):
```

```
    ...
```

```
    o = (tf.keras.layers.Conv2DTranspose( n_classes, kernel_size=(4,4),  
                                           strides=(2,2))(o)
```

```
    o = tf.keras.layers.Cropping2D(cropping=(1, 1))(o)
```

```
    o2 = ( tf.keras.layers.Conv2D(n_classes,(1,1), activation='relu',  
                                   padding='same'))(f3)
```

```
    o = tf.keras.layers.Add()([o, o2])
```

```
    ...
```

```
def fcn8_decoder(convs, n_classes):  
    ...  
    o = (tf.keras.layers.Conv2DTranspose( n_classes, kernel_size=(4,4),  
                                           strides=(2,2)))(o)
```

```
o = tf.keras.layers.Cropping2D(cropping=(1, 1))(o)
```

```
o2 = ( tf.keras.layers.Conv2D(n_classes,(1,1), activation='relu',  
                               padding='same'))(f3)
```

```
o = tf.keras.layers.Add()( [o, o2])
```

```
...
```

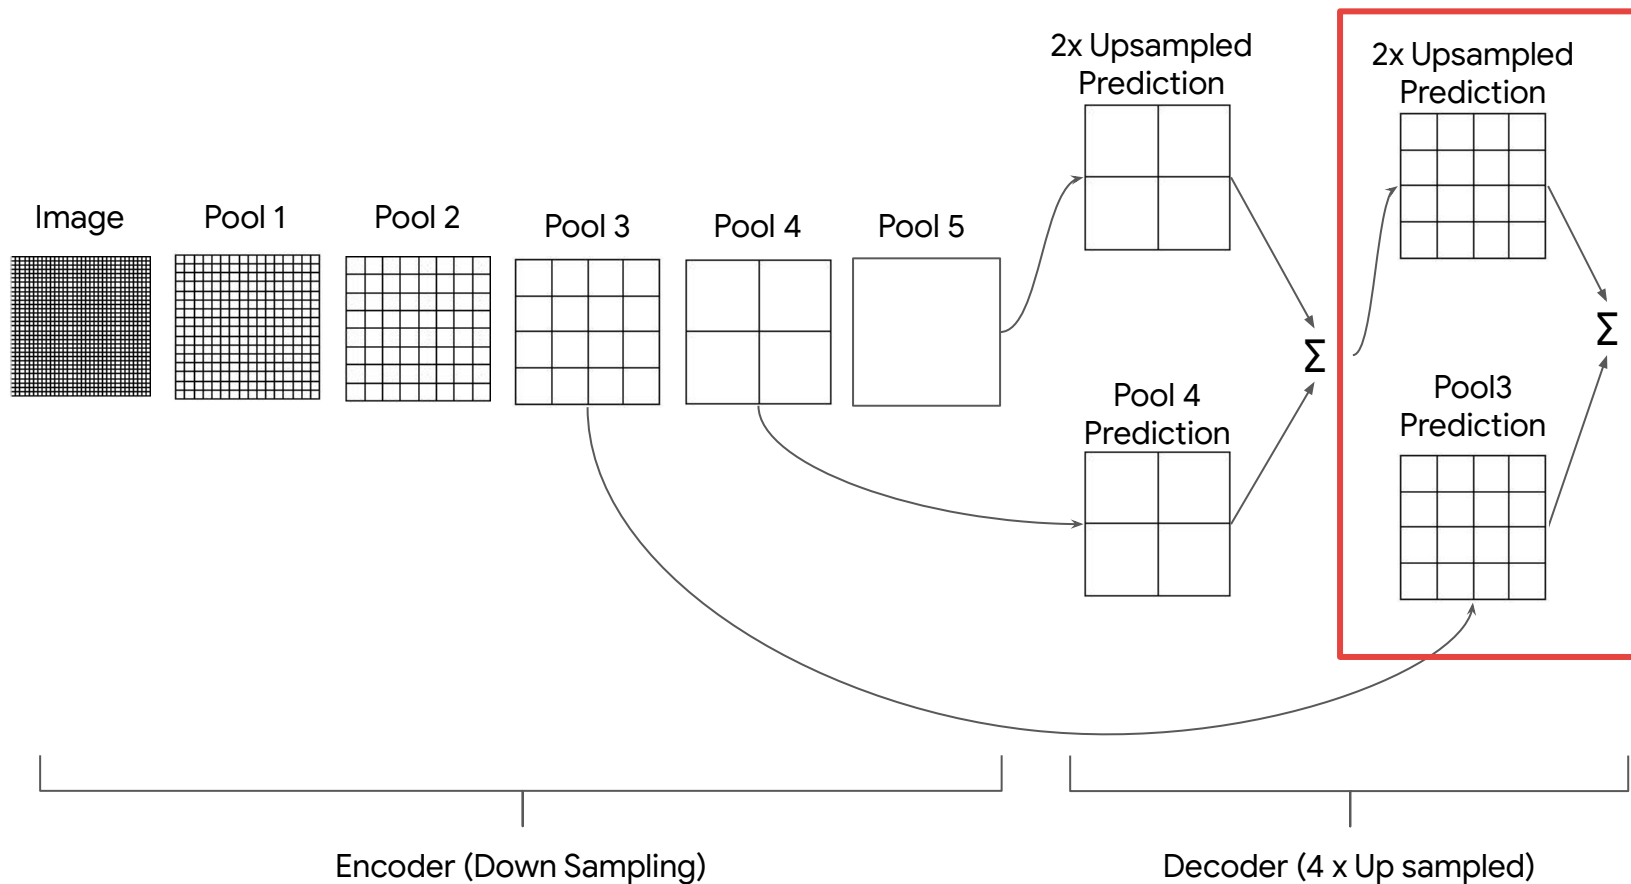
```
def fcn8_decoder(convs, n_classes):  
    ...  
    o = (tf.keras.layers.Conv2DTranspose( n_classes, kernel_size=(4,4),  
                                           strides=(2,2)))(o)  
  
    o = tf.keras.layers.Cropping2D(cropping=(1, 1))(o)  
  
    o2 = ( tf.keras.layers.Conv2D(n_classes,(1,1), activation='relu',  
                                   padding='same'))(f3)  
  
    o = tf.keras.layers.Add()( [o, o2])  
  
    ...
```

```
def fcn8_decoder(convs, n_classes):  
    ...  
    o = (tf.keras.layers.Conv2DTranspose( n_classes, kernel_size=(4,4),  
                                           strides=(2,2)))(o)  
  
    o = tf.keras.layers.Cropping2D(cropping=(1, 1))(o)  
  
    o2 = ( tf.keras.layers.Conv2D(n_classes,(1,1), activation='relu',  
                                   padding='same'))(f3)
```

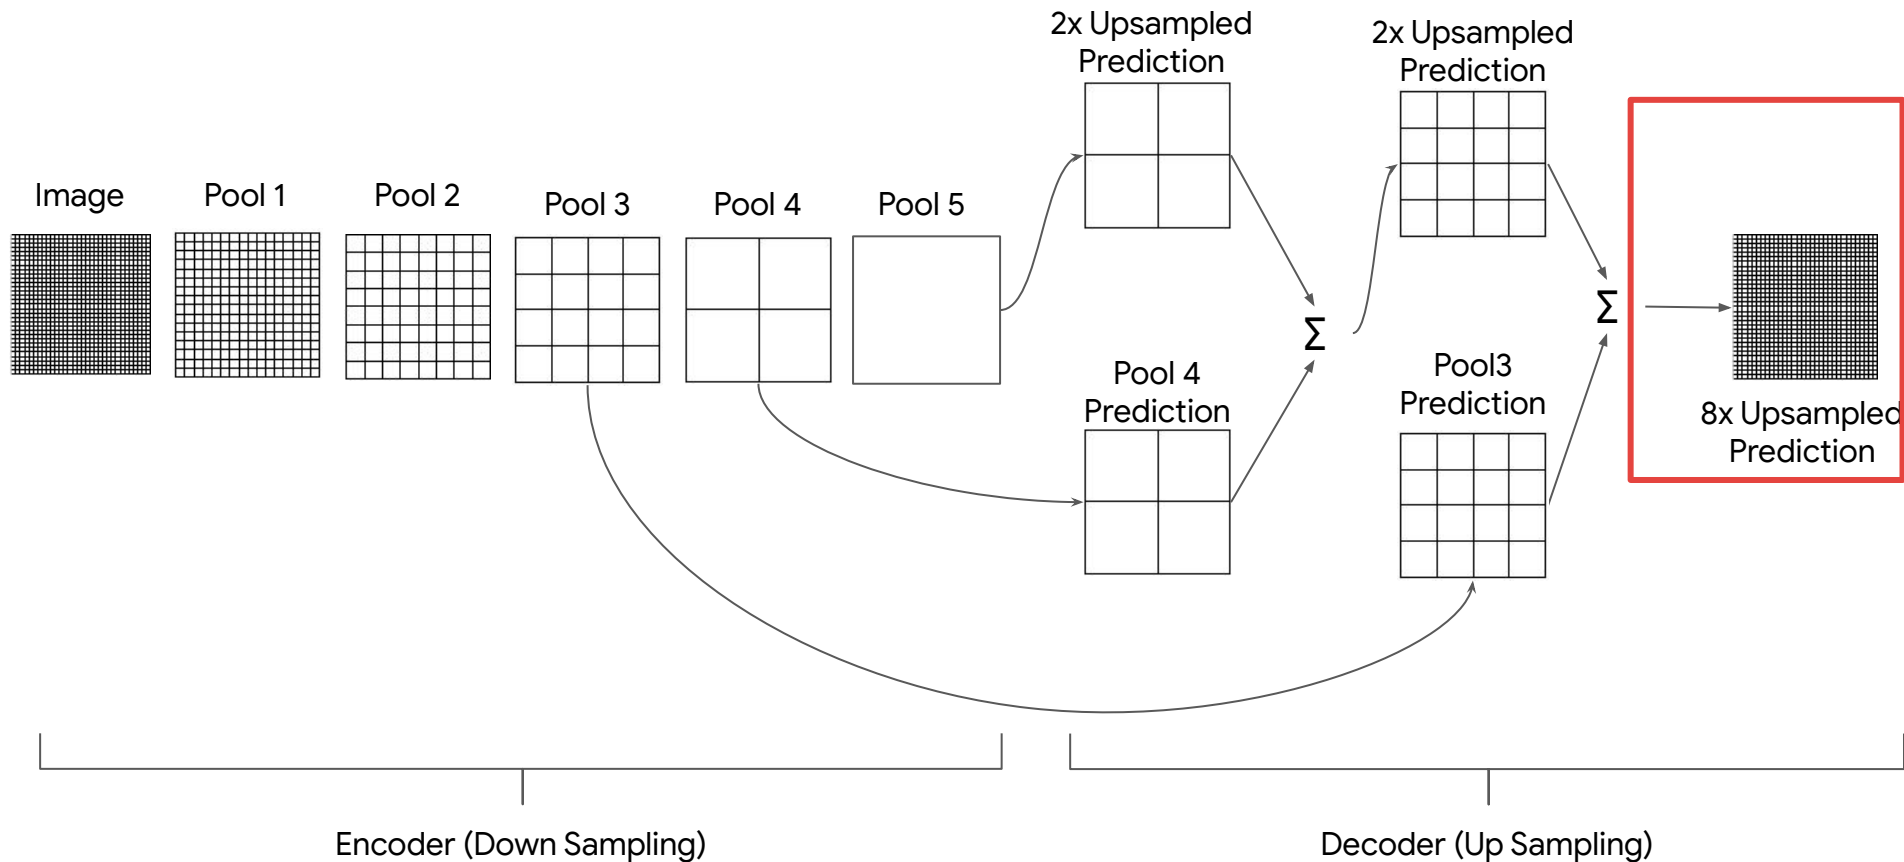
```
o = tf.keras.layers.Add()([o, o2])
```

```
...
```

2 x 2 Upsampled



2 x 2 x 8 Up Sampled



Define Decoder

```
def fcn8_decoder(convs, n_classes):  
    ...  
  
    o = tf.keras.layers.Conv2DTranspose(n_classes , kernel_size=(8,8),  
                                         strides=(8,8))(o)  
  
    o = (tf.keras.layers.Activation('softmax'))(o)  
  
    return o
```


Define Decoder

```
def fcn8_decoder(convs, n_classes):  
    ...  
  
    o = tf.keras.layers.Conv2DTranspose(n_classes, kernel_size=(8,8),  
                                         strides=(8,8))(o)  
  
    o = (tf.keras.layers.Activation('softmax'))(o)  
  
    return o
```

Define Decoder

```
def fcn8_decoder(convs, n_classes):  
    ...  
  
    o = tf.keras.layers.Conv2DTranspose(n_classes , kernel_size=(8,8),  
                                         strides=(8,8))(o)  
  
    o = (tf.keras.layers.Activation('softmax'))(o)  
  
    return o
```

Define Final Model

```
def segmentation_model():  
    inputs = tf.keras.layers.Input(shape=(224, 224, 3, ))  
    convs = VGG_16(image_input=inputs)  
    outputs = fcn8_decoder(convs, 12)  
    model = tf.keras.Model(inputs=inputs, outputs=outputs)  
    return model  
  
model = segmentation_model()
```

Define Final Model

```
def segmentation_model():  
    inputs = tf.keras.layers.Input(shape=(224, 224, 3, ))  
    convs = VGG_16(image_input=inputs)  
    outputs = fcn8_decoder(convs, 12)  
    model = tf.keras.Model(inputs=inputs, outputs=outputs)  
    return model  
  
model = segmentation_model()
```

Define Final Model


```
def segmentation_model():  
    inputs = tf.keras.layers.Input(shape=(224, 224, 3, ))  
    convs = VGG_16(image_input=inputs)  
    outputs = fcn8_decoder(convs, 12)  
    model = tf.keras.Model(inputs=inputs, outputs=outputs)  
    return model  
  
model = segmentation_model()
```

Define Final Model


```
def segmentation_model():  
    inputs = tf.keras.layers.Input(shape=(224, 224, 3, ))  
    convs = VGG_16(image_input=inputs)  
    outputs = fcn8_decoder(convs, 12)  
    model = tf.keras.Model(inputs=inputs, outputs=outputs)  
    return model
```

```
model = segmentation_model()
```


Sample Visualization of Predicted Segments



Original Image




Predicted Segments



Ground Truth Segments


Area of Overlap

Area of Overlap = sum(True Positives)




Combined Area

Combined Area = Total Pixels in predicted segmentation mask + Total Pixels in True Segmentation mask



Area of Union

Area of Union = Total Pixels in predicted segmentation mask + Total Pixels in True Segmentation mask - Area of Overlap



Calculate Areas

```
def class_wise_metrics(y_true, y_pred):  
    ...  
    smoothing_factor = 0.00001  
  
    for i in range(n_classes):  
        intersection = np.sum((y_pred == i) * (y_true == i))  
        y_true_area = np.sum((y_true == i))  
        y_pred_area = np.sum((y_pred == i))  
        combined_area = y_true_area + y_pred_area  
        union_area = combined_area - intersection  
    ...
```

Calculate Areas


```
def class_wise_metrics(y_true, y_pred):  
    ...  
    smoothening_factor = 0.00001  
  
    for i in range(n_classes):  
        intersection = np.sum((y_pred == i) * (y_true == i))  
        y_true_area = np.sum((y_true == i))  
        y_pred_area = np.sum((y_pred == i))  
        combined_area = y_true_area + y_pred_area  
        union_area = combined_area - intersection  
    ...
```

Calculate Areas

```
def class_wise_metrics(y_true, y_pred):  
    ...  
    smoothening_factor = 0.00001  
  
    for i in range(n_classes):  
        intersection = np.sum((y_pred == i) * (y_true == i))  
        y_true_area = np.sum((y_true == i))  
        y_pred_area = np.sum((y_pred == i))  
        combined_area = y_true_area + y_pred_area  
        union_area = combined_area - intersection  
    ...
```

Intersection Over Union

$$\text{IOU} = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$




Calculate IOU

```
def class_wise_metrics(y_true, y_pred):  
    class_wise_iou = []  
  
    ...  
    for i in range(n_classes):  
        ...  
        iou = (intersection) / (union_area)  
        class_wise_iou.append(iou)  
  
    return class_wise_iou
```



IOU Results

| | |
|---------------|------------------------|
| sky | 0.8779669959482955 |
| building | 0.7570989578412737 |
| column/pole | 4.57875457665808e-10 |
| road | 0.915543155822588 |
| side walk | 0.7235628237658467 |
| vegetation | 0.7664541807647628 |
| traffic light | 3.0202657798187055e-05 |
| fence | 0.006380242448568188 |
| vehicle | 0.2950299461448835 |
| pedestrian | 0.0001264333276608086 |
| bicyclist | 0.023621930993270864 |
| void | 0.16456276759816527 |





IOU Results

| | |
|---------------|------------------------|
| sky | 0.8779669959482955 |
| building | 0.7570989578412737 |
| column/pole | 4.57875457665808e-10 |
| road | 0.915543155822588 |
| side walk | 0.7235628237658467 |
| vegetation | 0.7664541807647628 |
| traffic light | 3.0202657798187055e-05 |
| fence | 0.006380242448568188 |
| vehicle | 0.2950299461448835 |
| pedestrian | 0.0001264333276608086 |
| bicyclist | 0.023621930993270864 |
| void | 0.16456276759816527 |




IOU Results

| | |
|---------------|------------------------|
| sky | 0.8779669959482955 |
| building | 0.7570989578412737 |
| column/pole | 4.57875457665808e-10 |
| road | 0.915543155822588 |
| side walk | 0.7235628237658467 |
| vegetation | 0.7664541807647628 |
| traffic light | 3.0202657798187055e-05 |
| fence | 0.006380242448568188 |
| vehicle | 0.2950299461448835 |
| pedestrian | 0.0001264333276608086 |
| bicyclist | 0.023621930993270864 |
| void | 0.16456276759816527 |



Dice Score

$$\text{Dice Score} = 2 \times \frac{\text{Area of Overlap}}{\text{Combined Area}}$$





Calculate Dice Score

```
def class_wise_metrics(y_true, y_pred):  
    class_wise_dice_score = []  
  
    ...  
    for i in range(n_classes):  
        ...  
        dice_score = 2 * (intersection) / (combined_area)  
        class_wise_dice_score.append(dice_score)  
  
    return class_wise_dice_score
```



Dice Score Results

| | |
|---------------|-----------------------|
| sky | 0.9350185576821789 |
| building | 0.861760180856767 |
| column/pole | 9.15750915331616e-10 |
| road | 0.9559097147402678 |
| side walk | 0.8396129387346395 |
| vegetation | 0.8677883515092748 |
| traffic light | 6.040349126864078e-05 |
| fence | 0.012679586065167121 |
| vehicle | 0.45563416820437186 |
| pedestrian | 0.0002528346886971326 |
| byciclist | 0.04615362426586659 |
| void | 0.2826172572009191 |



Dice Score Results

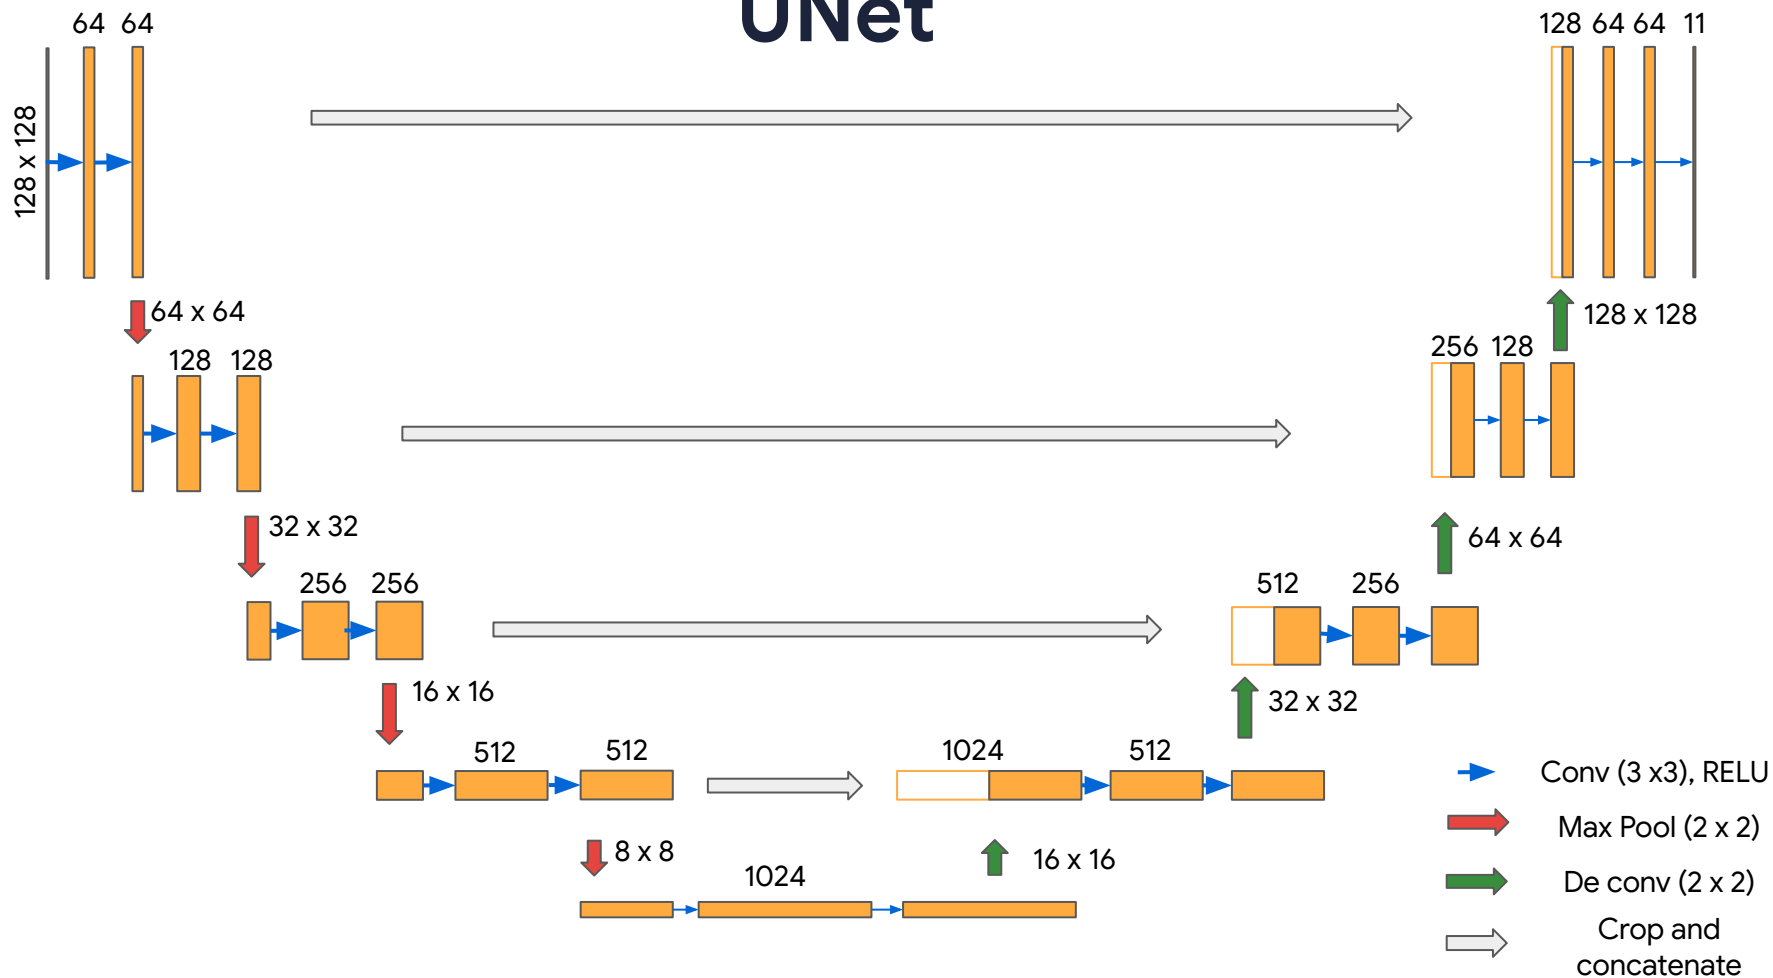
| | |
|---------------|-----------------------|
| sky | 0.9350185576821789 |
| building | 0.861760180856767 |
| column/pole | 9.15750915331616e-10 |
| road | 0.9559097147402678 |
| side walk | 0.8396129387346395 |
| vegetation | 0.8677883515092748 |
| traffic light | 6.040349126864078e-05 |
| fence | 0.012679586065167121 |
| vehicle | 0.45563416820437186 |
| pedestrian | 0.0002528346886971326 |
| bicyclist | 0.04615362426586659 |
| void | 0.2826172572009191 |



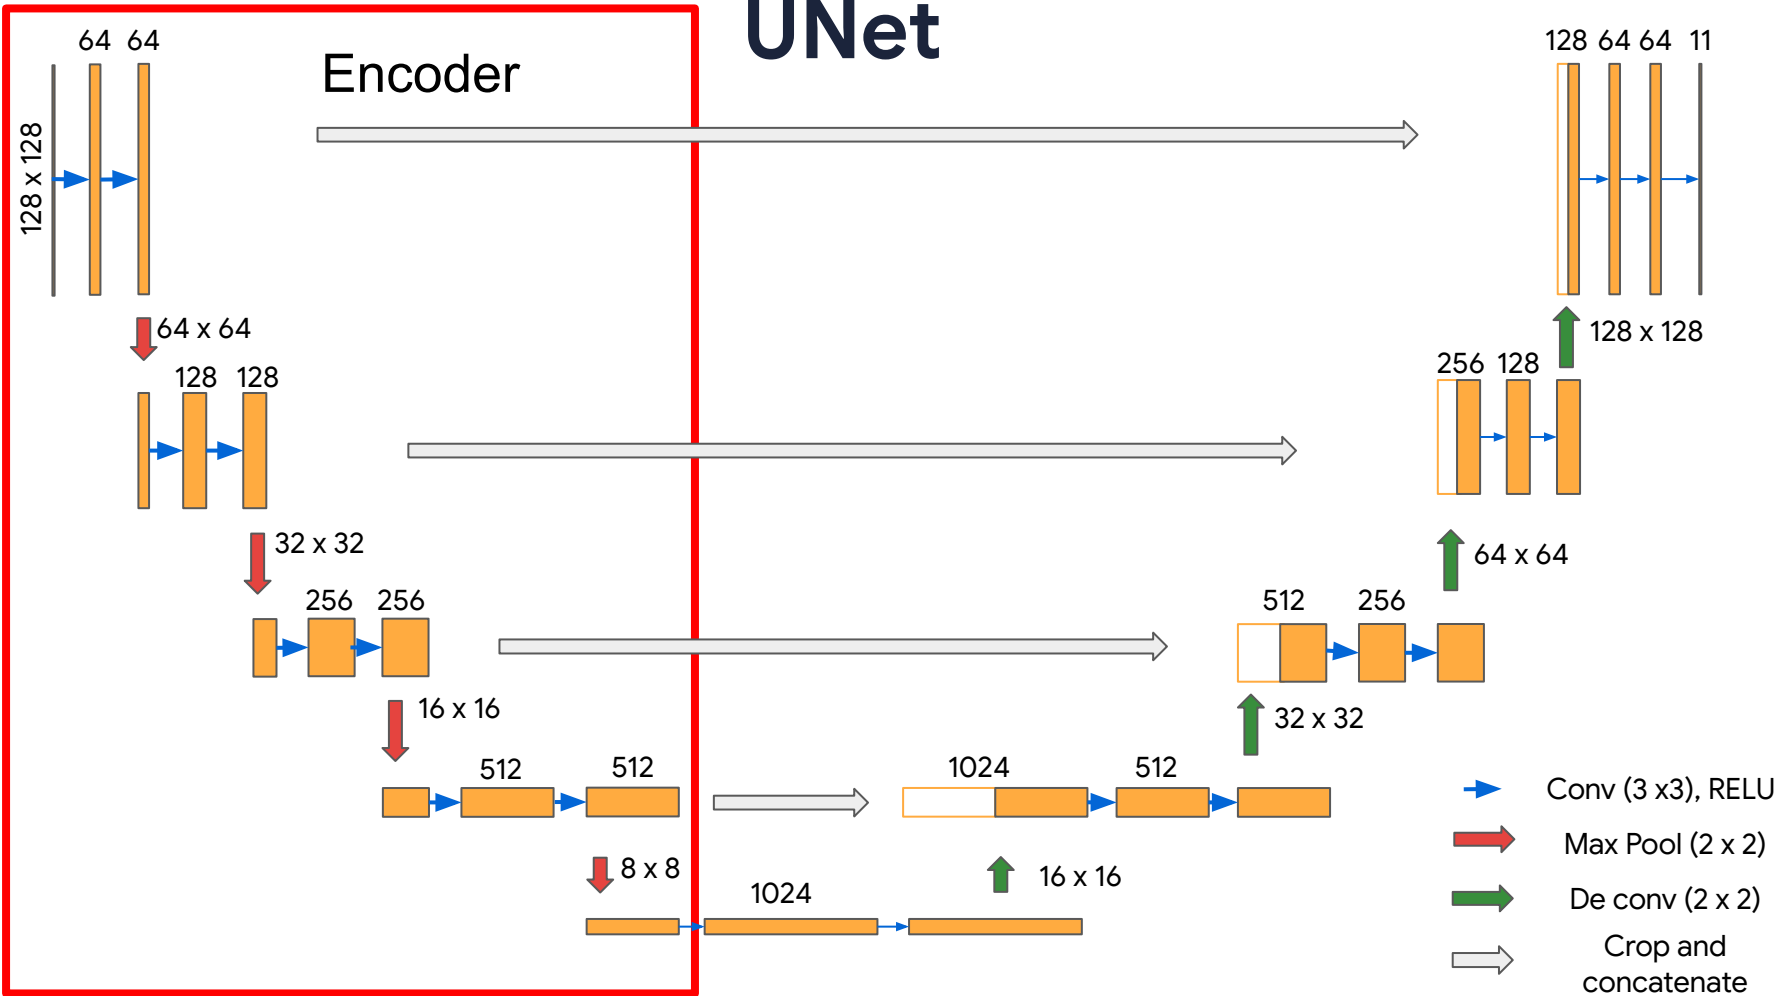
U-Net: Convolutional Networks for Biomedical Image Segmentation

Olaf Ronneberger, Philipp Fischer, Thomas Brox
<https://arxiv.org/abs/1505.04597>

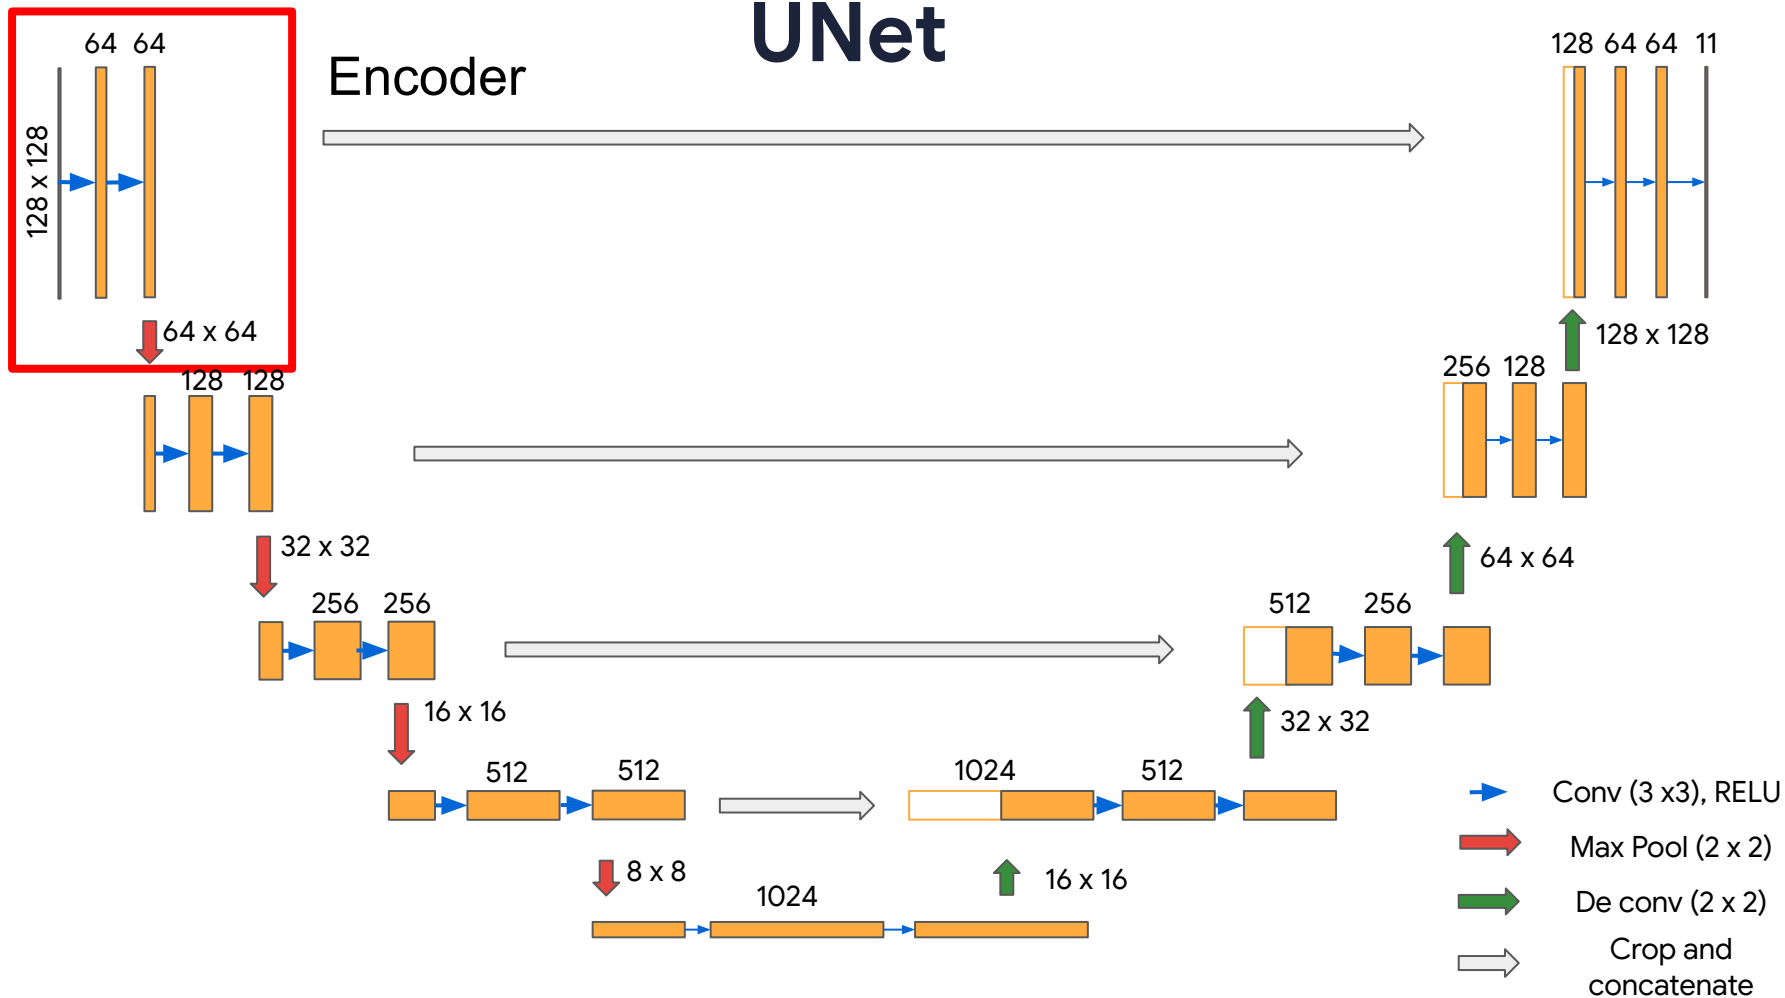
UNet



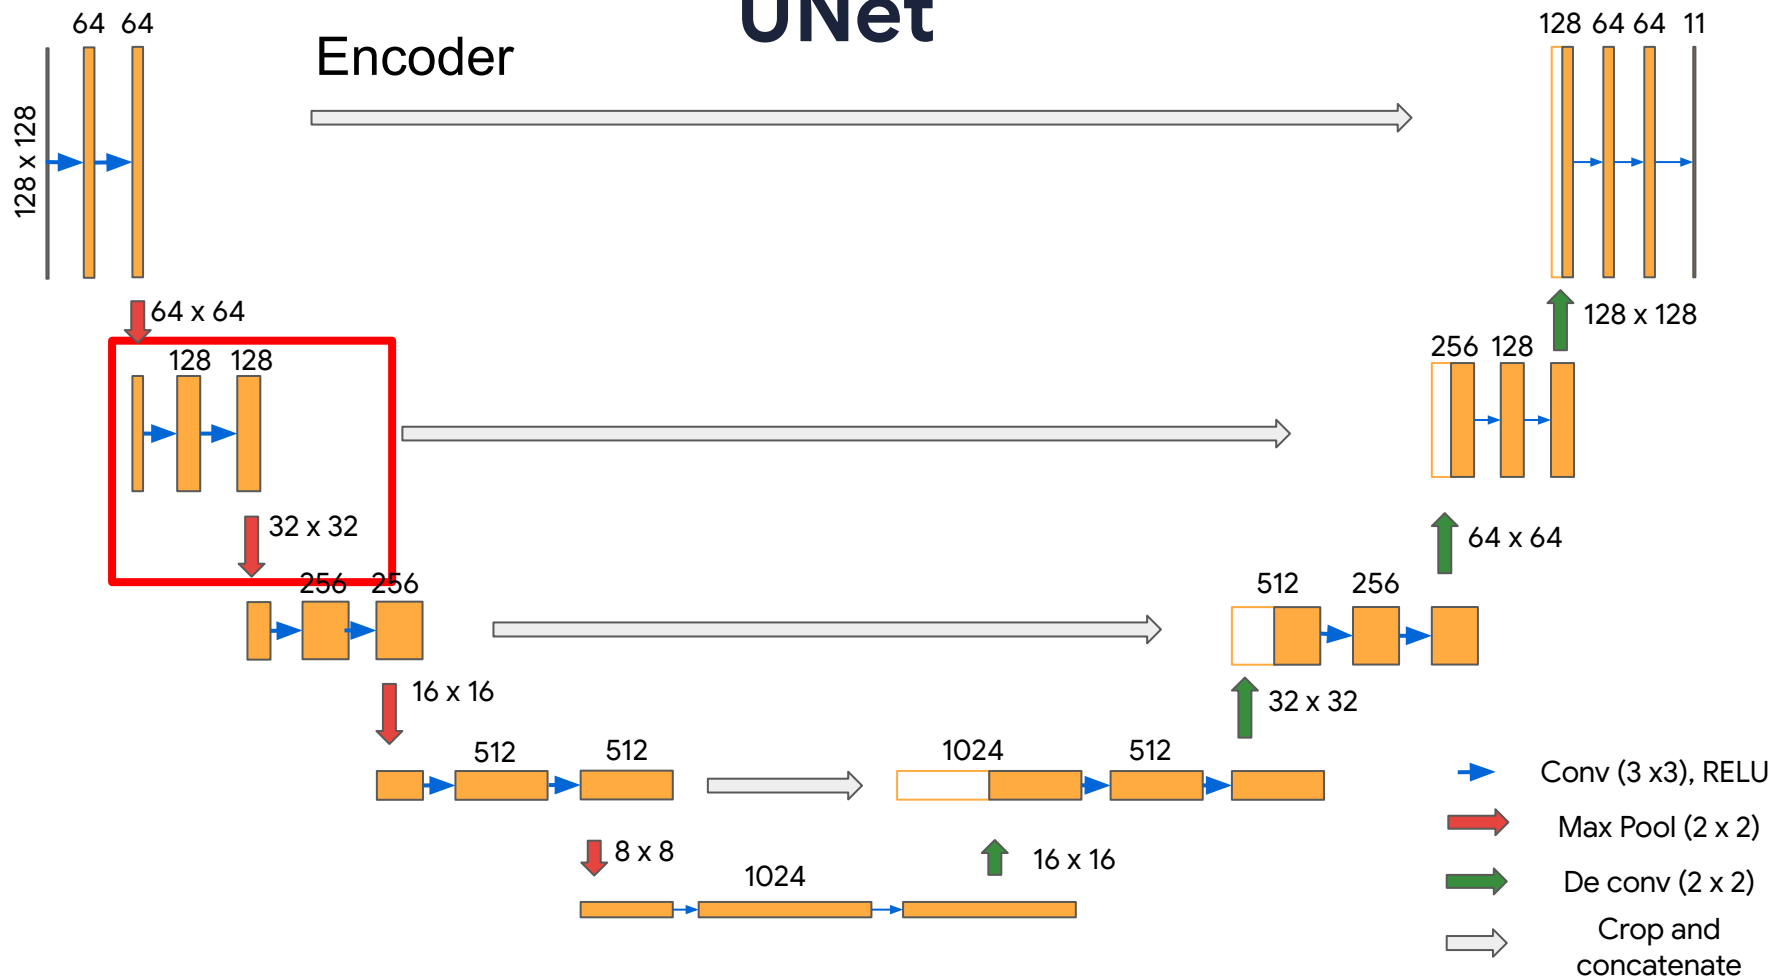
UNet



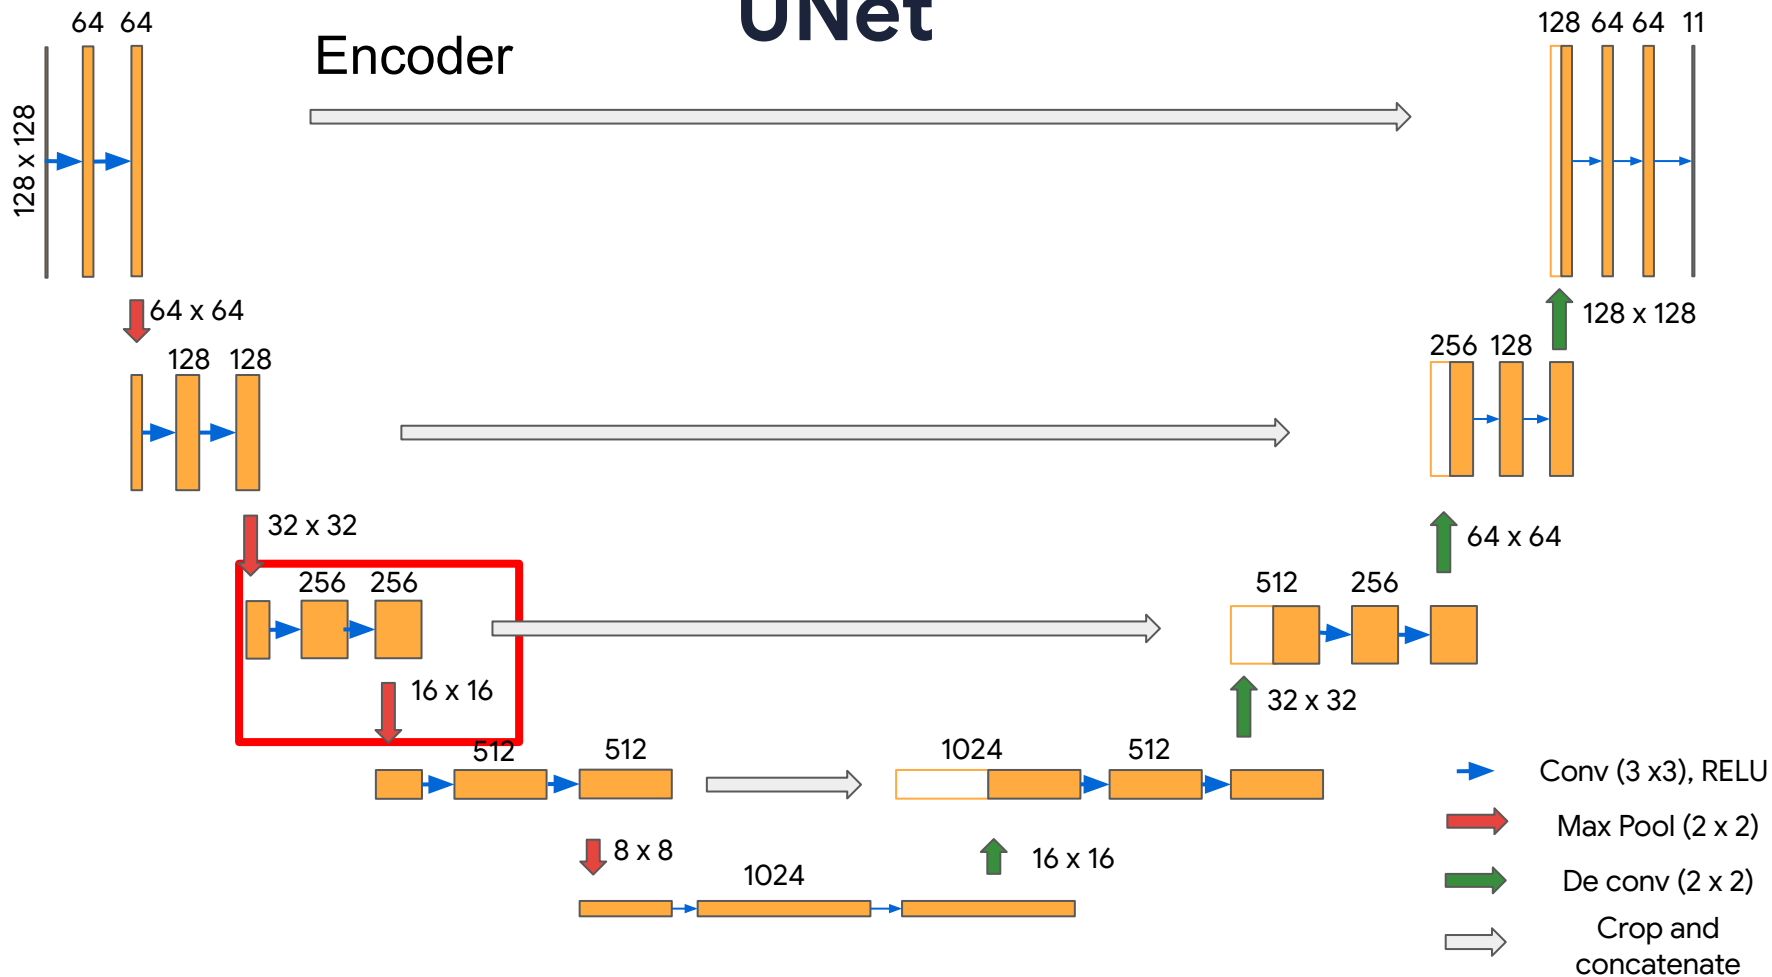
UNet




UNet



Encoder




UNet






UNet




UNet

Decoder




UNet

Decoder




UNet


Decoder



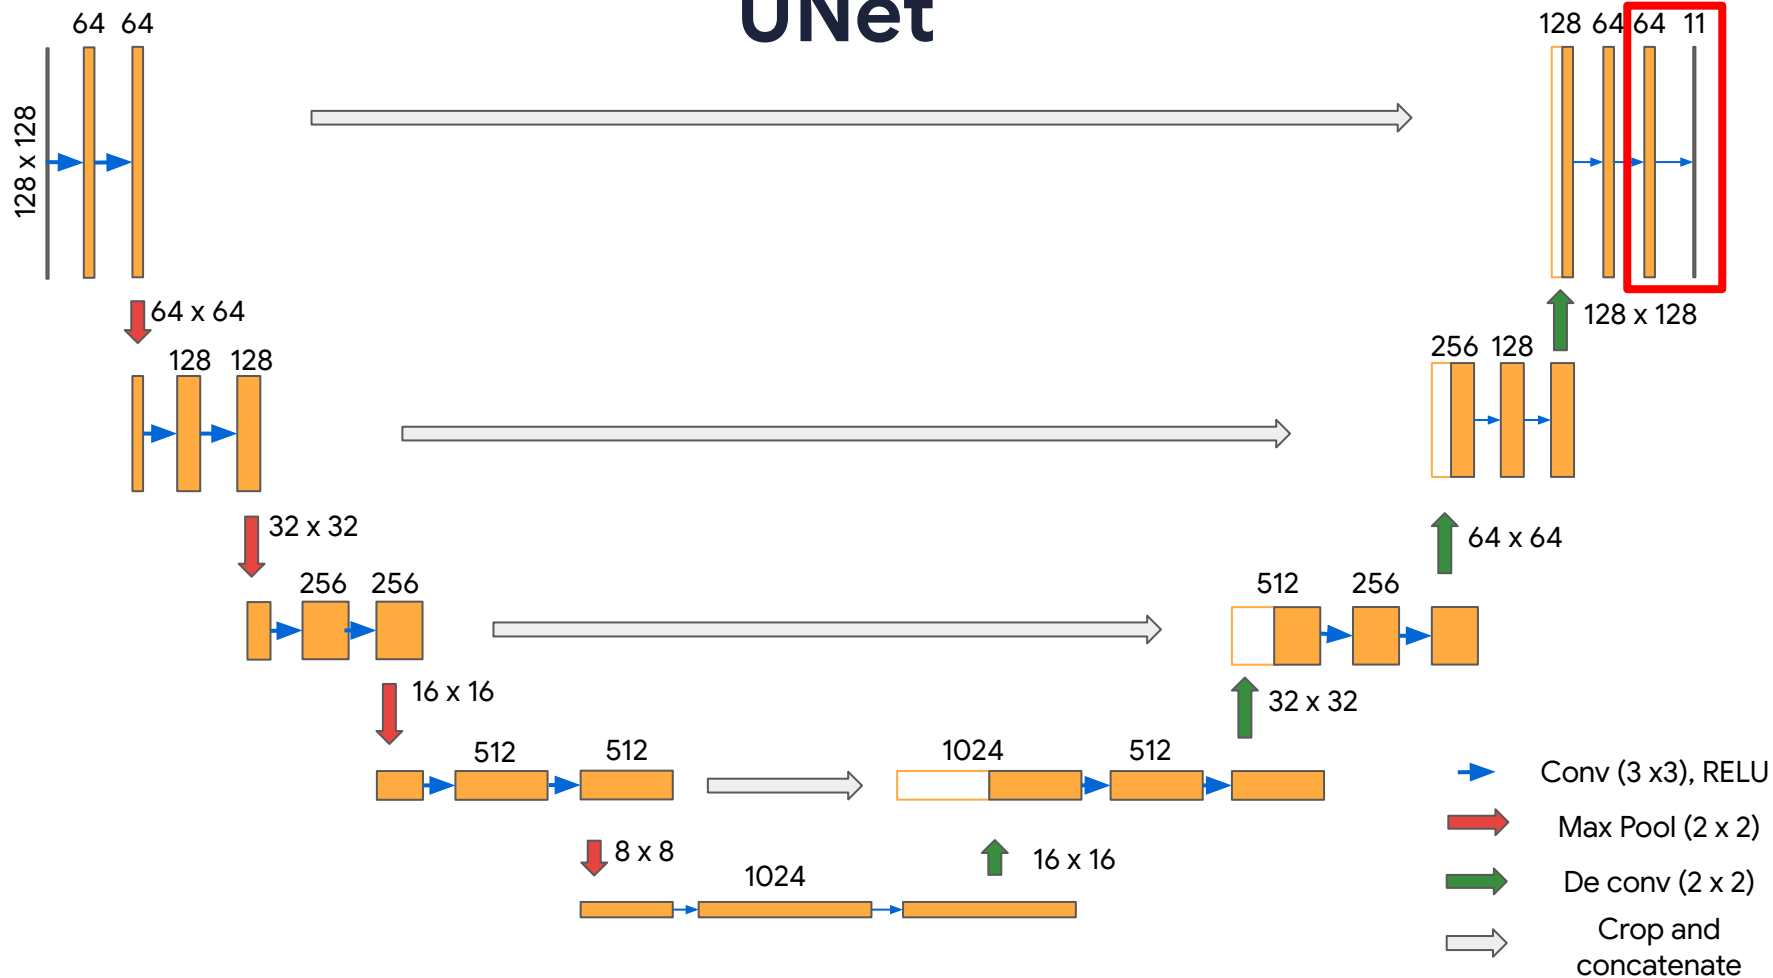
UNet



UNet




UNet



Oxford IIT Pets Dataset

- The dataset can be found at <https://www.robots.ox.ac.uk/~vgg/data/pets/>
- 37 category pet dataset.
- Dataset has images, associated breed, segmentation masks and head ROI.
- Segmentation Masks are labelled either {1, 2, 3}
 - 1 - Pet
 - 2 - Outline
 - 3 - Background


Oxford IIT Pets Dataset




Images

Segments

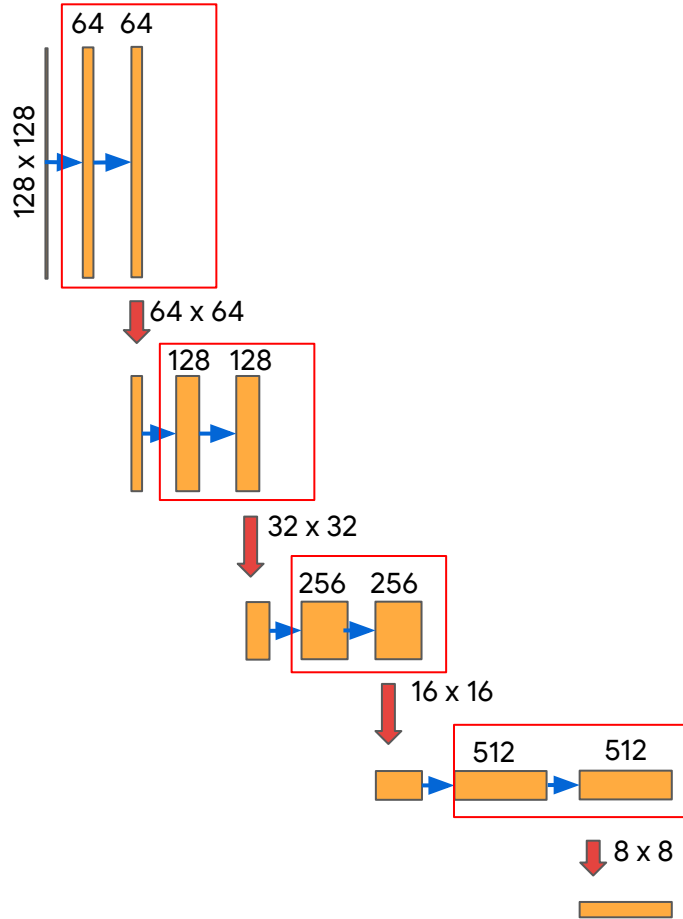
Architecture - UNet



Encoder



Encoder




```
def conv2d_block(input_tensor, n_filters, kernel_size = 3):  
    x = input_tensor  
  
    for i in range(2):  
        x = tf.keras.layers.Conv2D(filters = n_filters,  
                                     kernel_size = (kernel_size, kernel_size))(x)  
  
        x = tf.keras.layers.Activation('relu')(x)  
  
    return x
```

```
def conv2d_block(input_tensor, n_filters, kernel_size = 3):  
    x = input_tensor  
  
    for i in range(2):  
        x = tf.keras.layers.Conv2D(filters = n_filters,  
                                     kernel_size = (kernel_size, kernel_size))(x)  
  
        x = tf.keras.layers.Activation('relu')(x)  
  
    return x
```

```
def conv2d_block(input_tensor, n_filters, kernel_size = 3):  
    x = input_tensor  
  
    for i in range(2):  
        x = tf.keras.layers.Conv2D(filters = n_filters,  
                                     kernel_size = (kernel_size, kernel_size))(x)  
  
        x = tf.keras.layers.Activation('relu')(x)  
  
    return x
```

```
def conv2d_block(input_tensor, n_filters, kernel_size = 3):  
    x = input_tensor  
  
    for i in range(2):  
        x = tf.keras.layers.Conv2D(filters = n_filters,  
                                     kernel_size = (kernel_size, kernel_size))(x)  
  
        x = tf.keras.layers.Activation('relu')(x)  
  
    return x
```

Encoder



```
def encoder_block(inputs, n_filters, pool_size, dropout):  
    f = conv2d_block(inputs, n_filters=n_filters)  
    p = tf.keras.layers.MaxPooling2D(pool_size)(f)  
    p = tf.keras.layers.Dropout(dropout)(p)  
  
    return f, p
```



```
def encoder_block(inputs, n_filters, pool_size, dropout):  
    f = conv2d_block(inputs, n_filters=n_filters)  
    p = tf.keras.layers.MaxPooling2D(pool_size)(f)  
    p = tf.keras.layers.Dropout(dropout)(p)  
  
    return f, p
```

```
def encoder_block(inputs, n_filters, pool_size, dropout):  
    f = conv2d_block(inputs, n_filters=n_filters)  
    p = tf.keras.layers.MaxPooling2D(pool_size)(f)  
    p = tf.keras.layers.Dropout(dropout)(p)  
  
    return f, p
```

```
def encoder_block(inputs, n_filters, pool_size, dropout):  
    f = conv2d_block(inputs, n_filters=n_filters)  
    p = tf.keras.layers.MaxPooling2D(pool_size)(f)  
    p = tf.keras.layers.Dropout(dropout)(p)  
  
    return f, p
```

```
def encoder_block(inputs, n_filters, pool_size, dropout):  
    f = conv2d_block(inputs, n_filters=n_filters)  
    p = tf.keras.layers.MaxPooling2D(pool_size)(f)  
    p = tf.keras.layers.Dropout(dropout)(p)  
  
    return f, p
```

```
def encoder_block(inputs, n_filters, pool_size, dropout):  
    f = conv2d_block(inputs, n_filters=n_filters)  
    p = tf.keras.layers.MaxPooling2D(pool_size)(f)  
    p = tf.keras.layers.Dropout(dropout)(p)  
  
    return f, p
```

```
def encoder(inputs):
```

```
    f1, p1 = encoder_block(inputs, n_filters=64, pool_size=(2,2), dropout=0.3)
```

```
    f2, p2 = encoder_block(p1, n_filters=128, pool_size=(2,2), dropout=0.3)
```

```
    f3, p3 = encoder_block(p2, n_filters=256, pool_size=(2,2), dropout=0.3)
```

```
    f4, p4 = encoder_block(p3, n_filters=512, pool_size=(2,2), dropout=0.3)
```

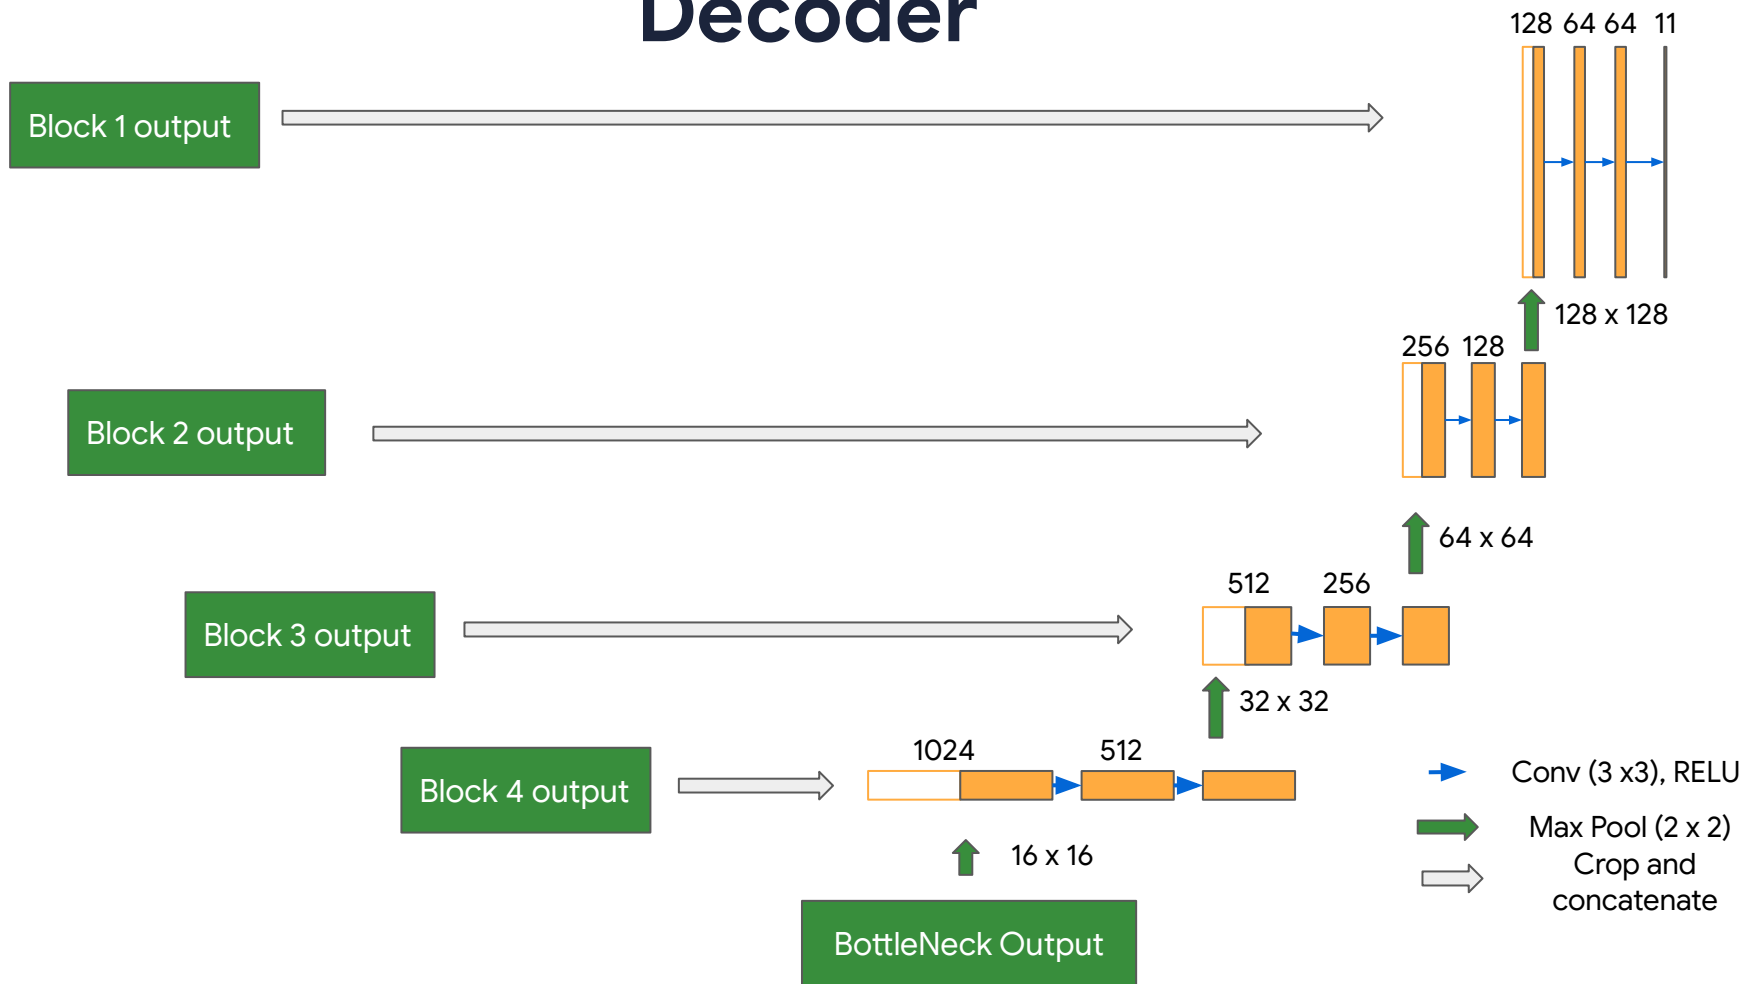
```
    return p4, (f1, f2, f3, f4)
```

Bottleneck



```
def bottleneck(inputs):  
    bottle_neck = conv2d_block(inputs, n_filters=1024)  
  
    return bottle_neck
```


Decoder



```
def decoder_block(inputs, conv_output, n_filters, kernel_size, strides, dropout):  
    u = tf.keras.layers.Conv2DTranspose(n_filters, kernel_size, strides = strides,  
                                         padding = 'same')(inputs)  
    c = tf.keras.layers.concatenate([u, conv_output])  
    c = tf.keras.layers.Dropout(dropout)(c)  
    c = conv2d_block(c, n_filters, kernel_size=3)  
  
    return c
```

```
def decoder_block(inputs, conv_output, n_filters, kernel_size, strides, dropout):  
    u = tf.keras.layers.Conv2DTranspose(n_filters, kernel_size, strides = strides,  
                                         padding = 'same')(inputs)  
  
    c = tf.keras.layers.concatenate([u, conv_output])  
  
    c = tf.keras.layers.Dropout(dropout)(c)  
  
    c = conv2d_block(c, n_filters, kernel_size=3)  
  
    return c
```

```
def decoder_block(inputs, conv_output, n_filters, kernel_size, strides, dropout):  
    u = tf.keras.layers.Conv2DTranspose(n_filters, kernel_size, strides = strides,  
                                         padding = 'same')(inputs)  
    c = tf.keras.layers.concatenate([u, conv_output])  
    c = tf.keras.layers.Dropout(dropout)(c)  
    c = conv2d_block(c, n_filters, kernel_size=3)  
  
    return c
```

```
def decoder_block(inputs, conv_output, n_filters, kernel_size, strides, dropout):  
    u = tf.keras.layers.Conv2DTranspose(n_filters, kernel_size, strides = strides,  
                                         padding = 'same')(inputs)  
  
    c = tf.keras.layers.concatenate([u, conv_output])  
  
    c = tf.keras.layers.Dropout(dropout)(c)  
  
    c = conv2d_block(c, n_filters, kernel_size=3)  
  
    return c
```

```
def decoder_block(inputs, conv_output, n_filters, kernel_size, strides, dropout):  
    u = tf.keras.layers.Conv2DTranspose(n_filters, kernel_size, strides = strides,  
                                         padding = 'same')(inputs)  
  
    c = tf.keras.layers.concatenate([u, conv_output])  
  
    c = tf.keras.layers.Dropout(dropout)(c)  
  
    c = conv2d_block(c, n_filters, kernel_size=3)  
  
    return c
```

```
def decoder_block(inputs, conv_output, n_filters, kernel_size, strides, dropout):  
    u = tf.keras.layers.Conv2DTranspose(n_filters, kernel_size, strides = strides,  
                                         padding = 'same')(inputs)  
    c = tf.keras.layers.concatenate([u, conv_output])  
    c = tf.keras.layers.Dropout(dropout)(c)  
    c = conv2d_block(c, n_filters, kernel_size=3)  
  
    return c
```

```
def decoder_block(inputs, conv_output, n_filters, kernel_size, strides, dropout):  
    u = tf.keras.layers.Conv2DTranspose(n_filters, kernel_size, strides = strides,  
                                         padding = 'same')(inputs)  
    c = tf.keras.layers.concatenate([u, conv_output])  
    c = tf.keras.layers.Dropout(dropout)(c)  
    c = conv2d_block(c, n_filters, kernel_size=3)  
  
    return c
```



```
def decoder_block(inputs, conv_output, n_filters, kernel_size, strides, dropout):  
    u = tf.keras.layers.Conv2DTranspose(n_filters, kernel_size, strides = strides,  
                                         padding = 'same')(inputs)  
    c = tf.keras.layers.concatenate([u, conv_output])  
    c = tf.keras.layers.Dropout(dropout)(c)  
    c = conv2d_block(c, n_filters, kernel_size=3)  
  
    return c
```

```
def decoder(inputs, convs):  
    f1, f2, f3, f4 = convs  
  
    c6 = decoder_block(inputs, f4, n_filters=512, kernel_size=(3,3), strides=(2,2), dropout=0.3)  
    c7 = decoder_block(c6, f3, n_filters=256, kernel_size=(3,3), strides=(2,2), dropout=0.3)  
    c8 = decoder_block(c7, f2, n_filters=128, kernel_size=(3,3), strides=(2,2), dropout=0.3)  
    c9 = decoder_block(c8, f1, n_filters=64, kernel_size=(3,3), strides=(2,2), dropout=0.3)  
  
    outputs = tf.keras.layers.Conv2D(3, (1, 1), activation='softmax')(c9)  
  
    return outputs
```

```
def unet():  
    inputs = tf.keras.layers.Input(shape=(128,128,3))  
    encoder_output, convs = encoder(inputs)  
  
    bottle_neck = bottleneck(encoder_output)  
  
    outputs = decoder(bottle_neck, convs)  
    model = tf.keras.Model(inputs=inputs, outputs=outputs)  
  
    return model
```

```
def unet():
```

```
    inputs = tf.keras.layers.Input(shape=(128,128,3))
```

```
    encoder_output, convs = encoder(inputs)
```

```
    bottle_neck = bottleneck(encoder_output)
```

```
    outputs = decoder(bottle_neck, convs)
```

```
    model = tf.keras.Model(inputs=inputs, outputs=outputs)
```

```
    return model
```

```
def unet():  
    inputs = tf.keras.layers.Input(shape=(128,128,3))  
    encoder_output, convs = encoder(inputs)  
  
    bottle_neck = bottleneck(encoder_output)  
  
    outputs = decoder(bottle_neck, convs)  
    model = tf.keras.Model(inputs=inputs, outputs=outputs)  
  
    return model
```

```
def unet():  
    inputs = tf.keras.layers.Input(shape=(128,128,3))  
    encoder_output, convs = encoder(inputs)  
  
    bottle_neck = bottleneck(encoder_output)  
  
    outputs = decoder(bottle_neck, convs)  
    model = tf.keras.Model(inputs=inputs, outputs=outputs)  
  
    return model
```

```
def unet():  
    inputs = tf.keras.layers.Input(shape=(128,128,3))  
    encoder_output, convs = encoder(inputs)  
  
    bottle_neck = bottleneck(encoder_output)  
  
    outputs = decoder(bottle_neck, convs)  
    model = tf.keras.Model(inputs=inputs, outputs=outputs)  
  
    return model
```


```
def unet():  
    inputs = tf.keras.layers.Input(shape=(128,128,3))  
    encoder_output, convs = encoder(inputs)  
  
    bottle_neck = bottleneck(encoder_output)  
  
    outputs = decoder(bottle_neck, convs)  
    model = tf.keras.Model(inputs=inputs, outputs=outputs)  
  
    return model
```


Sample Visualization of Predicted Segments

Image



Predicted Mask



True Mask




pet: IOU: 0.9084117322372148 Dice Score: 0.9520081199787308

background: IOU: 0.8237179491525691 Dice Score: 0.9033391930467228

outline: IOU: 0.5384615399787589 Dice Score: 0.7000000027777779



Mask R-CNN




Mask R-CNN (Facebook AI Research)


By: Kaiming He, Georgia Gkioxari, Piotr Dollár, Ross Girshick

<https://arxiv.org/abs/1703.06870>

Mask R-CNN



Mask R-CNN



Running Mask R-CNN on TPUs



```
[ ] 1 max_boxes_to_draw = 50  #@param {type:"  
2 min_score_thresh = 0.1  #@param {type:"  
3  
4 image_with_detections = visualization_ut  
5   np_image,  
6   detection_boxes,  
7   detection_classes,  
8   detection_scores,  
9   category_index,  
10  instance_masks=segmentations,  
11  use_normalized_coordinates=False,  
12  max_boxes_to_draw=max_boxes_to_draw,  
13  min_score_thresh=min_score_thresh)  
14 output_image_path = 'test_results.jpg'  
15 Image.fromarray(image_with_detections.as  
16 display.display(display.Image(output_ima
```

max_boxes_to_draw: 50

min_score_thresh: 0.1

```
def dict_to_tf_example(data, label_map_dict, image_dir):  
    # Convert segmentation masks to run length encoded masks  
    run_len_encoding = rle(data['segmentation'], image_height, image_width)  
    ...  
    return example
```

```
def dict_to_tf_example(data, label_map_dict, image_dir):  
    # Convert segmentation masks to run length encoded masks  
    run_len_encoding = rle(data['segmentation'], image_height, image_width)  
    ...  
    return example
```