

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>

# Style Transfer

Style Image



Wassily Kandinsky (wikipedia)

[https://upload.wikimedia.org/wikipedia/commons/b/b4/Vassily\\_Kandinsky%2C\\_1913\\_-\\_Composition\\_7.jpg](https://upload.wikimedia.org/wikipedia/commons/b/b4/Vassily_Kandinsky%2C_1913_-_Composition_7.jpg)

Content Image



[https://cdn.pixabay.com/photo/2017/02/28/23/00/swan-2107052\\_1280.jpg](https://cdn.pixabay.com/photo/2017/02/28/23/00/swan-2107052_1280.jpg)

Stylized Image



# Approaches to Style Transfer

1. Supervised Learning.

2. Neural Style Transfer

3. Fast Neural Style Transfer

# Approaches to Style Transfer

1. Supervised Learning.
2. Neural Style Transfer
3. Fast Neural Style Transfer

# Approaches to Style Transfer

1. Supervised Learning.
2. Neural Style Transfer
3. Fast Neural Style Transfer

# 1. Supervised Learning

- Pairs: original & stylized image.
  - Need lots of pairs!

# 1. Supervised Learning

- Pairs: original & stylized image.
  - Need lots of pairs!

## 2. Neural Style Transfer

- pre-trained model
- Inputs: a single pair of images
  - Extract style from image 1
  - Extract content from image 2
- Generate image to match style and content.
  - In a loop: minimize loss



## 2. Neural Style Transfer

- pre-trained model
- Inputs: a single pair of images
  - Extract style from image 1
  - Extract content from image 2
- Generate image to match style and content.
  - In a loop: minimize loss

## 2. Neural Style Transfer

- pre-trained model
- Inputs: a single pair of images
  - Extract style from image 1
  - Extract content from image 2
- Generate image to match style and content.
  - In a loop: minimize loss

## 2. Neural Style Transfer

- pre-trained model
- Inputs: a single pair of images
  - Extract style from image 1
  - Extract content from image 2
- Generate image to match style and content.
  - In a loop: minimize loss

## 2. Neural Style Transfer

- pre-trained model
- Inputs: a single pair of images
  - Extract style from image 1
  - Extract content from image 2
- Generate image to match style and content.
  - In a loop: minimize loss

## 2. Neural Style Transfer

- pre-trained model
- Inputs: a single pair of images
  - Extract style from image 1
  - Extract content from image 2
- Generate image to match style and content.
  - In a loop: minimize loss

# A Neural Algorithm of Artistic Style

Leon Gatys  
Alexander Ecker  
Matthias Bethge





Figure 2: Images that combine the content of a photograph with the style of several well-known artworks. The images were created by finding an image that simultaneously matches the content representation of the photograph and the style representation of the artwork (see Methods). The

Update Input image with gradients wrt. Loss using optimizer



Content Image




Pre trained CNN (VGG - 19)

Content Loss




Generated Image




Pre trained CNN (VGG - 19)

Style Loss



Style Image

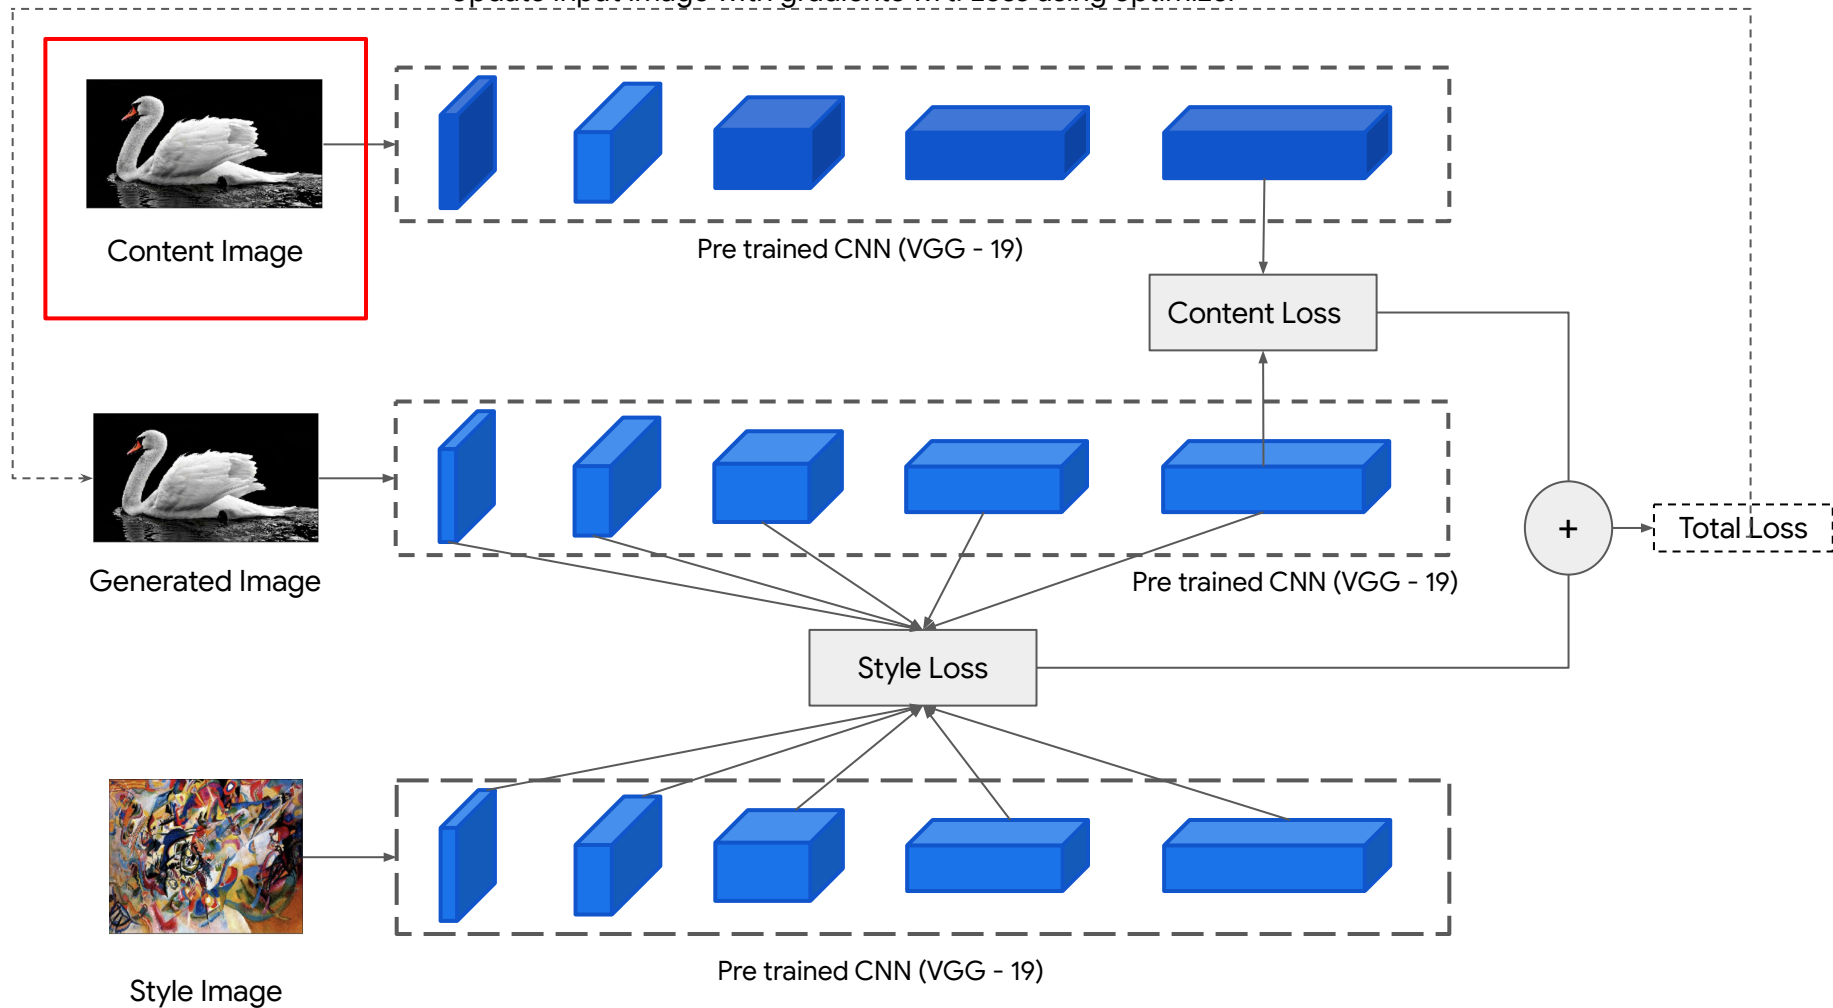


Pre trained CNN (VGG - 19)

+

Total Loss

Update Input image with gradients wrt. Loss using optimizer






Update Input image with gradients wrt. Loss using optimizer



Content Image




Pre trained CNN (VGG - 19)

Content Loss



Generated Image




Pre trained CNN (VGG - 19)

Style Loss



Style Image

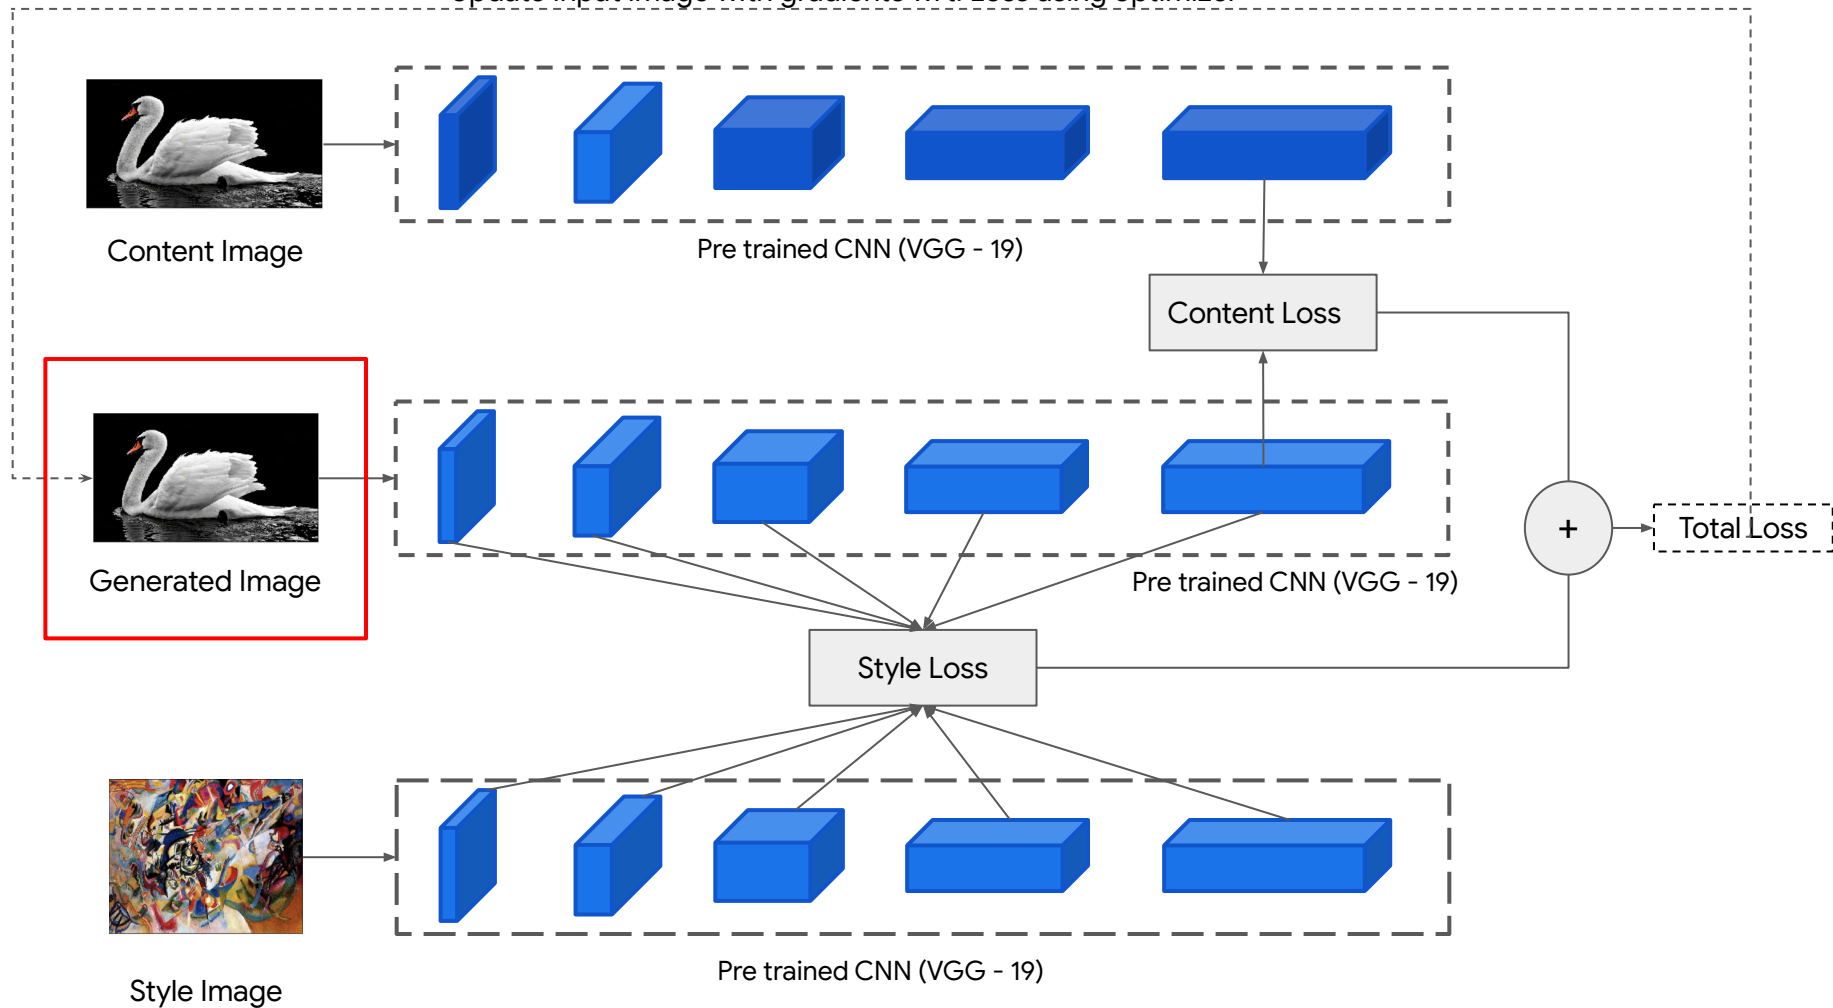


Pre trained CNN (VGG - 19)

+

Total Loss


Update Input image with gradients wrt. Loss using optimizer




Update Input image with gradients wrt. Loss using optimizer




Content Image



Generated Image



Style Image



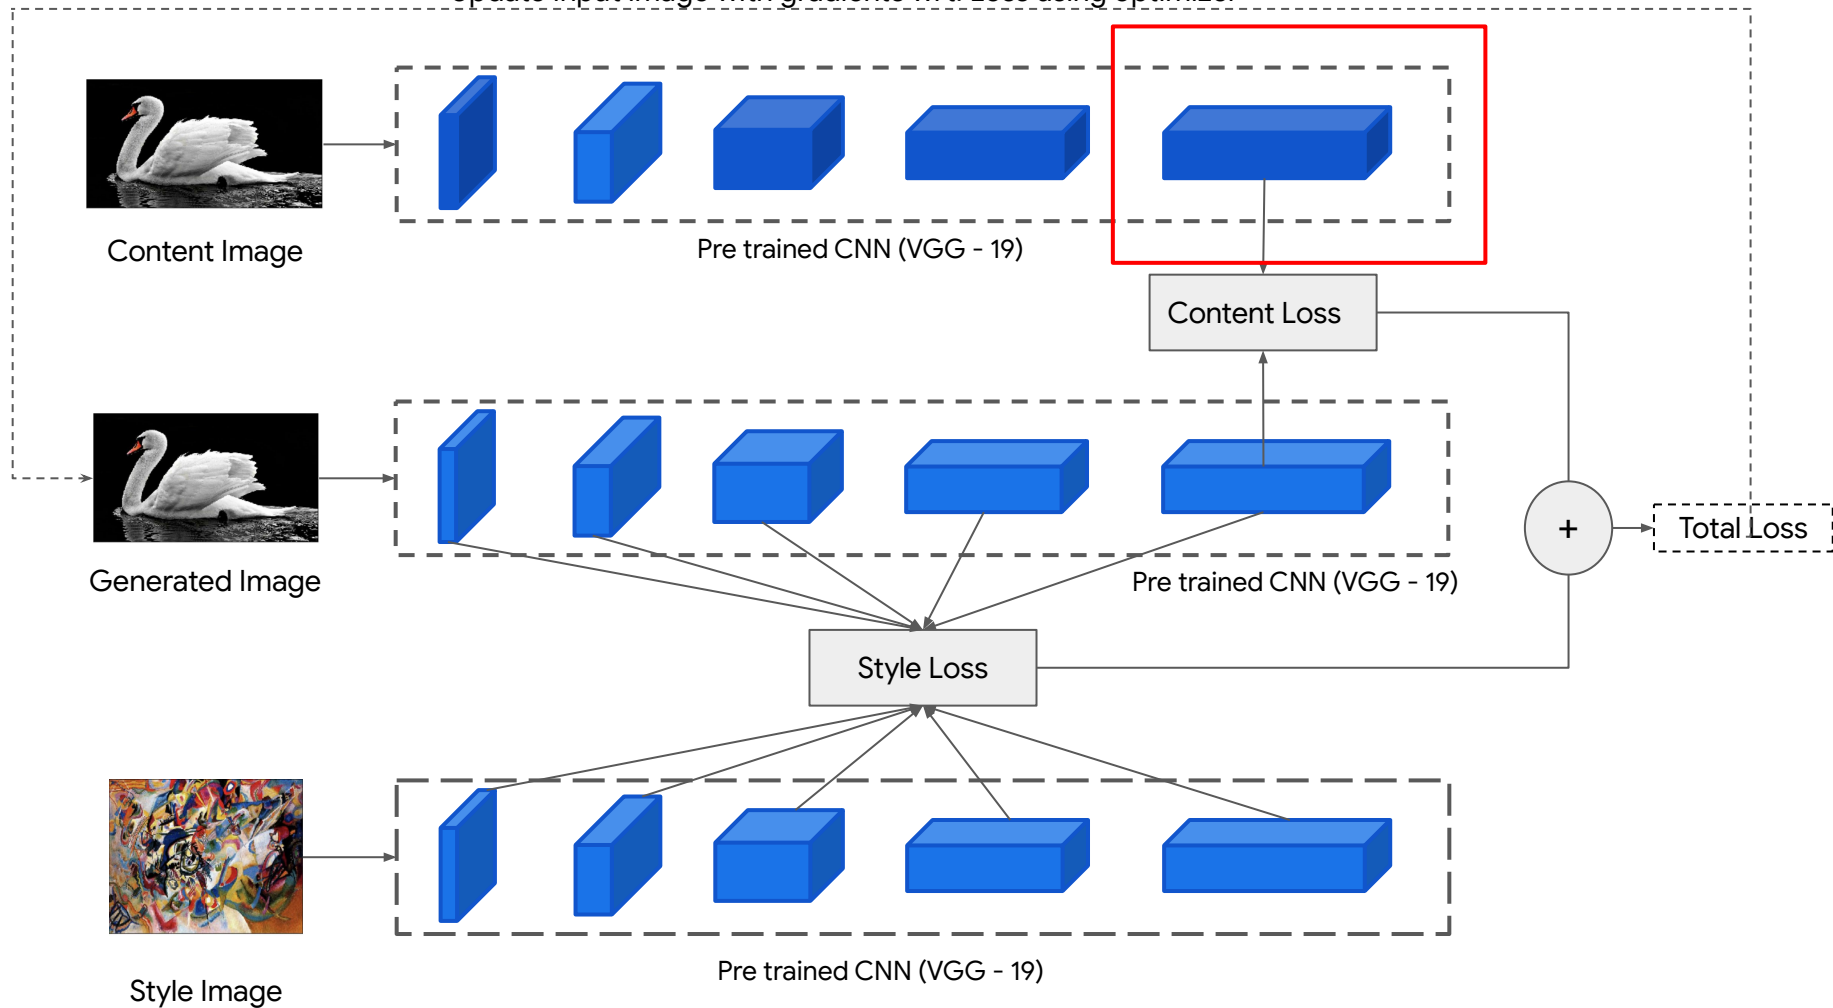
Style Loss

Content Loss

+

Total Loss


Update Input image with gradients wrt. Loss using optimizer



Update Input image with gradients wrt. Loss using optimizer



Content Image




Pre trained CNN (VGG - 19)

Content Loss




Generated Image




Pre trained CNN (VGG - 19)

Style Loss



Style Image

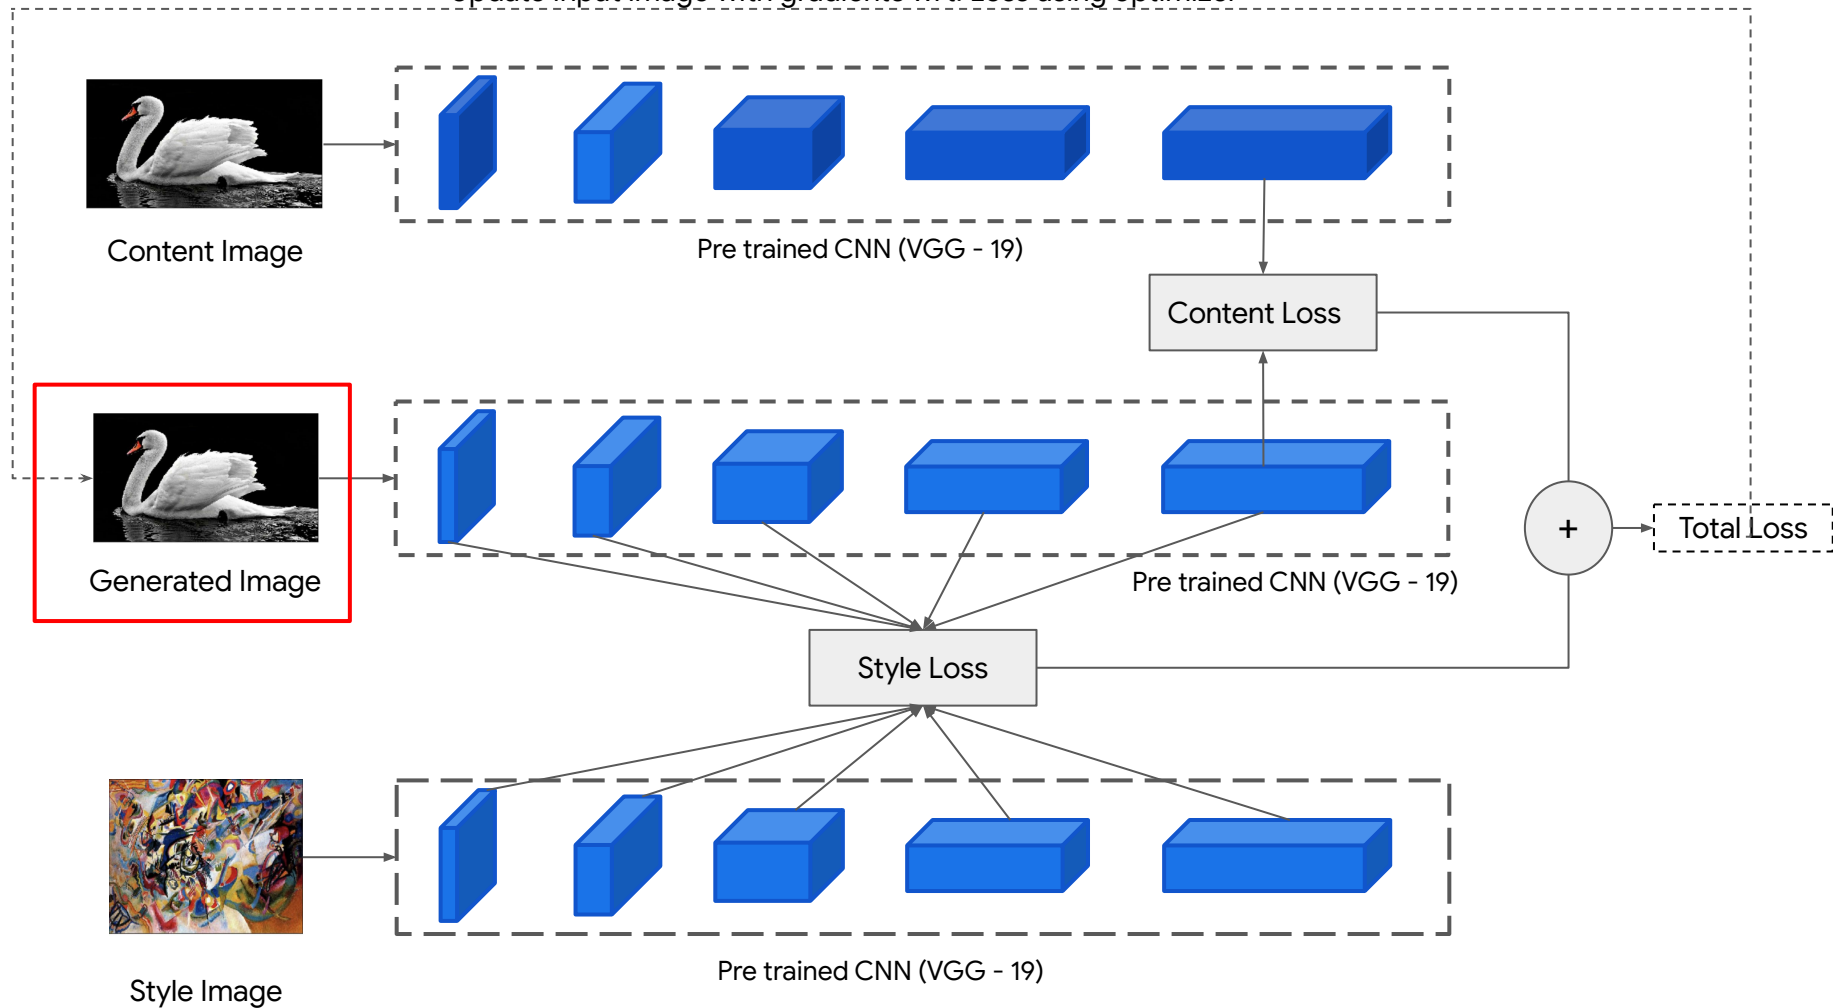


Pre trained CNN (VGG - 19)

+

Total Loss


Update Input image with gradients wrt. Loss using optimizer



Update Input image with gradients wrt. Loss using optimizer




Content Image



Pre trained CNN (VGG - 19)




Generated Image




Pre trained CNN (VGG - 19)

Content Loss

Style Loss




Style Image



Pre trained CNN (VGG - 19)




Total Loss



Update Input image with gradients wrt. Loss using optimizer



Content Image




Pre trained CNN (VGG - 19)

Content Loss



Generated Image




Pre trained CNN (VGG - 19)

Style Loss



Style Image



Pre trained CNN (VGG - 19)

+


Total Loss



Update Input image with gradients wrt. Loss using optimizer



Content Image




Pre trained CNN (VGG - 19)

Content Loss




Generated Image




Pre trained CNN (VGG - 19)

Style Loss



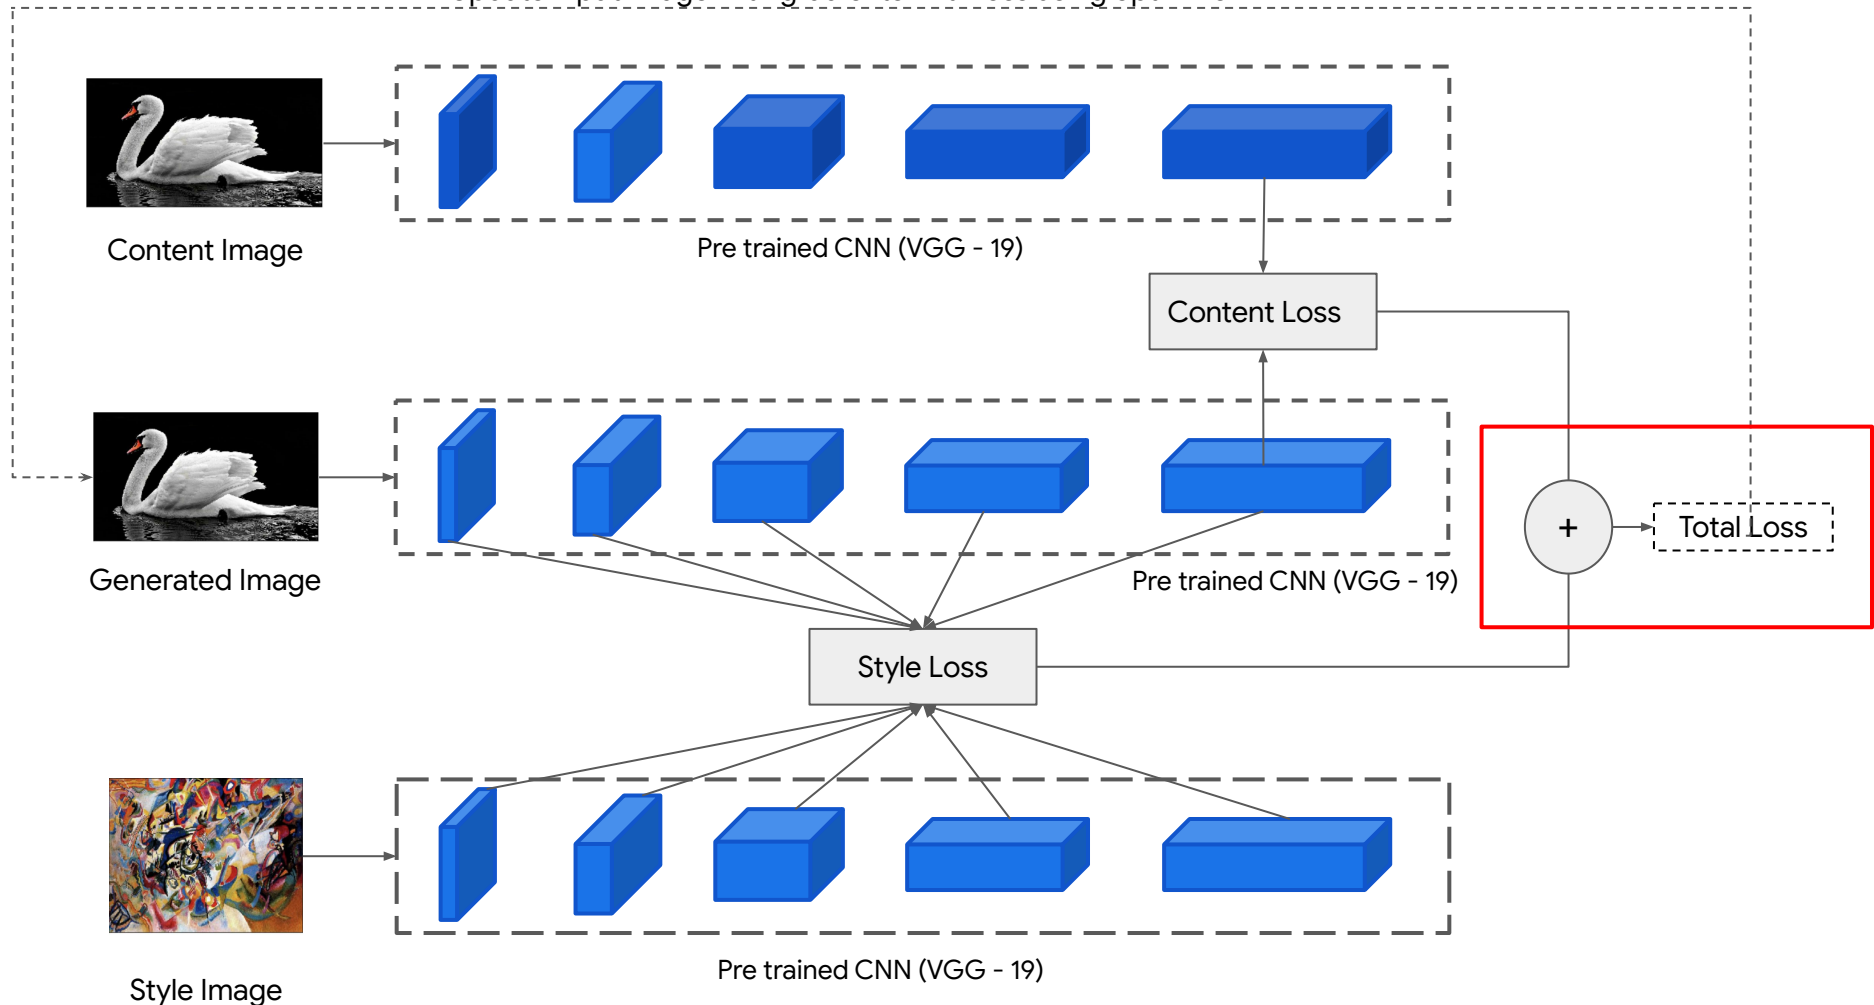
Style Image



Pre trained CNN (VGG - 19)

+


Total Loss



Update Input image with gradients wrt. Loss using optimizer




Content Image



Pre trained CNN (VGG - 19)



Generated Image




Pre trained CNN (VGG - 19)

Style Loss



Style Image



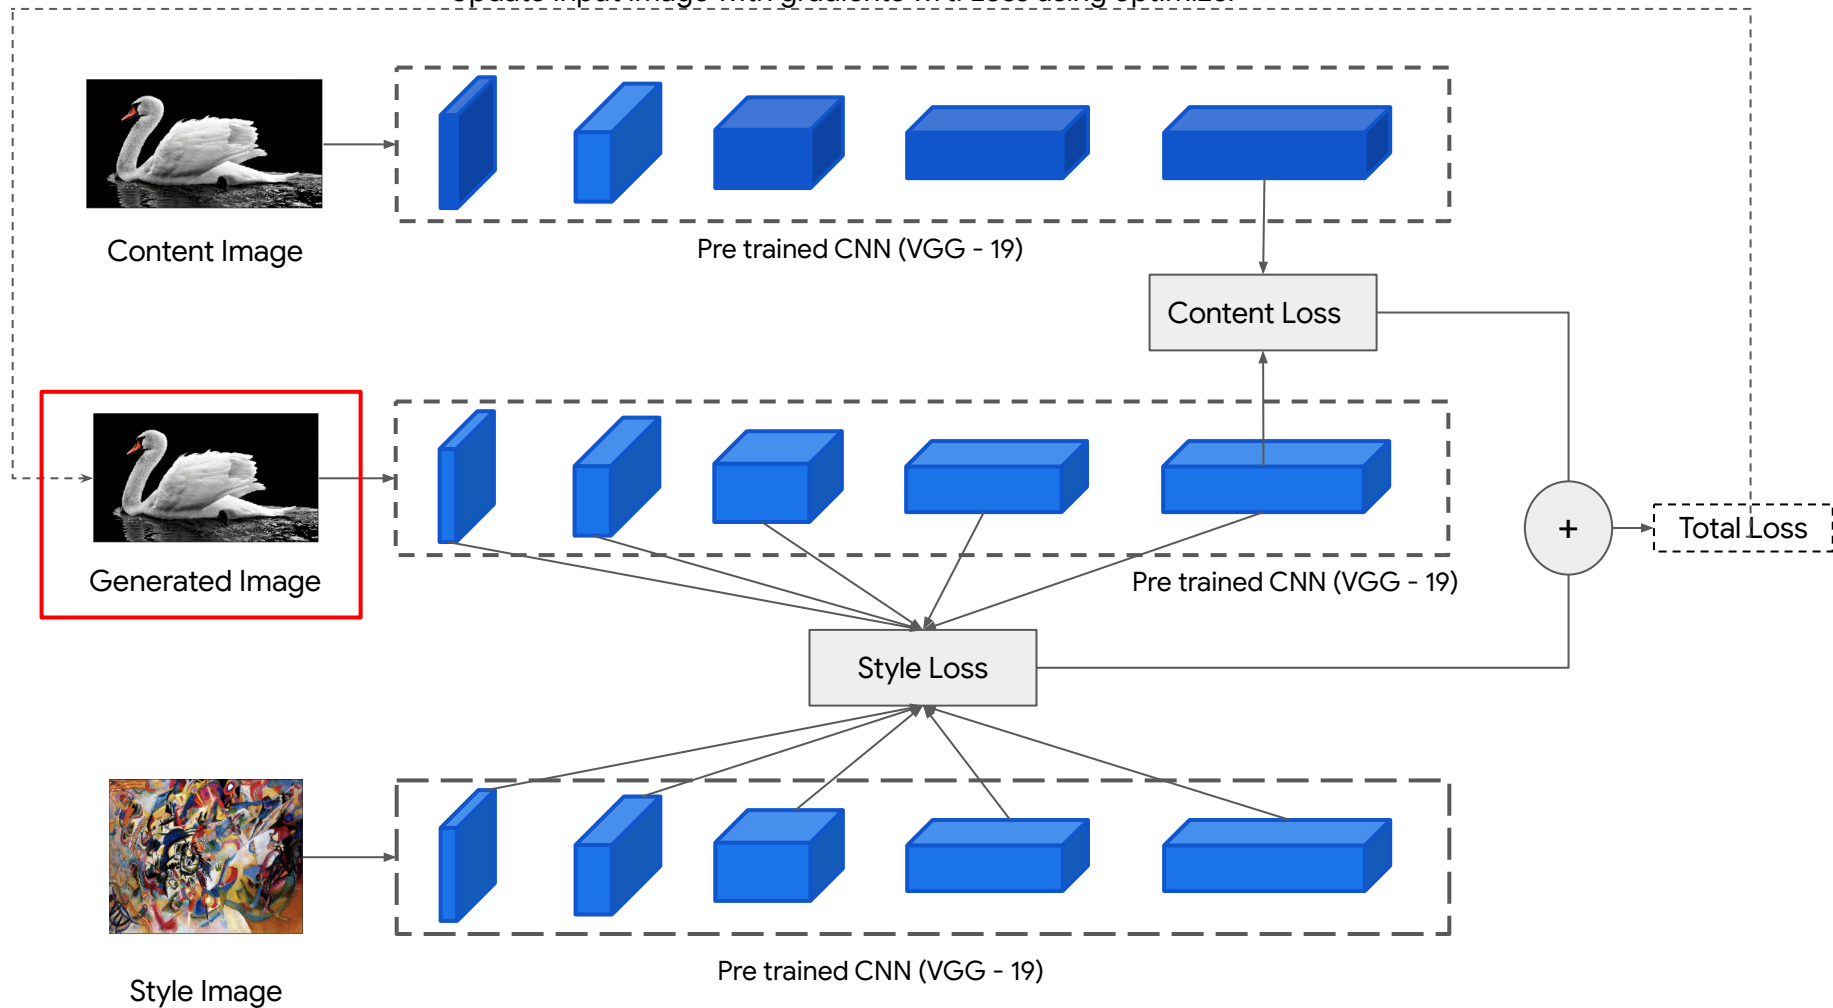
Pre trained CNN (VGG - 19)

Content Loss

+

Total Loss

Update Input image with gradients wrt. Loss using optimizer



## 2. Neural Style Transfer

### Advantages

- Requires only single pair of content image and style image.

### Drawbacks

- Requires many passes for stylizing each image.
- Takes time in stylizing every time an input is given.

# Perceptual Losses for Real-Time Style Transfer and Super Resolution

Justin Johnson / Alexandre Alahi / Fei Fei Li

<https://cs.stanford.edu/people/jcjohns/eccv16/>

Perceptual Losses for Real-Time Style Transfer and Super-Resolution

11




**Fig. 7.** Example results for style transfer on  $512 \times 512$  images by applying models trained on  $256 \times 256$  images. The style images are the same as Figure 6.

Update Input image with gradients wrt. Loss using optimizer



Content Image




Pre trained CNN (VGG - 19)

Content Loss




Generated Image




Pre trained CNN (VGG - 19)

Style Loss



Style Image

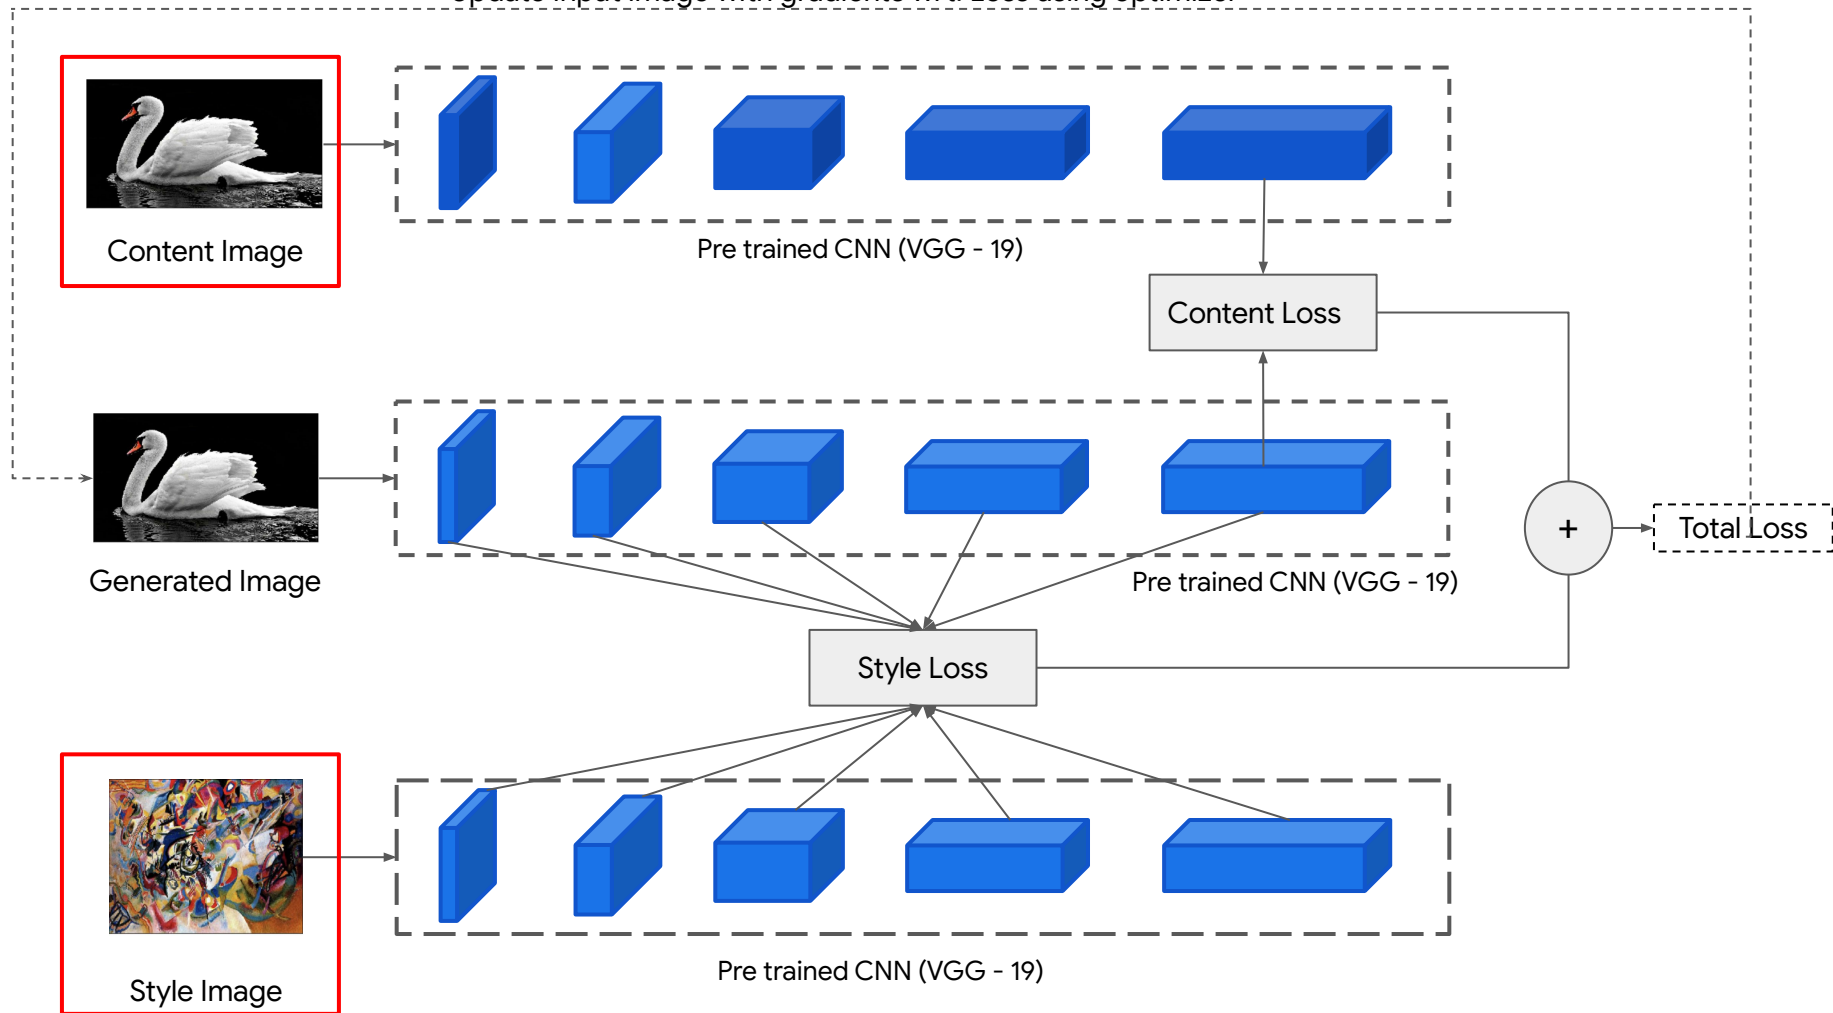


Pre trained CNN (VGG - 19)

+

Total Loss


Update Input image with gradients wrt. Loss using optimizer



Update Input image with gradients wrt. Loss using optimizer




Content Image



Pre trained CNN (VGG - 19)



Generated Image




Pre trained CNN (VGG - 19)

Style Loss


Content Loss



Total Loss



Style Image




Pre trained CNN (VGG - 19)



Update Input image with gradients wrt. Loss using optimizer




Content Image



Pre trained CNN (VGG - 19)




Generated Image





Pre trained CNN (VGG - 19)



Style Image



Pre trained CNN (VGG - 19)




Total Loss

+

Update Input image with gradients wrt. Loss using optimizer



Content Image




Pre trained CNN (VGG - 19)

Content Loss



Generated Image




Pre trained CNN (VGG - 19)


Style Loss



Style Image



Pre trained CNN (VGG - 19)




Total Loss

Update Input image with gradients wrt. Loss using optimizer




Content Image



Pre trained CNN (VGG - 19)




Generated Image




Pre trained CNN (VGG - 19)

Content Loss

Style Loss



Style Image



Pre trained CNN (VGG - 19)

+

Total Loss

# Steps to Develop Neural Style Transfer

1. Preprocess, content & style images
2. Load pre trained model, define loss functions
3. Loop:
  - a. Optimize and generate new image
  - b. Visualize Outputs

# Steps to Develop Neural Style Transfer

1. Preprocess, content & style images
2. Load pre trained model, define loss functions
3. Loop:
  - a. Optimize and generate new image
  - b. Visualize Outputs

# Steps to Develop Neural Style Transfer

1. Preprocess, content & style images
2. Load pre trained model, define loss functions
3. Loop:
  - a. Optimize and generate new image
  - b. Visualize Outputs

# Steps to Develop Neural Style Transfer

1. Preprocess, content & style images
2. Load pre trained model, define loss functions
3. Loop:
  - a. Optimize and generate new image
  - b. Visualize Outputs

# Steps to Develop Neural Style Transfer

1. Preprocess, content & style images
2. Load pre trained model, define loss functions
3. Loop:
  - a. Optimize and generate new image
  - b. Visualize Outputs



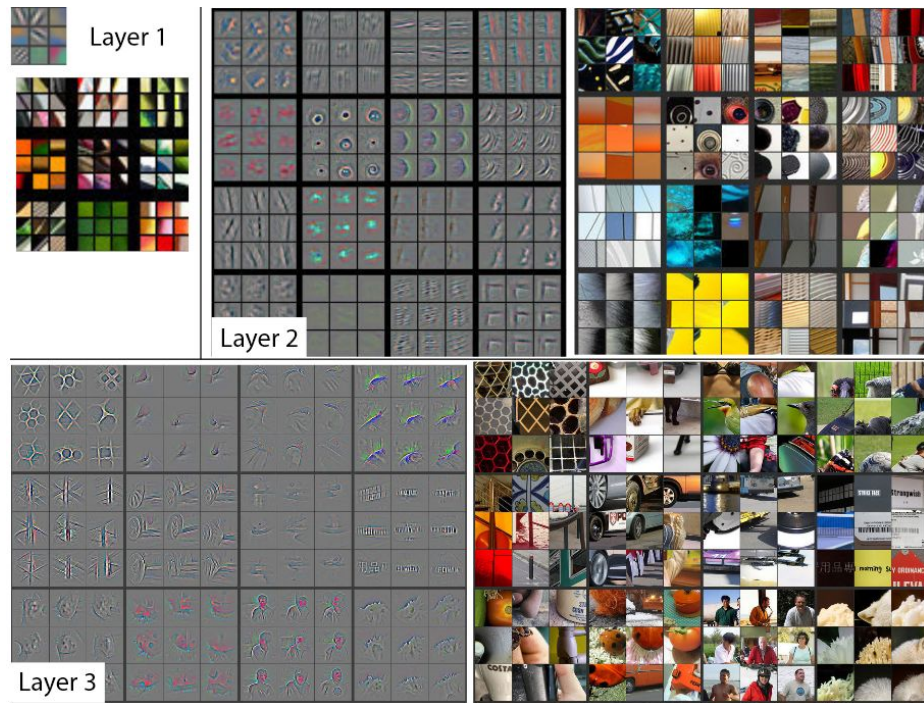
# Load and Preprocess Images

- VGG-19 pre trained on 'imagenet'
- Accepts pixels values [0...255]
  - Don't normalize to [0 ... 1]
- Expects centered pixel values

# Keras Preprocess\_input

```
def preprocess_image(image):  
    ...  
    image = tf.keras.applications.vgg19.preprocess_input(image)  
  
    return image
```

# How CNN visualizes images?



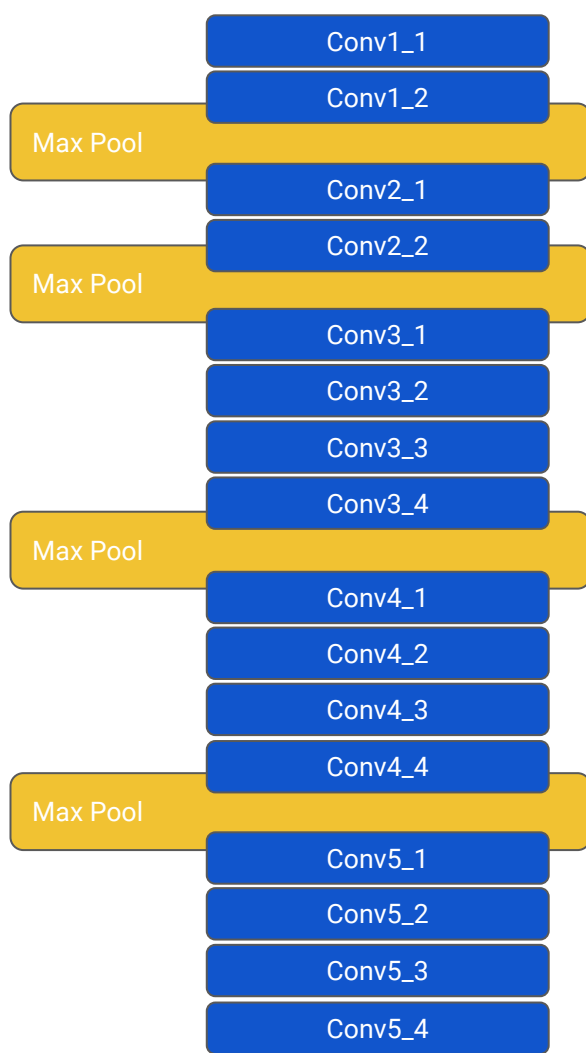
<https://arxiv.org/pdf/1311.2901.pdf>



# VGG-19 for style extraction

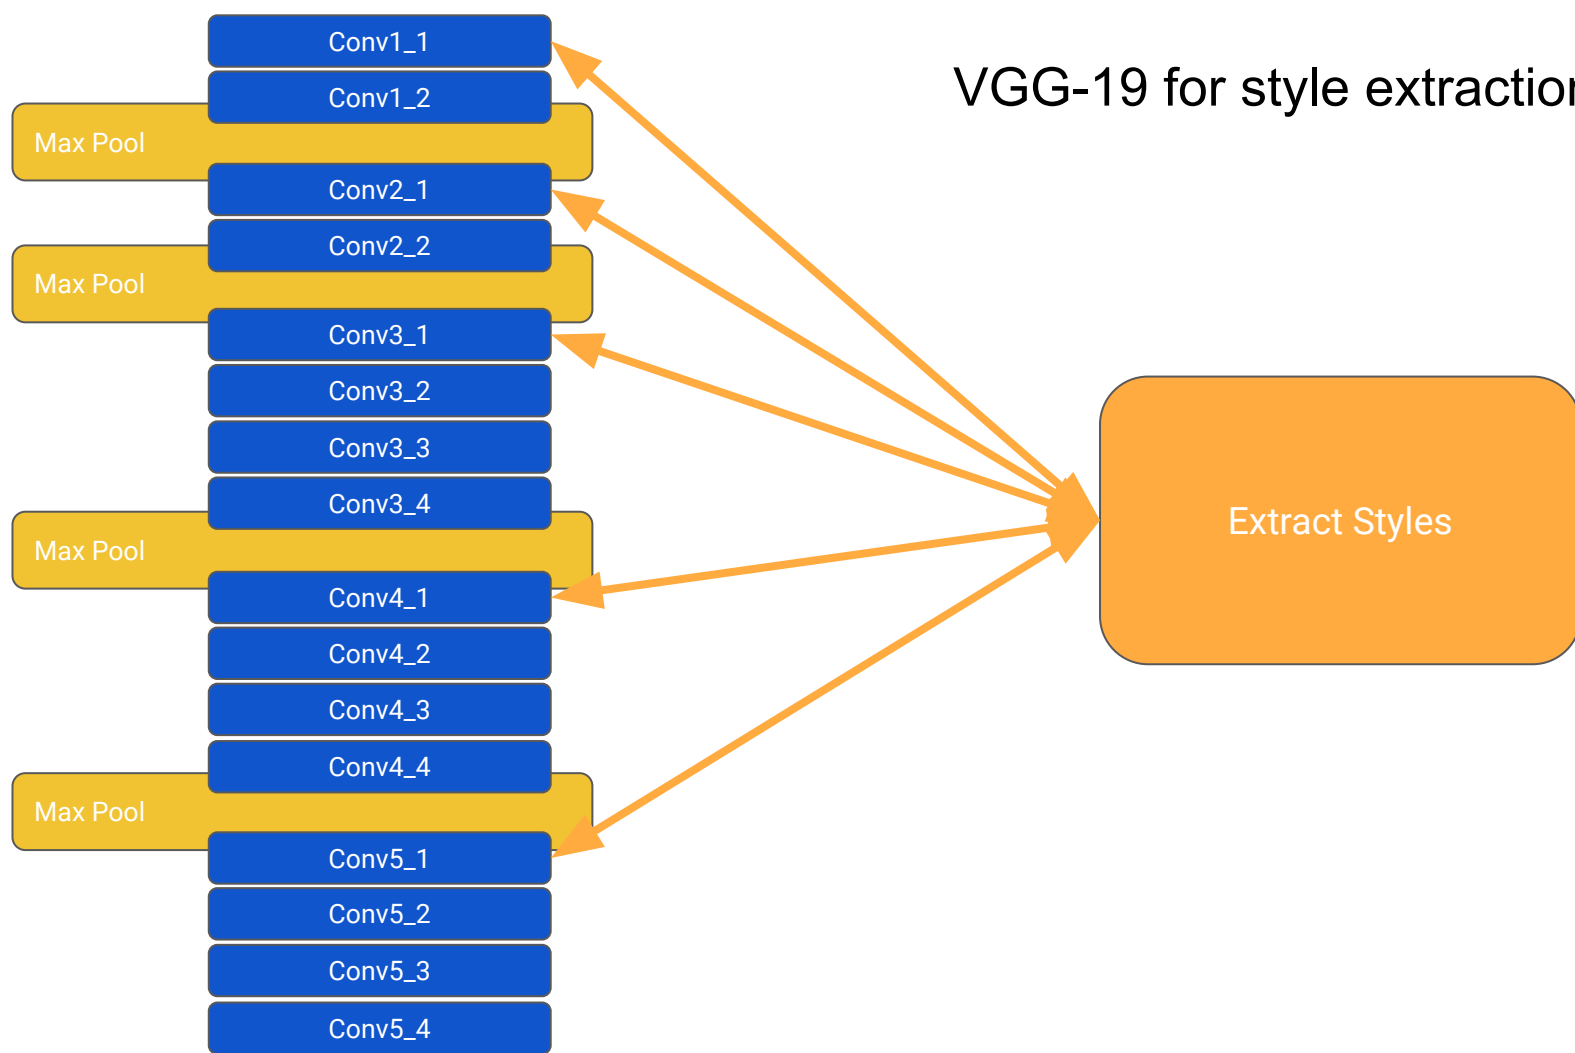


# VGG-19 for style extraction

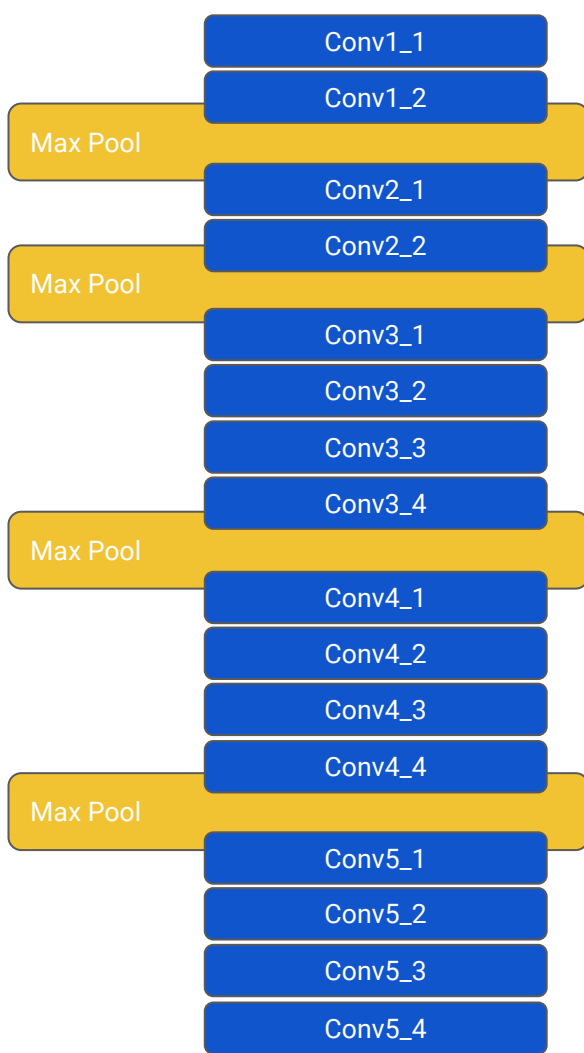


Extract Styles

# VGG-19 for style extraction



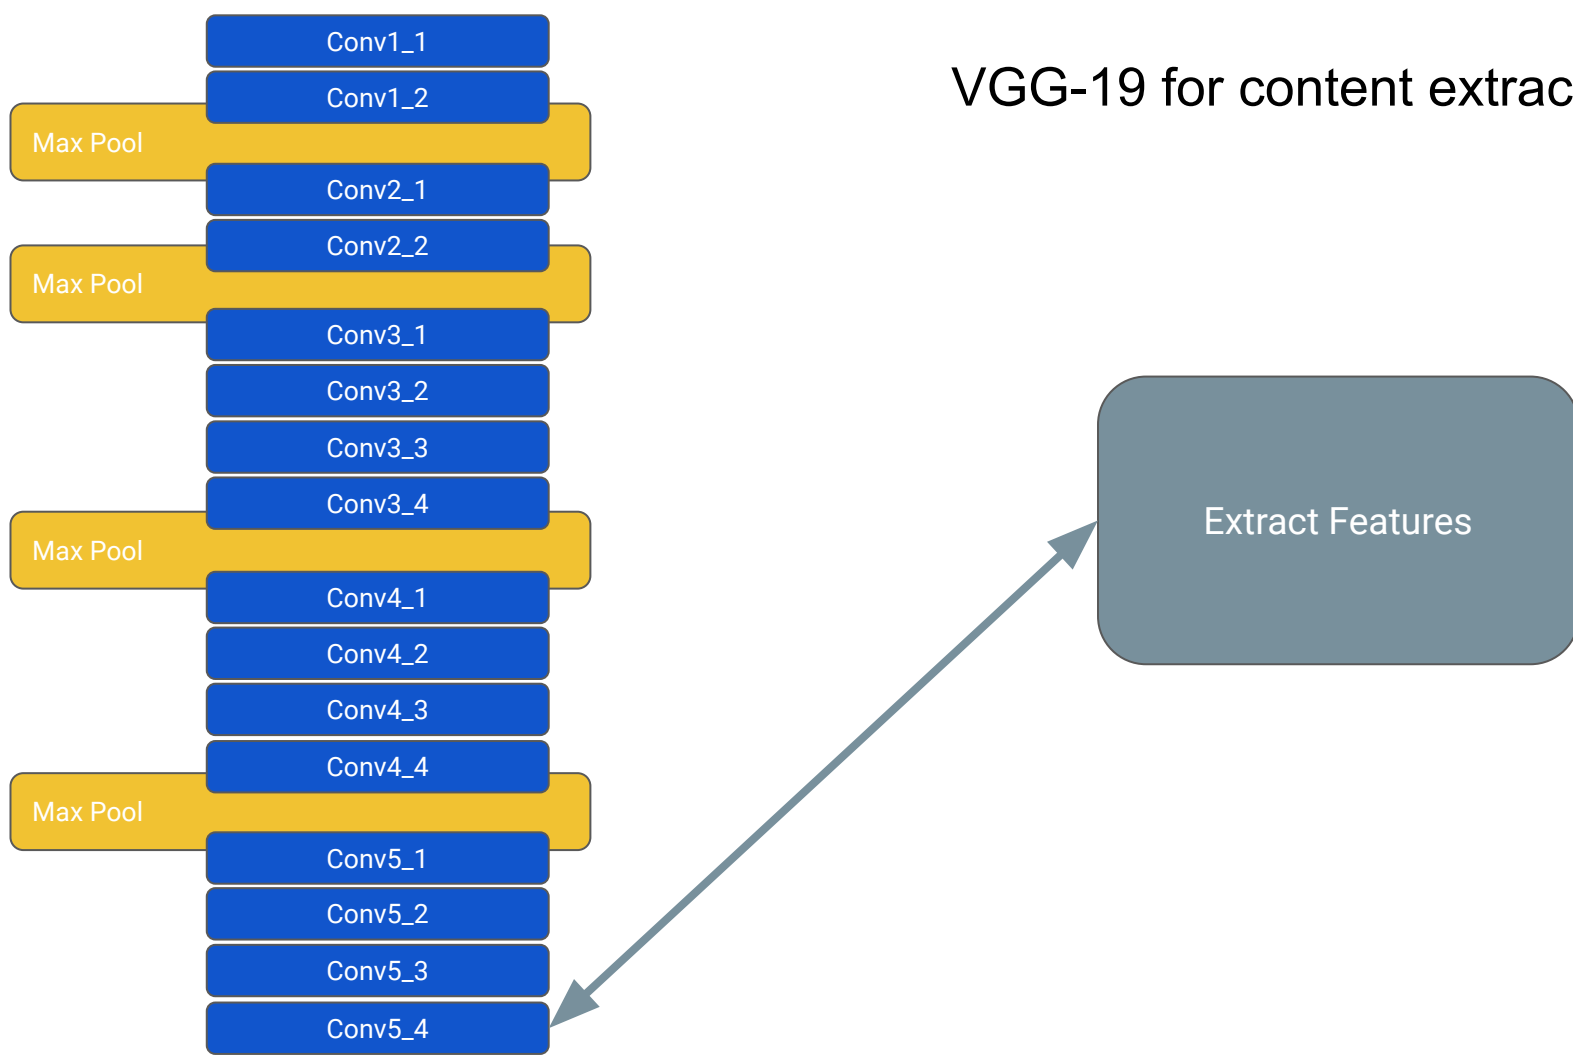
# VGG-19 for content extraction

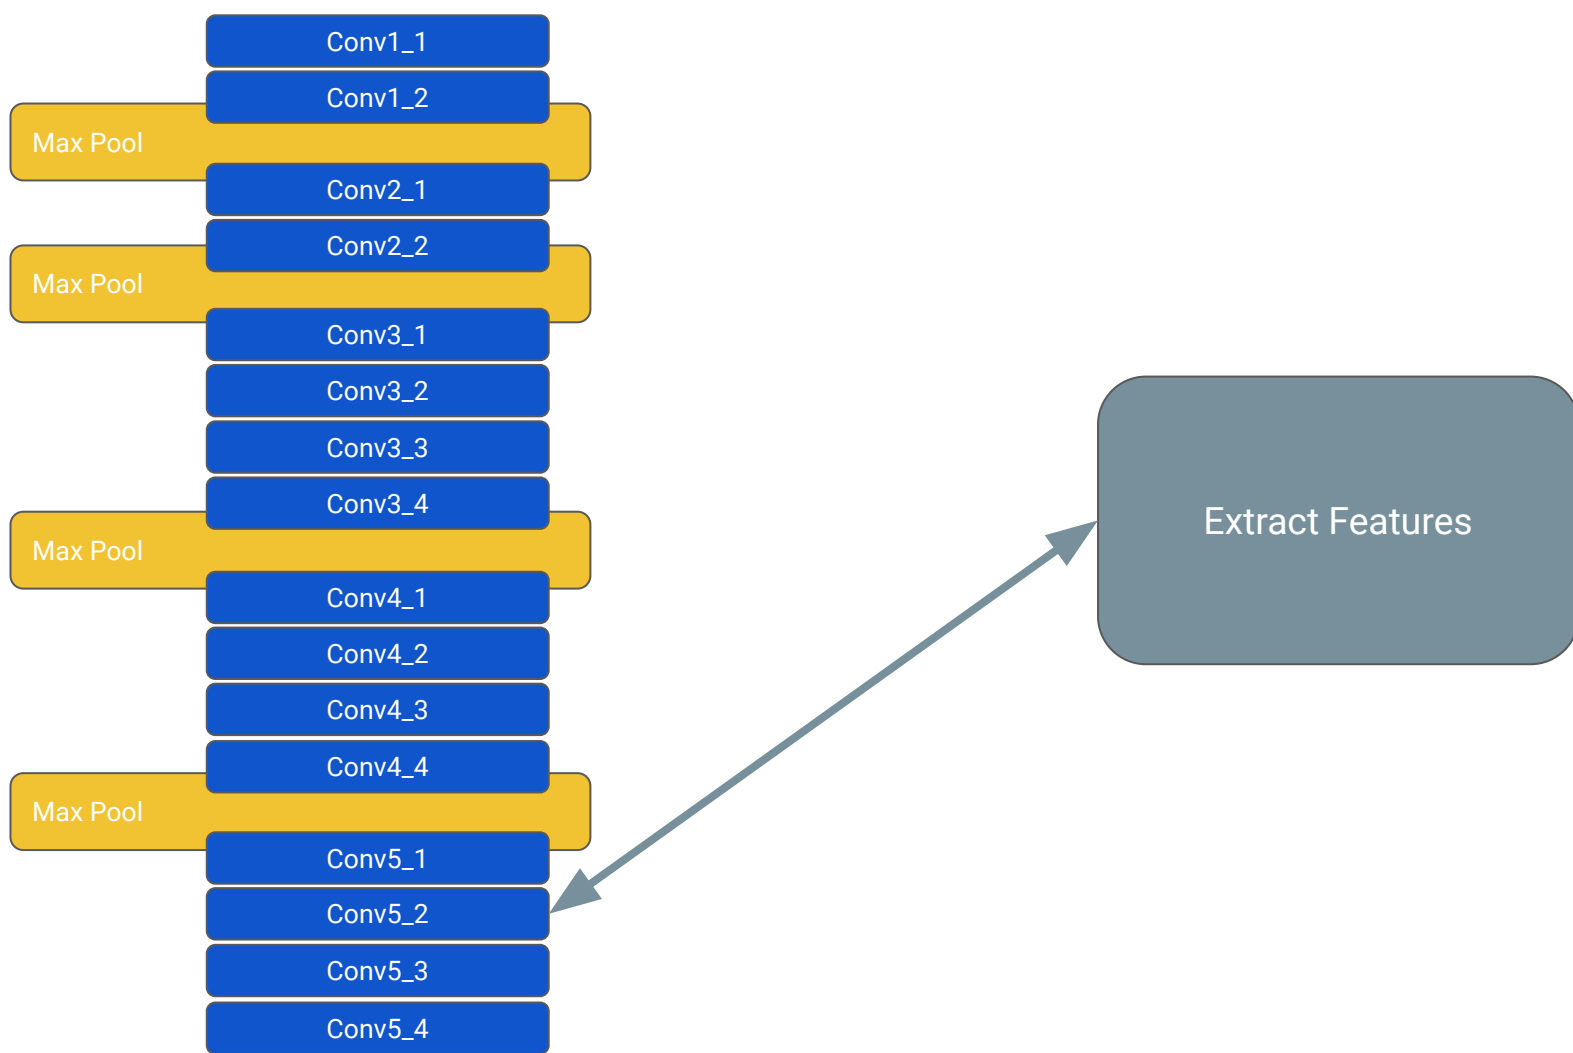


Extract Features



# VGG-19 for content extraction





```
# Content layer where will pull our feature maps
```

```
content_layers = ['block5_conv2']
```

```
# Style layer of interest
```

```
style_layers = ['block1_conv1',  
               'block2_conv1',  
               'block3_conv1',  
               'block4_conv1',  
               'Block5_conv1']
```

```
layer_names = content_layers + style_layers
```

```
# Content layer where will pull our feature maps
```

```
content_layers = ['block5_conv2']
```

```
# Style layer of interest
```

```
style_layers = ['block1_conv1',  
               'block2_conv1',  
               'block3_conv1',  
               'block4_conv1',  
               'Block5_conv1']
```

```
layer_names = content_layers + style_layers
```

```
# Content layer where will pull our feature maps
```

```
content_layers = ['block5_conv2']
```

```
# Style layer of interest
```

```
style_layers = ['block1_conv1',  
               'block2_conv1',  
               'block3_conv1',  
               'block4_conv1',  
               'Block5_conv1']
```

```
layer_names = content_layers + style_layers
```

```
# Content layer where we will pull our feature maps
```

```
content_layers = ['block5_conv2']
```

```
# Style layer of interest
```

```
style_layers = ['block1_conv1',
```

```
                'block2_conv1',
```

```
                'block3_conv1',
```

```
                'block4_conv1',
```

```
                'block5_conv1']
```

```
layer_names = content_layers + style_layers
```

```
def vgg_model(layer_names):  
    """ Creates a vgg model that returns a list of intermediate output values."""  
    vgg = tf.keras.applications.vgg19.VGG19(include_top=False, weights='imagenet')  
    vgg.trainable = False  
  
    outputs = [vgg.get_layer(name).output for name in layer_names]  
  
    model = tf.keras.Model(inputs=vgg.input, outputs=outputs)  
  
    return model
```

```
def vgg_model(layer_names):  
    """ Creates a vgg model that returns a list of intermediate output values."""  
    vgg = tf.keras.applications.vgg19.VGG19(include_top=False, weights='imagenet')  
    vgg.trainable = False  
  
    outputs = [vgg.get_layer(name).output for name in layer_names]  
  
    model = tf.keras.Model(inputs=vgg.input, outputs=outputs)  
  
    return model
```



```
def vgg_model(layer_names):  
    """ Creates a vgg model that returns a list of intermediate output values."""  
    vgg = tf.keras.applications.vgg19.VGG19(include_top=False, weights='imagenet')  
    vgg.trainable = False  
  
    outputs = [vgg.get_layer(name).output for name in layer_names]  
  
    model = tf.keras.Model(inputs=vgg.input, outputs=outputs)  
  
    return model
```

```
def vgg_model(layer_names):  
    """ Creates a vgg model that returns a list of intermediate output values."""  
    vgg = tf.keras.applications.vgg19.VGG19(include_top=False, weights='imagenet')  
    vgg.trainable = False  
  
    outputs = [vgg.get_layer(name).output for name in layer_names]  
  
    model = tf.keras.Model(inputs=vgg.input, outputs=outputs)  
  
    return model
```

```
def vgg_model(layer_names):  
    """ Creates a vgg model that returns a list of intermediate output values."""  
    vgg = tf.keras.applications.vgg19.VGG19(include_top=False, weights='imagenet')  
    vgg.trainable = False  
  
    outputs = [vgg.get_layer(name).output for name in layer_names]  
  
    model = tf.keras.Model(inputs=vgg.input, outputs=outputs)  
  
    return model
```


```
def vgg_model(layer_names):  
    """ Creates a vgg model that returns a list of intermediate output values."""  
    vgg = tf.keras.applications.vgg19.VGG19(include_top=False, weights='imagenet')  
    vgg.trainable = False  
  
    outputs = [vgg.get_layer(name).output for name in layer_names]  
  
    model = tf.keras.Model(inputs=vgg.input, outputs=outputs)  
  
    return model
```

```
def vgg_model(layer_names):  
    """ Creates a vgg model that returns a list of intermediate output values."""  
    vgg = tf.keras.applications.vgg19.VGG19(include_top=False, weights='imagenet')  
    vgg.trainable = False  
  
    outputs = [vgg.get_layer(name).output for name in layer_names]  
  
    model = tf.keras.Model(inputs=vgg.input, outputs=outputs)  
  
    return model
```

Update Input image with gradients wrt. Loss using optimizer



Content Image




Pre trained CNN (VGG - 19)

Content Loss




Generated Image




Pre trained CNN (VGG - 19)

Style Loss



Style Image



Pre trained CNN (VGG - 19)


+

Total Loss

Update Input image with gradients wrt. Loss using optimizer




Content Image



Pre trained CNN (VGG - 19)



Generated Image




Pre trained CNN (VGG - 19)


Content Loss

Content Loss

Style Loss



Style Image



Pre trained CNN (VGG - 19)




Total Loss

Update Input image with gradients wrt. Loss using optimizer



Content Image




Pre trained CNN (VGG - 19)

Content Loss




Generated Image




Pre trained CNN (VGG - 19)

Style Loss



Style Image



Pre trained CNN (VGG - 19)

+


Total Loss



Update Input image with gradients wrt. Loss using optimizer



Content Image




Pre trained CNN (VGG - 19)

Content Loss




Generated Image




Pre trained CNN (VGG - 19)


Style Loss



Style Image



Pre trained CNN (VGG - 19)



# Total Loss

$$L_{total} = L_{content} + L_{style}$$

# Total Loss

$$L_{total} = \alpha L_{content} + \beta L_{style}$$

**$\alpha$**  : Content Weight

**$\beta$**  : Style Weight

# Total Loss

$$L_{total}(\overrightarrow{p}, \overrightarrow{a}, \overrightarrow{x}) = \alpha L_{content}(\overrightarrow{p}, \overrightarrow{x}) + \beta L_{style}(\overrightarrow{a}, \overrightarrow{x})$$

**$\alpha$**  : Content Weight

**$\beta$**  : Style Weight

# Total Loss

$$L_{total}(\overrightarrow{p}, \overrightarrow{a}, \overrightarrow{x}) = \alpha L_{content}(\overrightarrow{p}, \overrightarrow{x}) + \beta L_{style}(\overrightarrow{a}, \overrightarrow{x})$$

$\overrightarrow{p}$  : Content Image (Original Photograph)

$\overrightarrow{x}$  : Generated Image initialized to Input Image

$\alpha$  : Content Weight

$\beta$  : Style Weight

# Total Loss

$$L_{total}(\vec{p}, \vec{a}, \vec{x}) = \alpha L_{content}(\vec{p}, \vec{x}) + \beta L_{style}(\vec{a}, \vec{x})$$

$\vec{a}$  : Style Image

$\vec{x}$  : Generated Image initialized to Input Image

$\alpha$  : Content Weight

$\beta$  : Style Weight

# Total Loss

$$L_{total}(\vec{p}, \vec{a}, \vec{x}) = \alpha L_{content}(\vec{p}, \vec{x}) + \beta L_{style}(\vec{a}, \vec{x})$$

$\vec{p}$  : Content Image (Original Photograph)

$\vec{a}$  : Style Image

$\vec{x}$  : Generated Image initialized to Input Image

$\alpha$  : Content Weight

$\beta$  : Style Weight

# Content Loss



# Content Loss

Generated image

1	2
3	4

—

Content image

2	2
2	2

# Content Loss

Generated image

1	2
3	4

−

Content image

2	2
2	2

=

1-2	2-2
3-2	4-2

Element-wise  
subtraction

# Content Loss

Generated image

1	2
3	4

−

Content image

2	2
2	2

=

1-2	2-2
3-2	4-2

Element-wise  
subtraction

$(1-2)^2$	$(2-2)^2$
$(3-2)^2$	$(4-2)^2$

Element-wise square

# Content Loss

Generated image

1	2
3	4

−

Content image

2	2
2	2

=

1-2	2-2
3-2	4-2

Element-wise subtraction

$(1-2)^2$	$(2-2)^2$
$(3-2)^2$	$(4-2)^2$

Element-wise square

$$1^2 + 0^2 + 1^2 + 2^2 =$$

6

Reduce sum

# Content Loss

Generated image

1	2
3	4

−

Content image

2	2
2	2

=

1-2	2-2
3-2	4-2

Element-wise subtraction

$(1-2)^2$	$(2-2)^2$
$(3-2)^2$	$(4-2)^2$

Element-wise square

$$1^2 + 0^2 + 1^2 + 2^2 =$$

6

Reduce sum

$$(\frac{1}{2}) 6 =$$

3

weight

# Content Loss

$$L_{content}(\vec{p}, \vec{x}, l) = \frac{1}{2} \sum_{i,j} (F_{i,j}^l - P_{i,j}^l)^2$$

$l$ : layer  $l$

$\mathbf{F}_{ij}^l$ : Content representation of generated image  $x$  in layer  $l$   
(Activation of  $i$ th feature map at position  $j$  in layer  $l$  of  
*generated image*).

$\mathbf{P}_{ij}^l$ : Content representation of content image  $p$  in layer  $l$ .  
(Activation of the  $i$ th feature map at position  $j$  in layer  $l$  of  
content image.)

# Content Loss - In Practice

```
def get_content_loss(features, targets):  
    return 0.5 * tf.reduce_sum(tf.square(features - targets))
```

# Style Loss

Generated image

1	2
3	4

-

Style image

2	2
2	2

=

1-2	2-2
3-2	4-2

Element-wise subtraction

$(1-2)^2$	$(2-2)^2$
$(3-2)^2$	$(4-2)^2$

Element-wise square

$$1^2 + 0^2 + 1^2 + 2^2 =$$

6

Reduce sum

$$(\frac{?}{?}) 6 =$$

?

weight



# Style Loss

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G_{i,j}^l - A_{ij}^l)^2$$

$l$ : layer  $l$

$\mathbf{A}_{ij}^l$ : Style Representation(Gram Matrix) of style image  $a$ .

$\mathbf{G}_{ij}^l$ : Style Representation(Gram Matrix) of generated image  $x$

# Style Loss

$$\boxed{E_l} = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G_{i,j}^l - A_{ij}^l)^2$$

$l$ : layer  $l$

$\mathbf{A}_{ij}^l$ : Style Representation(Gram Matrix) of style image  $a$ .

$\mathbf{G}_{ij}^l$ : Style Representation(Gram Matrix) of generated image  $x$

# Style Loss

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G_{i,j}^l - A_{ij}^l)^2$$

$l$ : layer  $l$

$\mathbf{A}_{ij}^l$ : Style Representation(Gram Matrix) of style image  $a$ .

$\mathbf{G}_{ij}^l$ : Style Representation(Gram Matrix) of generated image  $x$

# Gram Matrix

$$G_{ij}^l = \sum_k F_{ik}^l F_{jk}^l$$

$G_{ii}^l$ : inner product between vectorized feature maps  $i$  and  $j$  in layer  $l$ .

# Gram Matrix - In Practice

```
def gram_matrix(input_tensor):  
    result = tf.linalg.einsum('bijk,bijd->bcd', input_tensorT, input_tensor)  
    input_shape = tf.shape(input_tensor)  
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)  
    return result/(num_locations)
```

<https://numpy.org/doc/stable/reference/generated/numpy.einsum.html>

# Get Style Feature Representation at Layers

```
def get_style_image_features(image):  
    preprocessed_style_image = preprocess_image(image)  
    style_outputs = vgg(preprocessed_style_image)  
    gram_style_features =  
        [gram_matrix(style_layer) for style_layer in style_outputs[:num_style_layers]]  
  
    return gram_style_features
```

# Get Style Feature Representation at Layers

```
def get_style_image_features(image):  
    preprocessed_style_image = preprocess_image(image)  
    style_outputs = vgg(preprocessed_style_image)  
    gram_style_features =  
        [gram_matrix(style_layer) for style_layer in style_outputs[:num_style_layers]]  
  
    return gram_style_features
```

# Get Style Feature Representation at Layers

```
def get_style_image_features(image):  
    preprocessed_style_image = preprocess_image(image)  
    style_outputs = vgg(preprocessed_style_image)  
    gram_style_features =  
        [gram_matrix(style_layer) for style_layer in style_outputs[:num_style_layers]]  
  
    return gram_style_features
```



# Get Style Feature Representation at Layers

```
def get_style_image_features(image):  
    preprocessed_style_image = preprocess_image(image)  
    style_outputs = vgg(preprocessed_style_image)  
    gram_style_features =  
        [gram_matrix(style_layer) for style_layer in style_outputs[:num_style_layers]]  
  
    return gram_style_features
```

# Get Style Feature Representation at Layers

```
def get_style_loss(features, targets):  
    return tf.reduce_mean(tf.square(features - targets))
```

# Total Loss

$$L_{total}(\overrightarrow{p}, \overrightarrow{a}, \overrightarrow{x}) = \alpha L_{content}(\overrightarrow{p}, \overrightarrow{x}) + \beta L_{style}(\overrightarrow{a}, \overrightarrow{x})$$

```
def get_style_content_loss(style_targets, style_outputs,
                           content_targets, content_outputs,
                           style_weight, content_weight):

    style_loss = tf.add_n([get_style_loss(style_output, style_target)
                           for style_output, style_target in zip(style_outputs, style_targets)])

    style_loss *= style_weight / num_style_layers

    content_loss = tf.add_n([get_content_loss(content_output, content_target)
                             for content_output, content_target in zip(content_outputs, content_targets)])

    content_loss *= content_weight / num_content_layers

    loss = style_loss + content_loss
    return loss
```

```
def get_style_content_loss(style_targets, style_outputs,
                           content_targets, content_outputs,
                           style_weight, content_weight):

    style_loss = tf.add_n([get_style_loss(style_output, style_target)
                           for style_output, style_target in zip(style_outputs, style_targets)])

    style_loss *= style_weight / num_style_layers

    content_loss = tf.add_n([get_content_loss(content_output, content_target)
                             for content_output, content_target in zip(content_outputs, content_targets)])
    content_loss *= content_weight / num_content_layers

    loss = style_loss + content_loss
    return loss
```

```
def get_style_content_loss(style_targets, style_outputs,
                           content_targets, content_outputs,
                           style_weight, content_weight):

    style_loss = tf.add_n([get_style_loss(style_output, style_target)
                           for style_output, style_target in zip(style_outputs, style_targets)])

    style_loss *= style_weight / num_style_layers

    content_loss = tf.add_n([get_content_loss(content_output, content_target)
                             for content_output, content_target in zip(content_outputs, content_targets)])
    content_loss *= content_weight / num_content_layers

    loss = style_loss + content_loss
    return loss
```

```
def get_style_content_loss(style_targets, style_outputs,
                           content_targets, content_outputs,
                           style_weight, content_weight):

    style_loss = tf.add_n([get_style_loss(style_output, style_target)
                           for style_output, style_target in zip(style_outputs, style_targets)])

    style_loss *= style_weight / num_style_layers

    content_loss = tf.add_n([get_content_loss(content_output, content_target)
                             for content_output, content_target in zip(content_outputs, content_targets)])

    content_loss *= content_weight / num_content_layers

    loss = style_loss + content_loss
    return loss
```

```
def get_style_content_loss(style_targets, style_outputs,
                           content_targets, content_outputs,
                           style_weight, content_weight):

    style_loss = tf.add_n([get_style_loss(style_output, style_target)
                           for style_output, style_target in zip(style_outputs, style_targets)])

    style_loss *= style_weight / num_style_layers

    content_loss = tf.add_n([get_content_loss(content_output, content_target)
                             for content_output, content_target in zip(content_outputs, content_targets)])
    content_loss *= content_weight / num_content_layers

    loss = style_loss + content_loss
    return loss
```



```
def get_style_content_loss(style_targets, style_outputs,
                           content_targets, content_outputs,
                           style_weight, content_weight):

    style_loss = tf.add_n([get_style_loss(style_output, style_target)
                           for style_output, style_target in zip(style_outputs, style_targets)])

    style_loss *= style_weight / num_style_layers

    content_loss = tf.add_n([get_content_loss(content_output, content_target)
                              for content_output, content_target in zip(content_outputs, content_targets)])

    content_loss *= content_weight / num_content_layers

    loss = style_loss + content_loss
    return loss
```

```
def calculate_gradients(image, content_targets,
                        style_targets, style_weight,
                        content_weight, with_regularization=False):

    with tf.GradientTape() as tape:
        style_features = get_style_image_features(image)
        content_features = get_content_image_features(image)
        loss = get_style_content_loss(style_targets, style_features,
                                      content_targets, content_features,
                                      style_weight, content_weight)

    gradients = tape.gradient(loss, image)

    return gradients
```

```
def calculate_gradients(image, content_targets,  
                        style_targets, style_weight,  
                        content_weight, with_regularization=False):  
  
    with tf.GradientTape() as tape:  
        style_features = get_style_image_features(image)  
        content_features = get_content_image_features(image)  
        loss = get_style_content_loss(style_targets, style_features,  
                                      content_targets, content_features,  
                                      style_weight, content_weight)  
  
    gradients = tape.gradient(loss, image)  
  
    return gradients
```

```
def calculate_gradients(image, content_targets,
                        style_targets, style_weight,
                        content_weight, with_regularization=False):

    with tf.GradientTape() as tape:
        style_features = get_style_image_features(image)
        content_features = get_content_image_features(image)
        loss = get_style_content_loss(style_targets, style_features,
                                      content_targets, content_features,
                                      style_weight, content_weight)

        gradients = tape.gradient(loss, image)

    return gradients
```

```
def update_image_with_style(image, content_targets, style_targets,  
                             optimizer, style_weight, content_weight,  
                             with_regularization=False):  
  
    gradients = calculate_gradients(image, content_targets,  
                                    style_targets, style_weight,  
                                    content_weight, with_regularization)  
  
    optimizer.apply_gradients([(gradients, image)]) #Apply gradients on image
```

```
def update_image_with_style(image, content_targets, style_targets,  
                             optimizer, style_weight, content_weight,  
                             with_regularization=False):
```

```
    gradients = calculate_gradients(image, content_targets,  
                                    style_targets, style_weight,  
                                    content_weight, with_regularization)
```

```
    optimizer.apply_gradients([(gradients, image)]) #Apply gradients on image
```

```
def update_image_with_style(image, content_targets, style_targets,  
                             optimizer, style_weight, content_weight,  
                             with_regularization=False):  
  
    gradients = calculate_gradients(image, content_targets,  
                                    style_targets, style_weight,  
                                    content_weight, with_regularization)  
  
    optimizer.apply_gradients([(gradients, image)]) #Apply gradients on image
```

```
def fit_style_transfer(input_image, style_image, optimizer,
                       epochs=1, steps_per_epoch=1,
                       with_regularization=False, style_weight = 0.01):

    for n in range(epochs):
        for m in range(steps_per_epoch):
            update_image_with_style(input_image, content_targets,
                                    style_targets, optimizer,
                                    style_weight, content_weight,
                                    with_regularization=with_regularization)

    return input_image, images
```



# Visualizing Outputs - Final Results

Style Image



[https://upload.wikimedia.org/wikipedia/commons/b/b4/Vassily\\_Kandinsky%2C\\_1913\\_-\\_Composition\\_7.jpg](https://upload.wikimedia.org/wikipedia/commons/b/b4/Vassily_Kandinsky%2C_1913_-_Composition_7.jpg)

Content Image



[https://cdn.pixabay.com/photo/2017/02/28/23/00/swan-2107052\\_1280.jpg](https://cdn.pixabay.com/photo/2017/02/28/23/00/swan-2107052_1280.jpg)

Stylized Image



# Style Loss

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G_{i,j}^l - A_{ij}^l)^2$$

$l$ : layer  $l$

$\mathbf{A}_{ij}^l$ : Style Representation(Gram Matrix) of style image  $a$ .

$\mathbf{G}_{ij}^l$ : Style Representation(Gram Matrix) of generated image  $x$

# Gram Matrix

$$G_{ij}^l = \sum_k F_{ik}^l F_{jk}^l$$

$G_{ij}^l$ : inner product between vectorized feature maps  $i$  and  $j$  in layer  $l$ .

[http://mlwiki.org/index.php/Gram\\_Matrices](http://mlwiki.org/index.php/Gram_Matrices)

# Style Loss (one layer)

Style layer

H

W

F

# Style Loss (one layer)

Style layer

$H = 2$

$W = 2$

$F = 2$

1	2		
3	4	5	6
		7	8

# Style Loss (one layer)

Style layer

$H = 2$

$W = 2$

$F = 2$

1	2	5	6
3	4		
		7	8

1
2
3
4

5
6
7
8


# Style Loss (one layer)

Style layer

$H = 2$

$W = 2$

$F = 2$



$$A = \left[ \begin{array}{c} \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 3 \\ \hline 4 \\ \hline \end{array} \\ a_1 \end{array} \right. \left. \begin{array}{c} \begin{array}{|c|} \hline 5 \\ \hline 6 \\ \hline 7 \\ \hline 8 \\ \hline \end{array} \\ a_2 \end{array} \right]$$


# Style Loss (one layer)

Style layer

$H = 2$

$W = 2$

$F = 2$



$$A = \begin{bmatrix} \begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix} & \begin{bmatrix} 5 \\ 6 \\ 7 \\ 8 \end{bmatrix} \end{bmatrix} G = \begin{bmatrix} \phantom{\begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix}} & \phantom{\begin{bmatrix} 5 \\ 6 \\ 7 \\ 8 \end{bmatrix}} \end{bmatrix}$$

$a_1$        $a_2$




# Style Loss (one layer)

Style layer

$H = 2$

$W = 2$

$F = 2$



$$A = \left[ \begin{array}{c} \boxed{1} \\ \boxed{2} \\ \boxed{3} \\ \boxed{4} \end{array} \quad \begin{array}{c} \boxed{5} \\ \boxed{6} \\ \boxed{7} \\ \boxed{8} \end{array} \right] G = \left[ \begin{array}{c} a_1 \cdot a_1 \end{array} \right]$$

The diagram illustrates the calculation of the style loss for a single layer. The tensor  $A$  is represented as a 2x2x2 volume. The first two dimensions (height and width) are highlighted with blue boxes, and the third dimension (feature) is highlighted with green boxes. The values are arranged as follows:

1	2	5	6
3	4	7	8

The tensor  $A$  is then flattened into a vector  $a_1$  (blue box) and a vector  $a_2$  (green box). The style loss  $G$  is calculated as the dot product of  $a_1$  and  $a_1$ , resulting in  $a_1 \cdot a_1$ .

# Style Loss (one layer)

Style layer

$H = 2$

$W = 2$

$F = 2$

1	2		
3	4	5	6
		7	8

$$A = \left[ \begin{array}{c} \boxed{1} \\ \boxed{2} \\ \boxed{3} \\ \boxed{4} \end{array} \quad \begin{array}{c} \boxed{5} \\ \boxed{6} \\ \boxed{7} \\ \boxed{8} \end{array} \right] G = \left[ \begin{array}{cc} a_1 \cdot a_1 & a_1 \cdot a_2 \end{array} \right]$$

$a_1$   $a_2$

# Style Loss (one layer)

Style layer

$H = 2$

$W = 2$

$F = 2$

1	2		
3	4	5	6
		7	8

$$A = \begin{bmatrix} \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 3 \\ \hline 4 \\ \hline \end{array} & \begin{array}{|c|} \hline 5 \\ \hline 6 \\ \hline 7 \\ \hline 8 \\ \hline \end{array} \end{bmatrix} \begin{matrix} a_1 \\ a_2 \end{matrix} G = \begin{bmatrix} a_1 \cdot a_1 & a_1 \cdot a_2 \\ a_2 \cdot a_1 & \end{bmatrix}$$

# Style Loss (one layer)

Style layer

$H = 2$

$W = 2$

$F = 2$

1	2		
3	4	5	6
		7	8

$$A = \left[ \begin{array}{c} \boxed{1} \\ \boxed{2} \\ \boxed{3} \\ \boxed{4} \end{array} \quad \begin{array}{c} \boxed{5} \\ \boxed{6} \\ \boxed{7} \\ \boxed{8} \end{array} \right] G = \left[ \begin{array}{cc} a_1 \cdot a_1 & a_1 \cdot a_2 \\ a_2 \cdot a_1 & a_2 \cdot a_2 \end{array} \right]$$

$a_1$                    $a_2$


# Style Loss (one layer)

Style layer

$H = 2$

$W = 2$

$F = 2$



Gram matrix

$$A = \begin{bmatrix} \begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix} & \begin{bmatrix} 5 \\ 6 \\ 7 \\ 8 \end{bmatrix} \\ a_1 & a_2 \end{bmatrix} \quad G = \begin{bmatrix} a_1 \cdot a_1 & a_1 \cdot a_2 \\ a_2 \cdot a_1 & a_2 \cdot a_2 \end{bmatrix}$$

# Style Loss (one layer)

Style layer

$H = 2$

$W = 2$

$F = 2$

1	2		
3	4		
		5	6
		7	8

Gram matrix

$$A = \begin{bmatrix} \begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix} & \begin{bmatrix} 5 \\ 6 \\ 7 \\ 8 \end{bmatrix} \\ a_1 & a_2 \end{bmatrix} \quad G = \begin{bmatrix} a_1 \cdot a_1 & a_1 \cdot a_2 \\ a_2 \cdot a_1 & a_2 \cdot a_2 \end{bmatrix} = A^T A$$


# Style Loss (one layer)

Style layer

$H = 2$

$W = 2$

$F = 2$



Gram matrix

$$A = \begin{bmatrix} \begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix} & \begin{bmatrix} 5 \\ 6 \\ 7 \\ 8 \end{bmatrix} \end{bmatrix} \quad G = \begin{bmatrix} a_1 \cdot a_1 & a_1 \cdot a_2 \\ a_2 \cdot a_1 & a_2 \cdot a_2 \end{bmatrix} = A^T A$$

$a_1$                    $a_2$

Style feature

# Style Loss (code)

```
style_layer = tf.constant([1,2,3,4,5,6,7,8],  
                           shape=(2,2,2))
```



# Style Loss (code)

```
style_layer = tf.constant([1,2,3,4,5,6,7,8],  
                           shape=(2,2,2))
```

[[[1, 2],  
 [3, 4]],  
 [[5, 6],  
 [7, 8]]]

# Style Loss (code)

```
style_layer = tf.constant([1,2,3,4,5,6,7,8],  
                           shape=(2,2,2))
```

```
[[[1, 2],  
  [3, 4]],
```

```
 [[5, 6],  
  [7, 8]]]
```

```
A = tf.transpose(  
    tf.reshape(style_layer,  
               shape=(2,4)))
```

# Style Loss (code)

```
style_layer = tf.constant([1,2,3,4,5,6,7,8],  
                           shape=(2,2,2))
```

```
A = tf.transpose(  
    tf.reshape(style_layer,  
               shape=(2,4)))
```

```
[[[1, 2],  
  [3, 4]],
```

```
 [[5, 6],  
  [7, 8]]]
```

```
[[1, 5],  
 [2, 6],  
 [3, 7],  
 [4, 8]]
```

# Style Loss (code)

```
style_layer = tf.constant([1,2,3,4,5,6,7,8],  
                           shape=(2,2,2))
```

```
[[[1, 2],  
   [3, 4]],
```

```
[[5, 6],  
 [7, 8]]]
```

```
A = tf.transpose(  
    tf.reshape(style_layer,  
               shape=(2,4)))
```

```
[[1, 5],  
 [2, 6],  
 [3, 7],  
 [4, 8]]
```

```
AT = tf.transpose(A)
```

# Style Loss (code)

```
style_layer = tf.constant([1,2,3,4,5,6,7,8],  
                           shape=(2,2,2))
```

```
[[[1, 2],  
  [3, 4]],
```

```
 [[5, 6],  
  [7, 8]]]
```

```
A = tf.transpose(  
    tf.reshape(style_layer,  
               shape=(2,4)))
```

```
[[1, 5],  
 [2, 6],  
 [3, 7],  
 [4, 8]]
```

```
AT = tf.transpose(A)
```

```
[[1, 2, 3, 4],  
 [5, 6, 7, 8]]
```

# Style Loss (code)

```
style_layer = tf.constant([1,2,3,4,5,6,7,8],  
                           shape=(2,2,2))
```

```
[[[1, 2],  
  [3, 4]],
```

```
A = tf.transpose(  
    tf.reshape(style_layer,  
               shape=(2,4)))
```

```
[[5, 6],  
 [7, 8]]]
```

```
[[1, 5],  
 [2, 6],  
 [3, 7],  
 [4, 8]]
```

```
AT = tf.transpose(A)
```

```
[[1, 2, 3, 4],  
 [5, 6, 7, 8]]
```

```
G = tf.matmul(AT,A)
```

# Style Loss (code)

```
style_layer = tf.constant([1,2,3,4,5,6,7,8],  
                           shape=(2,2,2))
```

```
[[[1, 2],  
  [3, 4]],
```

```
A = tf.transpose(  
    tf.reshape(style_layer,  
               shape=(2,4)))
```

```
[[5, 6],  
 [7, 8]]]
```

```
[[1, 5],  
 [2, 6],  
 [3, 7],  
 [4, 8]]
```

```
AT = tf.transpose(A)
```

```
[[1, 2, 3, 4],  
 [5, 6, 7, 8]]
```

```
G = tf.matmul(AT,A)
```

```
[[ 30,  70],  
 [ 70, 174]]
```

# Style Loss (code)

```
style_layer = tf.constant([1,2,3,4,5,6,7,8],  
                           shape=(2,2,2))
```

```
[[[1, 2],  
  [3, 4]],
```

```
A = tf.transpose(  
    tf.reshape(style_layer,  
               shape=(2,4)))
```

```
[[5, 6],  
 [7, 8]]]
```

```
[[1, 5],  
 [2, 6],  
 [3, 7],  
 [4, 8]]
```

```
AT = tf.transpose(A)
```

```
[[1, 2, 3, 4],  
 [5, 6, 7, 8]]
```

```
G = tf.matmul(AT,A)
```

```
[[ 30,  70],  
 [ 70, 174]]
```

```
G = tf.linalg.einsum('cij,dij->cd',  
                     style_layer,  
                     style_layer)
```



# Einstein notation

`tf.linalg.einsum`

# Einstein notation

```
x = tf.constant([1,2,3,4], shape=(2,2))
```

# Einstein notation

```
x = tf.constant([1,2,3,4], shape=(2,2))
```

```
 $\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$ 
```

# Einstein notation

```
x = tf.constant([1,2,3,4], shape=(2,2))
```

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$

```
xT = tf.transpose(x)
```

# Einstein notation

```
x = tf.constant([1,2,3,4], shape=(2,2))
```

$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$

```
xT = tf.transpose(x)
```

$\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$

# Einstein notation

```
x = tf.constant([1,2,3,4], shape=(2,2))
```

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$

```
xT = tf.transpose(x)
```

$$\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$$

```
tf.linalg.einsum('ij,ij->ij', x, x)
```

# Einstein notation

```
x = tf.constant([1,2,3,4], shape=(2,2))  
  
xT = tf.transpose(x)  
  
tf.linalg.einsum('ij,ij->ij', x, x)
```

$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$

$\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$

$\begin{bmatrix} 1 & 4 \\ 9 & 16 \end{bmatrix}$

# Einstein notation

```
x = tf.constant([1,2,3,4], shape=(2,2))  
  
xT = tf.transpose(x)  
  
tf.linalg.einsum('ij,ij->ij', x, x)  
  
tf.square(x)
```

$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$

$\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$

$\begin{bmatrix} 1 & 4 \\ 9 & 16 \end{bmatrix}$





# Einstein notation

```
x = tf.constant([1,2,3,4], shape=(2,2))  
  
xT = tf.transpose(x)  
  
tf.linalg.einsum('ij,ij->ij', x, x)  
  
tf.square(x)
```

$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$

$\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$

$\begin{bmatrix} 1 & 4 \\ 9 & 16 \end{bmatrix}$





# Einstein notation

```
x = tf.constant([1,2,3,4], shape=(2,2))  
  
xT = tf.transpose(x)  
  
tf.linalg.einsum('ij,ij->ij', x, x)  
  
tf.square(x)
```

$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$

$\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$

$\begin{bmatrix} 1 & 4 \\ 9 & 16 \end{bmatrix}$




# Einstein notation

```
x = tf.constant([1,2,3,4], shape=(2,2))  
  
xT = tf.transpose(x)  
  
tf.linalg.einsum('ij,ij->ij', x, x)  
  
tf.square(x)
```

$\begin{bmatrix} 1, & 2 \\ 3, & 4 \end{bmatrix},$

$\begin{bmatrix} 1, & 3 \\ 2, & 4 \end{bmatrix}$

$\begin{bmatrix} 1, & 4 \\ 9, & 16 \end{bmatrix},$



# Einstein notation

```
x = tf.constant([1,2,3,4], shape=(2,2))  
  
xT = tf.transpose(x)  
  
tf.linalg.einsum('ij,ij->ij', x, x)  
  
tf.square(x)  
  
tf.linalg.einsum('ij,jk->ik', xT, x)
```

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix},$$
$$\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$$
$$\begin{bmatrix} 1 & 4 \\ 9 & 16 \end{bmatrix}$$

# Einstein notation

<code>x = tf.constant([1,2,3,4], shape=(2,2))</code>	$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$
<code>xT = tf.transpose(x)</code>	$\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$
<code>tf.linalg.einsum('ij,ij-&gt;ij', x, x)</code>	$\begin{bmatrix} 1 & 4 \\ 9 & 16 \end{bmatrix}$
<code>tf.square(x)</code>	
<code>tf.linalg.einsum('ij,jk-&gt;ik', xT, x)</code>	$\begin{bmatrix} 10 & 14 \\ 14 & 20 \end{bmatrix}$

# Einstein notation

<code>x = tf.constant([1,2,3,4], shape=(2,2))</code>	$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$
<code>xT = tf.transpose(x)</code>	$\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$
<code>tf.linalg.einsum('ij,ij-&gt;ij', x, x)</code>	$\begin{bmatrix} 1 & 4 \\ 9 & 16 \end{bmatrix}$
<code>tf.square(x)</code>	
<code>tf.linalg.einsum('ij,jk-&gt;ik', xT, x)</code>	$\begin{bmatrix} 10 & 14 \\ 14 & 20 \end{bmatrix}$
<code>tf.matmul(xT, x)</code>	

# Einstein notation

```
x = tf.constant([1,2,3,4], shape=(2,2))
```

$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$

```
xT = tf.transpose(x)
```

$\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$

```
tf.linalg.einsum('ij,ij->ij', x, x)
```


$\begin{bmatrix} 1 & 4 \\ 9 & 16 \end{bmatrix}$

```
tf.square(x)
```

```
tf.linalg.einsum('ij,jk->ik', xT, x)
```

$\begin{bmatrix} 10 & 14 \\ 14 & 20 \end{bmatrix}$

```
tf.matmul(xT, x)
```



# Einstein notation


```
x = tf.constant([1,2,3,4], shape=(2,2))  
  
xT = tf.transpose(x)  
  
tf.linalg.einsum('ij,ij->ij', x, x)  
  
tf.square(x)  
  
tf.linalg.einsum('ij,jk->ik', xT, x)  
  
tf.matmul(xT, x)
```

$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$

$\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$

$\begin{bmatrix} 1 & 4 \\ 9 & 16 \end{bmatrix}$

$\begin{bmatrix} 10 & 14 \\ 14 & 20 \end{bmatrix}$






# Einstein notation

```
x = tf.constant([1,2,3,4], shape=(2,2))  
  
xT = tf.transpose(x)  
  
tf.linalg.einsum('ij,ij->ij', x, x)  
  
tf.square(x)  
  
tf.linalg.einsum('ij,jk->ik', xT, x)
```

 $\begin{bmatrix} 1, & 2 \\ 3, & 4 \end{bmatrix}$  $\begin{bmatrix} 1, & 3 \\ 2, & 4 \end{bmatrix}$  $\begin{bmatrix} 1, & 4 \\ 9, & 16 \end{bmatrix}$  $\begin{bmatrix} 10, & 14 \\ 14, & 20 \end{bmatrix}$ 

```
tf.matmul(xT, x)
```



# Einstein notation


```
x = tf.constant([1,2,3,4], shape=(2,2))  
  
xT = tf.transpose(x)  
  
tf.linalg.einsum('ij,ij->ij', x, x)  
  
tf.square(x)  
  
tf.linalg.einsum('ij,jk->ik', xT, x)  
  
tf.matmul(xT, x)
```

$\begin{bmatrix} 1, & 2 \\ 3, & 4 \end{bmatrix}$

$\begin{bmatrix} 1, & 3 \\ 2, & 4 \end{bmatrix}$

$\begin{bmatrix} 1, & 4 \\ 9, & 16 \end{bmatrix}$

$\begin{bmatrix} 10, & 14 \\ 14, & 20 \end{bmatrix}$



# Einstein notation

```
x = tf.constant([1,2,3,4], shape=(2,2))
```

 $\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$ 

```
xT = tf.transpose(x)
```

 $\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$ 

```
tf.linalg.einsum('ij,ij->ij', x, x)
```


 $\begin{bmatrix} 1 & 4 \\ 9 & 16 \end{bmatrix}$ 

```
tf.square(x)
```

```
tf.linalg.einsum('ij,jk->ik', xT, x)
```

 $\begin{bmatrix} 10 & 14 \\ 14 & 20 \end{bmatrix}$ 

```
tf.matmul(xT, x)
```



# Reduce sum

# Reduce sum


```
tf.linalg.einsum('ij,jk->ik', xT, x)
```

```
[10, 14],  
[14, 20]
```

# Reduce sum

```
tf.linalg.einsum('ij,jk->ik', xT, x)
```

$\begin{bmatrix} 10, & 14 \\ 14, & 20 \end{bmatrix}$




```
tf.linalg.einsum('ij,jk->', xT, x)
```

# Reduce sum

```
tf.linalg.einsum('ij,jk->ik', xT, x)
```

$\begin{bmatrix} 10, & 14 \\ 14, & 20 \end{bmatrix}$




```
tf.linalg.einsum('ij,jk->', xT, x)
```

# Reduce sum

```
tf.linalg.einsum('ij,jk->ik', xT, x)
```

$\begin{bmatrix} 10, & 14 \\ 14, & 20 \end{bmatrix}$



```
tf.linalg.einsum('ij,jk->', xT, x)
```


$[58]$



# Reduce sum


```
tf.linalg.einsum('ij,jk->ik', xT, x)
```

$\begin{bmatrix} 10 & 14 \\ 14 & 20 \end{bmatrix}$



```
tf.linalg.einsum('ij,jk->', xT, x)
```

$[58]$



# Einstein notation

<code>x = tf.constant([1,2,3,4], shape=(2,2))</code>	$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$
<code>v = tf.reshape(x, shape=(4,1) )</code>	$\begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix}$
<code>vT = tf.transpose(v)</code>	$[1, 2, 3, 4]$
<code>tf.matmul(vT, v)</code>	$[30]$
<code>tf.linalg.einsum('ij,ij-&gt;', x, x)</code>	$[30]$

# Style Loss

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G_{i,j}^l - A_{ij}^l)^2$$

$l$ : layer  $l$

$\mathbf{A}_{ij}^l$ : Style Representation(Gram Matrix) of style image  $a$ .

$\mathbf{G}_{ij}^l$ : Style Representation(Gram Matrix) of generated image  $x$

# Weighting the style loss

Style layer (height, width, filters)

1	2	5	6
3	4	7	8

# Weighting the style loss

Style layer (height, width, filters)

1	2
3	4

5	6
7	8

Height = 2

Width = 2

Filters = 2

# Weighting the style loss

Style layer (height, width, filters)

1	2
3	4

5	6
7	8

Height = 2  
Width = 2  
Filters = 2

Style weight

$$\frac{1}{(2 \times h \times w \times f)^2}$$

# Weighting the style loss

Style layer (height, width, filters)

1	2
3	4

5	6
7	8

Height = 2  
Width = 2  
Filters = 2

Style weight

$$\frac{1}{(2 \times h \times w \times f)^2}$$

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G_{i,j}^l - A_{i,j}^l)^2$$

# Gram Matrix - In Practice

```
def gram_matrix(input_tensor):  
    input_tensorT = tf.transpose(input_tensor, perm[0,2,1,3])  
    result = tf.linalg.einsum('bijc,bijd->bcd' input_tensorT, input_tensor)  
    input_shape = tf.shape(input_tensor)  
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)  
    return result/(num_locations)
```

<https://numpy.org/doc/stable/reference/generated/numpy.einsum.html>



# Gram Matrix - In Practice

(batch, height, width, filters)

```
def gram_matrix(input_tensor):  
    input_tensorT = tf.transpose(input_tensor, perm[0,2,1,3])  
    result = tf.linalg.einsum('bijc,bijd->bcd', input_tensorT, input_tensor)  
    input_shape = tf.shape(input_tensor)  
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)  
    return result/(num_locations)
```

# Gram Matrix - In Practice

(batch, height, width, filters)

(batch, height, width, filters)

```
def gram_matrix(input_tensor):  
    input_tensorT = tf.transpose(input_tensor, perm[0,2,1,3])  
    result = tf.linalg.einsum('bijc,bijd->bcd', input_tensorT, input_tensor)  
    input_shape = tf.shape(input_tensor)  
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)  
    return result/(num_locations)
```

# Gram Matrix - In Practice

(batch, height, width, filters) (batch, height, width, filters)

b

b

```
def gram_matrix(input_tensor):  
    input_tensorT = tf.transpose(input_tensor, perm[0,2,1,3])  
    result = tf.linalg.einsum('bijc,bijd->bcd', input_tensorT, input_tensor)  
    input_shape = tf.shape(input_tensor)  
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)  
    return result/(num_locations)
```

# Gram Matrix - In Practice

(batch, height, width, filters) (batch, height, width, filters)

b

i

j

b

```
def gram_matrix(input_tensor):  
    input_tensorT = tf.transpose(input_tensor, perm[0,2,1,3])  
    result = tf.linalg.einsum('bijc,bijd->bcd', input_tensorT, input_tensor)  
    input_shape = tf.shape(input_tensor)  
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)  
    return result/(num_locations)
```

# Gram Matrix - In Practice

(batch, height, width, filters) (batch, height, width, filters)

b

i

j

b

i

j

```
def gram_matrix(input_tensor):  
    input_tensorT = tf.transpose(input_tensor, perm[0,2,1,3])  
    result = tf.linalg.einsum('bijc,bijd->bcd', input_tensorT, input_tensor)  
    input_shape = tf.shape(input_tensor)  
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)  
    return result/(num_locations)
```

# Gram Matrix - In Practice

(batch, height, width, filters) (batch, height, width, filters)

b

i

j

c

b

i

j

```
def gram_matrix(input_tensor):  
    input_tensorT = tf.transpose(input_tensor, perm[0,2,1,3])  
    result = tf.linalg.einsum('bijc,bijd->bcd', input_tensorT, input_tensor)  
    input_shape = tf.shape(input_tensor)  
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)  
    return result/(num_locations)
```

# Gram Matrix - In Practice

(batch, height, width, filters) (batch, height, width, filters)

b

i

j

c

b

i

j

d

```
def gram_matrix(input_tensor):  
    input_tensorT = tf.transpose(input_tensor, perm[0,2,1,3])  
    result = tf.linalg.einsum('bijc,bijd->bcd', input_tensorT, input_tensor)  
    input_shape = tf.shape(input_tensor)  
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)  
    return result/(num_locations)
```

# Gram Matrix - In Practice

(batch, height, width, filters)

b i j c

(batch, height, width, filters)

b i j d

(batch, filters, filters)

b c d

```
def gram_matrix(input_tensor):  
    input_tensorT = tf.transpose(input_tensor, perm[0,2,1,3])  
    result = tf.linalg.einsum('bijc,bijd->bcd', input_tensorT, input_tensor)  
    input_shape = tf.shape(input_tensor)  
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)  
    return result/(num_locations)
```



# Gram Matrix - In Practice

(batch, height, width, filters)

b i j c

(batch, height, width, filters)

b i j d

(batch, filters, filters)

b c d

```
def gram_matrix(input_tensor):  
    input_tensorT = tf.transpose(input_tensor, perm[0,2,1,3])  
    result = tf.linalg.einsum('bijc,bijd->bcd', input_tensorT, input_tensor)  
    input_shape = tf.shape(input_tensor)  
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)  
    return result/(num_locations)
```

<https://numpy.org/doc/stable/reference/generated/numpy.einsum.html>




# Extract High Frequency Components


```
def high_pass_x_y(image):  
    x_var = image[:, :, 1:, :] - image[:, :, :-1, :]  
    y_var = image[:, 1:, :, :] - image[:, :-1, :, :]  
  
    return x_var, y_var
```

# Visualization-High Frequency Artifacts


Horizontal Deltas: Original




Vertical Deltas: Original



Horizontal Deltas: Stylized Image



Vertical Deltas: Stylized Image



# Solution - Total Variation Loss

- Decrease high frequency artefacts.
- Explicit regularization term on high frequency components

# Factoring in Total Variation using TensorFlow

```
loss += total_variation_weight * tf.image.total_variation(image)
```

```
def calculate_gradients(image, content_targets,
                        style_targets, style_weight,
                        content_weight, with_regularization=False ):

    total_variation_weight = 30

    with tf.GradientTape() as tape:
        if with_regularization:
            loss += total_variation_weight*tf.image.total_variation(image)

    gradients = tape.gradient(loss, image)
    return gradients
```

```
def calculate_gradients(image, content_targets,  
                        style_targets, style_weight,  
                        content_weight, with_regularization=False ):  
  
    total_variation_weight = 30  
    with tf.GradientTape() as tape:  
        if with_regularization:  
            loss += total_variation_weight*tf.image.total_variation(image)  
  
    gradients = tape.gradient(loss, image)  
    return gradients
```



# Comparison

Without Variation Loss



With Variation Loss



[https://cdn.pixabay.com/photo/2017/02/28/23/00/swan-2107052\\_1280.jpg](https://cdn.pixabay.com/photo/2017/02/28/23/00/swan-2107052_1280.jpg)

## Exploring the structure of a real-time, arbitrary neural artistic stylization network

Golnaz Ghiasi Google Brain golnazg@google.com	Honglak Lee Google Brain honglak@google.com	Manjunath Kudlur Google Brain keveaman@google.com
Vincent Dumoulin MILA, Université de Montréal vi.dumoulin@gmail.com	Jonathon Shlens Google Brain shlens@google.com	

August 28, 2017

### Abstract

In this paper, we present a method which combines the flexibility of the neural algorithm of artistic style with the speed of fast style transfer networks to allow real-time stylization using any content/style image pair. We build upon recent work leveraging conditional instance normalization for multi-style transfer networks by learning to predict the conditional instance normalization parameters directly from a style image. The model is successfully trained on a corpus of roughly 80,000 paintings and is able to generalize to paintings previously unobserved. We demonstrate that the learned embedding space is smooth and contains a rich structure and organizes semantic information associated with paintings in an entirely unsupervised manner.

## 1 Introduction

*Elmyr de Hory* gained world-wide fame by forging thousands of pieces of artwork and selling them to art dealers and museums [13]. The forger's skill is a testament to the human talent and intelligence required to reproduce the artistic details of a diverse set of paintings. In computer vision, much work has been invested in teaching computers to likewise capture the artistic style of a painting with the goal of conferring this style in arbitrary photographs in a convincing manner.



# magenta/arbitrary-image-stylization-v1-256

Fast arbitrary image style transfer.

Publisher: **Google** Updated: 08/25/2020 License: **Apache-2.0**

Architecture: Dataset:

- Other
- Multiple

Overall usage data

133 Downloads

## Model formats

<

TF

TFLite (magenta/arbitrary-image-stylization-v1-256/fp16/prediction)

TFLite (magenta/arbitrary-image-stylization-v1-256/fp16/transfer)

TFLite (magenta/arbitrary-image-stylization-v1-256/fp16/quantized) >

Hub module (v2)

Usage data: 133 Downloads V2

Fine tunable: No License: [Apache-2.0](#) Last updated: 08/25/2020 Format: Hub module

...rbitrary-image-stylization-v1-256/2

Copy URL

```
import tensorflow_hub as hub
```

```
h_m = hub.load
```

```
( 'https://tfhub.dev/google/magenta/arbitrary-image-stylization-v1-256/2' )
```

```
import tensorflow_hub as hub

h_m = hub.load

( 'https://tfhub.dev/google/magenta/arbitrary-image-stylization-v1-256/2' )

content_image, style_image = load_images("content.jpg", "style.jpg")
```

```
import tensorflow_hub as hub

h_m = hub.load
      ('https://tfhub.dev/google/magenta/arbitrary-image-stylization-v1-256/2')

content_image, style_image = load_images("content.jpg", "style.jpg")

stylized_image = h_m(tf.image.convert_image_dtype(content_image, tf.float32),
                    tf.image.convert_image_dtype(style_image, tf.float32))[0]
```

```
import tensorflow_hub as hub

h_m = hub.load
      ('https://tfhub.dev/google/magenta/arbitrary-image-stylization-v1-256/2')

content_image, style_image = load_images("content.jpg", "style.jpg")

stylized_image = h_m(tf.image.convert_image_dtype(content_image, tf.float32),
                     tf.image.convert_image_dtype(style_image, tf.float32))[0]

tensor_to_image(stylized_image)
```

