

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>

What are AutoEncoders?

- Neural networks capable of learning dense representations of input data without supervision
 - Training data is not labelled
- Useful for dimensionality reduction and for visualization
- Can be used to generate new data that resembles input data
- In practice they
 - Copy input to output
 - They learn efficient ways to represent data

What are AutoEncoders?

- Neural networks capable of learning dense representations of input data without supervision
 - Training data is not labelled
- Useful for dimensionality reduction and for visualization
- Can be used to generate new data that resembles input data
- In practice they
 - Copy input to output
 - They learn efficient ways to represent data

What are AutoEncoders?

- Neural networks capable of learning dense representations of input data without supervision
 - Training data is not labelled
- Useful for dimensionality reduction and for visualization
- Can be used to generate new data that resembles input data
- In practice they
 - Copy input to output
 - They learn efficient ways to represent data

What are AutoEncoders?

- Neural networks capable of learning dense representations of input data without supervision
 - Training data is not labelled
- Useful for dimensionality reduction and for visualization
- Can be used to generate new data that resembles input data
- In practice they
 - Copy input to output
 - They learn efficient ways to represent data

What are AutoEncoders?

- Neural networks capable of learning dense representations of input data without supervision
 - Training data is not labelled
- Useful for dimensionality reduction and for visualization
- Can be used to generate new data that resembles input data
- In practice they
 - Copy input to output
 - They learn efficient ways to represent data

83 12 21 42 99 18 51


4 9 16 25 36 49 64 81 100 121 144 169


83 12 21 42 99 18 51


4 9 16 25 36 49 64 81 100 121 144 169


83 12 21 42 99 18 51


4 9 16 25 36 49 64 81 100 121 144 169










































```
encoder = keras.models.Sequential([keras.layers.Dense(2, input_shape=[3])])  
decoder = keras.models.Sequential([keras.layers.Dense(3, input_shape=[2])])  
autoencoder = keras.models.Sequential([encoder, decoder])  
autoencoder.compile(loss="mse", optimizer=keras.optimizers.SGD(lr=1.5))
```




```
encoder = keras.models.Sequential([keras.layers.Dense(2, input_shape=[3])])  
decoder = keras.models.Sequential([keras.layers.Dense(3, input_shape=[2])])  
autoencoder = keras.models.Sequential([encoder, decoder])  
autoencoder.compile(loss="mse", optimizer=keras.optimizers.SGD(lr=1.5))
```




```
encoder = keras.models.Sequential([keras.layers.Dense(2, input_shape=[3])])
```


```
decoder = keras.models.Sequential([keras.layers.Dense(3, input_shape=[2])])
```

```
autoencoder = keras.models.Sequential([encoder, decoder])
```

```
autoencoder.compile(loss="mse", optimizer=keras.optimizers.SGD(lr=1.5))
```




```
encoder = keras.models.Sequential([keras.layers.Dense(2, input_shape=[3])])  
decoder = keras.models.Sequential([keras.layers.Dense(3, input_shape=[2])])  
autoencoder = keras.models.Sequential([encoder, decoder])  
autoencoder.compile(loss="mse", optimizer=keras.optimizers.SGD(lr=1.5))
```



```
history = autoencoder.fit(X_train, X_train, epochs=200)
```


```
history = autoencoder.fit(X_train, X_train, epochs=200)
```

```
history = autoencoder.fit(X_train, X_train, epochs=200)
```




```
codings = encoder.predict(data)
```

```
decodings = decoder.predict(codings)
```





5 3 8 8 7 5 1 1 9 8




5 3 8 8 7 5 1 1 9 8

5 3 8 8 7 5 1 1 9 8



5 3 8 8 7 5 1 1 9 8

5 3 8 8 7 5 1 1 9 8



5 3 8 8 7 5 1 1 9 8

```
inputs = tf.keras.layers.Input(shape=(784,))
```

```
def simple_autoencoder():  
    encoder = tf.keras.layers.Dense(units=32, activation='relu')(inputs)  
    decoder = tf.keras.layers.Dense(units=784, activation='sigmoid')(encoder)  
    return encoder, decoder
```

```
encoder_output, decoder_output = simple_autoencoder()
```

```
encoder_model = tf.keras.Model(inputs=inputs, outputs=encoder_output)
```

```
autoencoder_model = tf.keras.Model(inputs=inputs, outputs=decoder_output)
```

```
inputs = tf.keras.layers.Input(shape=(784,))
```

```
def simple_autoencoder():  
    encoder = tf.keras.layers.Dense(units=32, activation='relu')(inputs)  
    decoder = tf.keras.layers.Dense(units=784, activation='sigmoid')(encoder)  
    return encoder, decoder
```

```
encoder_output, decoder_output = simple_autoencoder()
```

```
encoder_model = tf.keras.Model(inputs=inputs, outputs=encoder_output)
```

```
autoencoder_model = tf.keras.Model(inputs=inputs, outputs=decoder_output)
```

```
inputs = tf.keras.layers.Input(shape=(784,))
```

```
def simple_autoencoder():  
    encoder = tf.keras.layers.Dense(units=32, activation='relu')(inputs)  
    decoder = tf.keras.layers.Dense(units=784, activation='sigmoid')(encoder)  
    return encoder, decoder
```

```
encoder_output, decoder_output = simple_autoencoder()
```

```
encoder_model = tf.keras.Model(inputs=inputs, outputs=encoder_output)
```

```
autoencoder_model = tf.keras.Model(inputs=inputs, outputs=decoder_output)
```

```
inputs = tf.keras.layers.Input(shape=(784,))
```

```
def simple_autoencoder():  
    encoder = tf.keras.layers.Dense(units=32, activation='relu')(inputs)  
    decoder = tf.keras.layers.Dense(units=784, activation='sigmoid')(encoder)  
    return encoder, decoder
```

```
encoder_output, decoder_output = simple_autoencoder()
```

```
encoder_model = tf.keras.Model(inputs=inputs, outputs=encoder_output)
```

```
autoencoder_model = tf.keras.Model(inputs=inputs, outputs=decoder_output)
```



```
inputs = tf.keras.layers.Input(shape=(784,))
```

```
def simple_autoencoder():  
    encoder = tf.keras.layers.Dense(units=32, activation='relu')(inputs)  
    decoder = tf.keras.layers.Dense(units=784, activation='sigmoid')(encoder)  
    return encoder, decoder
```

```
encoder_output, decoder_output = simple_autoencoder()
```

```
encoder_model = tf.keras.Model(inputs=inputs, outputs=encoder_output)
```

```
autoencoder_model = tf.keras.Model(inputs=inputs, outputs=decoder_output)
```

```
inputs = tf.keras.layers.Input(shape=(784,))
```

```
def simple_autoencoder():  
    encoder = tf.keras.layers.Dense(units=32, activation='relu')(inputs)  
    decoder = tf.keras.layers.Dense(units=784, activation='sigmoid')(encoder)  
    return encoder, decoder
```


```
encoder_output, decoder_output = simple_autoencoder()
```

```
encoder_model = tf.keras.Model(inputs=inputs, outputs=encoder_output)
```

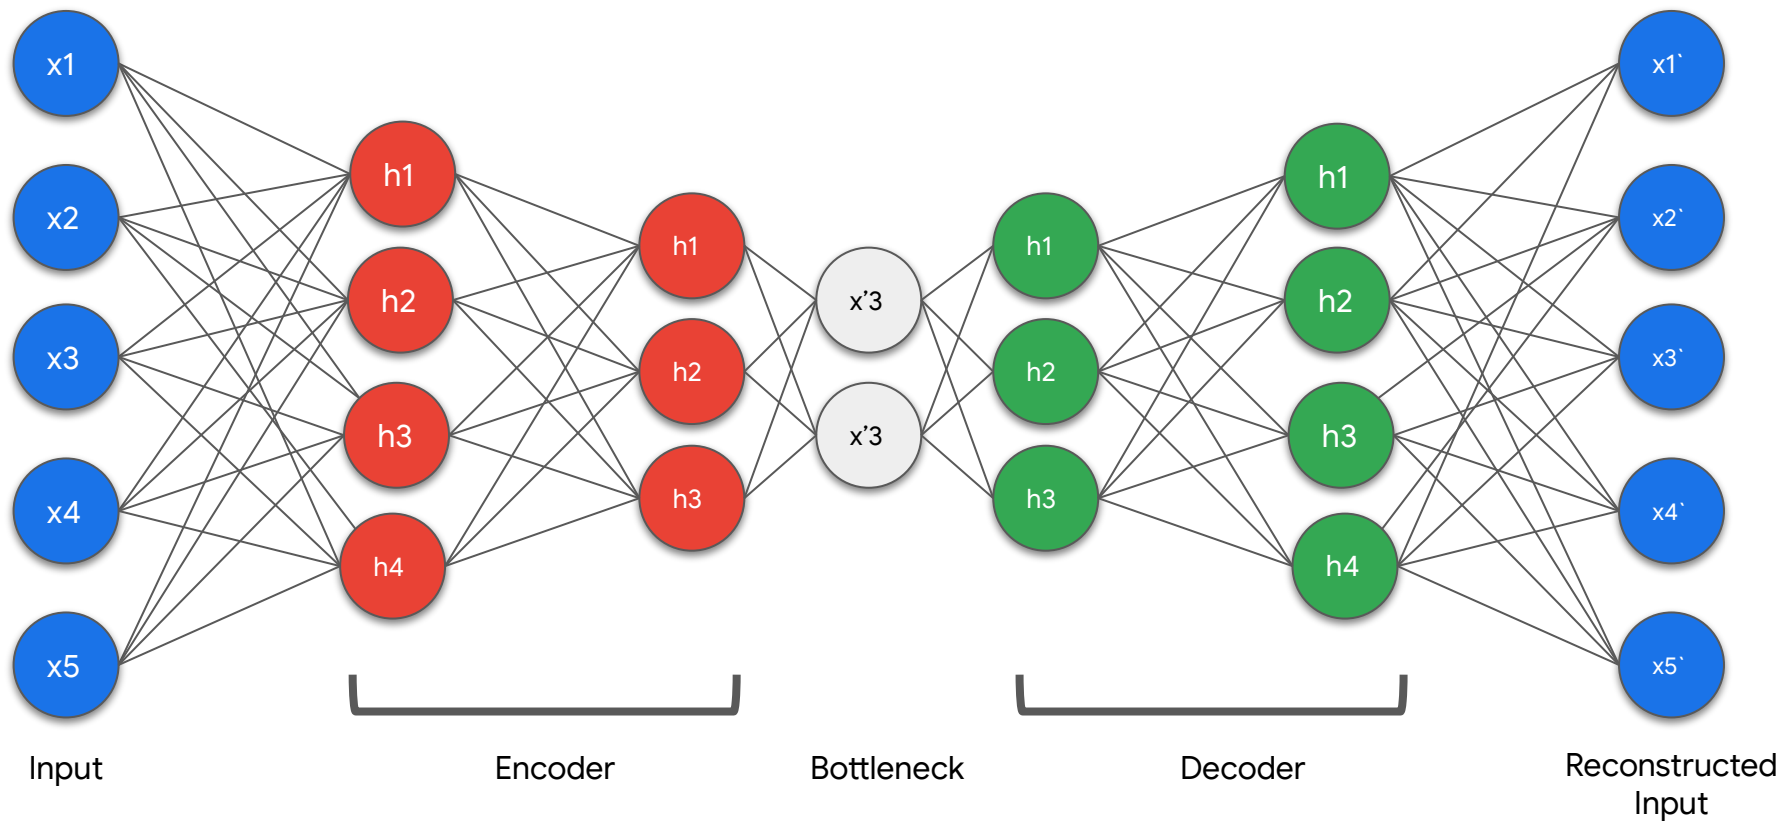
```
autoencoder_model = tf.keras.Model(inputs=inputs, outputs=decoder_output)
```

```
autoencoder_model.compile(  
    optimizer=tf.keras.optimizers.Adam(),  
    loss='binary_crossentropy')
```

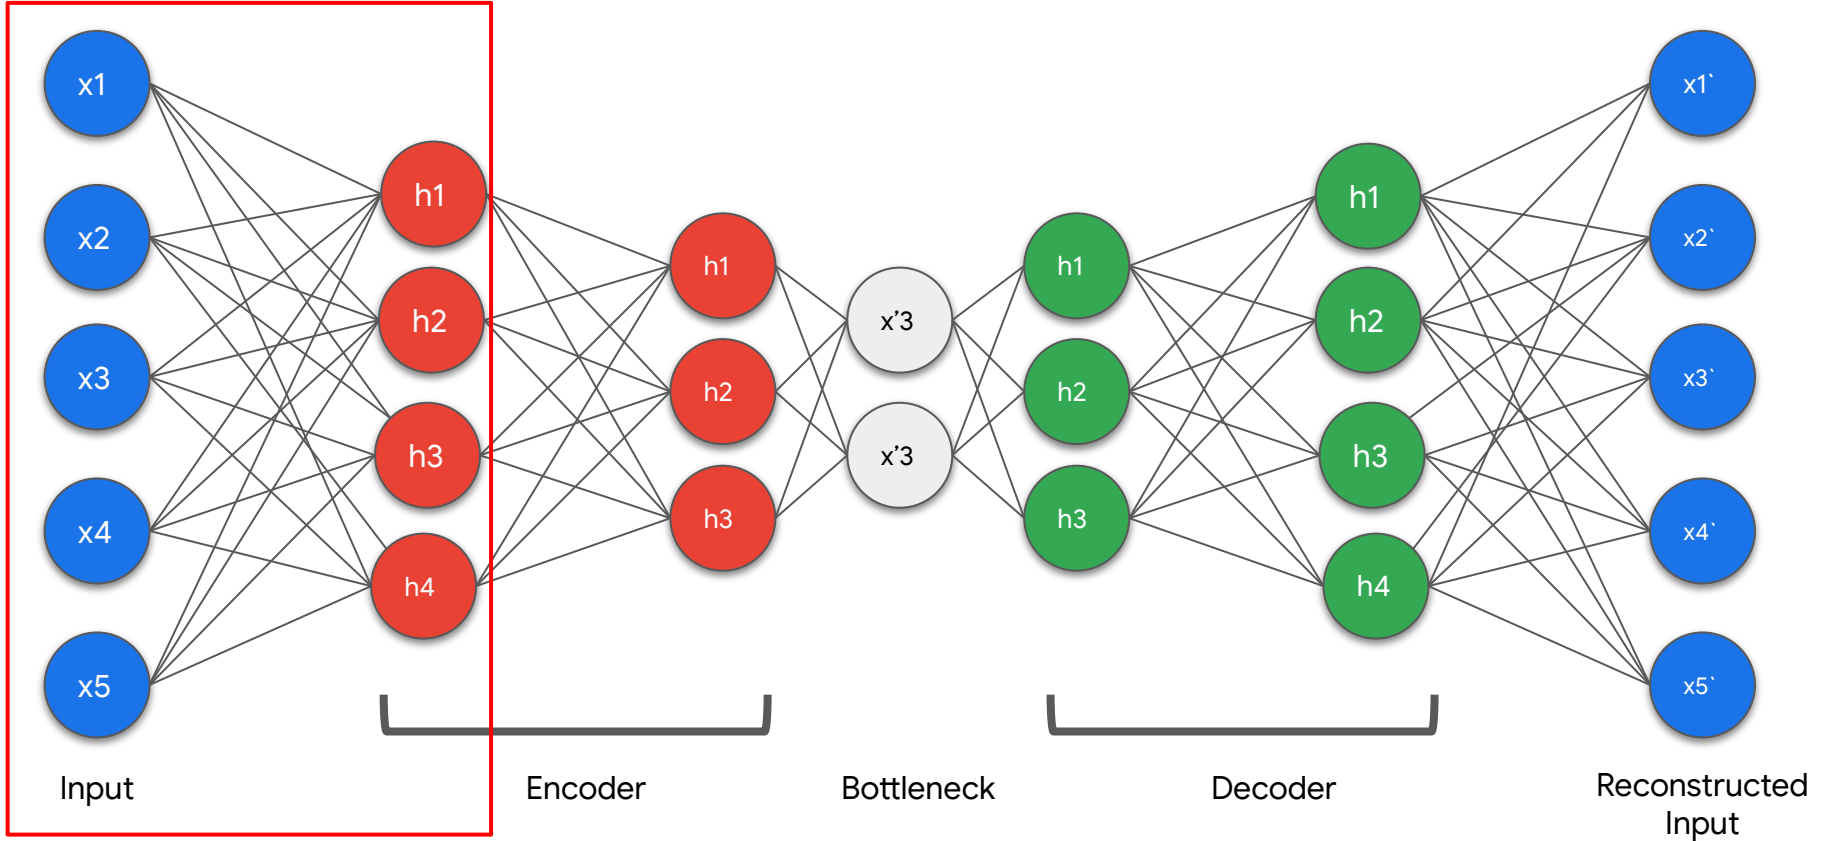
```
autoencoder_model.compile(  
    optimizer=tf.keras.optimizers.Adam(),  
    loss='binary_crossentropy')
```



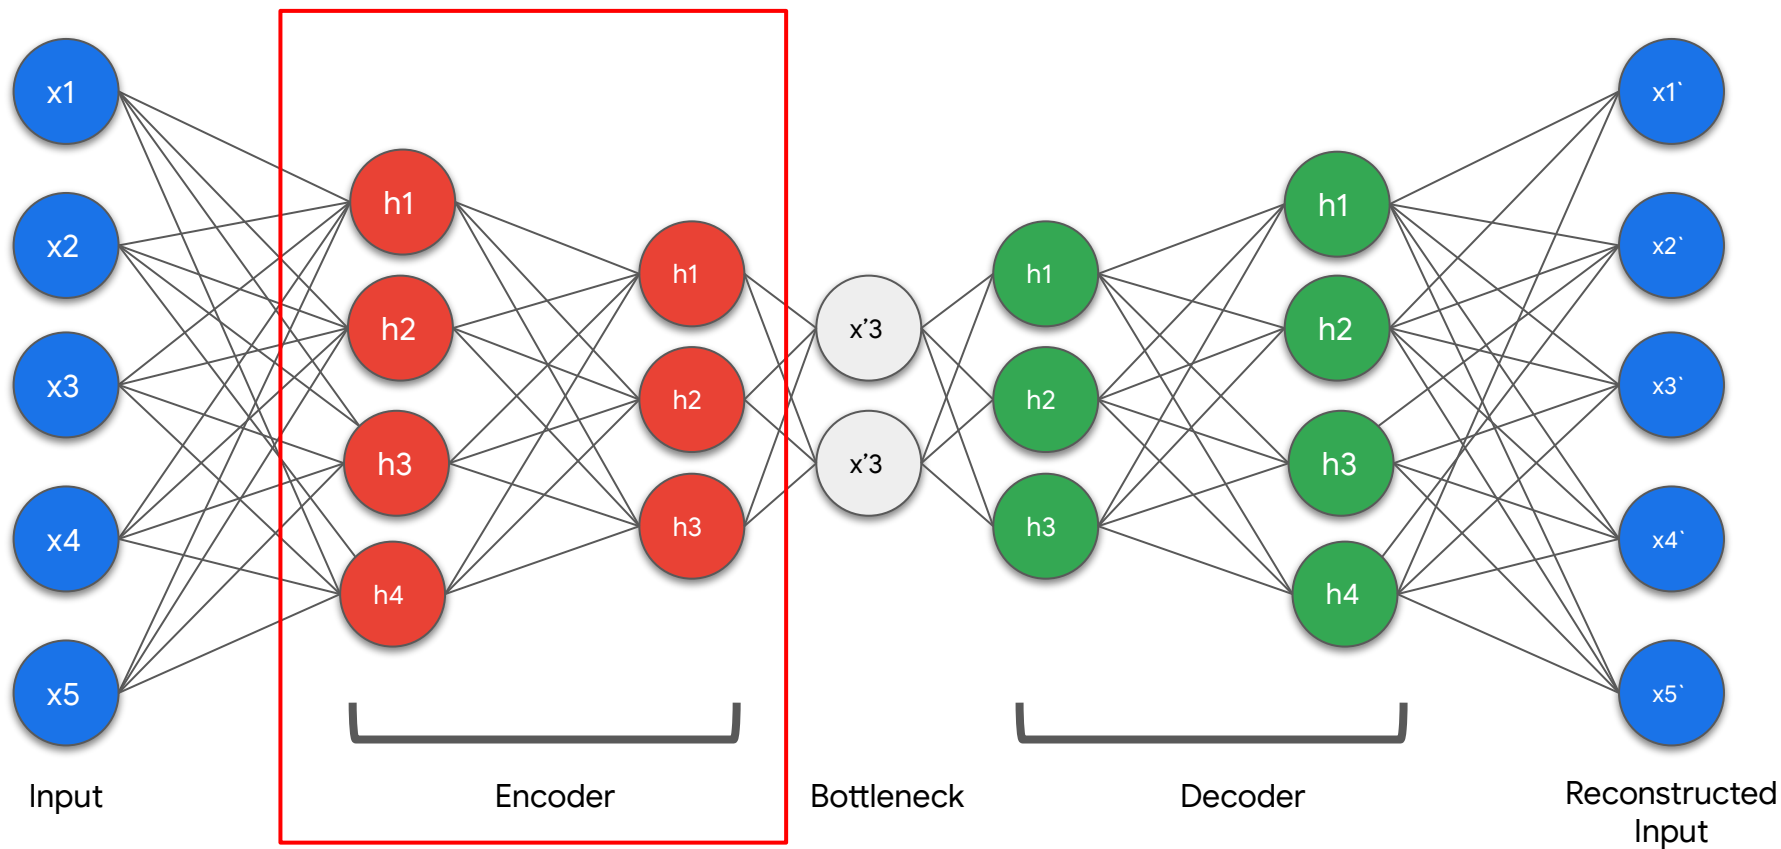
Stacked Auto-Encoders



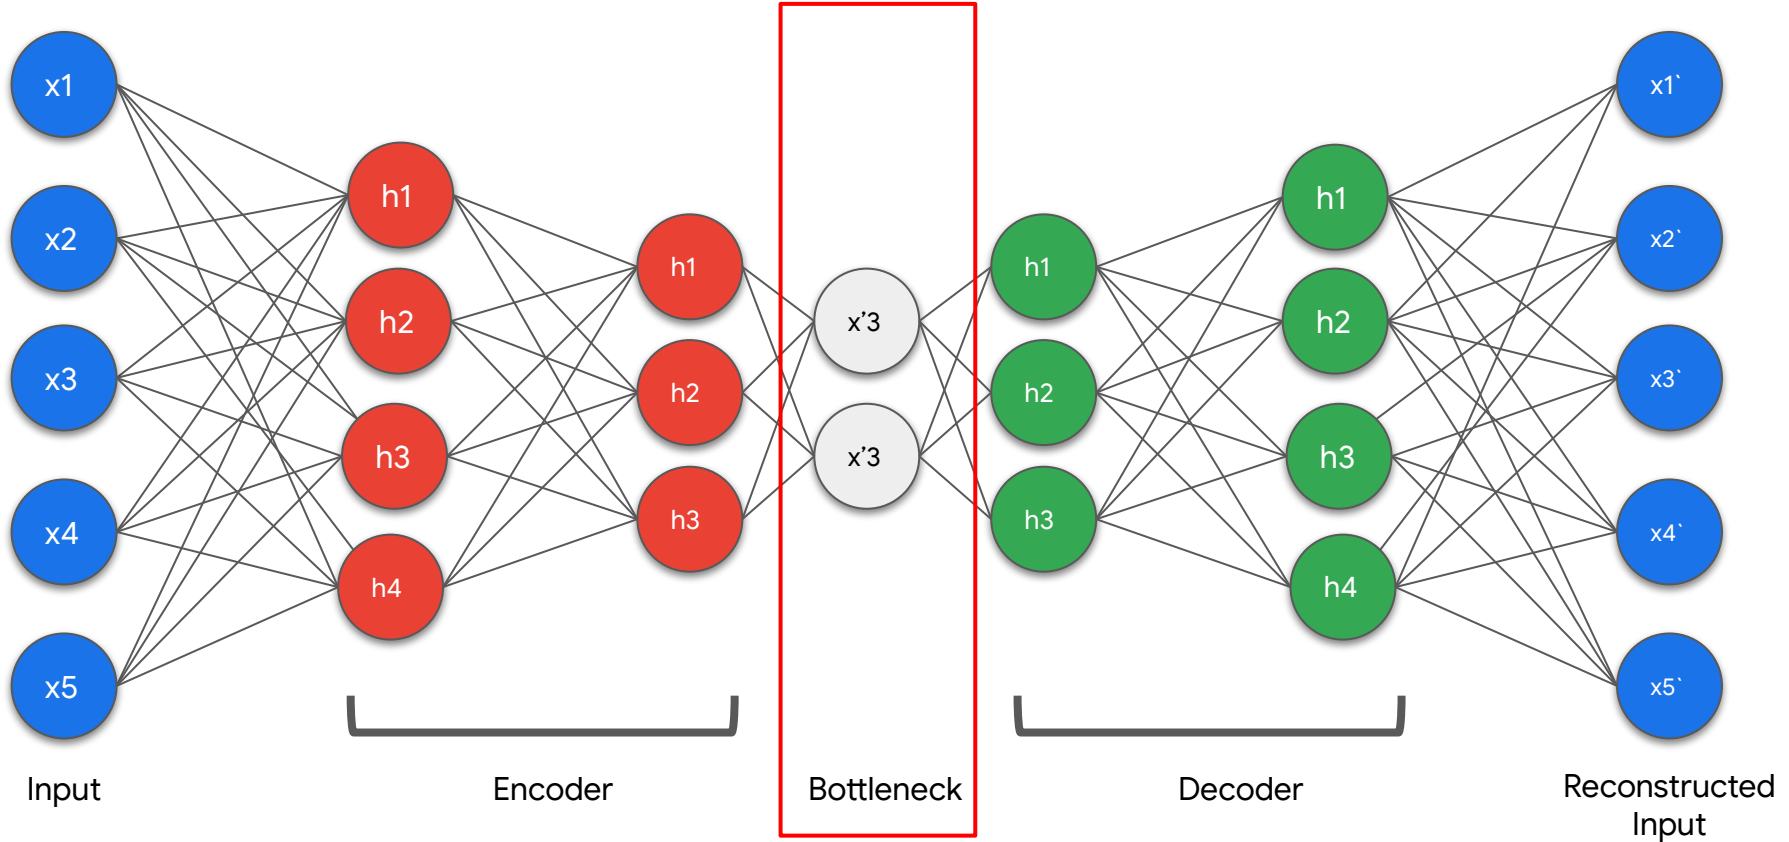
Stacked Auto-Encoders



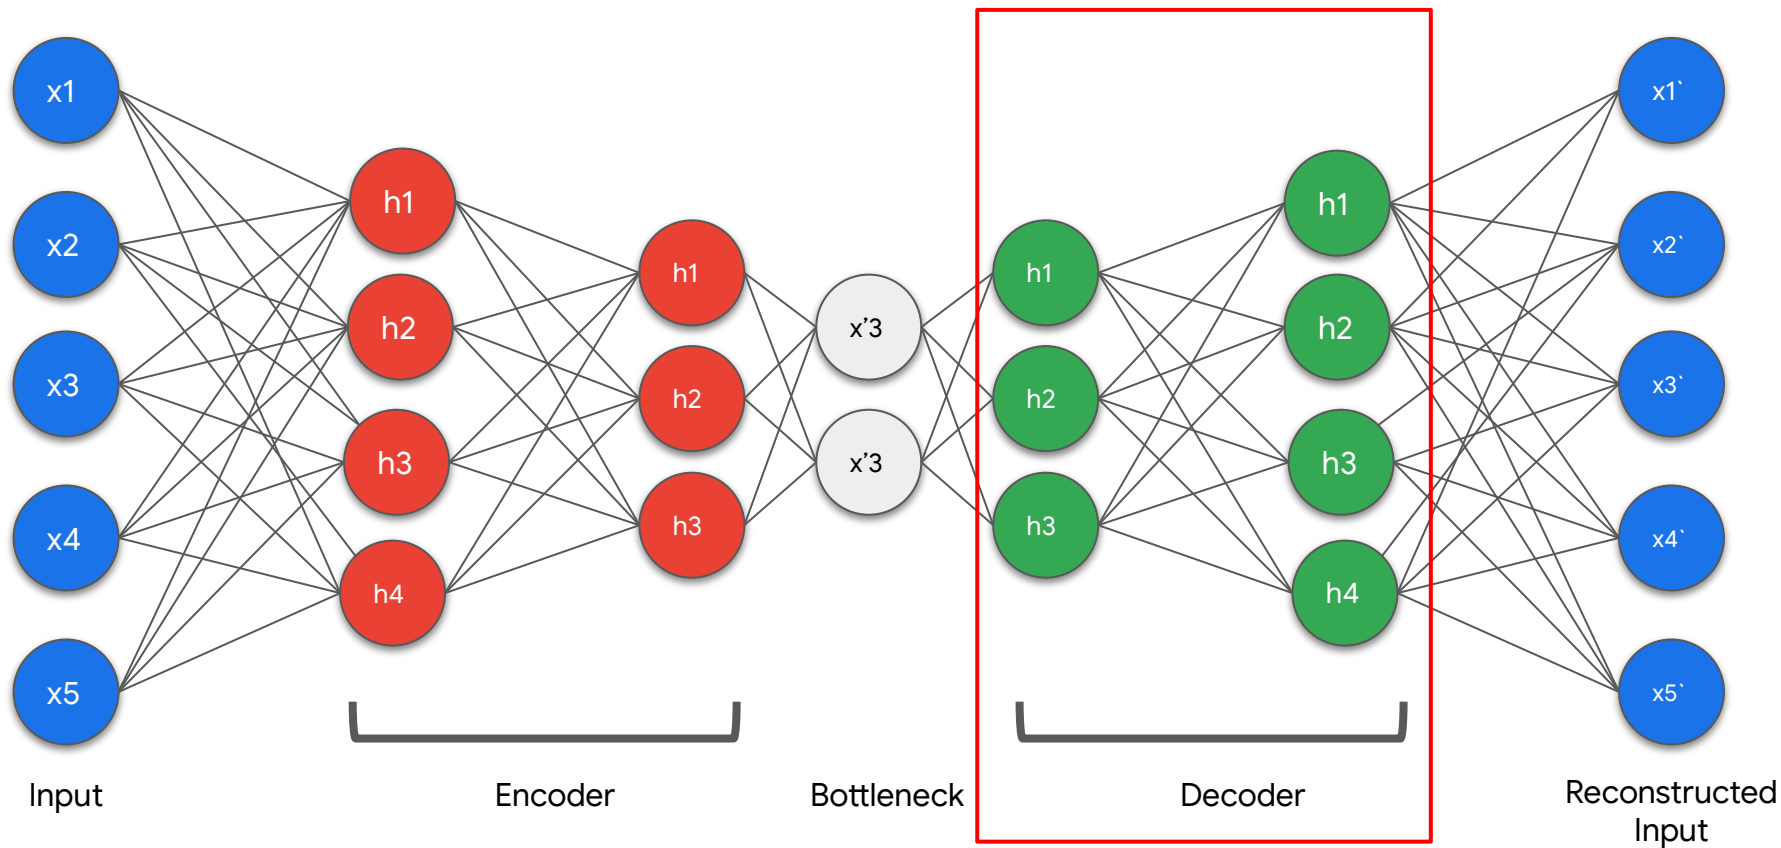
Stacked Auto-Encoders



Stacked Auto-Encoders



Stacked Auto-Encoders



```
inputs = tf.keras.layers.Input(shape=(784,))
```

```
def deep_autoencoder():
```

```
    encoder = tf.keras.layers.Dense(units=128, activation='relu')(inputs)
```

```
    encoder = tf.keras.layers.Dense(units=64, activation='relu')(encoder)
```

```
    encoder = tf.keras.layers.Dense(units=32, activation='relu')(encoder)
```

```
    decoder = tf.keras.layers.Dense(units=64, activation='relu')(encoder)
```

```
    decoder = tf.keras.layers.Dense(units=128, activation='relu')(decoder)
```

```
    decoder = tf.keras.layers.Dense(units=784, activation='sigmoid')(decoder)
```

```
    return encoder, decoder
```

```
deep_encoder_output, deep_autoencoder_output = deep_autoencoder()
```

```
deep_encoder_model = tf.keras.Model(inputs=inputs, outputs=deep_encoder_output)
```

```
deep_autoencoder_model = tf.keras.Model(inputs=inputs, outputs=deep_autoencoder_output)
```

```
def deep_autoencoder():  
    encoder = tf.keras.layers.Dense(units=128, activation='relu')(inputs)  
    encoder = tf.keras.layers.Dense(units=64, activation='relu')(encoder)  
    encoder = tf.keras.layers.Dense(units=32, activation='relu')(encoder)  
  
    decoder = tf.keras.layers.Dense(units=64, activation='relu')(encoder)  
    decoder = tf.keras.layers.Dense(units=128, activation='relu')(decoder)  
    decoder = tf.keras.layers.Dense(units=784, activation='sigmoid')(decoder)  
  
    return encoder, decoder
```

```
deep_encoder_output, deep_autoencoder_output = deep_autoencoder()
```

```
deep_encoder_model = tf.keras.Model(inputs=inputs, outputs=deep_encoder_output)  
deep_autoencoder_model = tf.keras.Model(inputs=inputs, outputs=deep_autoencoder_output)
```

```
def deep_autoencoder():  
    encoder = tf.keras.layers.Dense(units=128, activation='relu')(inputs)  
    encoder = tf.keras.layers.Dense(units=64, activation='relu')(encoder)  
    encoder = tf.keras.layers.Dense(units=32, activation='relu')(encoder)  
  
    decoder = tf.keras.layers.Dense(units=64, activation='relu')(encoder)  
    decoder = tf.keras.layers.Dense(units=128, activation='relu')(decoder)  
    decoder = tf.keras.layers.Dense(units=784, activation='sigmoid')(decoder)  
  
    return encoder, decoder
```


```
deep_encoder_output, deep_autoencoder_output = deep_autoencoder()
```

```
deep_encoder_model = tf.keras.Model(inputs=inputs, outputs=deep_encoder_output)  
deep_autoencoder_model = tf.keras.Model(inputs=inputs, outputs=deep_autoencoder_output)
```



```
def deep_autoencoder():  
    encoder = tf.keras.layers.Dense(units=128, activation='relu')(inputs)  
    encoder = tf.keras.layers.Dense(units=64, activation='relu')(encoder)  
    encoder = tf.keras.layers.Dense(units=32, activation='relu')(encoder)  
  
    decoder = tf.keras.layers.Dense(units=64, activation='relu')(encoder)  
    decoder = tf.keras.layers.Dense(units=128, activation='relu')(decoder)  
    decoder = tf.keras.layers.Dense(units=784, activation='sigmoid')(decoder)  
  
    return encoder, decoder
```


```
deep_encoder_output, deep_autoencoder_output = deep_autoencoder()  
  
deep_encoder_model = tf.keras.Model(inputs=inputs, outputs=deep_encoder_output)  
deep_autoencoder_model = tf.keras.Model(inputs=inputs, outputs=deep_autoencoder_output)
```



1 4 5 8 6 6 6 1 6 6







1 4 5 8 6 6 6 1 6 6























Convolutional Auto-Encoders




Convolutional Auto-Encoders



Convolutional Auto-Encoders



Convolutional Auto-Encoders



```
def encoder(inputs):  
    conv_1 = tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3),  
                                     activation='relu', padding='same')(inputs)  
  
    max_pool_1 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_1)  
  
    conv_2 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3),  
                                     activation='relu', padding='same')(max_pool_1)  
  
    max_pool_2 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_2)  
  
    return max_pool_2
```

```
def encoder(inputs):
```

```
    conv_1 = tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3),  
                                     activation='relu', padding='same')(inputs)
```

```
    max_pool_1 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_1)
```

```
    conv_2 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3),  
                                     activation='relu', padding='same')(max_pool_1)
```

```
    max_pool_2 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_2)
```

```
    return max_pool_2
```

```
def encoder(inputs):  
    conv_1 = tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3),  
                                     activation='relu', padding='same')(inputs)  
  
    max_pool_1 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_1)  
  
    conv_2 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3),  
                                     activation='relu', padding='same')(max_pool_1)  
  
    max_pool_2 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_2)  
  
    return max_pool_2
```

```
def encoder(inputs):  
    conv_1 = tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3),  
                                     activation='relu', padding='same')(inputs)  
  
    max_pool_1 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_1)  
  
    conv_2 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3),  
                                     activation='relu', padding='same')(max_pool_1)  
  
    max_pool_2 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_2)  
  
    return max_pool_2
```

```
def bottle_neck(inputs):  
    bottle_neck = tf.keras.layers.Conv2D(filters=256, kernel_size=(3,3),  
                                           activation='relu', padding='same')(inputs)  
  
    encoder_visualization = tf.keras.layers.Conv2D(filters=1, kernel_size=(3,3),  
                                                    activation='sigmoid',  
                                                    padding='same')(bottle_neck)  
  
    return bottle_neck, encoder_visualization
```

```
def bottle_neck(inputs):  
    bottle_neck = tf.keras.layers.Conv2D(filters=256, kernel_size=(3,3),  
                                           activation='relu', padding='same')(inputs)  
  
    encoder_visualization = tf.keras.layers.Conv2D(filters=1, kernel_size=(3,3),  
                                                    activation='sigmoid',  
                                                    padding='same')(bottle_neck)  
  
    return bottle_neck, encoder_visualization
```



```
def bottle_neck(inputs):  
    bottle_neck = tf.keras.layers.Conv2D(filters=256, kernel_size=(3,3),  
                                           activation='relu', padding='same')(inputs)  
  
    encoder_visualization = tf.keras.layers.Conv2D(filters=1, kernel_size=(3,3),  
                                                    activation='sigmoid',  
                                                    padding='same')(bottle_neck)  
  
    return bottle_neck, encoder_visualization
```

```
def bottle_neck(inputs):  
    bottle_neck = tf.keras.layers.Conv2D(filters=256, kernel_size=(3,3),  
                                           activation='relu', padding='same')(inputs)  
  
    encoder_visualization = tf.keras.layers.Conv2D(filters=1, kernel_size=(3,3),  
                                                    activation='sigmoid',  
                                                    padding='same')(bottle_neck)  
  
    return bottle_neck, encoder_visualization
```

```
def decoder(inputs):  
    conv_1 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3),  
                                     activation='relu', padding='same')(inputs)  
    up_sample_1 = tf.keras.layers.UpSampling2D(size=(2,2))(conv_1)  
  
    conv_2 = tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3),  
                                     activation='relu', padding='same')(up_sample_1)  
    up_sample_2 = tf.keras.layers.UpSampling2D(size=(2,2))(conv_2)  
  
    conv_3 = tf.keras.layers.Conv2D(filters=1, kernel_size=(3,3),  
                                     activation='sigmoid',  
                                     padding='same')(up_sample_2)  
  
    return conv_3
```

```
def decoder(inputs):
```

```
    conv_1 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3),  
                                     activation='relu', padding='same')(inputs)  
    up_sample_1 = tf.keras.layers.UpSampling2D(size=(2,2))(conv_1)
```

```
    conv_2 = tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3),  
                                     activation='relu', padding='same')(up_sample_1)  
    up_sample_2 = tf.keras.layers.UpSampling2D(size=(2,2))(conv_2)
```

```
    conv_3 = tf.keras.layers.Conv2D(filters=1, kernel_size=(3,3),  
                                     activation='sigmoid',  
                                     padding='same')(up_sample_2)
```

```
    return conv_3
```

```
def decoder(inputs):  
    conv_1 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3),  
                                     activation='relu', padding='same')(inputs)  
    up_sample_1 = tf.keras.layers.UpSampling2D(size=(2,2))(conv_1)  
  
    conv_2 = tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3),  
                                     activation='relu', padding='same')(up_sample_1)  
    up_sample_2 = tf.keras.layers.UpSampling2D(size=(2,2))(conv_2)  
  
    conv_3 = tf.keras.layers.Conv2D(filters=1, kernel_size=(3,3),  
                                     activation='sigmoid',  
                                     padding='same')(up_sample_2)  
  
    return conv_3
```

```
def decoder(inputs):  
    conv_1 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3),  
                                     activation='relu', padding='same')(inputs)  
    up_sample_1 = tf.keras.layers.UpSampling2D(size=(2,2))(conv_1)  
  
    conv_2 = tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3),  
                                     activation='relu', padding='same')(up_sample_1)  
    up_sample_2 = tf.keras.layers.UpSampling2D(size=(2,2))(conv_2)  
  
    conv_3 = tf.keras.layers.Conv2D(filters=1, kernel_size=(3,3),  
                                     activation='sigmoid',  
                                     padding='same')(up_sample_2)  
  
    return conv_3
```

```
def convolutional_auto_encoder():  
    inputs = tf.keras.layers.Input(shape=(28, 28, 1,))  
  
    encoder_output = encoder(inputs)  
  
    bottleneck_output, encoder_visualization = bottle_neck(encoder_output)  
  
    decoder_output = decoder(bottleneck_output)  
  
    model = tf.keras.Model(inputs =inputs, outputs=decoder_output)  
    encoder_model = tf.keras.Model(inputs=inputs, outputs=encoder_visualization)  
  
    return model, encoder_model
```


```
def convolutional_auto_encoder():  
    inputs = tf.keras.layers.Input(shape=(28, 28, 1,))  
  
    encoder_output = encoder(inputs)  
  
    bottleneck_output, encoder_visualization = bottle_neck(encoder_output)  
  
    decoder_output = decoder(bottleneck_output)  
  
    model = tf.keras.Model(inputs =inputs, outputs=decoder_output)  
    encoder_model = tf.keras.Model(inputs=inputs, outputs=encoder_visualization)  
  
    return model, encoder_model
```




```
def convolutional_auto_encoder():  
    inputs = tf.keras.layers.Input(shape=(28, 28, 1,))  
  
    encoder_output = encoder(inputs)  
  
    bottleneck_output, encoder_visualization = bottle_neck(encoder_output)  
  
    decoder_output = decoder(bottleneck_output)  
  
    model = tf.keras.Model(inputs=inputs, outputs=decoder_output)  
    encoder_model = tf.keras.Model(inputs=inputs, outputs=encoder_visualization)  
  
    return model, encoder_model
```

```
def convolutional_auto_encoder():  
    inputs = tf.keras.layers.Input(shape=(28, 28, 1,))  
  
    encoder_output = encoder(inputs)  
  
    bottleneck_output, encoder_visualization = bottle_neck(encoder_output)  
  
    decoder_output = decoder(bottleneck_output)  
  
    model = tf.keras.Model(inputs =inputs, outputs=decoder_output)  
    encoder_model = tf.keras.Model(inputs=inputs, outputs=encoder_visualization)  
  
    return model, encoder_model
```


Convolutional Auto-Encoders




Convolutional Auto-Encoders




Convolutional Auto-Encoders








Convolutional Auto-Encoders



Convolutional Auto-Encoders







```
def map_image_with_noise(image, label):  
    noise_factor = 0.5  
    image = tf.cast(image, dtype=tf.float32)  
    image = image / 255.0  
  
    factor = noise_factor * tf.random.normal(shape=image.shape)  
    image_noisy = image + factor  
    image_noisy = tf.clip_by_value(image_noisy, 0.0, 1.0)  
  
    return image_noisy, image
```

```
def map_image_with_noise(image, label):  
    noise_factor = 0.5  
    image = tf.cast(image, dtype=tf.float32)  
    image = image / 255.0  
  
    factor = noise_factor * tf.random.normal(shape=image.shape)  
    image_noisy = image + factor  
    image_noisy = tf.clip_by_value(image_noisy, 0.0, 1.0)  
  
    return image_noisy, image
```

```
def map_image_with_noise(image, label):  
    noise_factor = 0.5  
    image = tf.cast(image, dtype=tf.float32)  
    image = image / 255.0  
  
    factor = noise_factor * tf.random.normal(shape=image.shape)  
    image_noisy = image + factor  
    image_noisy = tf.clip_by_value(image_noisy, 0.0, 1.0)  
  
    return image_noisy, image
```

```
def map_image_with_noise(image, label):  
    noise_factor = 0.5  
    image = tf.cast(image, dtype=tf.float32)  
    image = image / 255.0  
  
    factor = noise_factor * tf.random.normal(shape=image.shape)  
    image_noisy = image + factor  
    image_noisy = tf.clip_by_value(image_noisy, 0.0, 1.0)  
  
    return image_noisy, image
```